

ACTOR PROGRAMMING FOR THE CLOUD

WHO AM I?

GUSTAVO PETRI

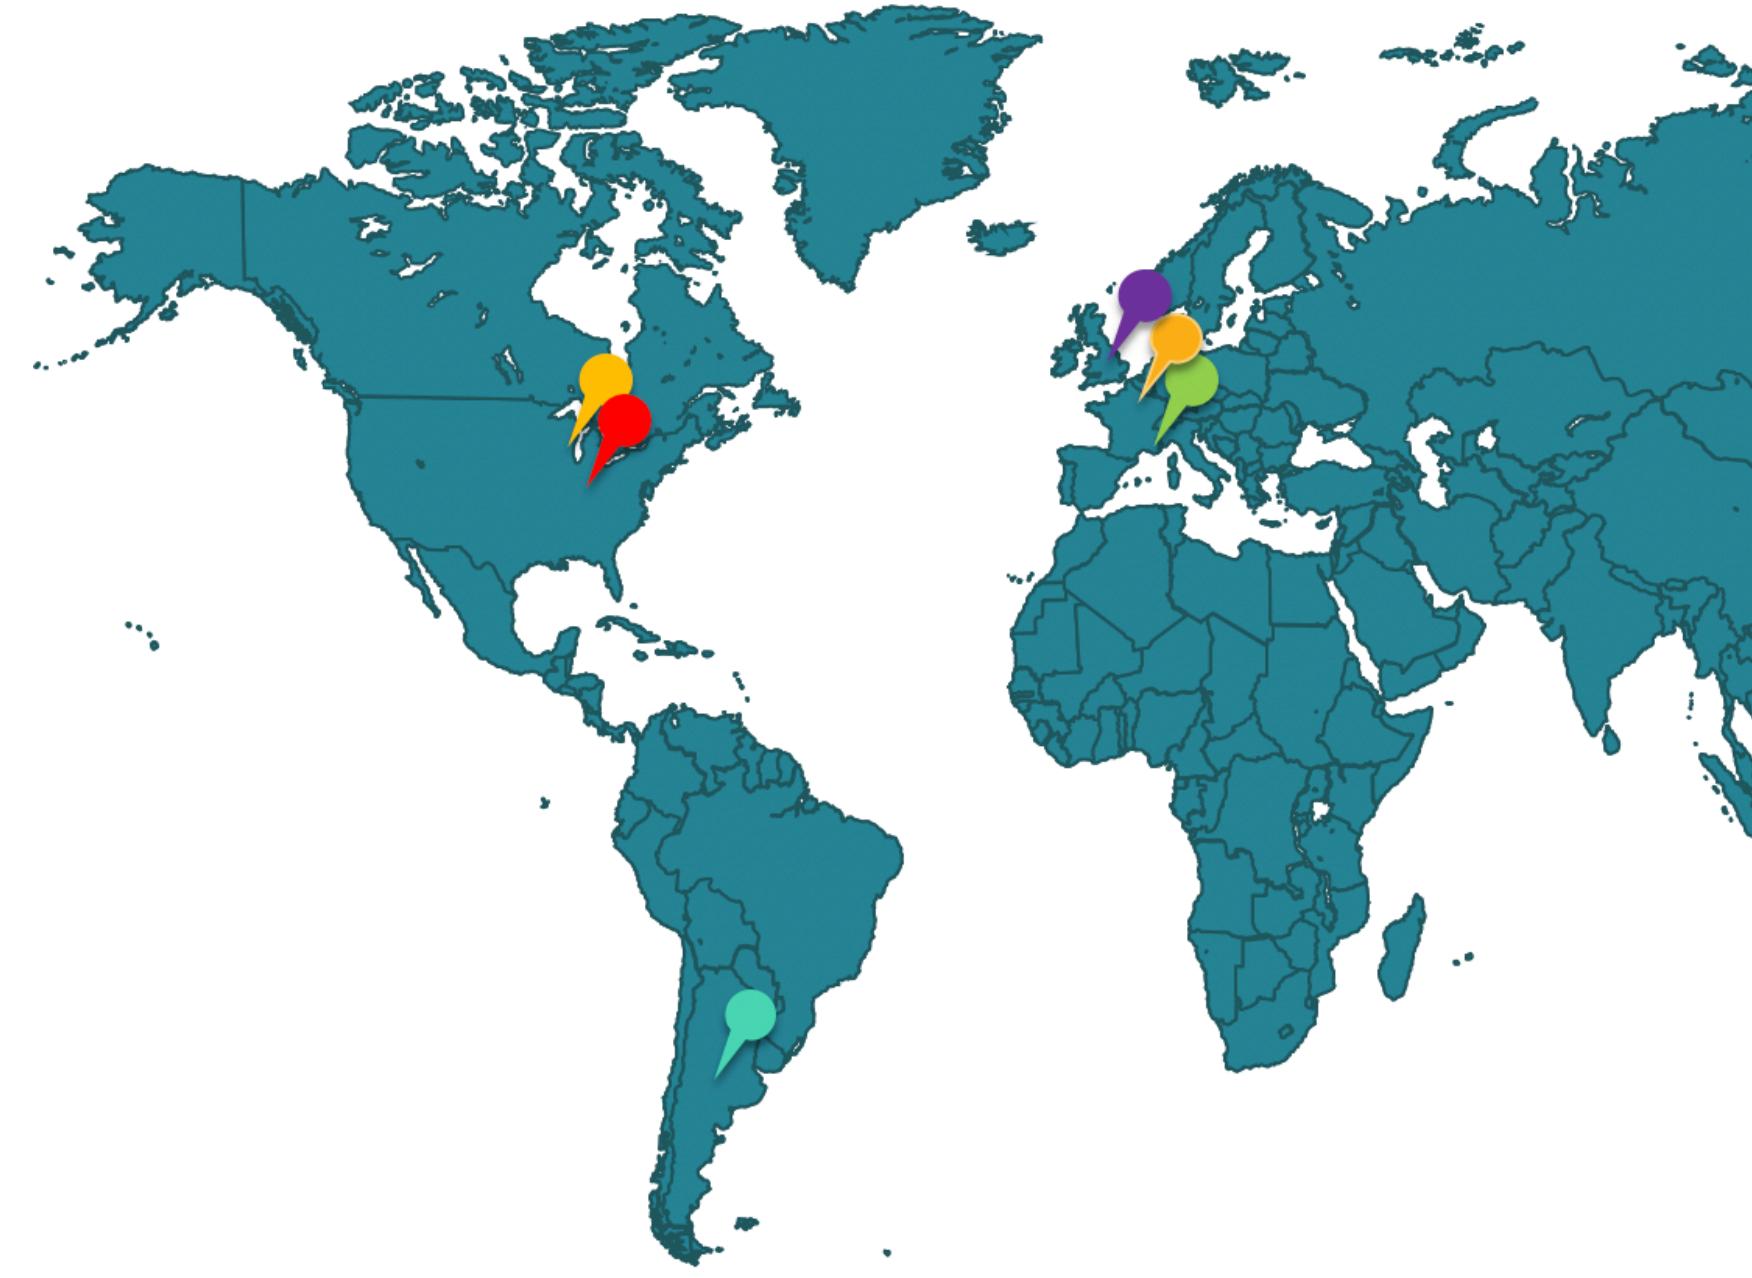
[2006 - 2010] PhD. Sophia Antipolis, *INRIA*, France

[2011 - 2012] PostDoc *DePaul University*, Chicago

[2012 - 2015] Visiting A. Professor *Purdue University*, Indiana

[2015 - 2018] Assistant Professor *IRIF - Universite Paris Diderot*, Paris

[2018 -] Formal Methods Researcher - *ARM Research*, Cambridge



WHO AM I?

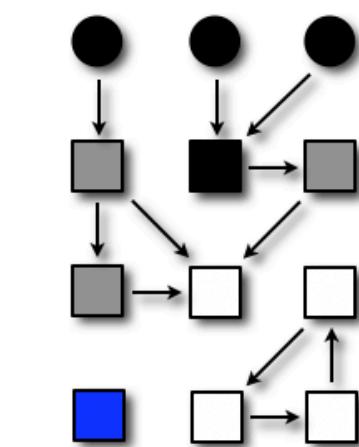
Topics

- ▶ Relaxed Memory Models
- ▶ Concurrent Program Verification
- ▶ Distributed Systems Verification
- ▶ Models for Distributed Computing
- ▶ and now ... Security

Tools

- ▶ Formal Semantics
- ▶ Formal Methods
- ▶ Programming Languages

Verifying a
Concurrent GC
(Rely/Guarantee)

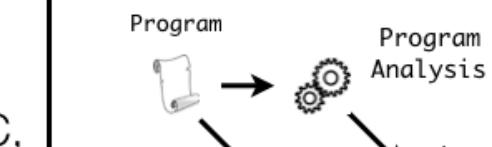


Theorem:
In any execution of the CGC,

Learning Shape
Specifications

$$\begin{aligned} \forall u, v, \nu : u \cup v \Rightarrow \\ (h1 : u \cup v \vee h1 : v \cup u \vee \\ h2 : u \cup v \vee h2 : v \cup u \vee \\ (h1 \rightarrow u \wedge h2 \rightarrow v) \vee \\ (h2 \rightarrow u \wedge h1 \rightarrow v)) \end{aligned}$$

CEGAR Loop



Elastic Cloud
Programming



Invariants for
Relaxed
Consistency

Invariant	Protocol
1Ob	Single Object
2Ob	Partial Order 2 objects
NOb	Equivalence Relations



WHENCE THIS COURSE?

- ▶ The way we write programs is changing
 - ▶ Move to mobile applications
 - ▶ Data availability
 - ▶ IoT, embedded, etc.
- ▶ Programming for the cloud
 - ▶ Distributed Systems
 - ▶ Elasticity
- ▶ Application Architecture



THE REACTIVE MANIFESTO (MICROSERVICES)

THE REACTIVE MANIFESTO

The Reactive Manifesto

Published on September 16 2014. (v2.0)

Organisations working in disparate domains are independently discovering patterns for building software that look the same. These systems are more robust, more resilient, more flexible and better positioned to meet modern demands.

These changes are happening because application requirements have changed dramatically in recent years. Only a few years ago a large application had tens of servers, seconds of response time, hours of offline maintenance and gigabytes of data. Today applications are deployed on everything from mobile devices to cloud-based clusters running thousands of multi-core processors. Users expect millisecond response times and 100% uptime. Data is measured in Petabytes. Today's demands are simply not met by yesterday's software architectures.

We believe that a coherent approach to systems architecture is needed, and we believe that all necessary aspects are already recognised individually: we want systems that are **Responsive**, **Resilient**, **Elastic** and **Message Driven**. We call these **Reactive Systems**.

Systems built as Reactive Systems are more flexible, loosely-coupled and scalable. This makes it easier to develop and amenable to change. They are significantly more tolerant of failure as they meet it with elegance rather than disaster. Reactive Systems provide effective interactive feedback.



THE REACTIVE MANIFESTO



Organisations working in disparate domains are independently discovering patterns of software that look the same. These systems are more robust, more resilient, more flexible and better positioned to meet modern demands.

These changes are happening because application requirements have changed dramatically in recent years. Only a few years ago a large application had tens of servers, seconds of response time, hours of offline maintenance and gigabytes of data. Today applications are deployed on everything from mobile devices to cloud-based clusters running thousands of multi-core processors. Users expect millisecond response times and 100% uptime. Data is measured in Petabytes. Today's demands are simply not met by yesterday's software architectures.

We believe that a coherent approach to systems architecture is needed, and we believe that all necessary aspects are already recognised individually: we want systems that are **Responsive**, **Resilient**, **Elastic** and **Message Driven**. We call these **Reactive Systems**.

Systems built as Reactive Systems are more flexible, loosely-coupled and scalable. This makes it easier to develop and amenable to change. They are significantly more tolerant of failure as they occur they meet it with elegance rather than disaster. Reactive Systems provide effective interactive feedback.



THE REACTIVE MANIFESTO

The Reactive Manifesto

System Requirements and Architecture for Modern Distributed Applications

Published on September 16, 2014. (v2.0)

Organisations working software that look the better positioned to n These changes are ha recent years. Only a time, hours of offli everything from m processors. Users Petabytes. Today' We believe that necessary aspect Resilient, Elast

Systems built as Reactive Systems are easier to develop and amenable to change. They occur they meet it with elegance rather than effective interactive feedback.

```
graph TD; A[Responsive] --> B[Resilient]; B --> C[MessageDriven]; C --> D[Elastic]; D --> A; A <--> B; C <--> D;
```



INFRASTRUCTURE THROUGH TIME

Typical infrastructure for a single application

INFRASTRUCTURE THROUGH TIME

Typical infrastructure for a single application

1970



INFRASTRUCTURE THROUGH TIME

Typical infrastructure for a single application

1970



1980



INFRASTRUCTURE THROUGH TIME

Typical infrastructure for a single application

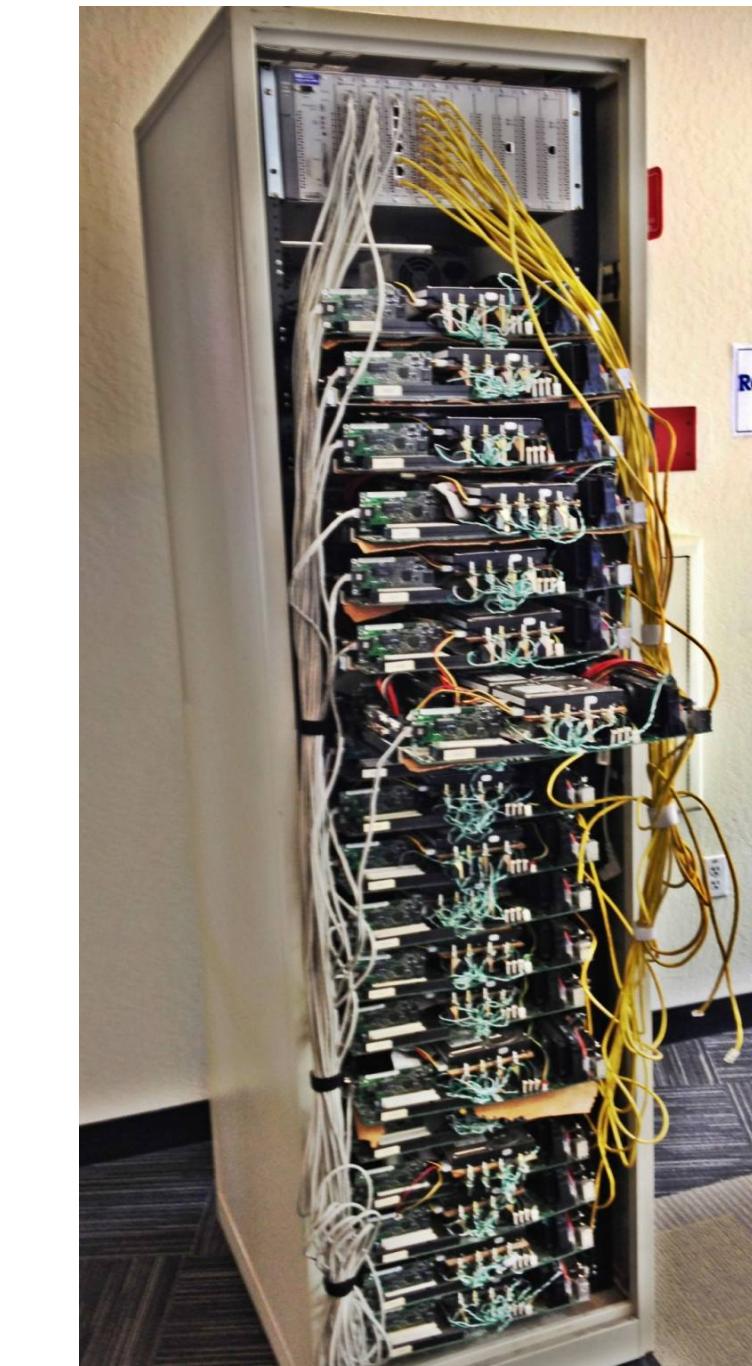
1970



1980



1990



INFRASTRUCTURE THROUGH TIME

Typical infrastructure for a single application

1970



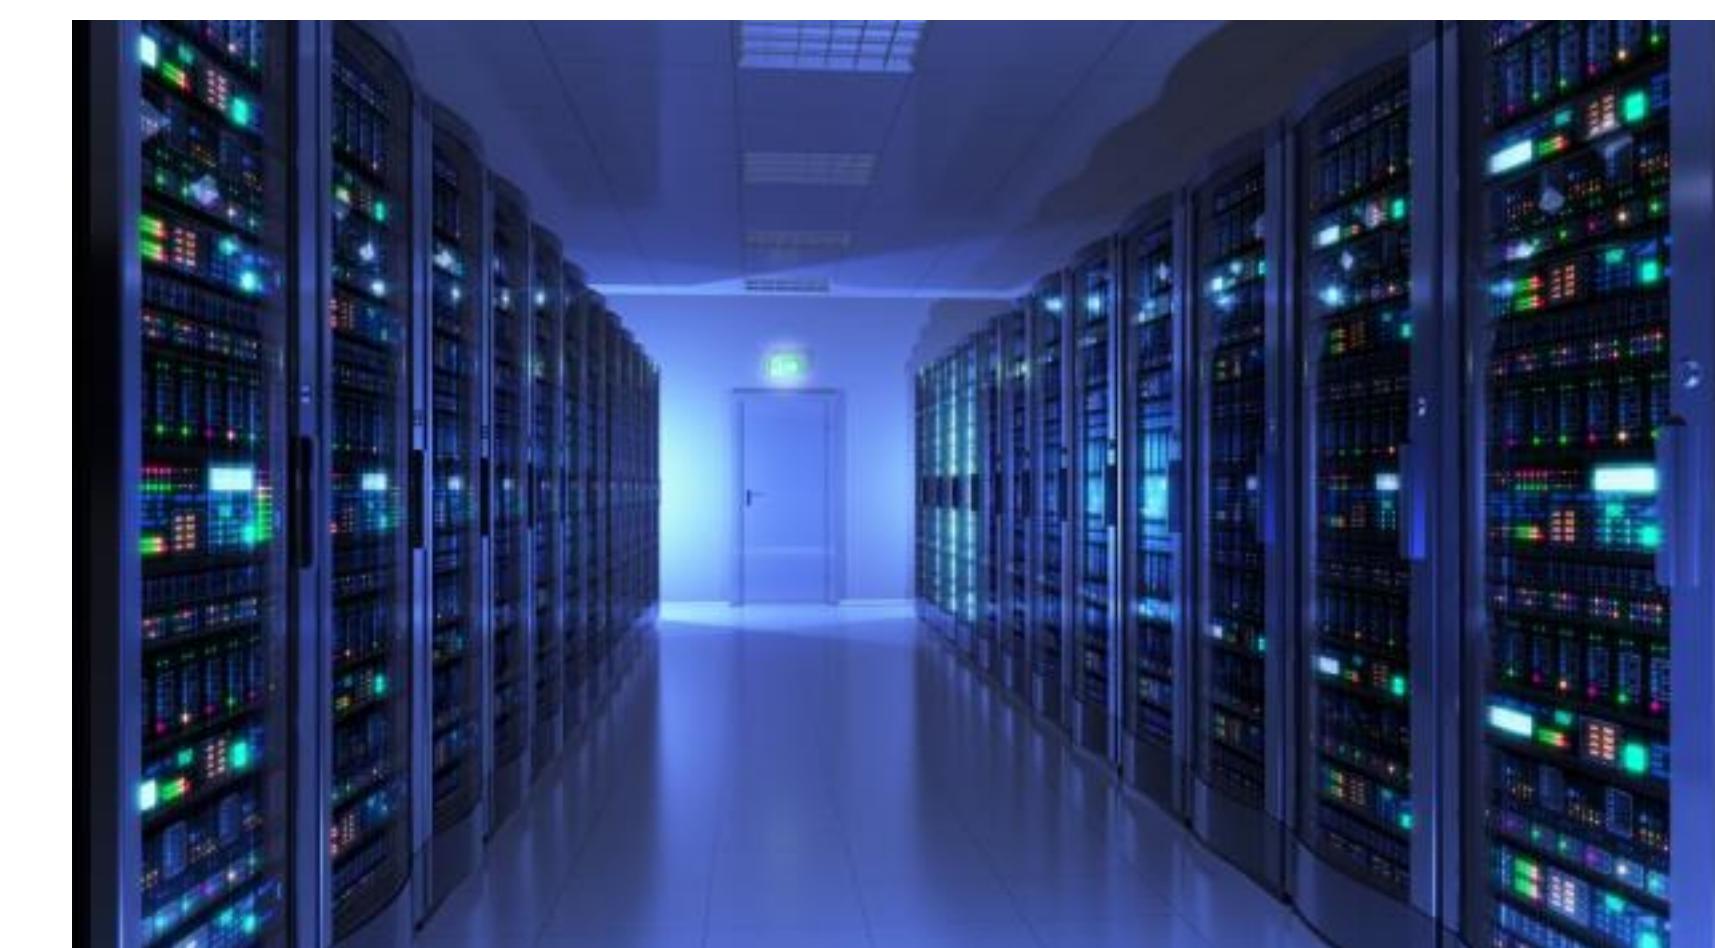
1980



1990



2010



THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible. Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible.

Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible. Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible. Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible. Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESPONSIVE

Responsive: The system responds in a timely manner if at all possible. Responsiveness is the cornerstone of usability and utility, but more than that, responsiveness means that problems may be detected quickly and dealt with effectively. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service. This consistent behaviour in turn simplifies error handling, builds end user confidence, and encourages further interaction.

THE REACTIVE MANIFESTO: RESILIENT

Resilient: The system stays responsive in the face of failure. This applies not only to highly-available, mission critical systems — any system that is not resilient will be unresponsive after a failure. Resilience is achieved by replication, containment, isolation and delegation. Failures are contained within each component, isolating components from each other and thereby ensuring that parts of the system can fail and recover without compromising the system as a whole. Recovery of each component is delegated to another (external) component and high-availability is ensured by replication where necessary. The client of a component is not burdened with handling its failures.

THE REACTIVE MANIFESTO: RESILIENT

Resilient: The system stays **responsive in the face of failure**. This applies not only to highly-available, mission critical systems — any system that is not resilient will be unresponsive after a failure. Resilience is achieved by replication, containment, isolation and delegation. Failures are contained within each component, isolating components from each other and thereby ensuring that parts of the system can fail and recover without compromising the system as a whole. Recovery of each component is delegated to another (external) component and high-availability is ensured by replication where necessary. The client of a component is not burdened with handling its failures.

THE REACTIVE MANIFESTO: RESILIENT

Resilient: The system stays **responsive in the face of failure**. This applies not only to highly-available, mission critical systems — any system that is not resilient will be unresponsive after a failure. Resilience is achieved by **replication, containment, isolation and delegation**. Failures are contained within each component, isolating components from each other and thereby ensuring that parts of the system can fail and recover without compromising the system as a whole. Recovery of each component is delegated to another (external) component and high-availability is ensured by replication where necessary. The client of a component is not burdened with handling its failures.

THE REACTIVE MANIFESTO: RESILIENT

Resilient: The system stays **responsive in the face of failure**. This applies not only to highly-available, mission critical systems — any system that is not resilient will be unresponsive after a failure. Resilience is achieved by **replication, containment, isolation and delegation**. Failures are **contained within each component**, isolating components from each other and thereby ensuring that parts of the system can fail and recover without compromising the system as a whole. Recovery of each component is delegated to another (external) component and high-availability is ensured by replication where necessary. The client of a component is not burdened with handling its failures.

THE REACTIVE MANIFESTO: ELASTIC

Elastic: The system stays responsive under varying workload. Reactive Systems can react to changes in the input rate by increasing or decreasing the resources allocated to service these inputs. This implies designs that have no contention points or central bottlenecks, resulting in the ability to shard or replicate components and distribute inputs among them. Reactive Systems support predictive, as well as Reactive, scaling algorithms by providing relevant live performance measures. They achieve elasticity in a cost-effective way on commodity hardware and software platforms.

THE REACTIVE MANIFESTO: ELASTIC

Elastic: The system stays **responsive under varying workload**. Reactive Systems can react to changes in the input rate by increasing or decreasing the resources allocated to service these inputs. This implies designs that have no contention points or central bottlenecks, resulting in the ability to shard or replicate components and distribute inputs among them. Reactive Systems support predictive, as well as Reactive, scaling algorithms by providing relevant live performance measures. They achieve elasticity in a cost-effective way on commodity hardware and software platforms.

THE REACTIVE MANIFESTO: ELASTIC

Elastic: The system stays responsive under varying workload. Reactive Systems can react to changes in the input rate by increasing or decreasing the resources allocated to service these inputs. This implies designs that have no contention points or central bottlenecks, resulting in the ability to shard or replicate components and distribute inputs among them. Reactive Systems support predictive, as well as Reactive, scaling algorithms by providing relevant live performance measures. They achieve elasticity in a cost-effective way on commodity hardware and software platforms.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying back-pressure when necessary. Location transparent messaging as a means of communication makes it possible for the management of failure to work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying back-pressure when necessary. Location transparent messaging as a means of communication makes it possible for the management of failure to work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying back-pressure when necessary. Location transparent messaging as a means of communication makes it possible for the management of failure to work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying back-pressure when necessary. Location transparent messaging as a means of communication makes it possible for the management of failure to work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying back-pressure when necessary. Location transparent messaging as a means of communication makes it possible for the management of failure to work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

THE REACTIVE MANIFESTO: MESSAGE DRIVEN

Message Driven: Reactive Systems rely on asynchronous message-passing to establish a boundary between components that ensures loose coupling, isolation, location transparency, and provides the means to delegate errors as messages.

Employing event-driven control by signal, back-pressure, and communication, we can work with the same constructs and semantics across a cluster or within a single host. Non-blocking communication allows recipients to only consume resources while active, leading to less system overhead.

How do we program for The Reactive Manifesto?

icity, and flow
and applying
means of
ork with the

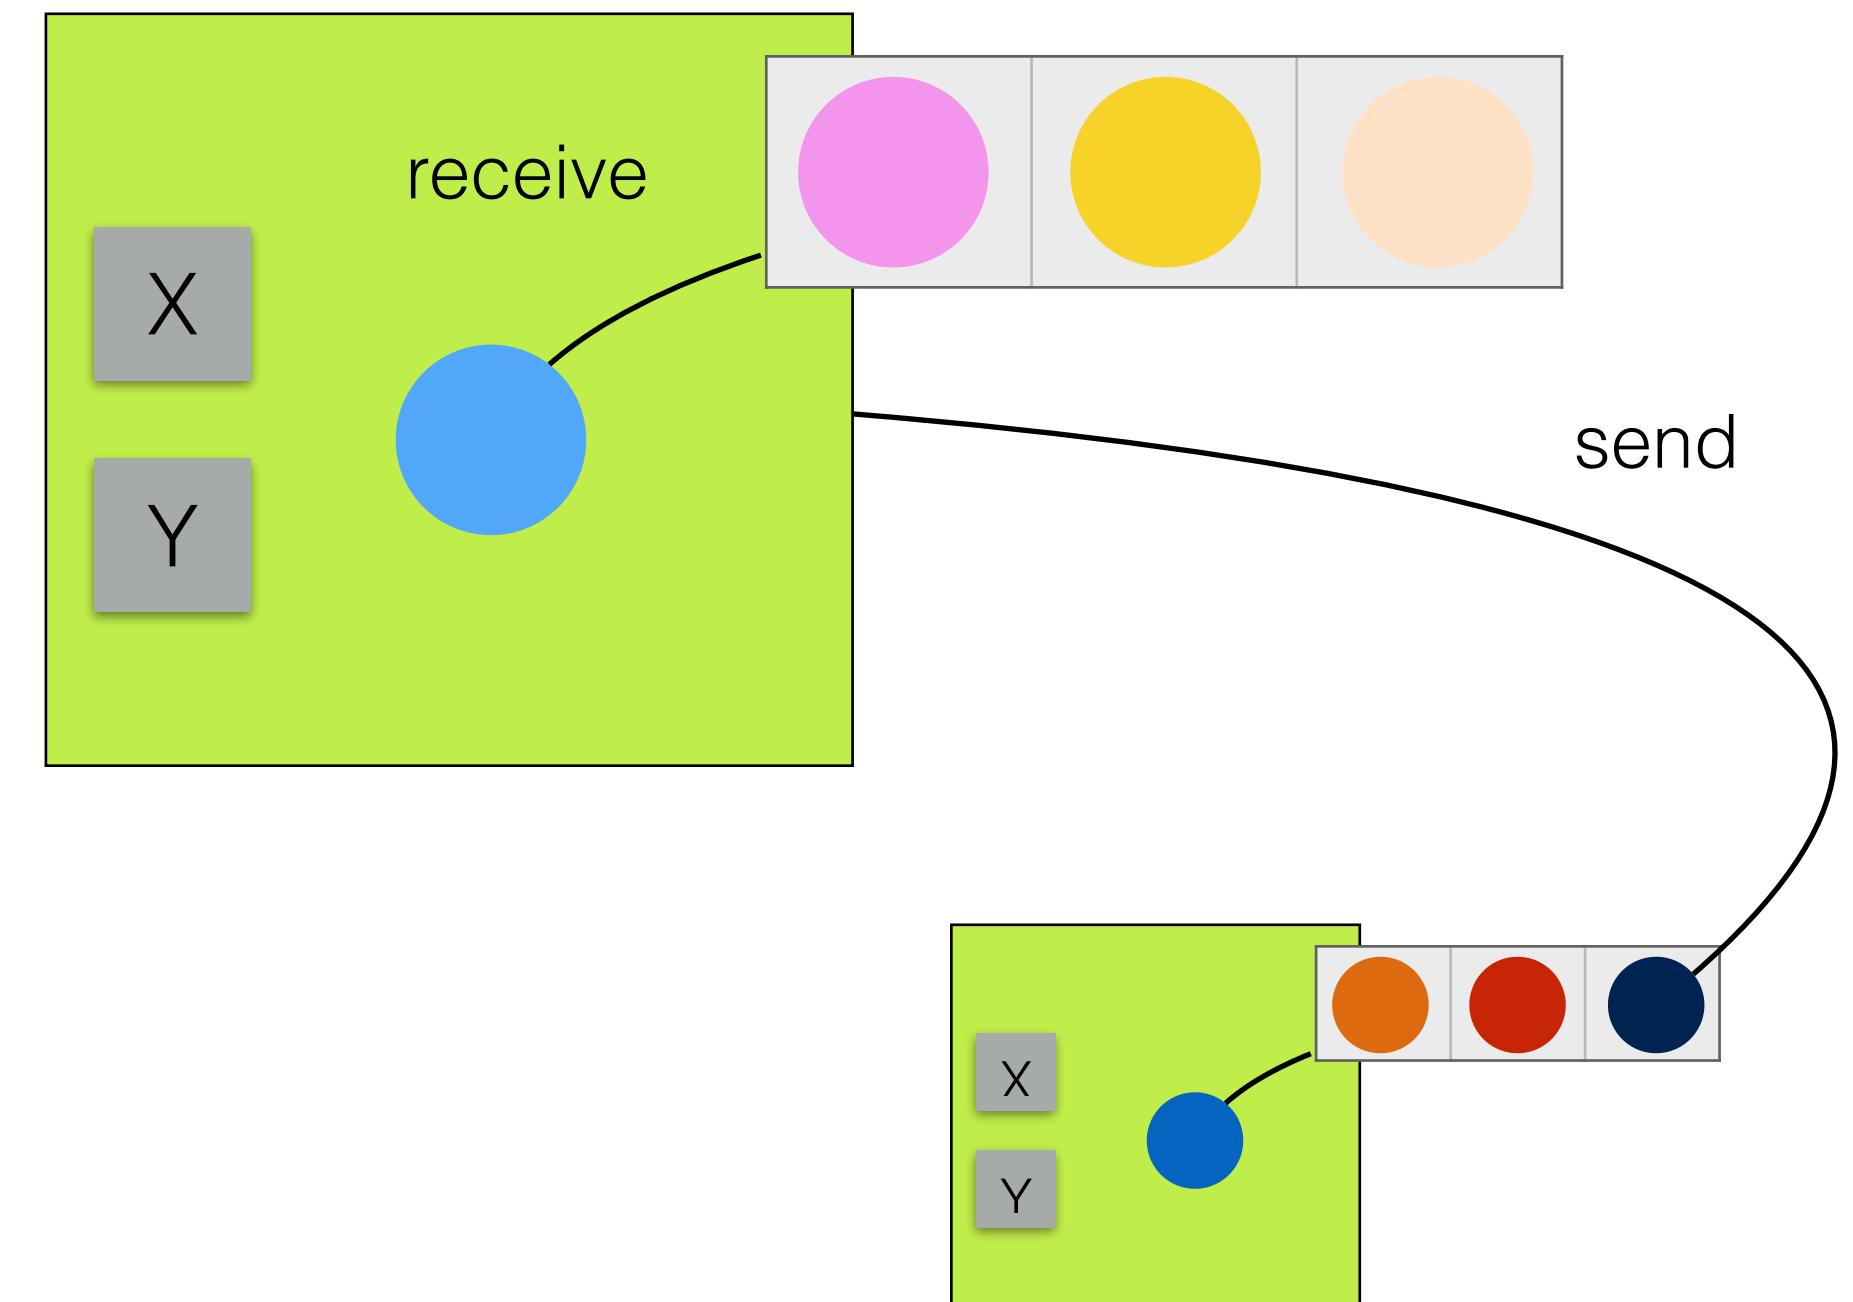
ACTOR PROGRAMMING

ACTOR EXECUTION MODEL

- ▶ Outline
 - ▶ What is an Actor
 - ▶ Actors vs. OO + Threads
 - ▶ Why Actors
 - ▶ Why Actors for Distributed Computing

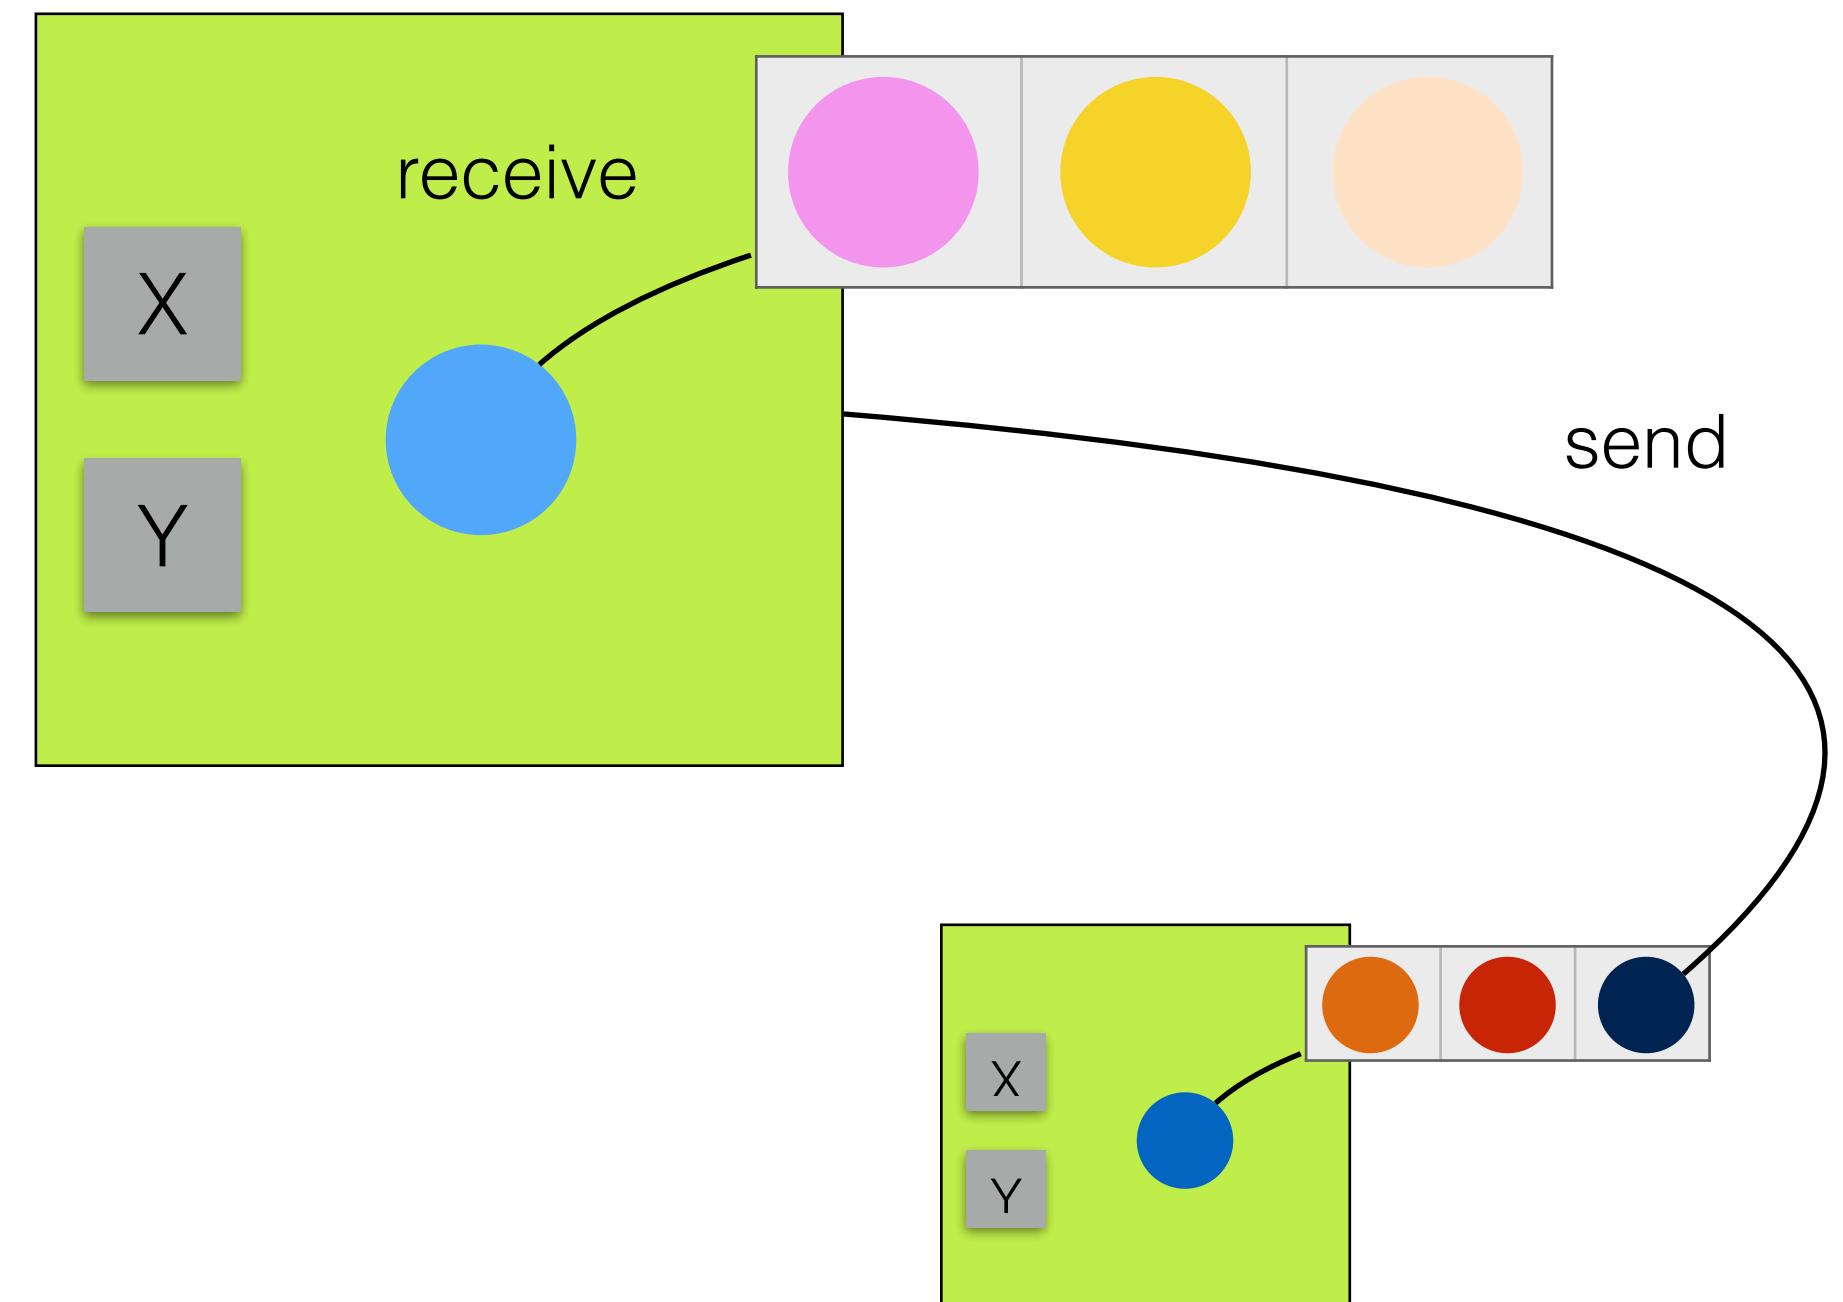
WHAT IS AN ACTOR

- ▶ Focus on Distribution (and parallelism)
- ▶ Sometimes Stateful
- ▶ Single-threaded
- ▶ Asynchronous
- ▶ Message Passing
- ▶ Messages go into a Mailbox



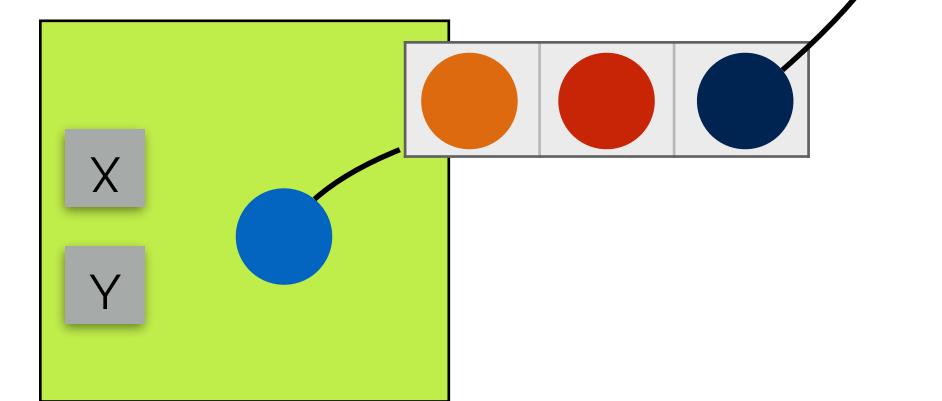
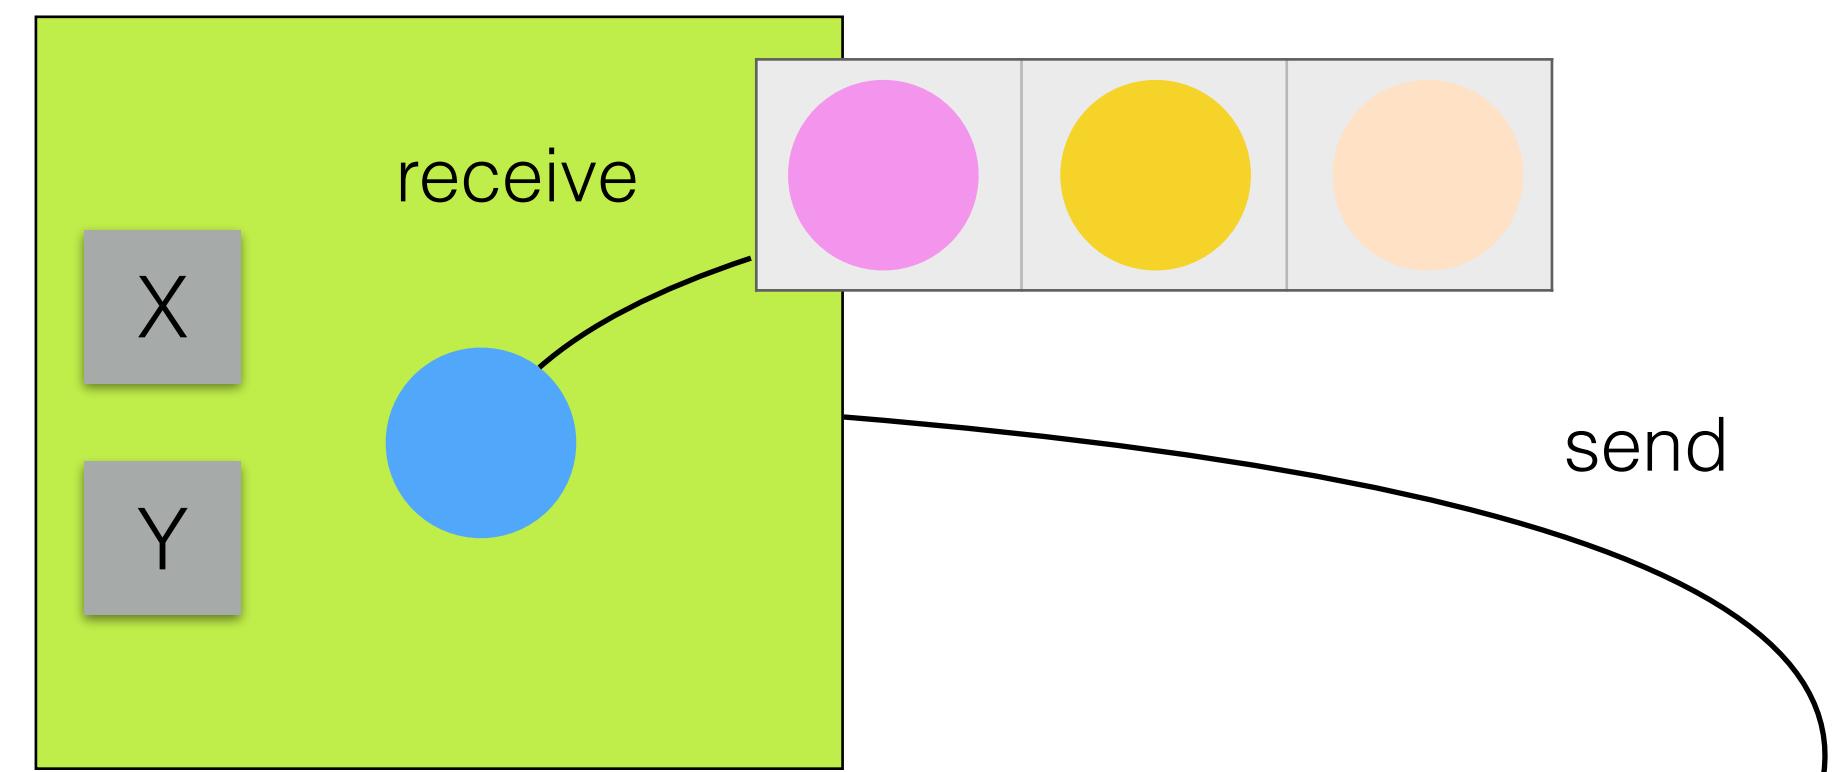
WHAT IS AN ACTOR

- ▶ Actor Reference
 - ▶ Information Hiding
 - ▶ Location Transparency
- ▶ Naming: ActorPath (public addresses)
- ▶ Local Actor State
 - ▶ No access from other actors
 - ▶ Interaction only through messages



WHAT IS AN ACTOR

- ▶ Behavior
 - ▶ Represented through message/response
 - ▶ Mailbox: FIFO? Priorities? Transactions?
 - ▶ Fault Tolerance: Supervision / Monitoring
 - ▶ Scalability
 - ▶ Elasticity
 - ▶ ...



ACTORS VS. OO

OO Programming

Actor Programming

ACTORS VS. OO

	OO Programming	Actor Programming
State	Public fields accessible to other objects	Encapsulated inside each actor

ACTORS VS. OO

	OO Programming	Actor Programming
State	Public fields accessible to other objects	Encapsulated inside each actor
Communication	Method invocation	Message passing

ACTORS VS. OO

	OO Programming	Actor Programming
State	Public fields accessible to other objects	Encapsulated inside each actor
Communication	Method invocation	Message passing
Scalability	Multi-threading Shared-Memory	Actors are independent processes

THOUGHT EXPERIMENT IN OO

- ▶ Consider developing an application that maintains state at runtime
- ▶ The application receives requests from clients on the internet

THOUGHT EXPERIMENT IN OO

- ▶ Consider developing an application that maintains state at runtime
- ▶ The application receives requests from clients on the internet
- ▶ But the workloads are unknown

THOUGHT EXPERIMENT IN OO

- ▶ Consider developing an application that maintains state at runtime
- ▶ The application receives requests from clients on the internet
- ▶ But the workloads are unknown
- ▶ Can we buy more machines to accommodate the workload?

THOUGHT EXPERIMENT IN OO

- ▶ Consider developing an application that maintains state at runtime
- ▶ The application receives requests from clients on the internet
- ▶ But the workloads are unknown
- ▶ Can we buy more machines to accommodate the workload?
- ▶ What happens if one machine fails?

THOUGHT EXPERIMENT IN OO

- ▶ Consider developing an application that maintains state at runtime
- ▶ The application receives requests from clients on the internet
- ▶ But the workloads are unknown
- ▶ Can we buy more machines to accommodate the workload?
- ▶ What happens if one machine fails?
- ▶ How can we architect this application?

OO + THREADS

- ▶ What is the underlying programming language Memory Model?

```
TS0.x = 0;
TS0.y = 0;
TS0.r1 = 0;
TS0.r2 = 0;

Thread wrXrdY =
    new Thread() {
        public void run() {
            // wait 1 millisecond
            try { Thread.sleep(1); }
            catch (Exception e) {System.exit(1);}
            // write x read y
            TS0.x = 1;
            TS0.r1 = TS0.y;
        }
};

Thread wrYrdX =
    new Thread() {
        public void run() {
            // wait 1 millisecond
            try { Thread.sleep(1); }
            catch (Exception e) {System.exit(1);}
            // write y read x
            TS0.y = 1;
            TS0.r2 = TS0.x;
        }
};
```

$$\begin{array}{ll} x = y = 0 \\ x := 1; & \parallel & y := 1; \\ r_0 = y & & r_1 = x \\ r_0 = r_1 = 0 \end{array}$$

Demo

ACTORS AS CONCURRENCY CONTROL

- ▶ Actors enforce a crude [and rather specific] form of Concurrency Control
 - ▶ Each actor's state is manipulated by strictly one process [conceptually]
 - ▶ No races are possible at the individual actor level
- ▶ Well designed actors can fail independently
 - ▶ “Let it crash” model (later)

PROGRAMMING WITH ACTORS

- ▶ How can we structure the application into Actors?

PROGRAMMING WITH ACTORS

- ▶ How can we structure the application into Actors?
- ▶ Each actor performs a small and specific function

PROGRAMMING WITH ACTORS

- ▶ How can we structure the application into Actors?
 - ▶ Each actor performs a small and specific function
 - ▶ Actors are created and terminated upon request

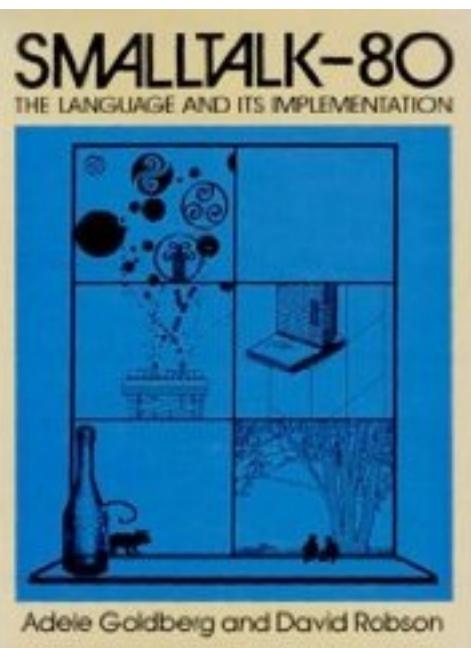
PROGRAMMING WITH ACTORS

- ▶ How can we structure the application into Actors?
 - ▶ Each actor performs a small and specific function
 - ▶ Actors are created and terminated upon request
 - ▶ Actors are supervised, and faults are treated where necessary

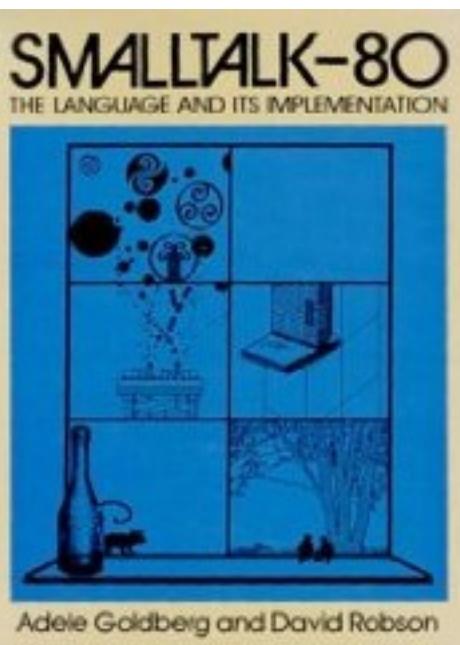
PROGRAMMING WITH ACTORS

- ▶ How can we structure the application into Actors?
 - ▶ Each actor performs a small and specific function
 - ▶ Actors are created and terminated upon request
 - ▶ Actors are supervised, and faults are treated where necessary
 - ▶ Some long-lived actors serve as services entry points

A NOTE ON SMALLTALK

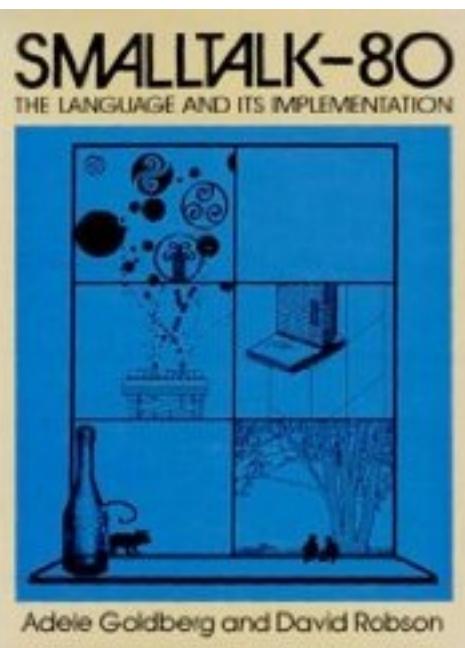


A NOTE ON SMALLTALK



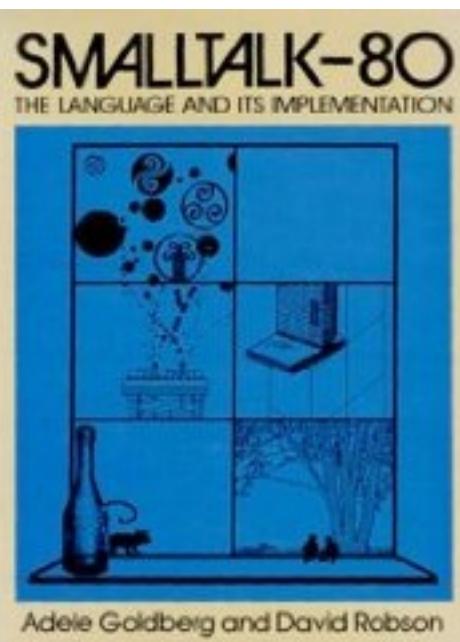
- ▶ Created in the '70 by Alan Kay and other
 - ▶ Introduces Objects
“abstract representation of a concrete entity”
 - ▶ Everything is an object
 - ▶ Interaction with objects is through message/response
 - ▶ Methods bind messages to functionality
 - ▶ Method binding is dynamic

A NOTE ON SMALLTALK



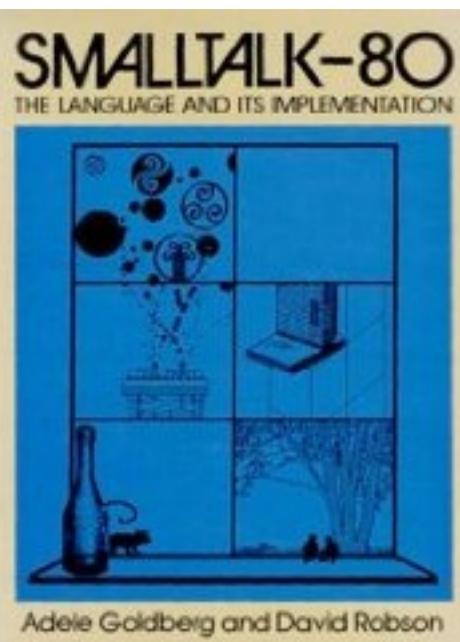
- ▶ Created in the '70 by Alan Kay and other
 - ▶ Introduces Objects
“abstract representation of a concrete entity”
 - ▶ Everything is an object
 - ▶ Interaction with objects is through message/response
 - ▶ Methods bind messages to functionality
 - ▶ Method binding is dynamic
- ▶ Classes define the representation of objects
 - ▶ Class Protocol
(messages that the object accepts – API)
 - ▶ Inheritance
 - ▶ Code reuse
 - ▶ Generalization
 - ▶ Specializatoin

A NOTE ON SMALLTALK



- ▶ Created in the '70 by Alan Kay and other
- ▶ Introduces Objects “abstract representation of a concrete entity”
- ▶ Everything is an object
- ▶ Interaction with objects is through message/response
- ▶ Methods bind messages to functionality
- ▶ Method binding is dynamic
- ▶ Classes define the representation of objects
- ▶ Class Protocol (messages that the object accepts – API)
- ▶ Inheritance
- ▶ Code reuse
- ▶ Generalization
- ▶ Specializatoin
- ▶ Purist Object Orientation
 - ▶ Everything is an object
 - ▶ No conditionals: expressed as messages ifTrue, ifFalse
 - ▶ The intent was to construct libraries for all expectable objects out there
 - ▶ Inspired the OO languages we know and love :-)

A NOTE ON SMALLTALK



- ▶ Created in the '70 by Alan Kay and other
- ▶ Introduces Objects “abstract representation of a concrete entity”
- ▶ Everything is an object
- ▶ Interaction with objects through message-response
- ▶ Classes define the representation of objects
- ▶ Class Protocol
- ▶ Code reuse
- ▶ Generalization
- ▶ Specializatoin
- ▶ Purist Object Orientation
 - ▶ Everything is an object
 - ▶ No conditionals:
 - ↳ ed as messages
 - ↳ False
 - ▶ was to construct for all foreseeable objects out there
 - ▶ Inspired the OO languages we know and love :-)

Actors borrow a lot from
this view of OO

A SNAPSHOT OF ERLANG



A SNAPSHOT OF ERLANG



- ▶ Actors due to Hewitt,
Bishop & Steiger '73

A SNAPSHOT OF ERLANG



- ▶ Actors due to Hewitt,
Bishop & Steiger '73
- ▶ Erlang: Joe Armstrong
(Ericson) '86
- ▶ Functional Language
- ▶ Pattern Matching
- ▶ Immutable State
- ▶ Everything is a Process

A SNAPSHOT OF ERLANG



- ▶ Actors due to Hewitt,
Bishop & Steiger '73
- ▶ Key characteristics
 - ▶ Distributed
 - ▶ Fault-tolerant
 - ▶ Soft real-time
 - ▶ Highly available
 - ▶ Hot swapping
- ▶ Erlang: Joe Armstrong
(Ericson) '86
- ▶ Functional Language
- ▶ Pattern Matching
- ▶ Immutable State
- ▶ Everything is a Process

A SNAPSHOT OF ERLANG



- ▶ Actors due to Hewitt, Bishop & Steiger '73
- ▶ Erlang: Joe Armstrong (Ericson) '86
- ▶ Functional Language
- ▶ Pattern Matching
- ▶ Immutable State
- ▶ Everything is a Process
- ▶ Key characteristics
 - ▶ Distributed
 - ▶ Fault-tolerant
 - ▶ Soft real-time
 - ▶ Highly available
 - ▶ Hot swapping
- ▶ Popularizes Actors:
 - ▶ No shared state
 - ▶ Lightweight processes
 - ▶ Asynch message-passing
 - ▶ Mailboxes messages

ERLANG'S HELLO WORLD



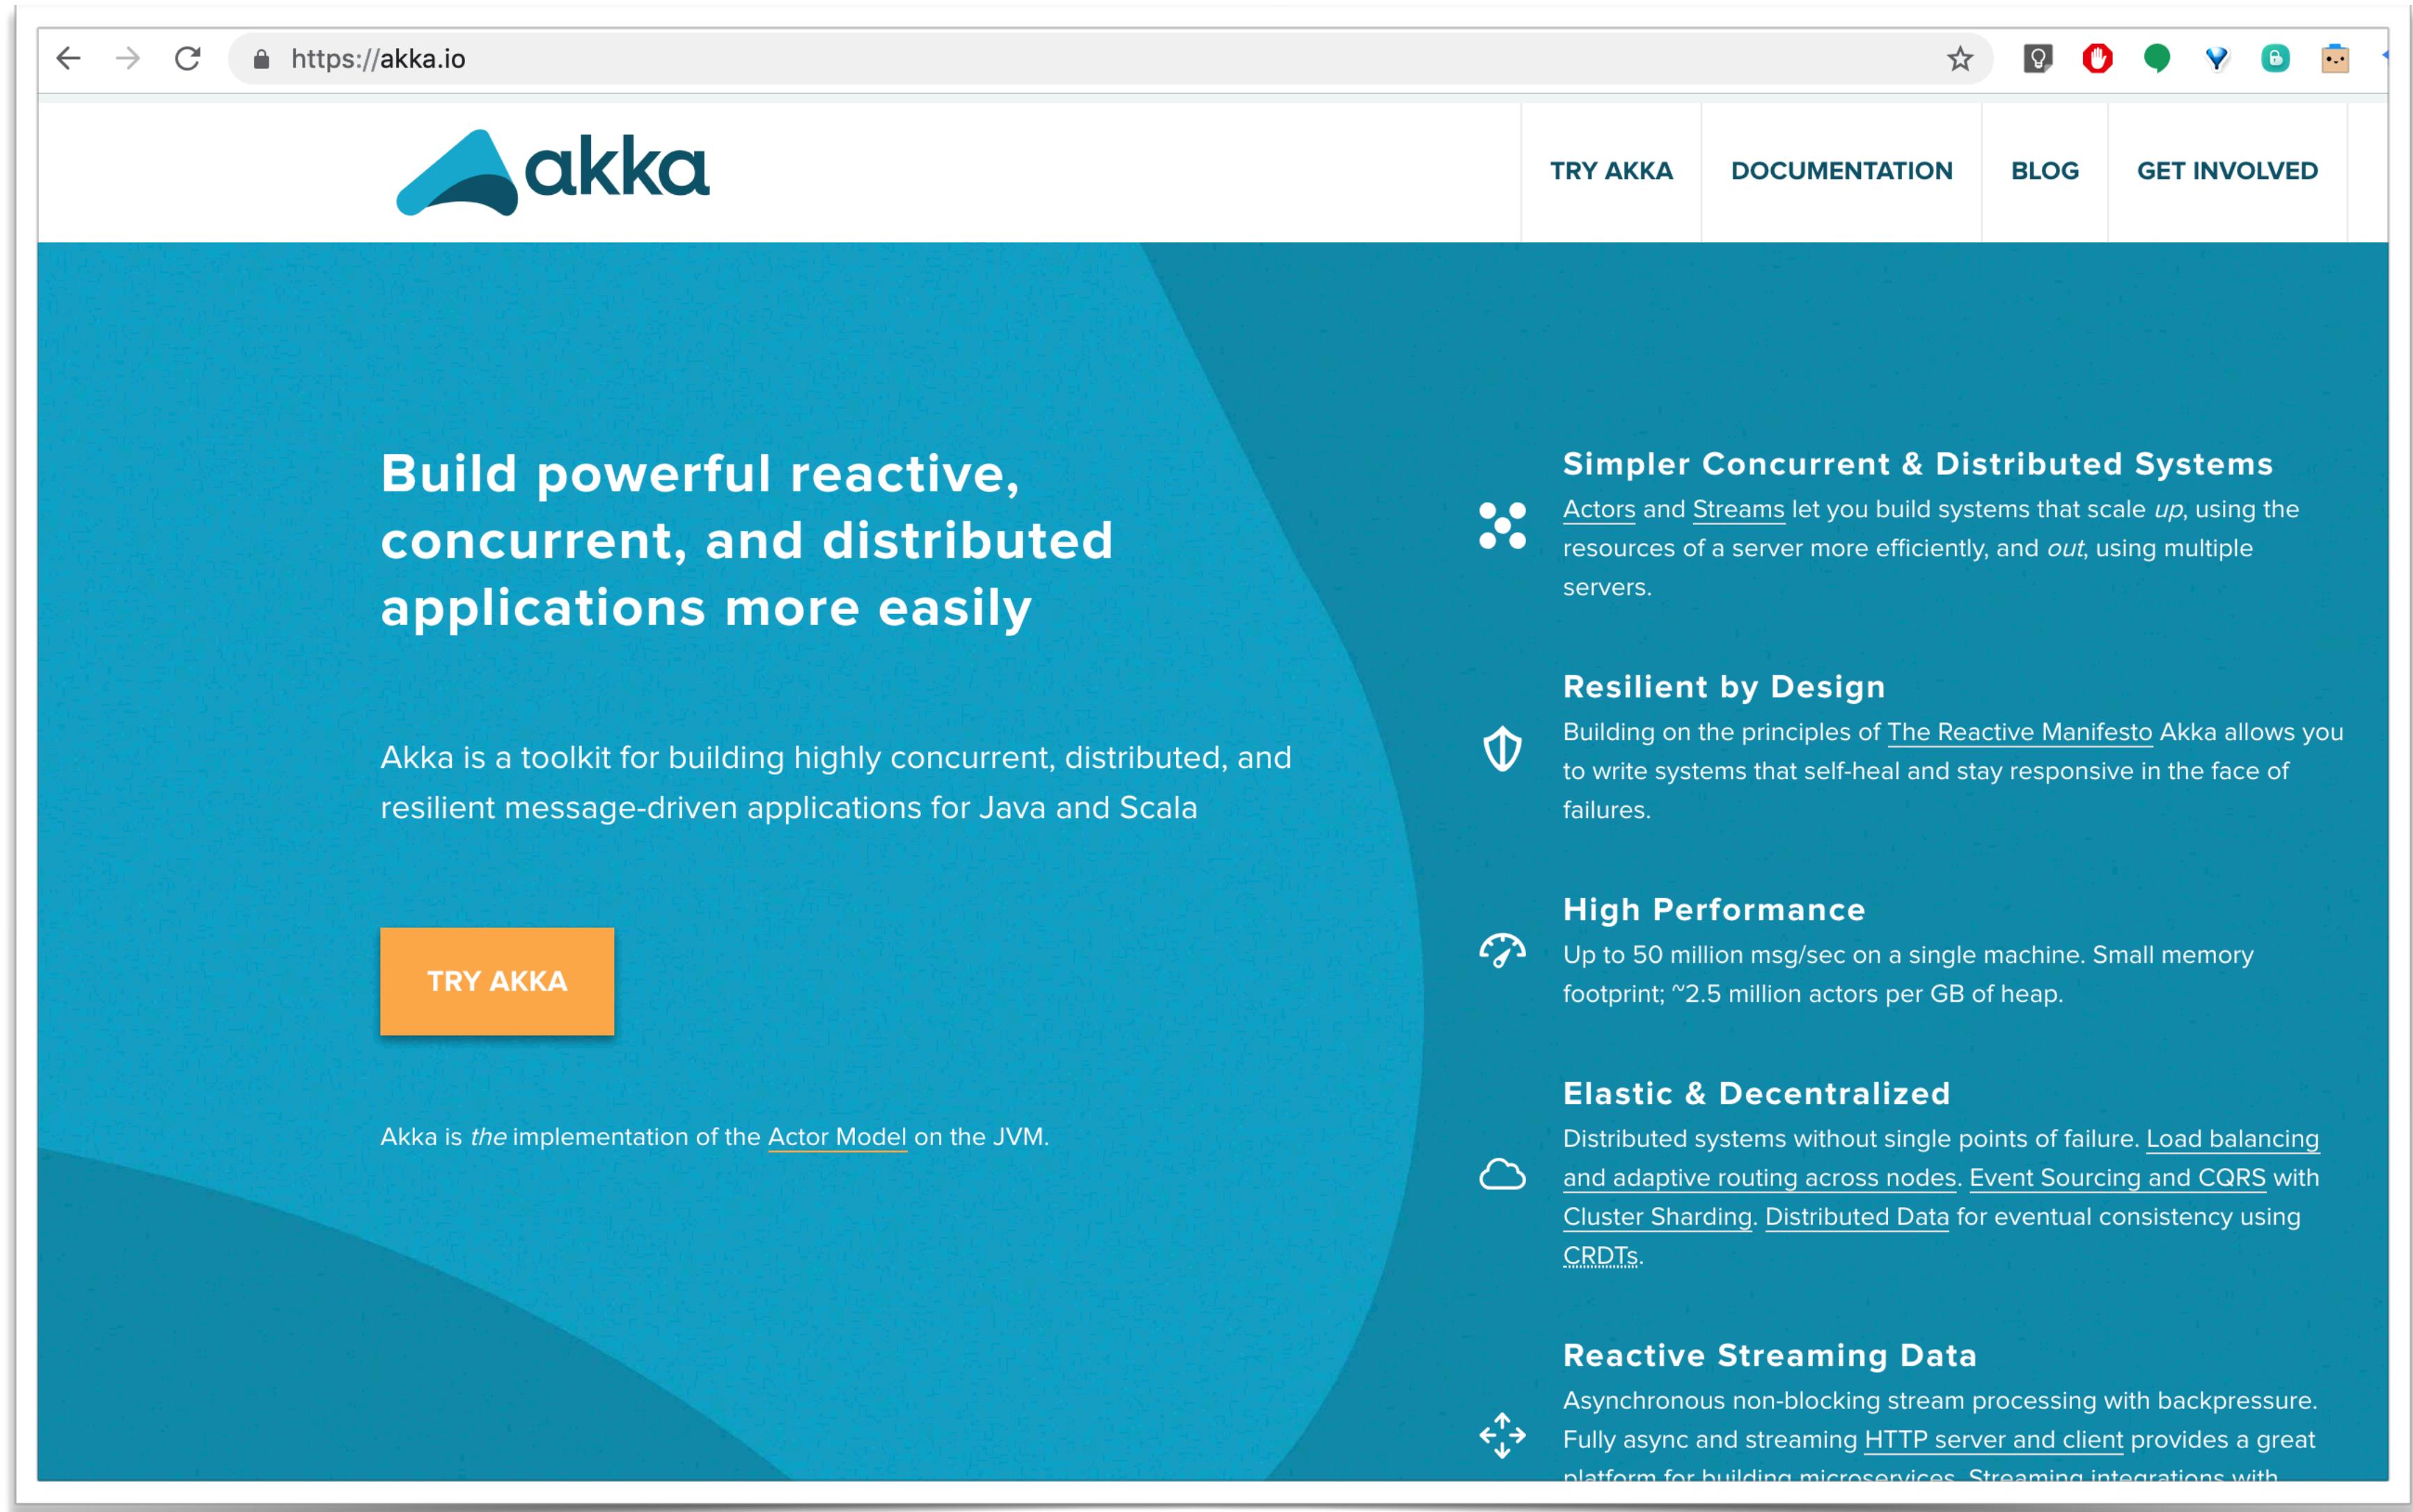
```
ping(0, PongNode) ->
    {pong, PongNode} ! finished,
    io:format("ping finished~n", []);
ping(N, PongNode) ->
    {pong, PongNode} ! {ping, self()},
    receive
        pong ->
            io:format("Ping received pong~n", [])
    end,
    ping(N - 1, PongNode).

pong() ->
    receive
        finished ->
            io:format("Pong finished~n", []);
        {ping, Ping_Node} ->
            io:format("Pong received ping~n", []),
            Ping_Node ! pong,
            pong()
    end.
```



Demo

AKKA ACTORS



The screenshot shows the official Akka website at <https://akka.io>. The page features a large blue header with the Akka logo and navigation links for TRY AKKA, DOCUMENTATION, BLOG, and GET INVOLVED. Below the header, a large white section contains the text "Build powerful reactive, concurrent, and distributed applications more easily". A smaller text below it states: "Akka is a toolkit for building highly concurrent, distributed, and resilient message-driven applications for Java and Scala". A prominent orange "TRY AKKA" button is located in the bottom left of this section. At the bottom of this white area, there's a note: "Akka is *the* implementation of the [Actor Model](#) on the JVM." To the right of this white section, there are five columns, each representing a key feature of Akka:

- Simpler Concurrent & Distributed Systems**: Accompanied by a gear icon, this section discusses how Akka's Actors and Streams allow for efficient scaling up and out.
- Resilient by Design**: Accompanied by a shield icon, this section explains how Akka systems self-heal and stay responsive in the face of failures, based on the Reactive Manifesto principles.
- High Performance**: Accompanied by a speedometer icon, this section highlights Akka's ability to handle up to 50 million msg/sec on a single machine with a small memory footprint.
- Elastic & Decentralized**: Accompanied by a cloud icon, this section describes Akka's distributed nature without single points of failure, mentioning load balancing, adaptive routing, event sourcing, CQRS, cluster sharding, and distributed data.
- Reactive Streaming Data**: Accompanied by a double-headed arrow icon, this section covers Akka's asynchronous stream processing with backpressure, its use as an HTTP server and client for microservices, and streaming integrations.

AKKA Actors (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies

AKKA Actors (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies
- ▶ Tuesday
 - ▶ Distribution
 - ▶ Scalability / Elasticity
 - ▶ Some Actor Patterns

AKKA ACTORS (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies
- ▶ Tuesday
 - ▶ Distribution
 - ▶ Scalability / Elasticity
 - ▶ Some Actor Patterns
- ▶ Wednesday
 - ▶ Cloud Computing
 - ▶ Virtualization
 - ▶ Akka on a cluster

AKKA ACTORS (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies
- ▶ Tuesday
 - ▶ Distribution
 - ▶ Scalability / Elasticity
 - ▶ Some Actor Patterns
- ▶ Wednesday
 - ▶ Cloud Computing
 - ▶ Virtualization
 - ▶ Akka on a cluster
- ▶ Thursday
 - ▶ Synchronization and State
 - ▶ Sharing Data
 - ▶ Cluster Orchestration

AKKA ACTORS (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies
- ▶ Tuesday
 - ▶ Distribution
 - ▶ Scalability / Elasticity
 - ▶ Some Actor Patterns
- ▶ Wednesday
 - ▶ Cloud Computing
 - ▶ Virtualization
 - ▶ Akka on a cluster
- ▶ Thursday
 - ▶ Synchronization and State
 - ▶ Sharing Data
 - ▶ Cluster Orchestration
- ▶ Friday
 - ▶ Orleans / AEON
 - ▶ Transactions
 - ▶ Lambda?

AKKA ACTORS (OUTLINE)

- ▶ Monday
 - ▶ Basics of Akka Programming
 - ▶ Fault Tolerance
 - ▶ Some Scala goodies
- ▶ Tuesday
 - ▶ Distribution
 - ▶ Scalability / Elasticity
 - ▶ Some Actor Patterns
- ▶ Wednesday
 - ▶ Cloud Computing
 - ▶ Virtualization
 - ▶ Akka on a cluster
- ▶ Thursday
 - ▶ Synchronization and State
 - ▶ Sharing Data
 - ▶ Cluster Orchestration
- ▶ Friday
 - ▶ Orleans / AEON
 - ▶ Transactions
 - ▶ Lambda?

AKKA ACTORS

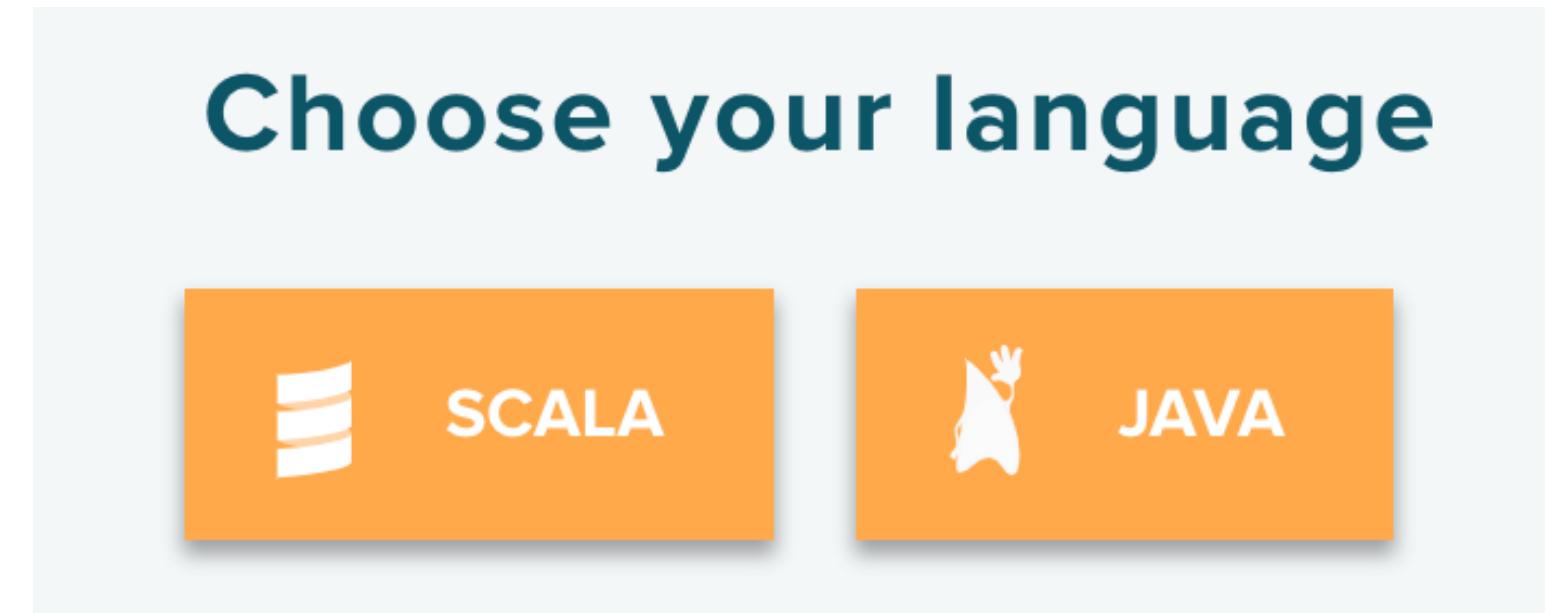


- ▶ We will be using Scala for the lectures, but ...

AKKA ACTORS



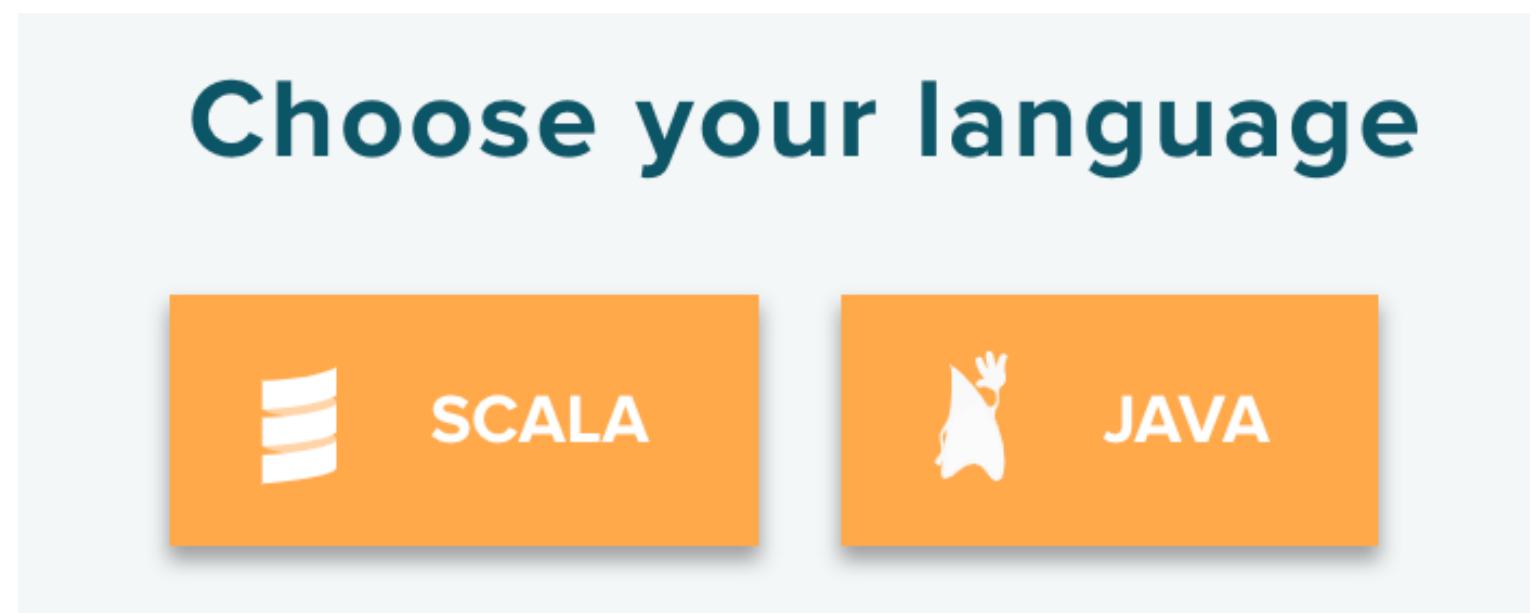
- We will be using Scala for the lectures, but ...



AKKA ACTORS



- ▶ We will be using Scala for the lectures, but ...

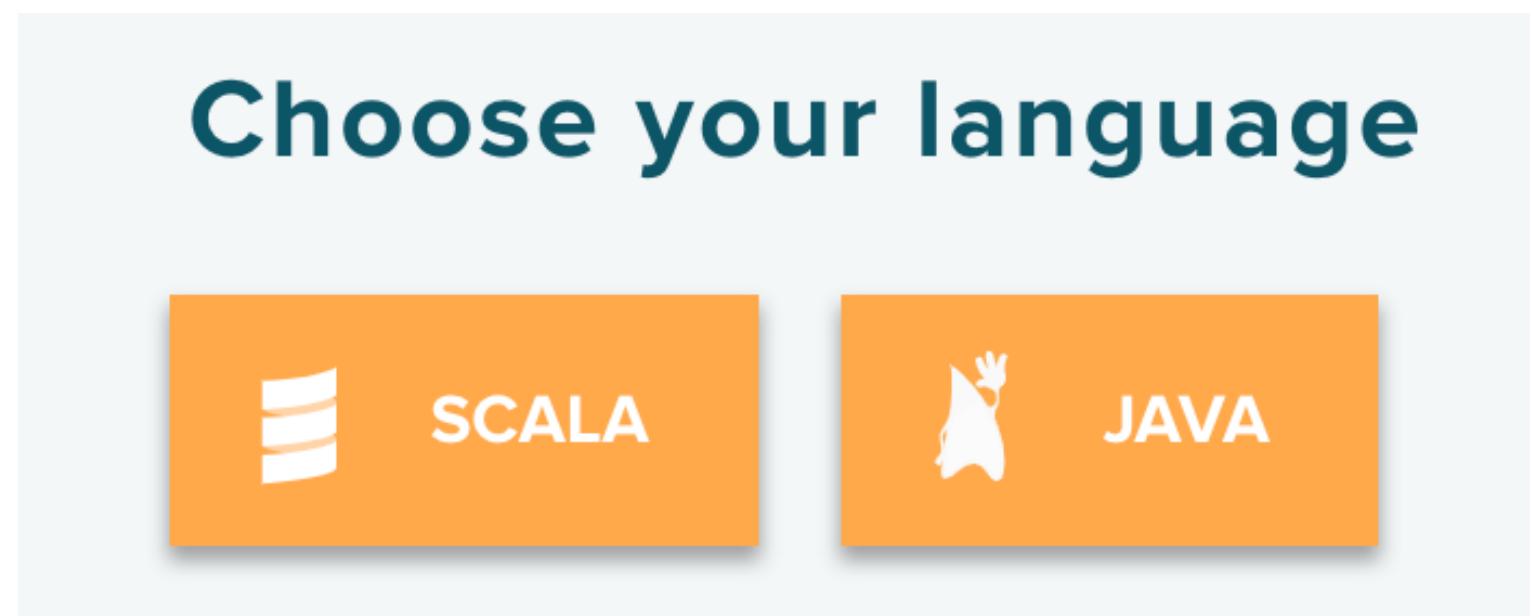


- ▶ The lectures are not Scala specific

AKKA ACTORS



- ▶ We will be using Scala for the lectures, but ...

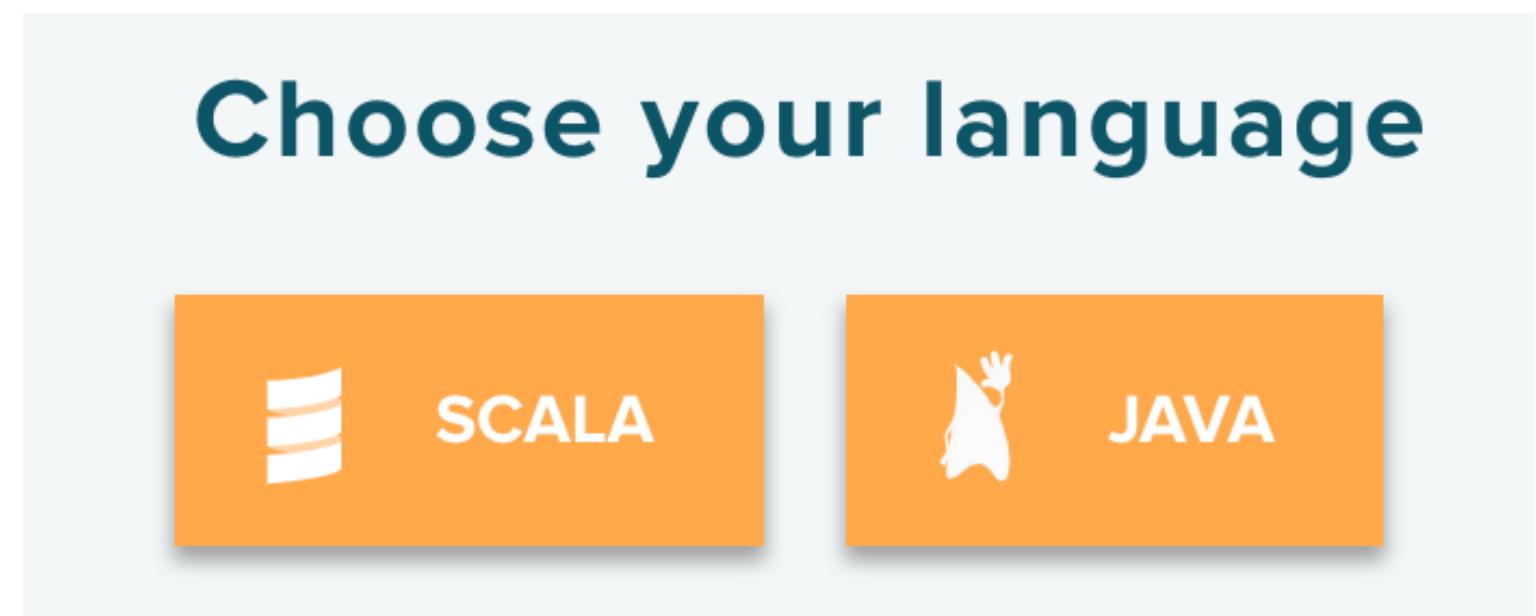


- ▶ The lectures are not Scala specific
- ▶ The documentation of Akka is great! Go try it out!

AKKA ACTORS



- ▶ We will be using Scala for the lectures, but ...

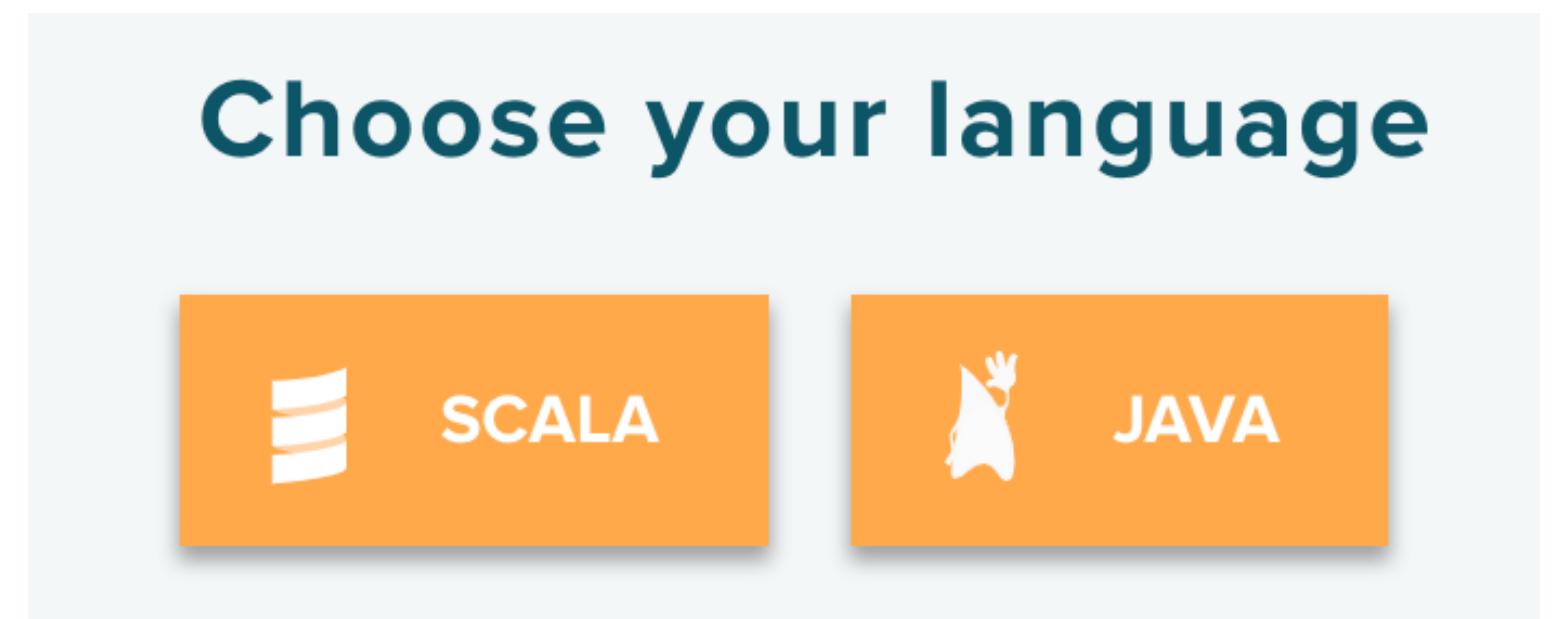


- ▶ The lectures are not Scala specific
- ▶ The documentation of Akka is great! Go try it out!
- ▶ There will be a small project at the end of the course

AKKA ACTORS



- ▶ We will be using Scala for the lectures, but ...

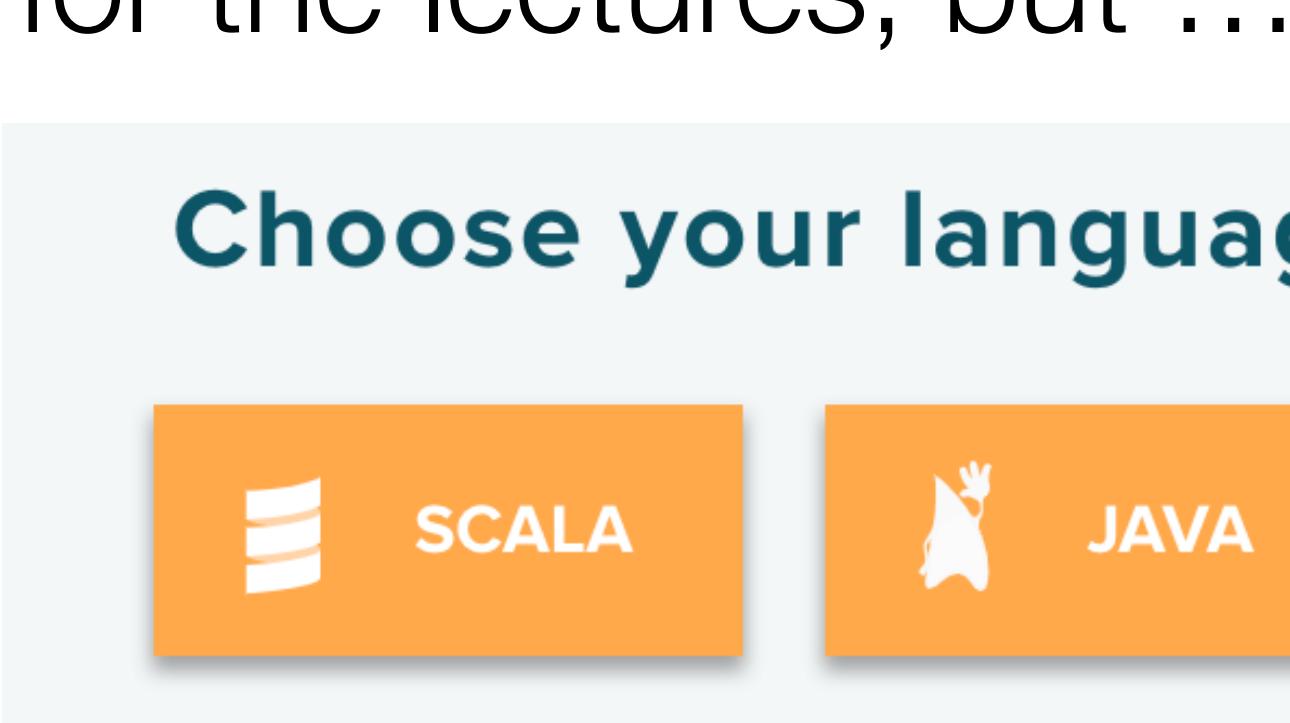


- ▶ The lectures are not Scala specific
- ▶ The documentation of Akka is great! Go try it out!
- ▶ There will be a small project at the end of the course
- ▶ Start reading the documentation of Akka soon!

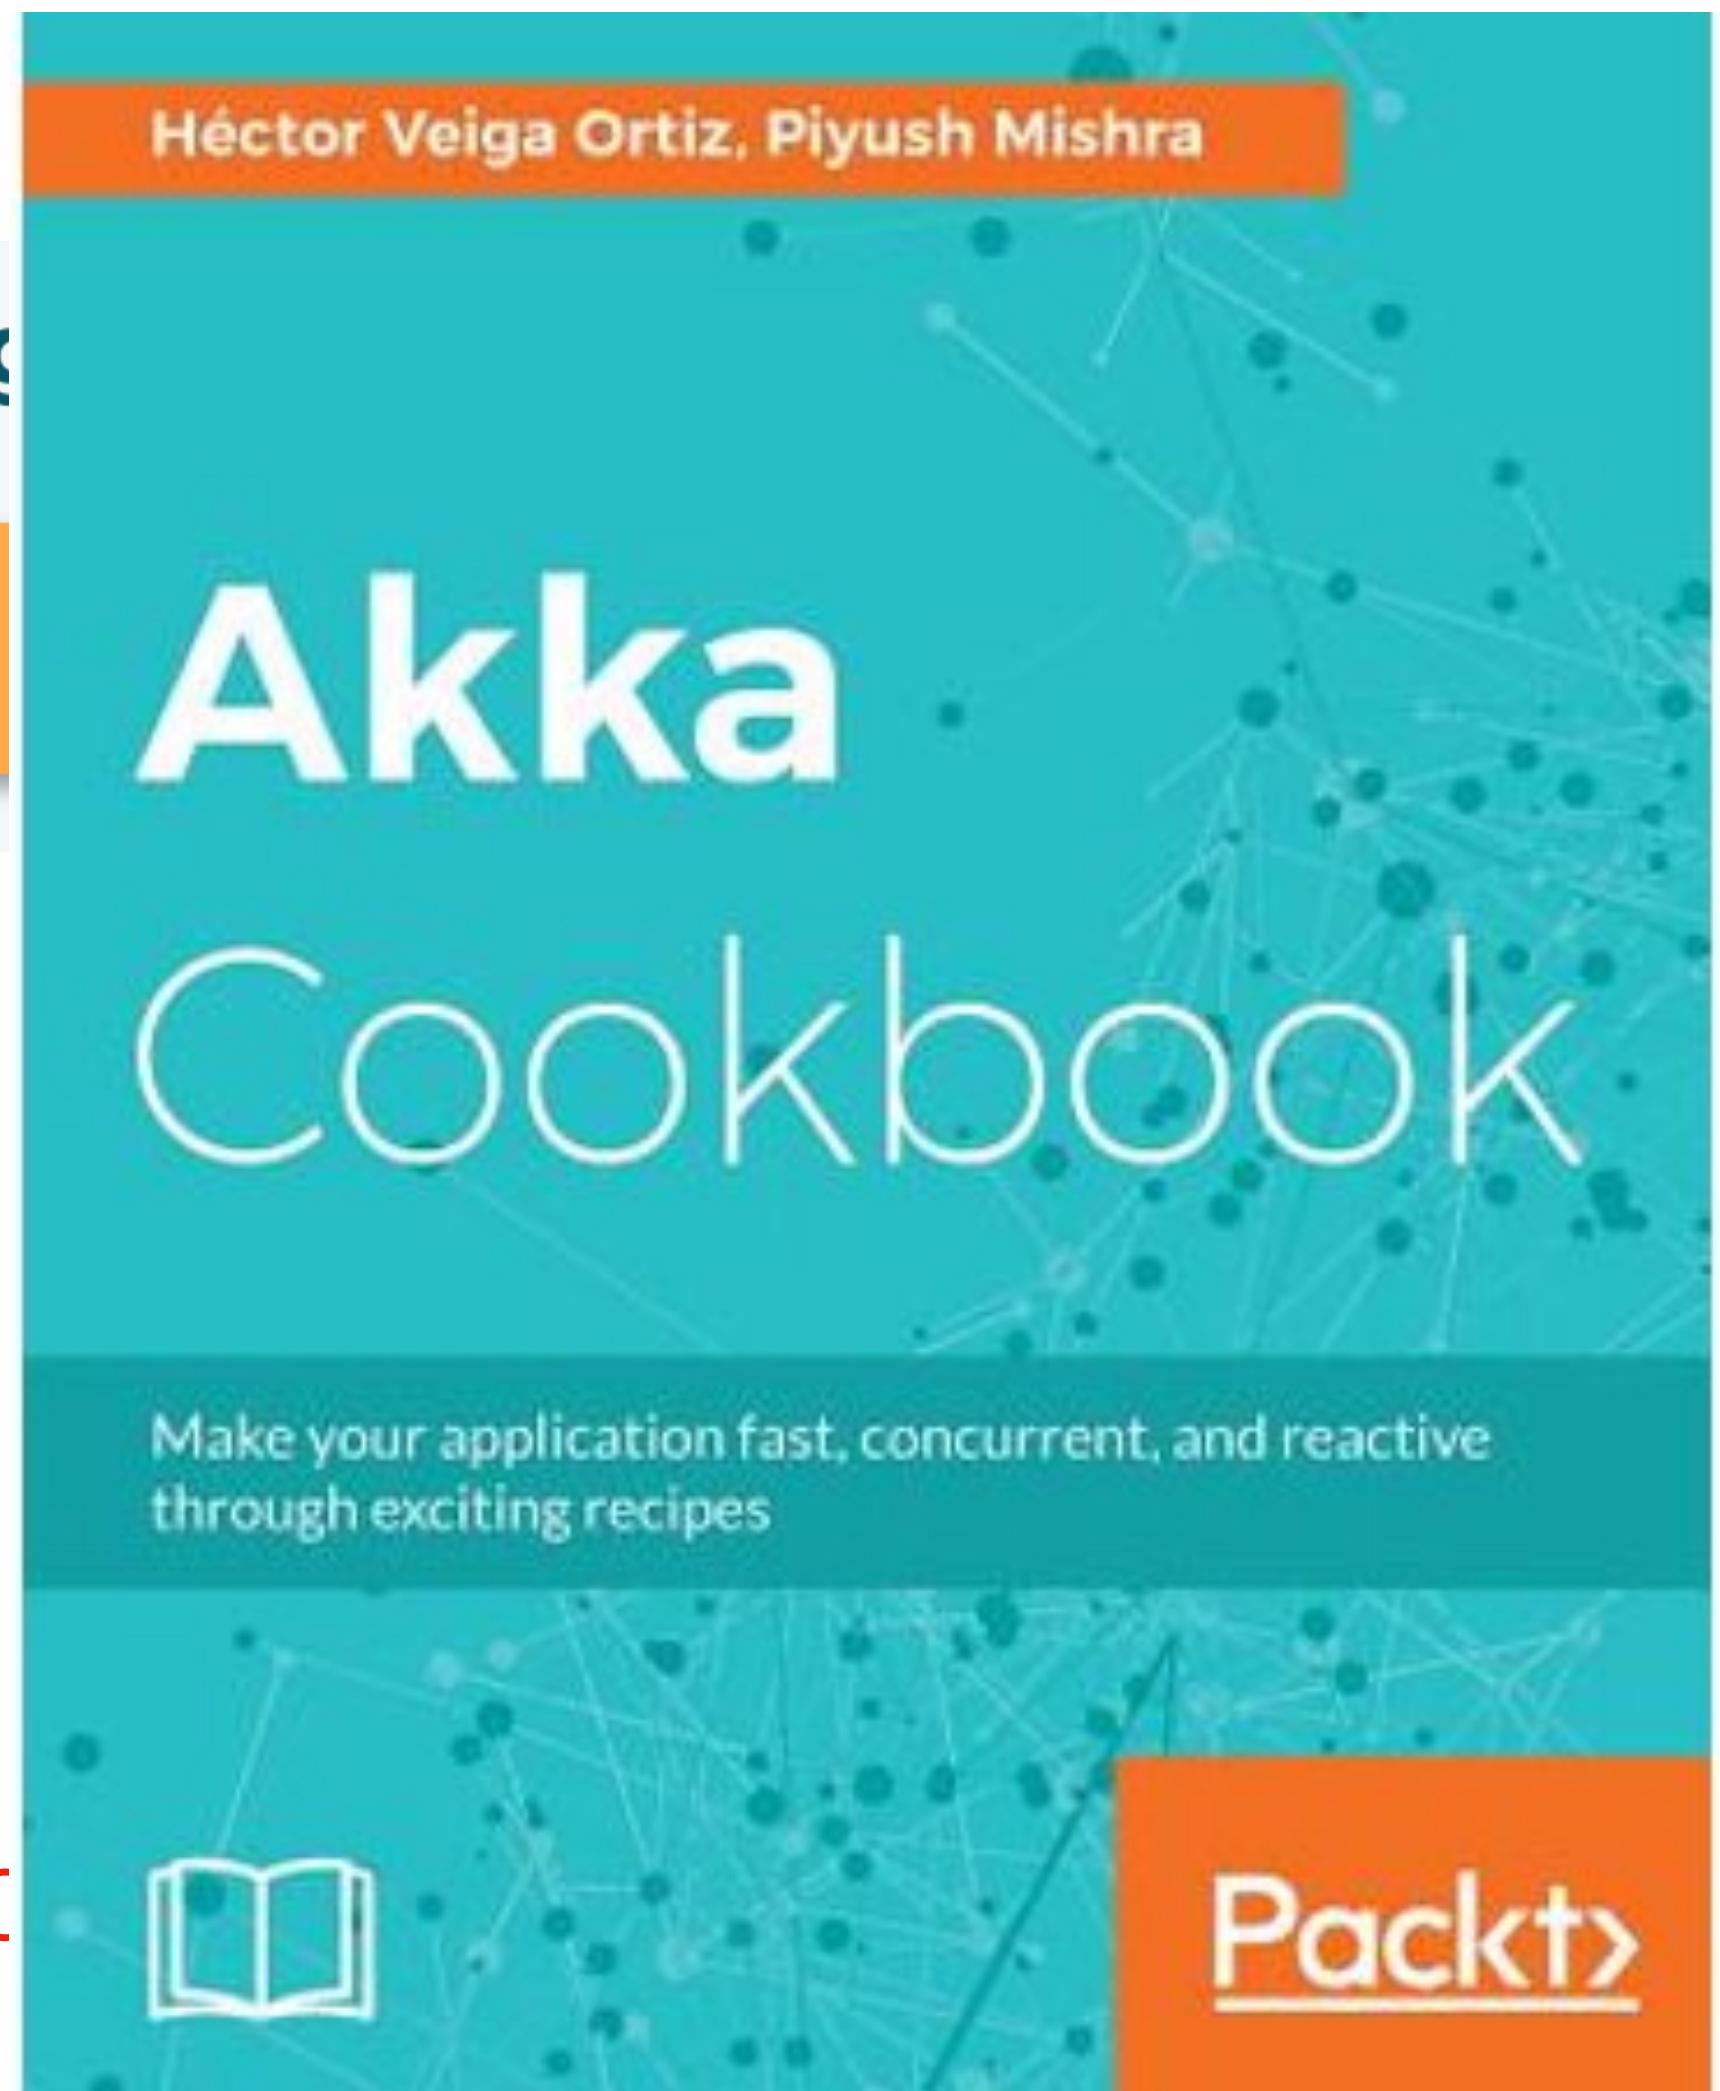
AKKA ACTORS



- ▶ We will be using Scala for the lectures, but ...



- ▶ The lectures are not Scala specific
- ▶ The documentation of Akka is great! Go try it
- ▶ There will be a small project at the end of the
- ▶ Start reading the documentation of Akka so

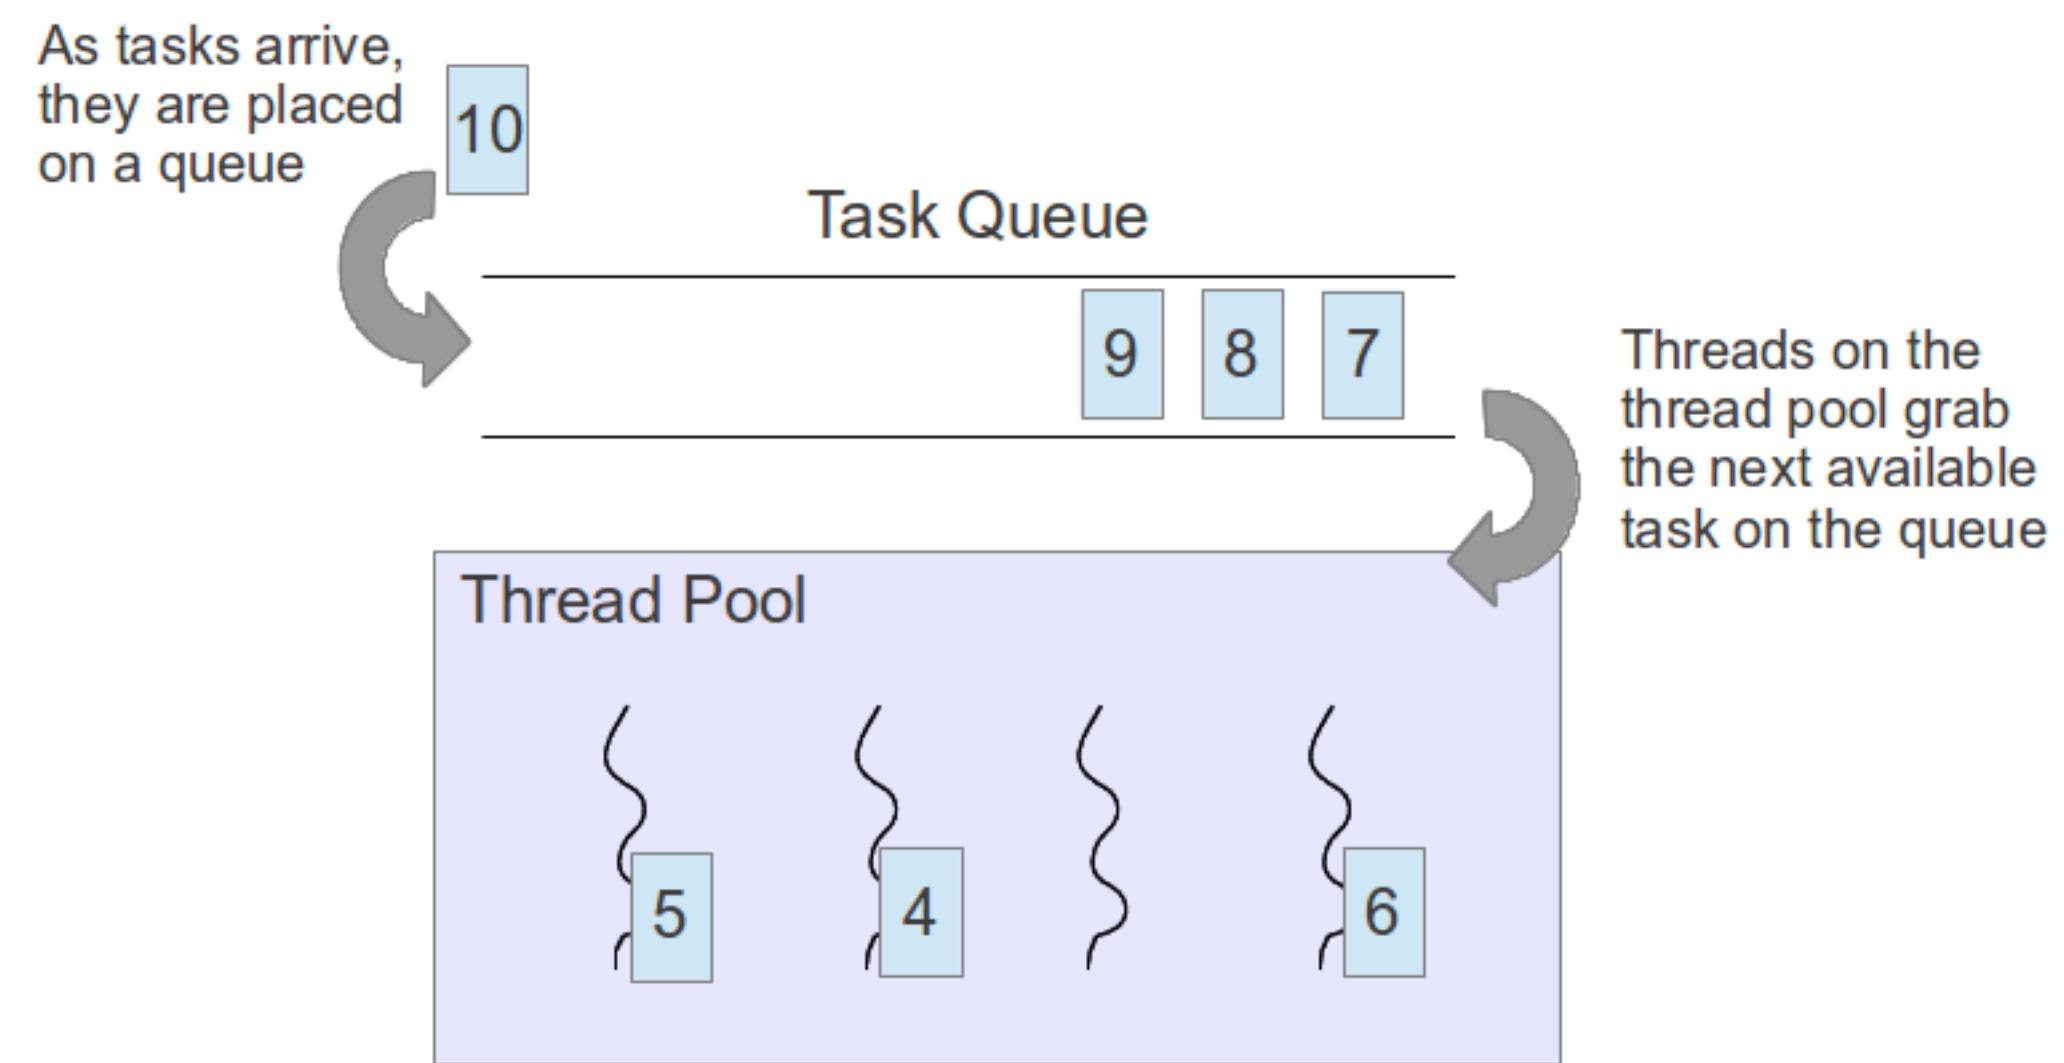


ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)

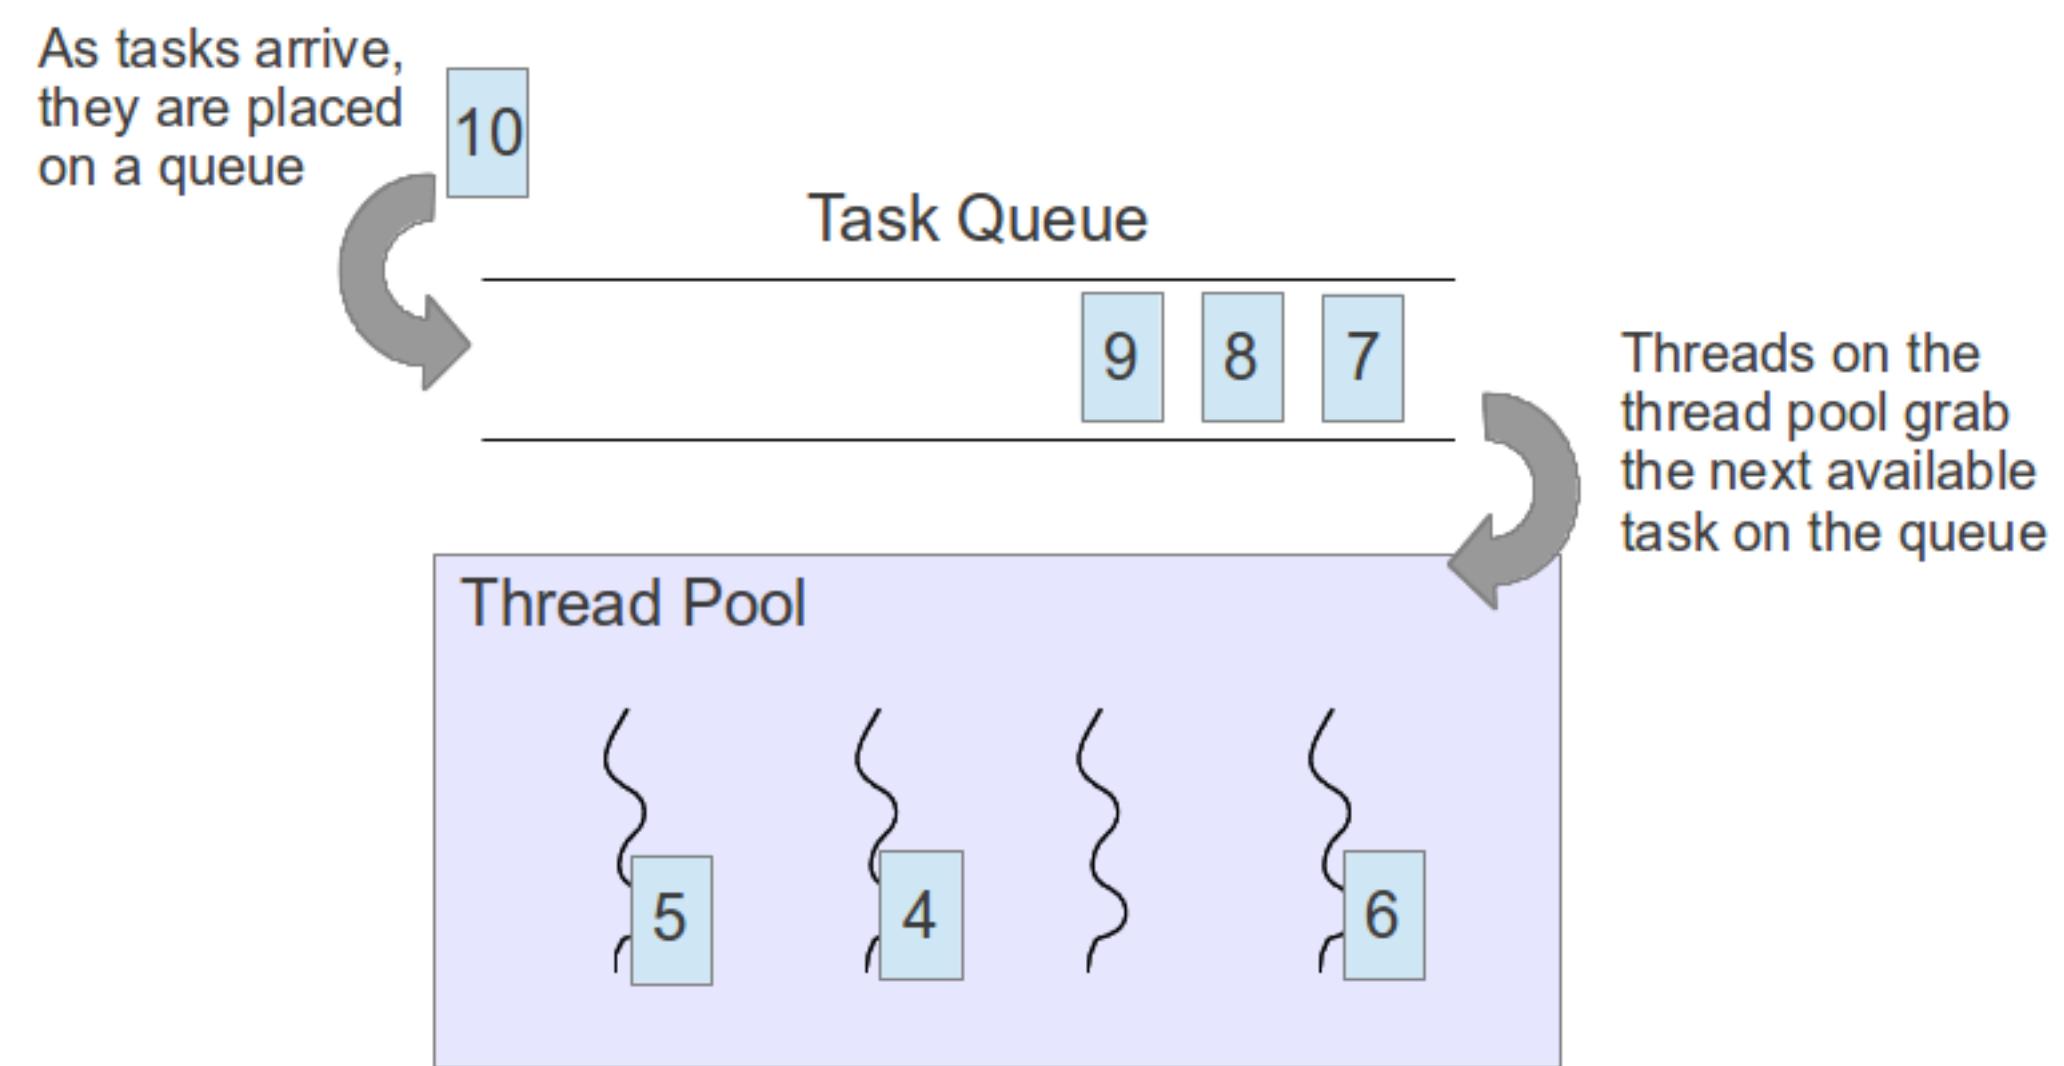
ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)



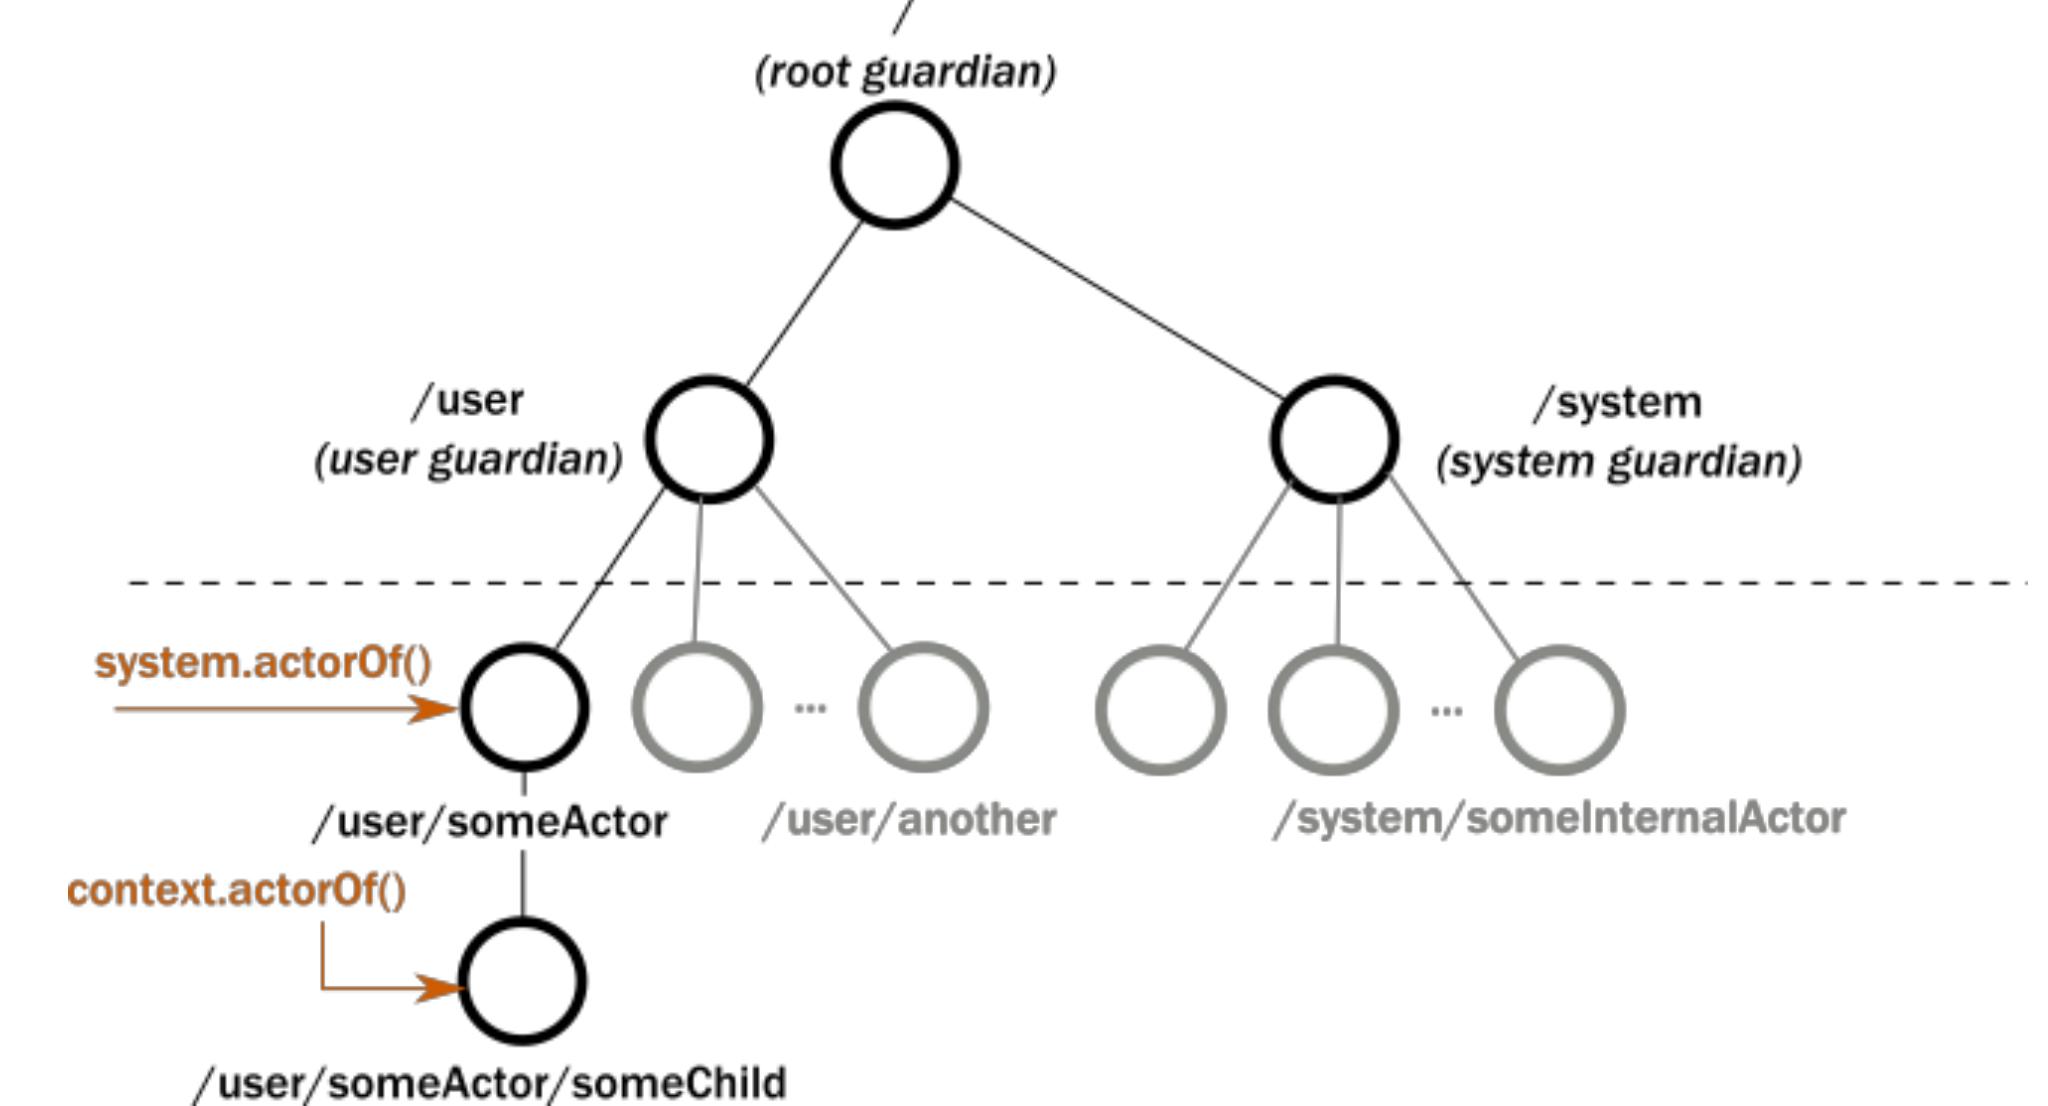
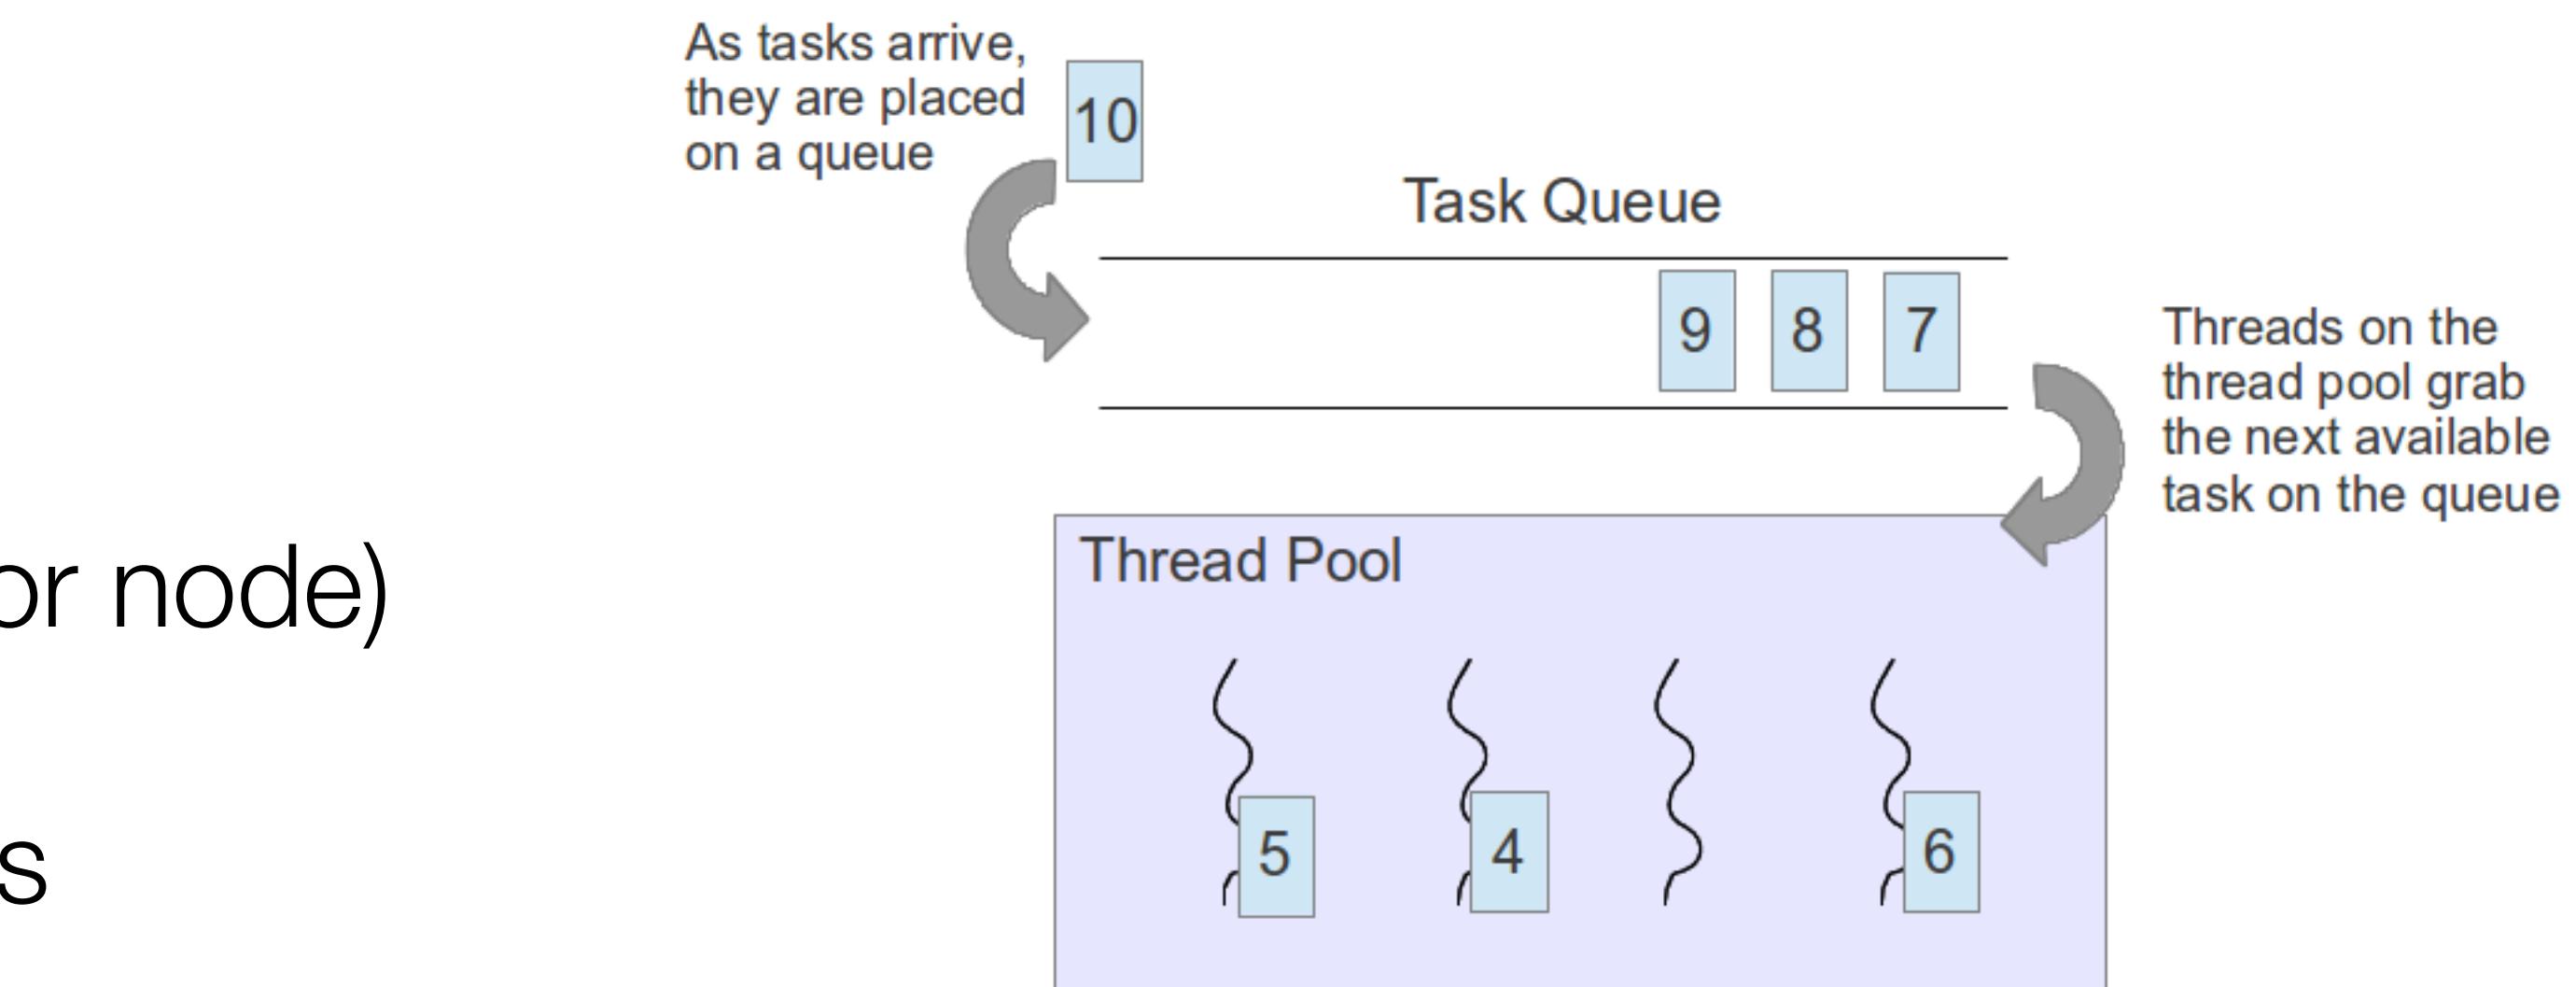
ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)
 - ▶ Takes care of the creation of Actors
 - ▶ Cares for the Actor lifecycle



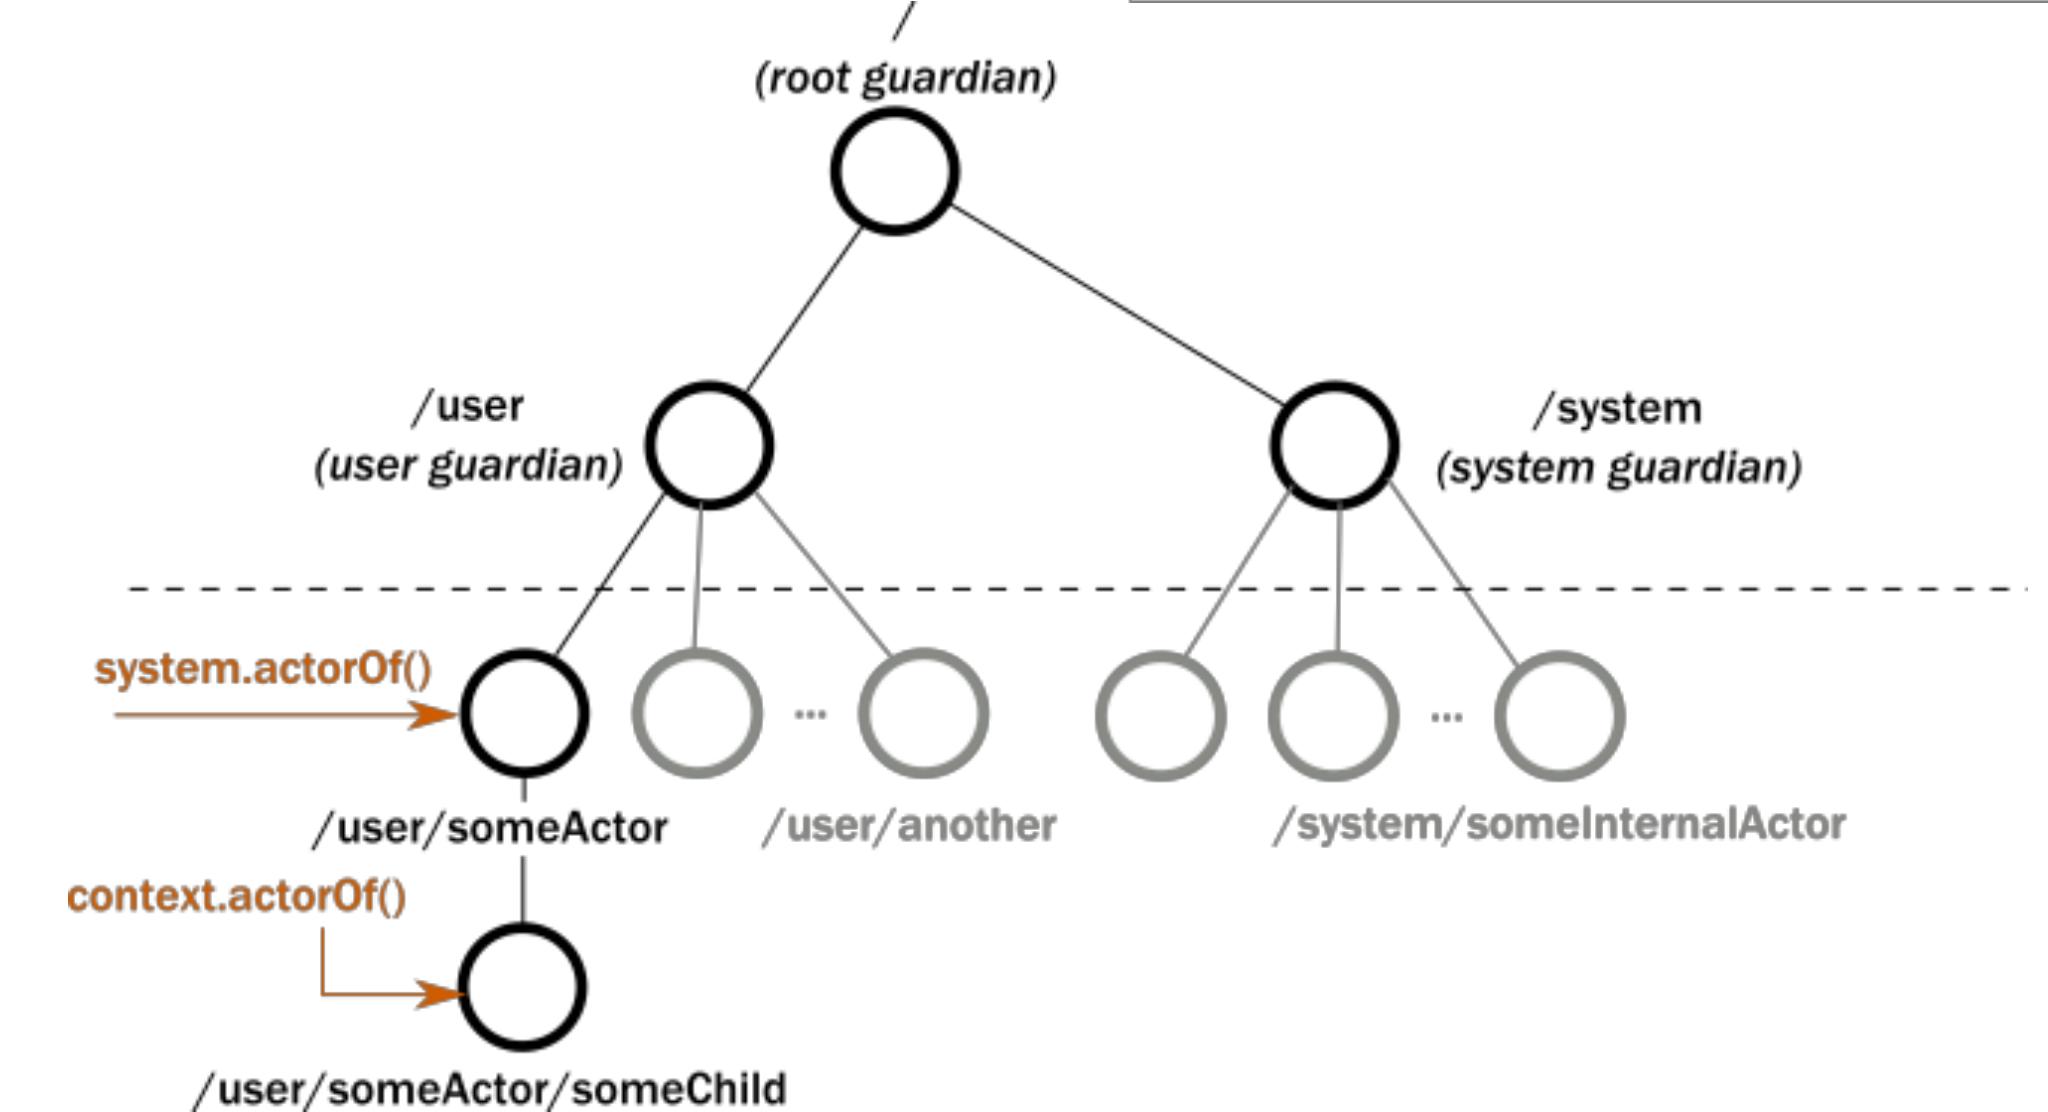
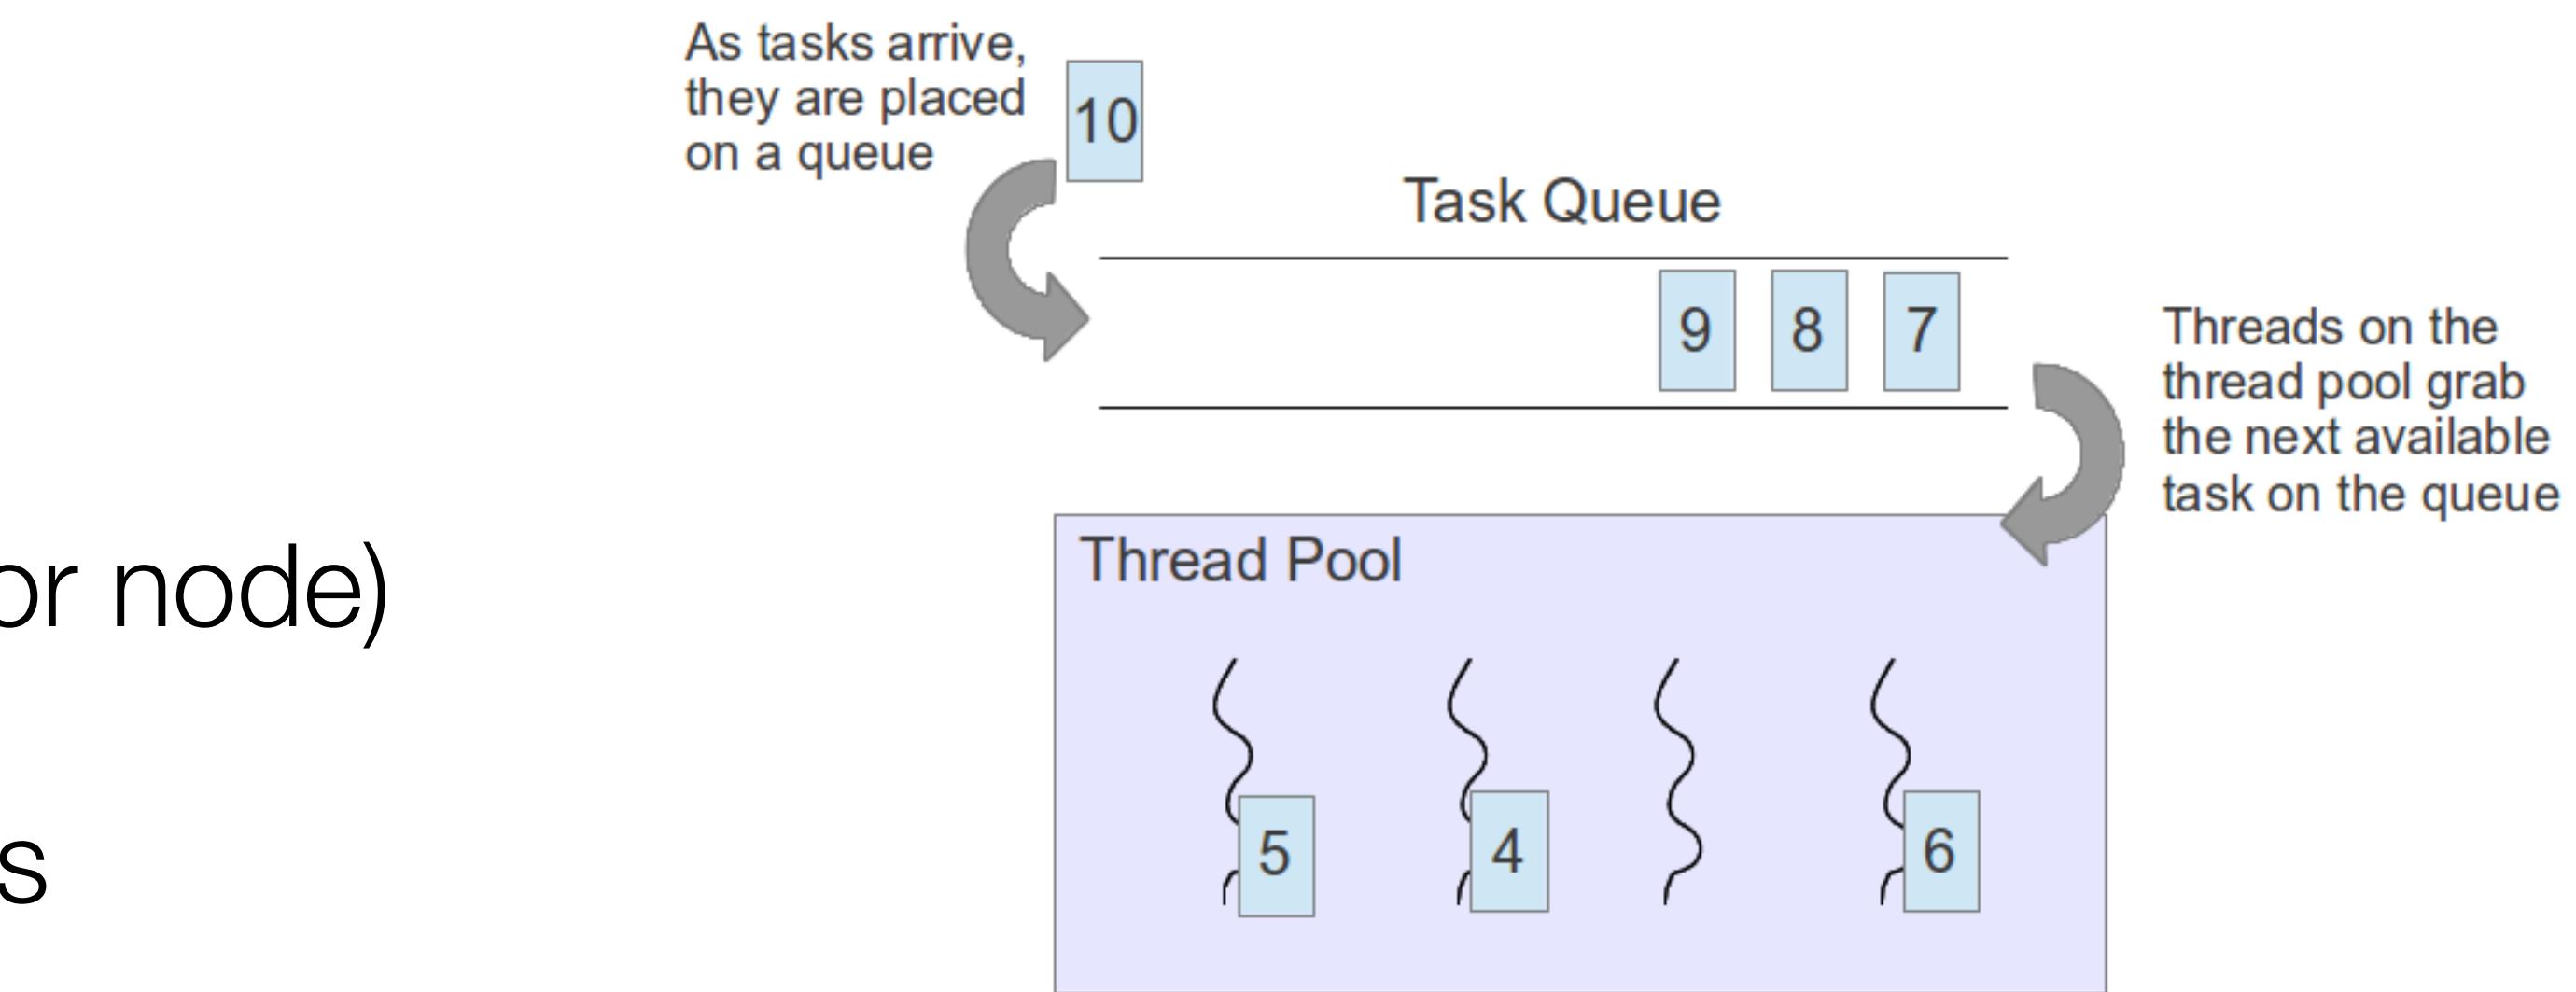
ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)
 - ▶ Takes care of the creation of Actors
 - ▶ Cares for the Actor lifecycle
 - ▶ Creates 3 actors on startup
 - ▶ /user, /system, / (the root)



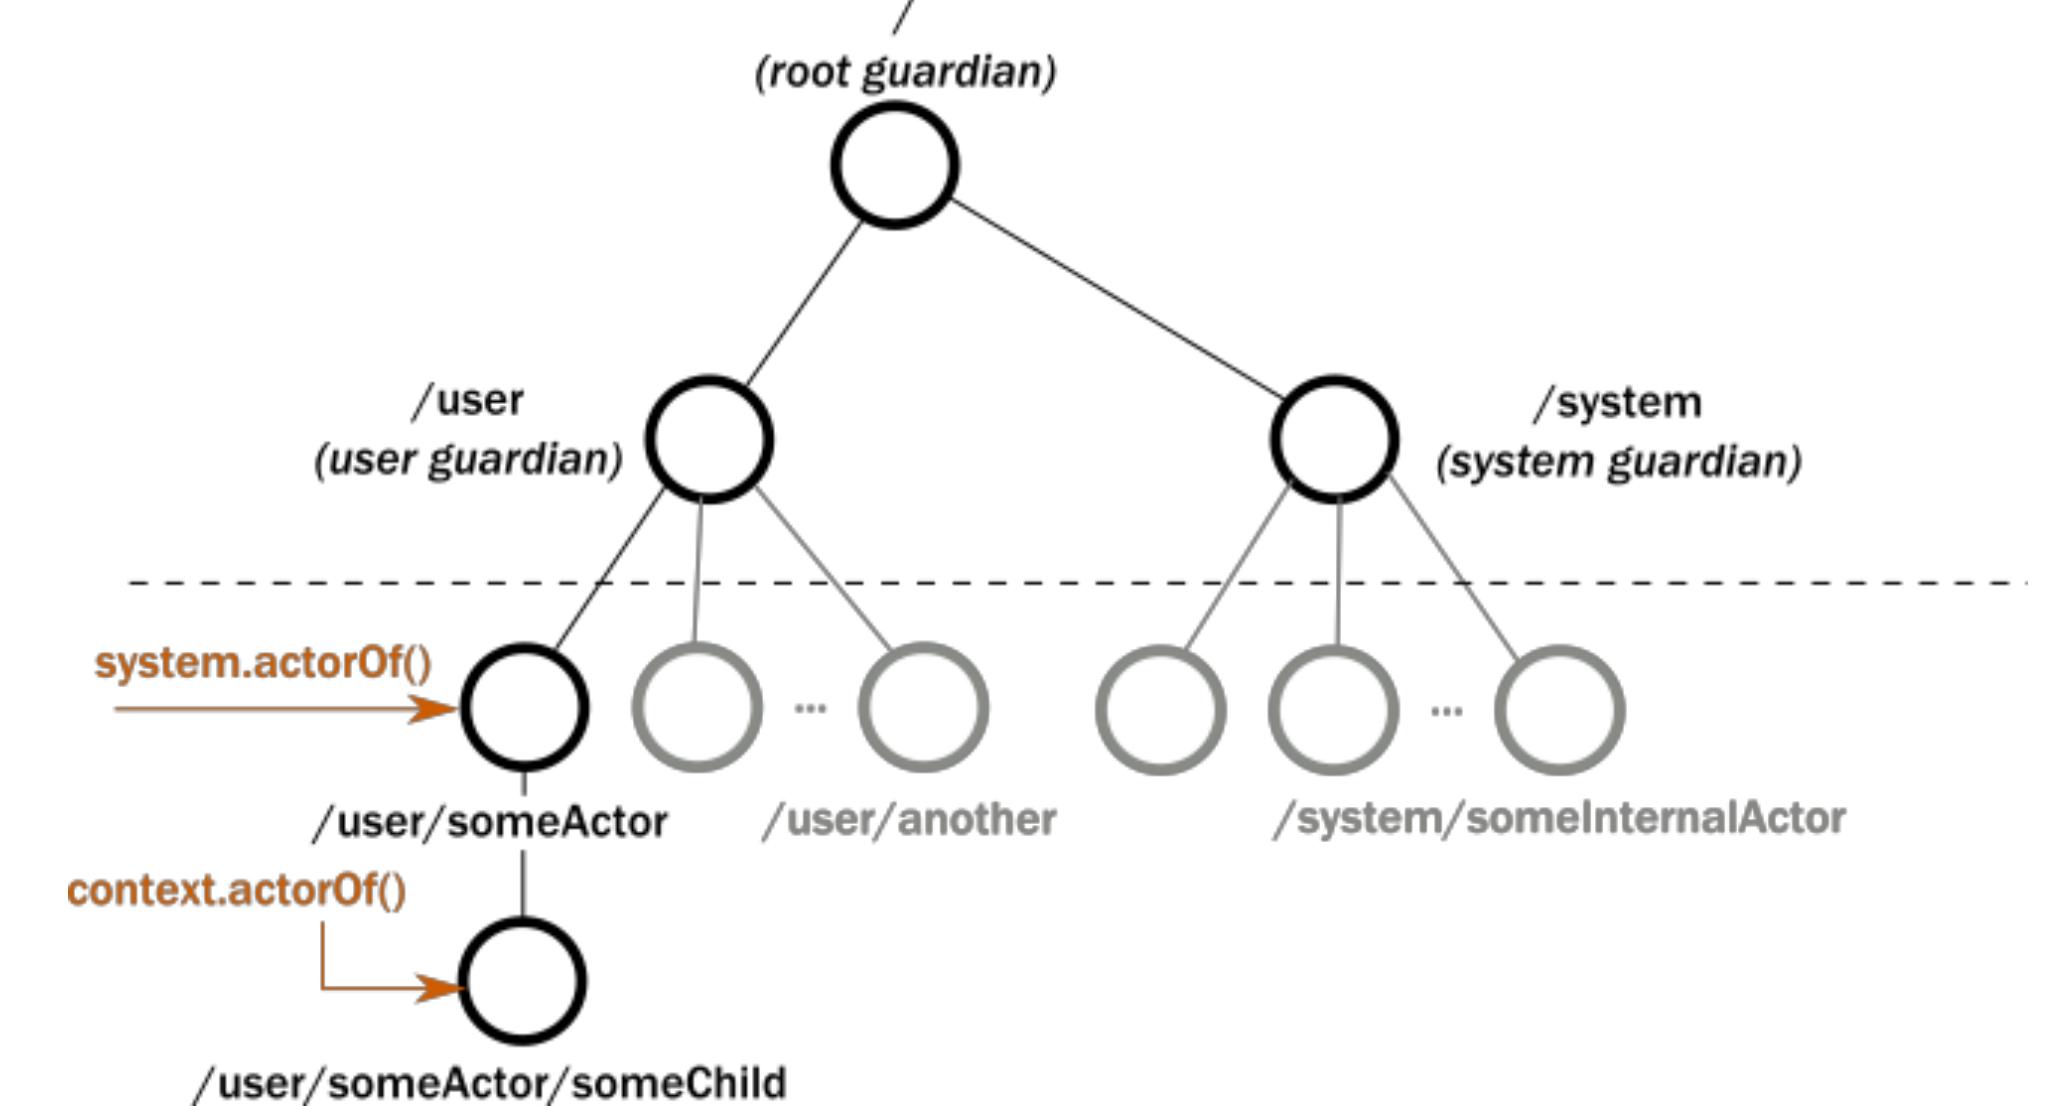
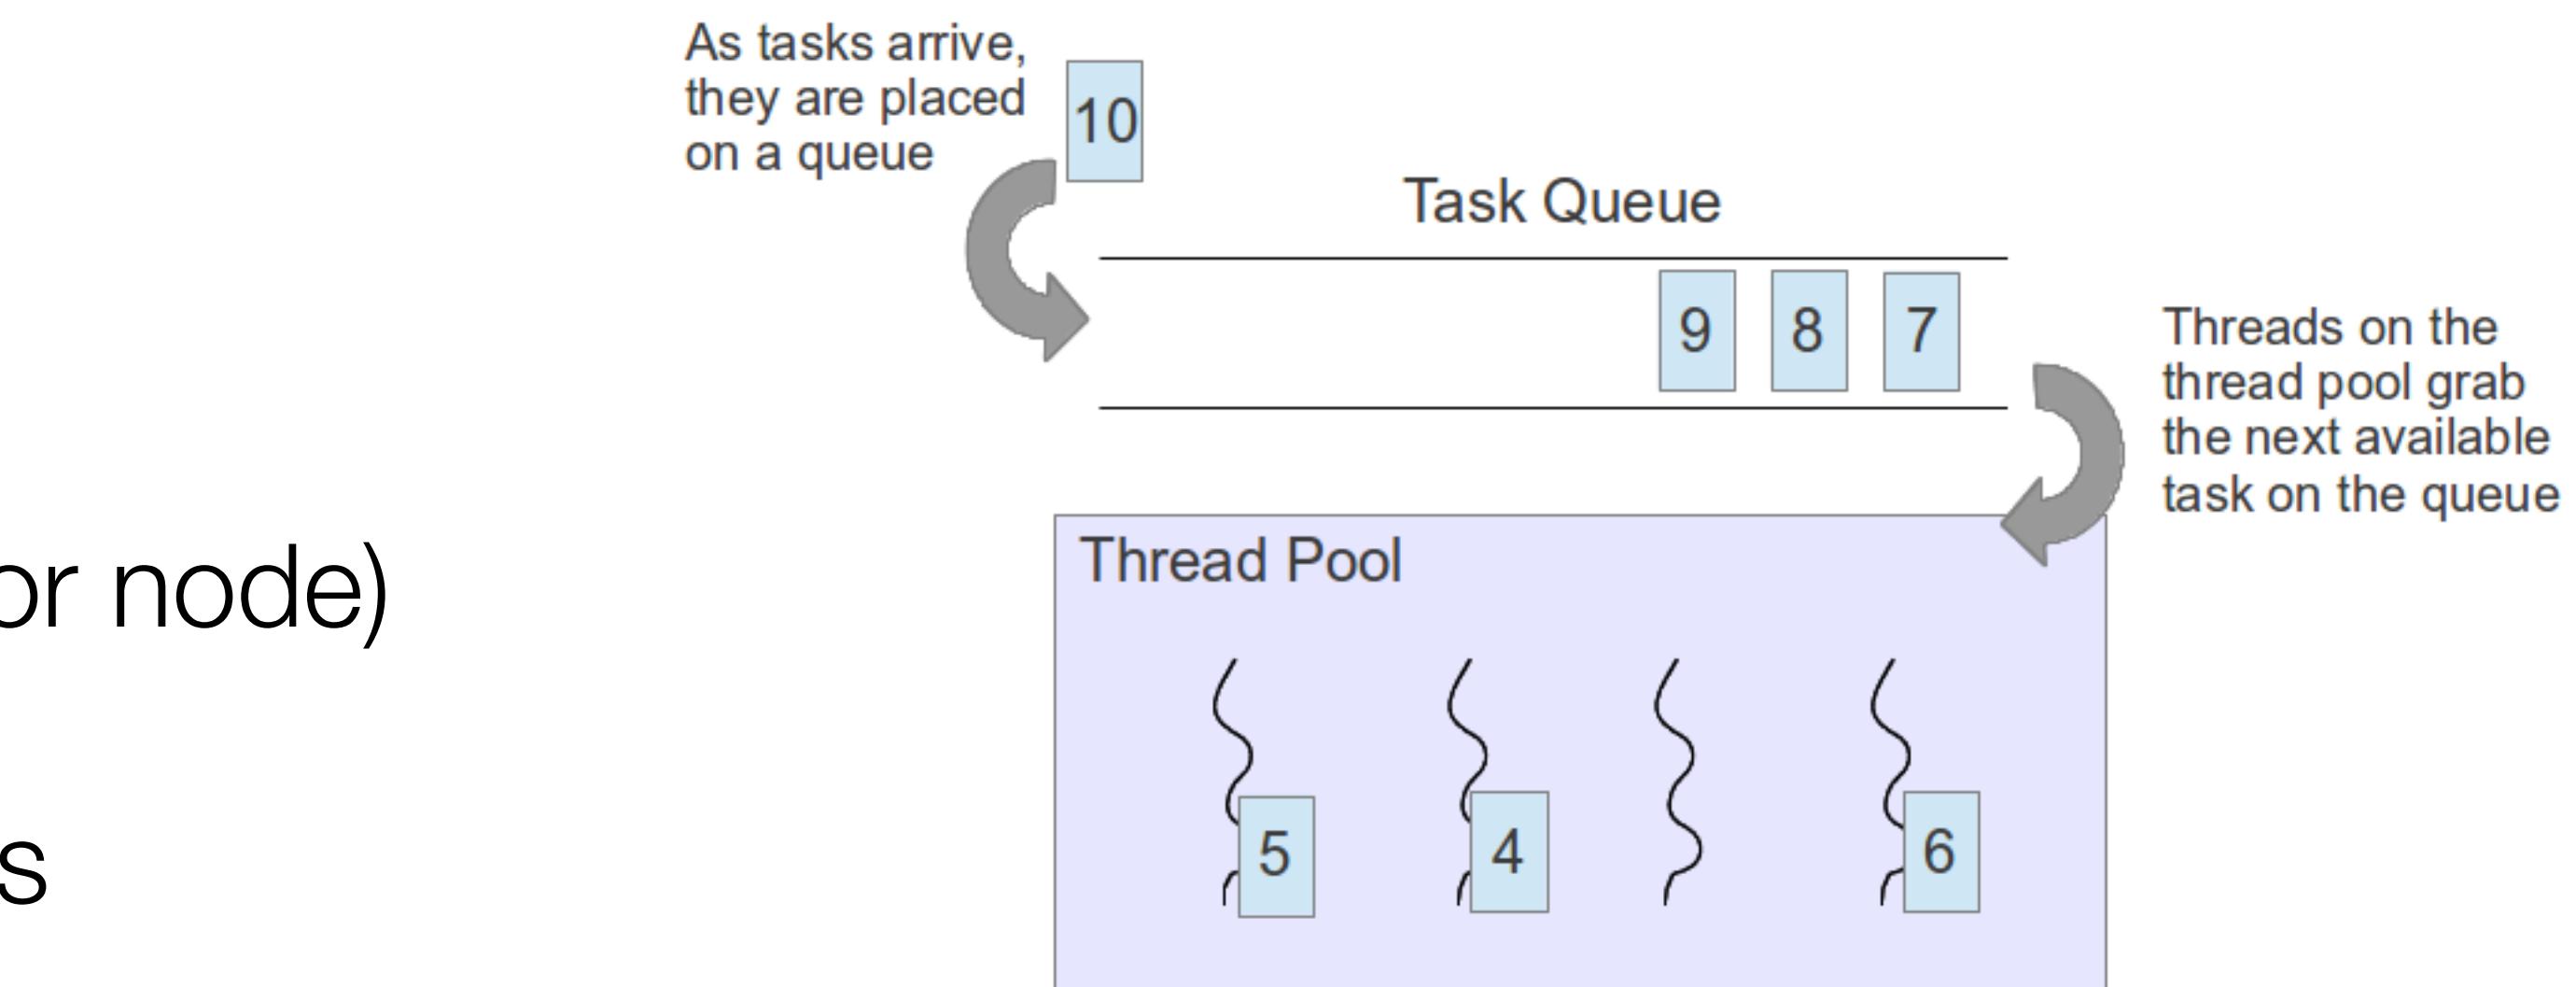
ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)
 - ▶ Takes care of the creation of Actors
 - ▶ Cares for the Actor lifecycle
 - ▶ Creates 3 actors on startup
 - ▶ /user, /system, / (the root)
 - ▶ AkkaActors.scala



ACTOR SYSTEM: CREATION

- ▶ It manages a thread pool
 - ▶ Most likely one per application (or node)
 - ▶ Takes care of the creation of Actors
 - ▶ Cares for the Actor lifecycle
 - ▶ Creates 3 actors on startup
 - ▶ /user, /system, / (the root)
 - ▶ AkkaActors.scala

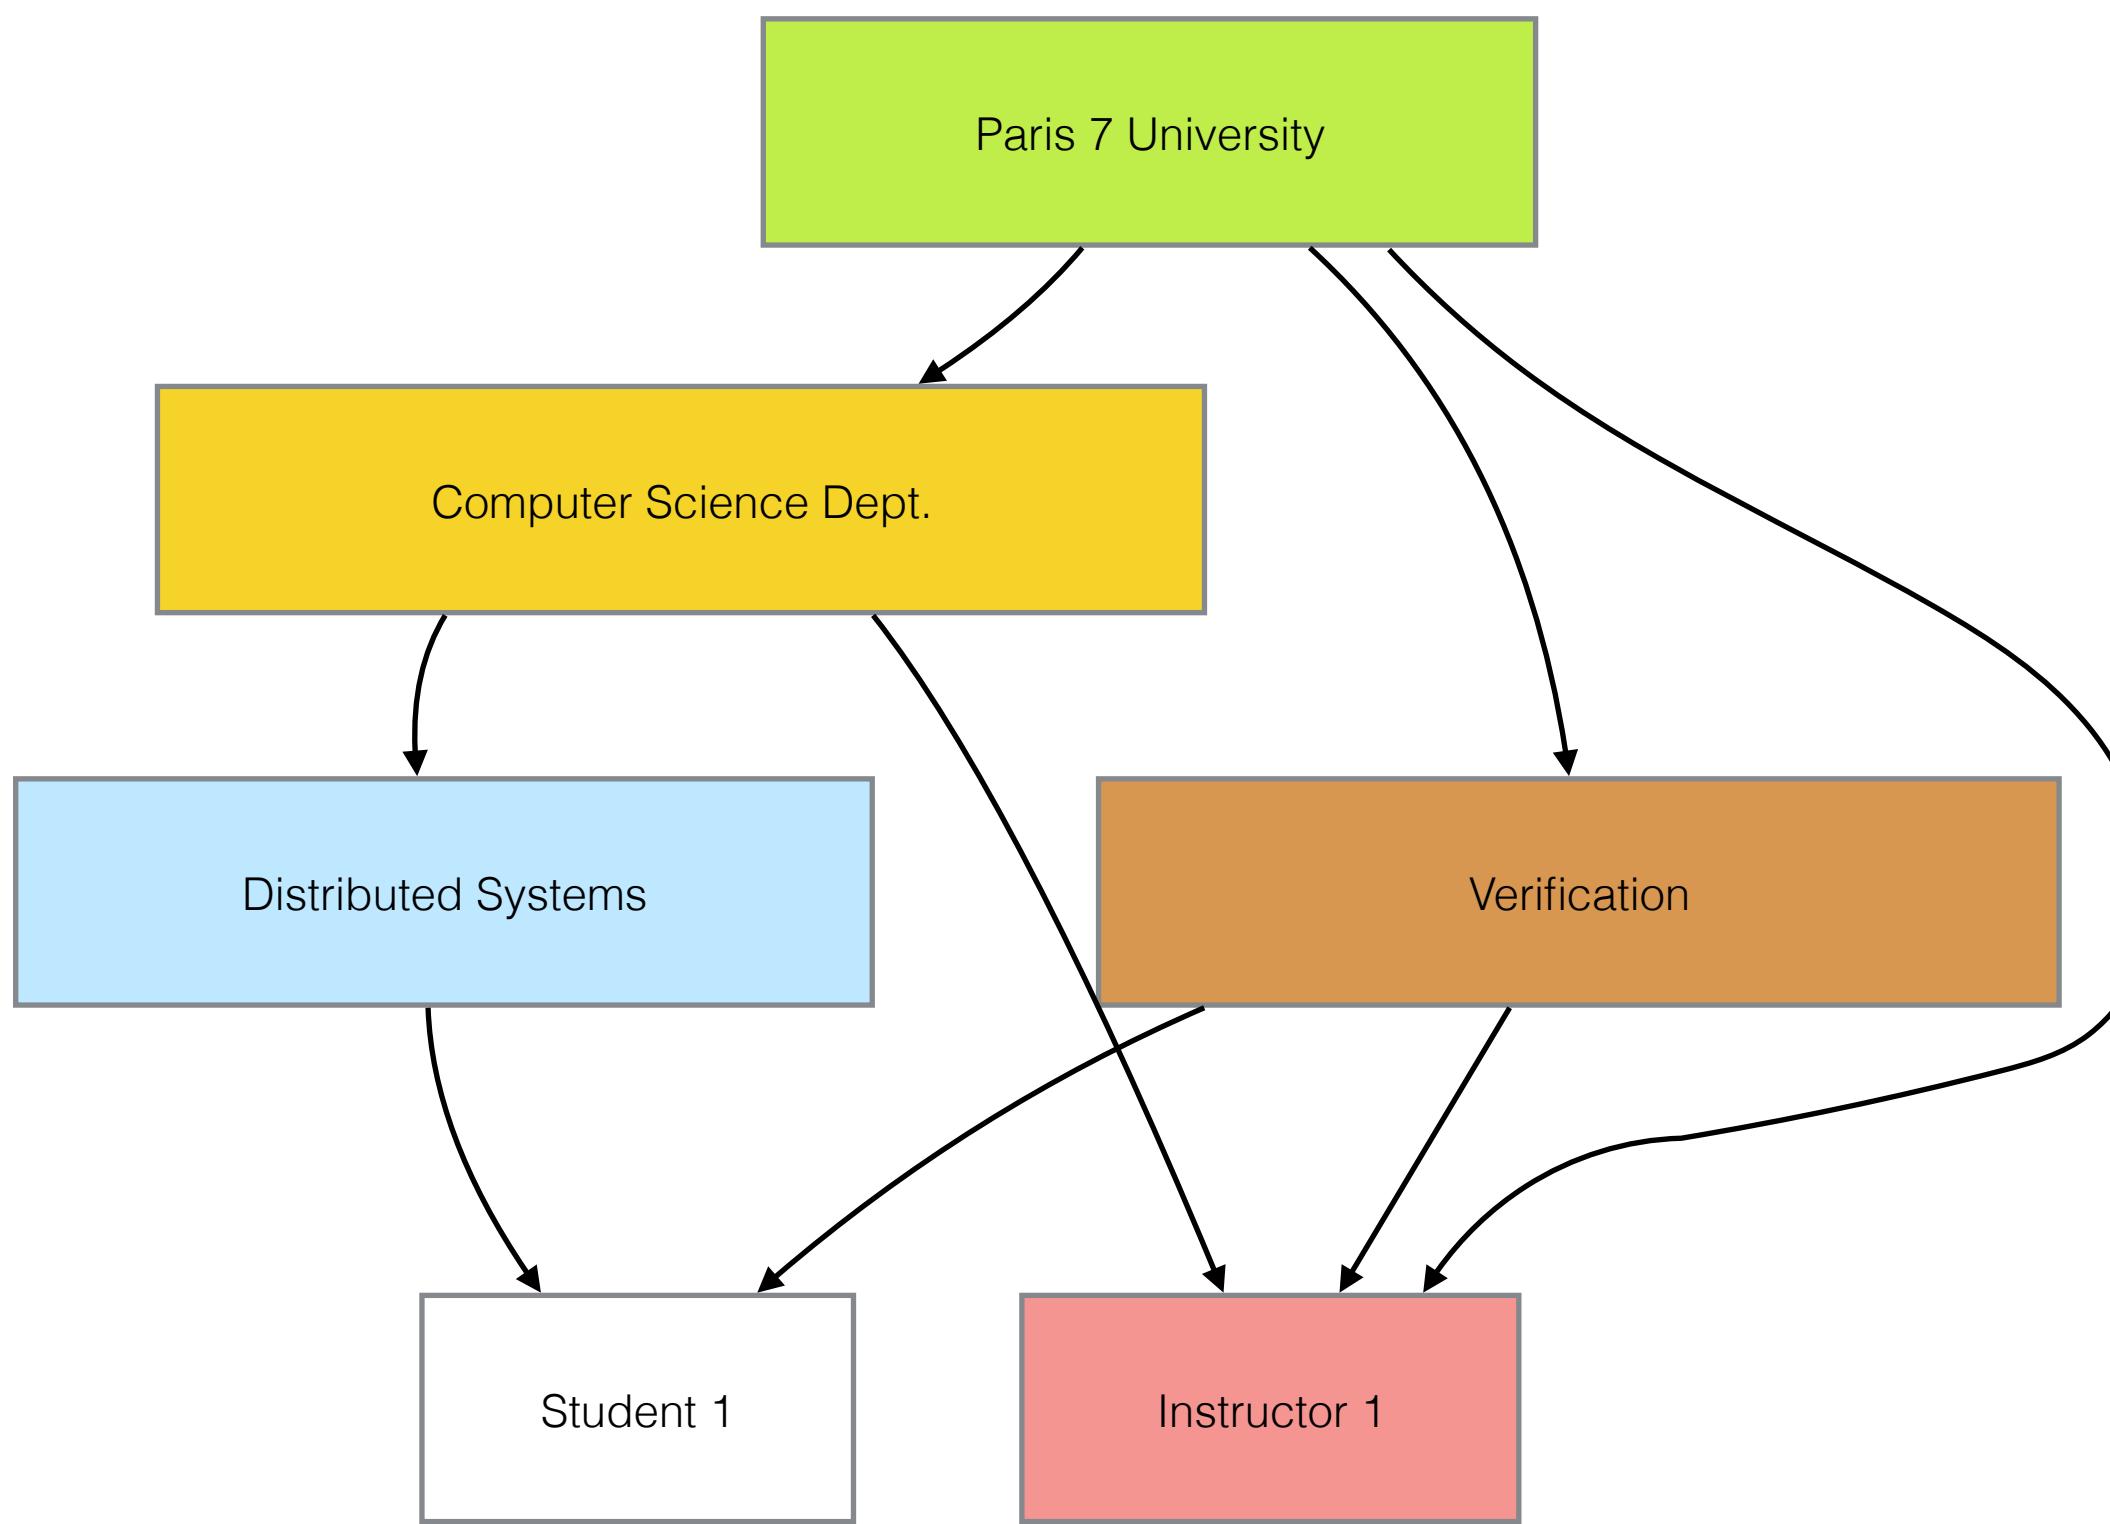


○ Exercise 1: Create an UM6P Actor System

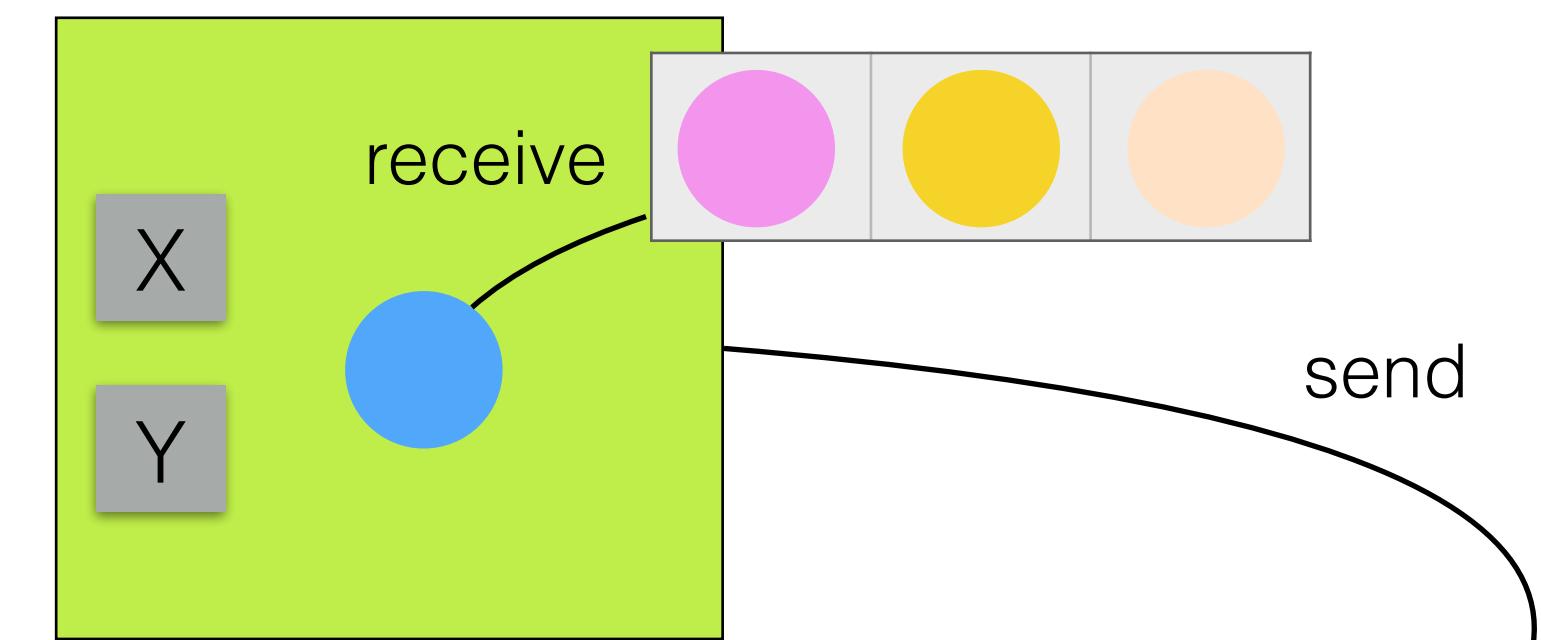
WHAT IS AN ACTOR?

► State

An actor system runtime



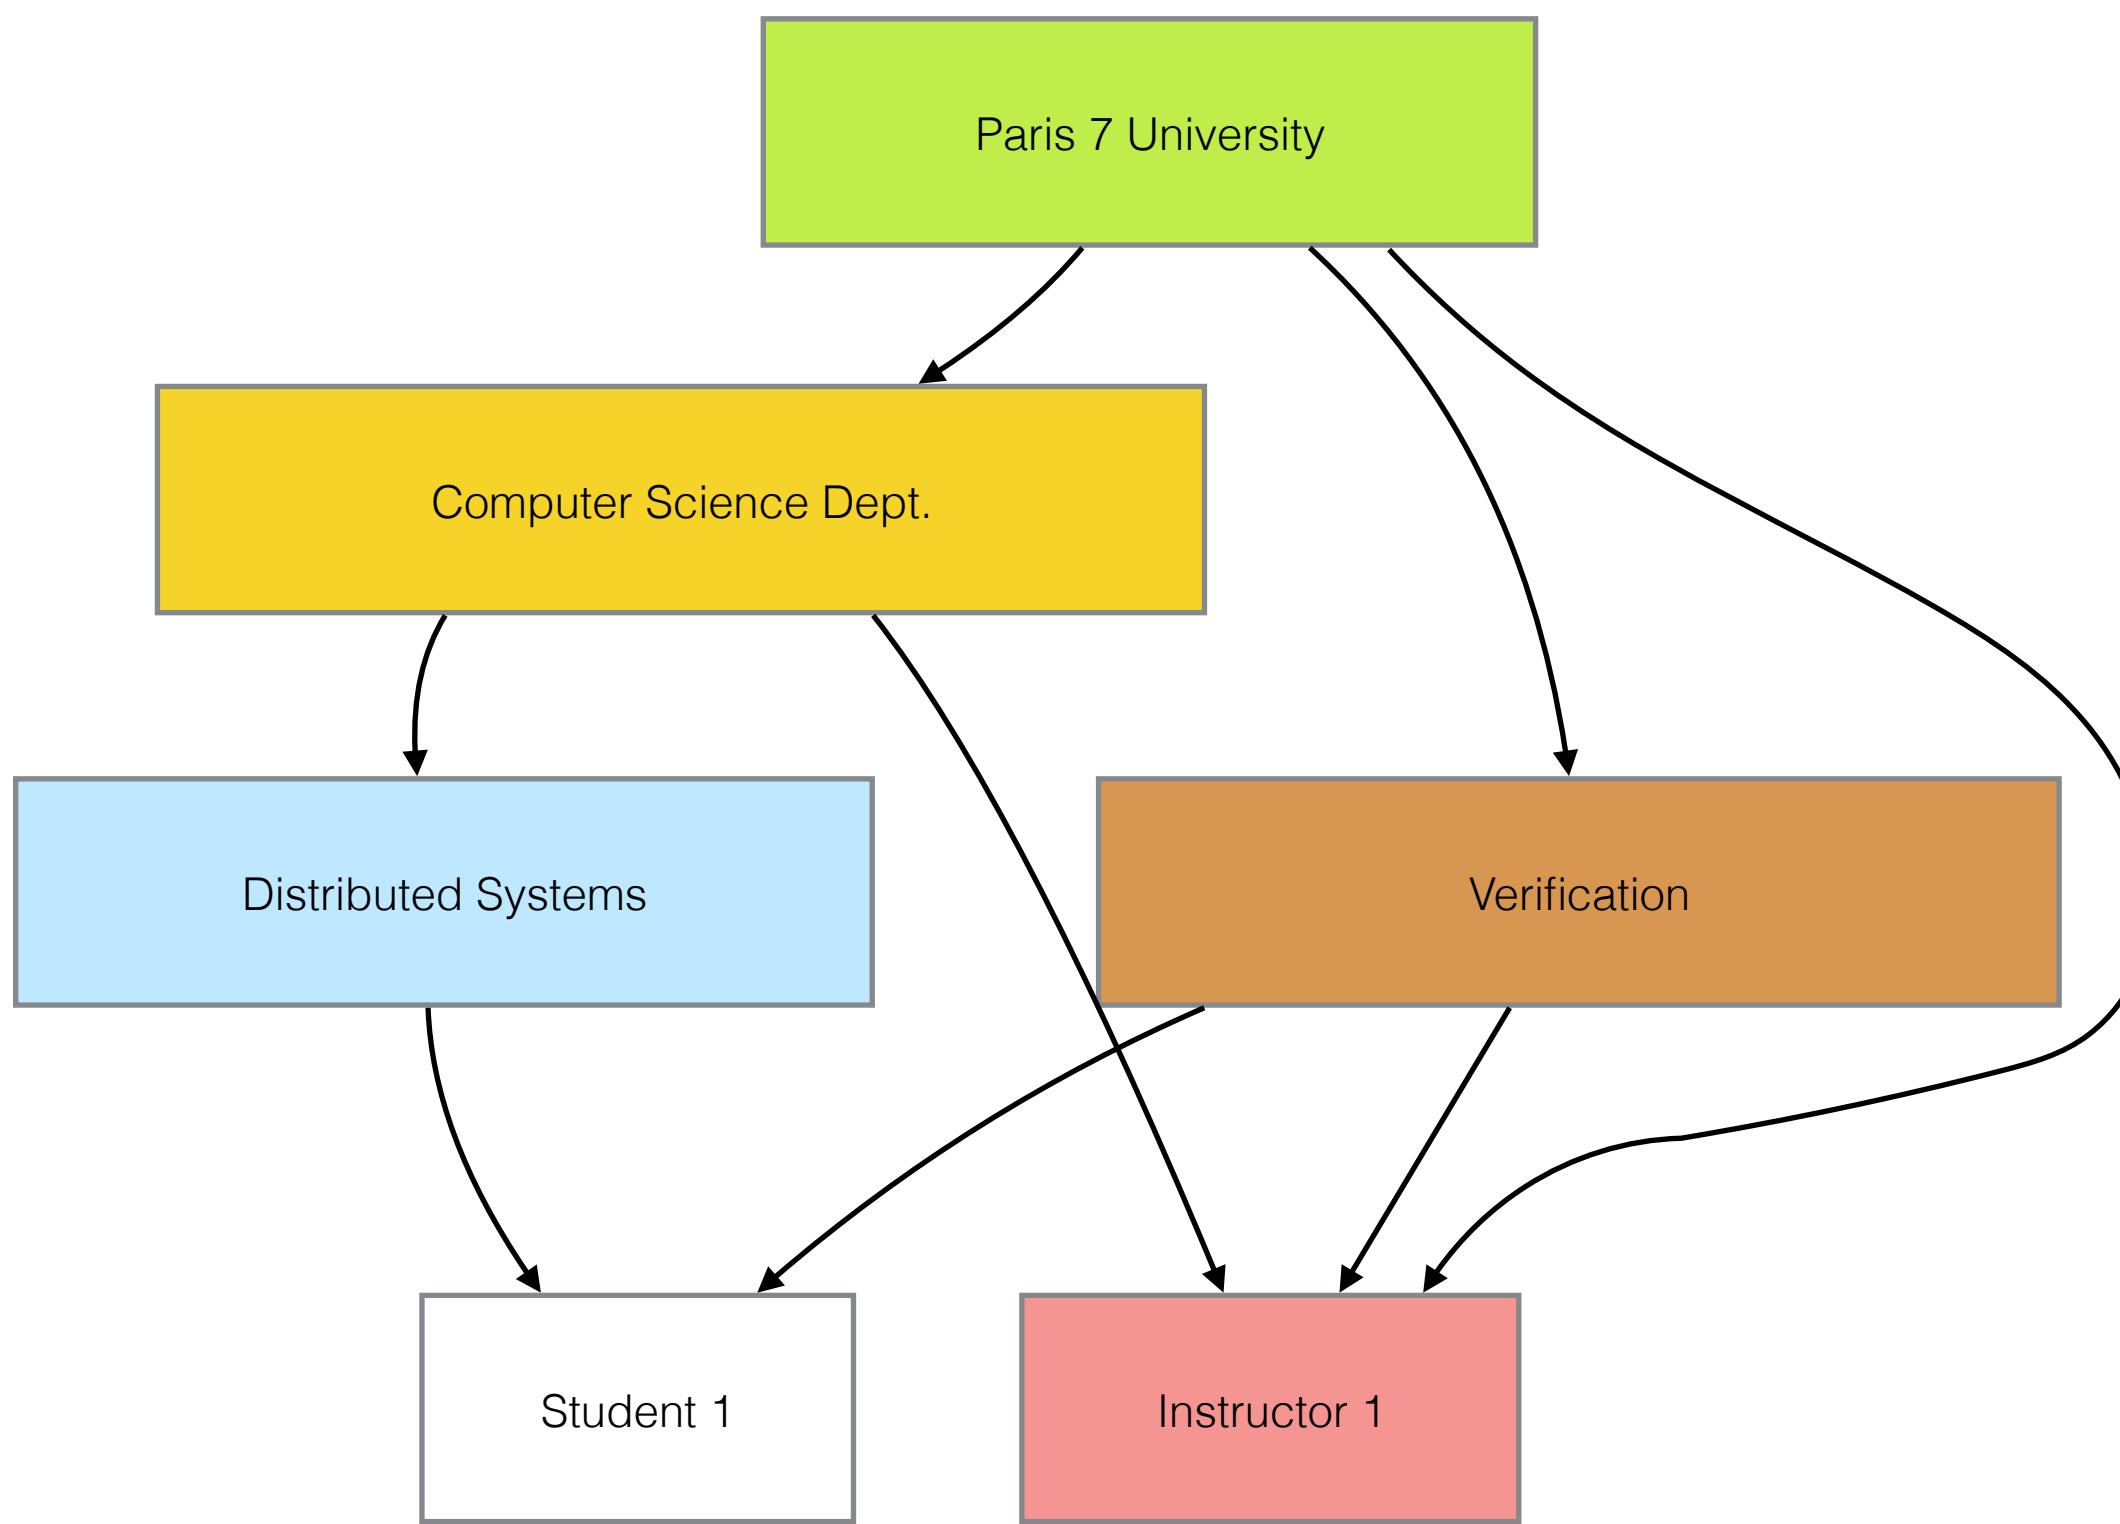
Two actors



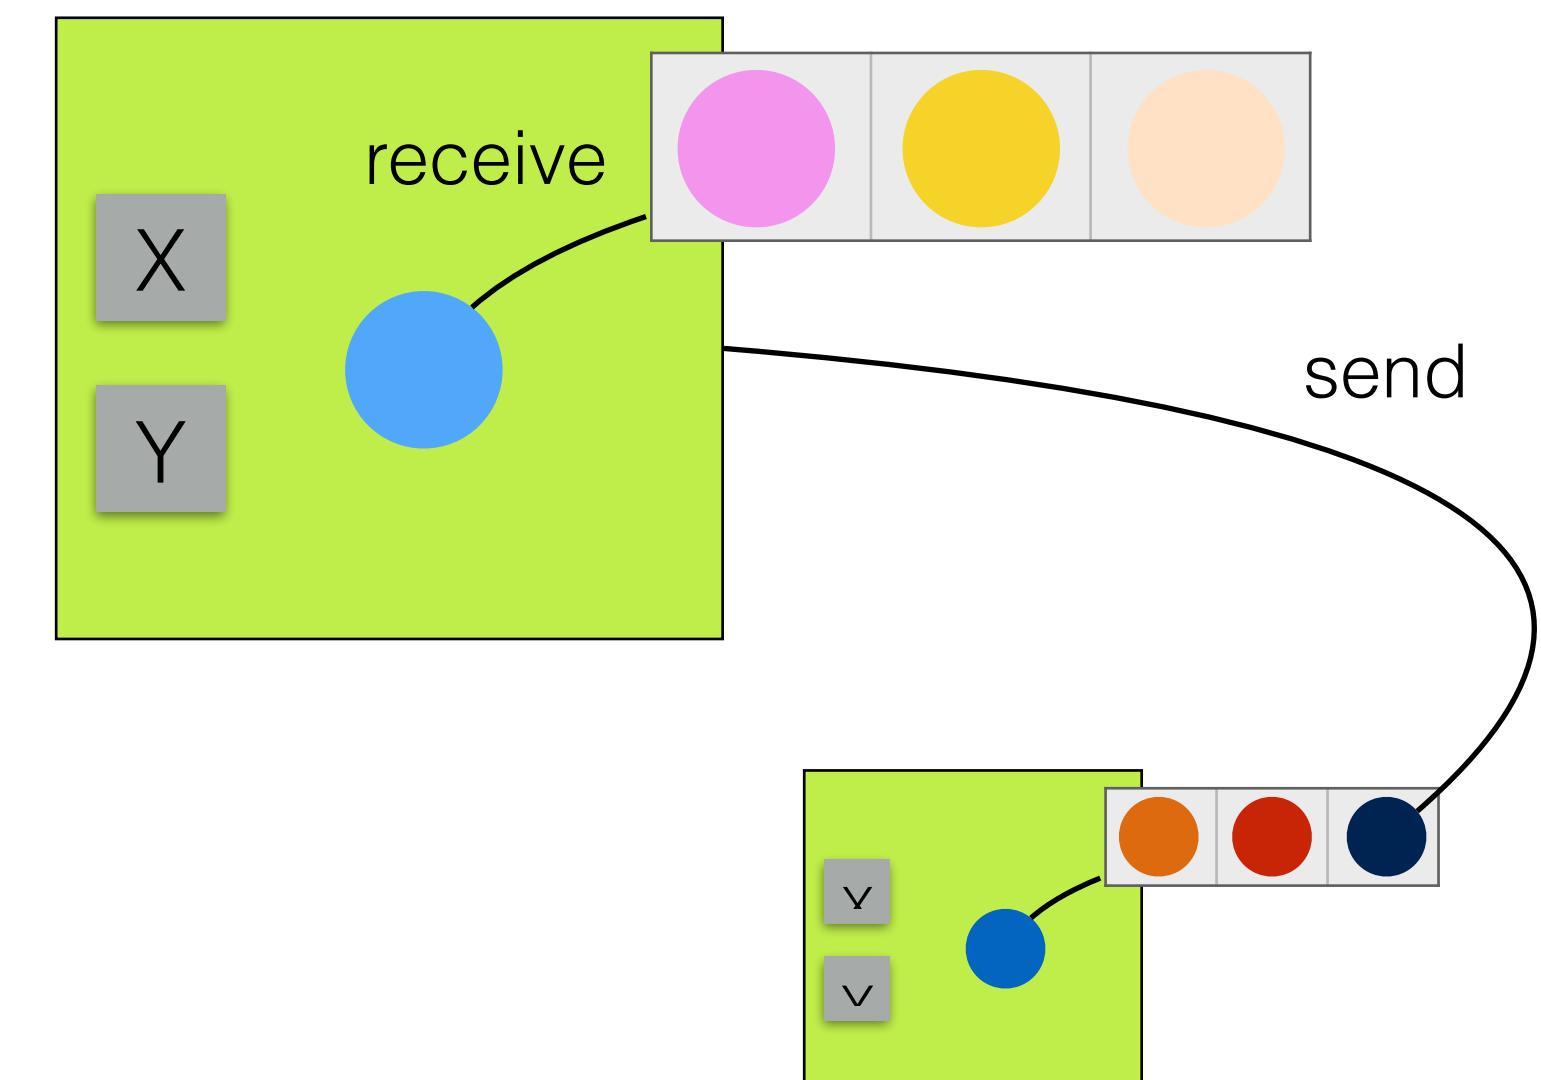
WHAT IS AN ACTOR?

- ▶ State
- ▶ Messages

An actor system runtime



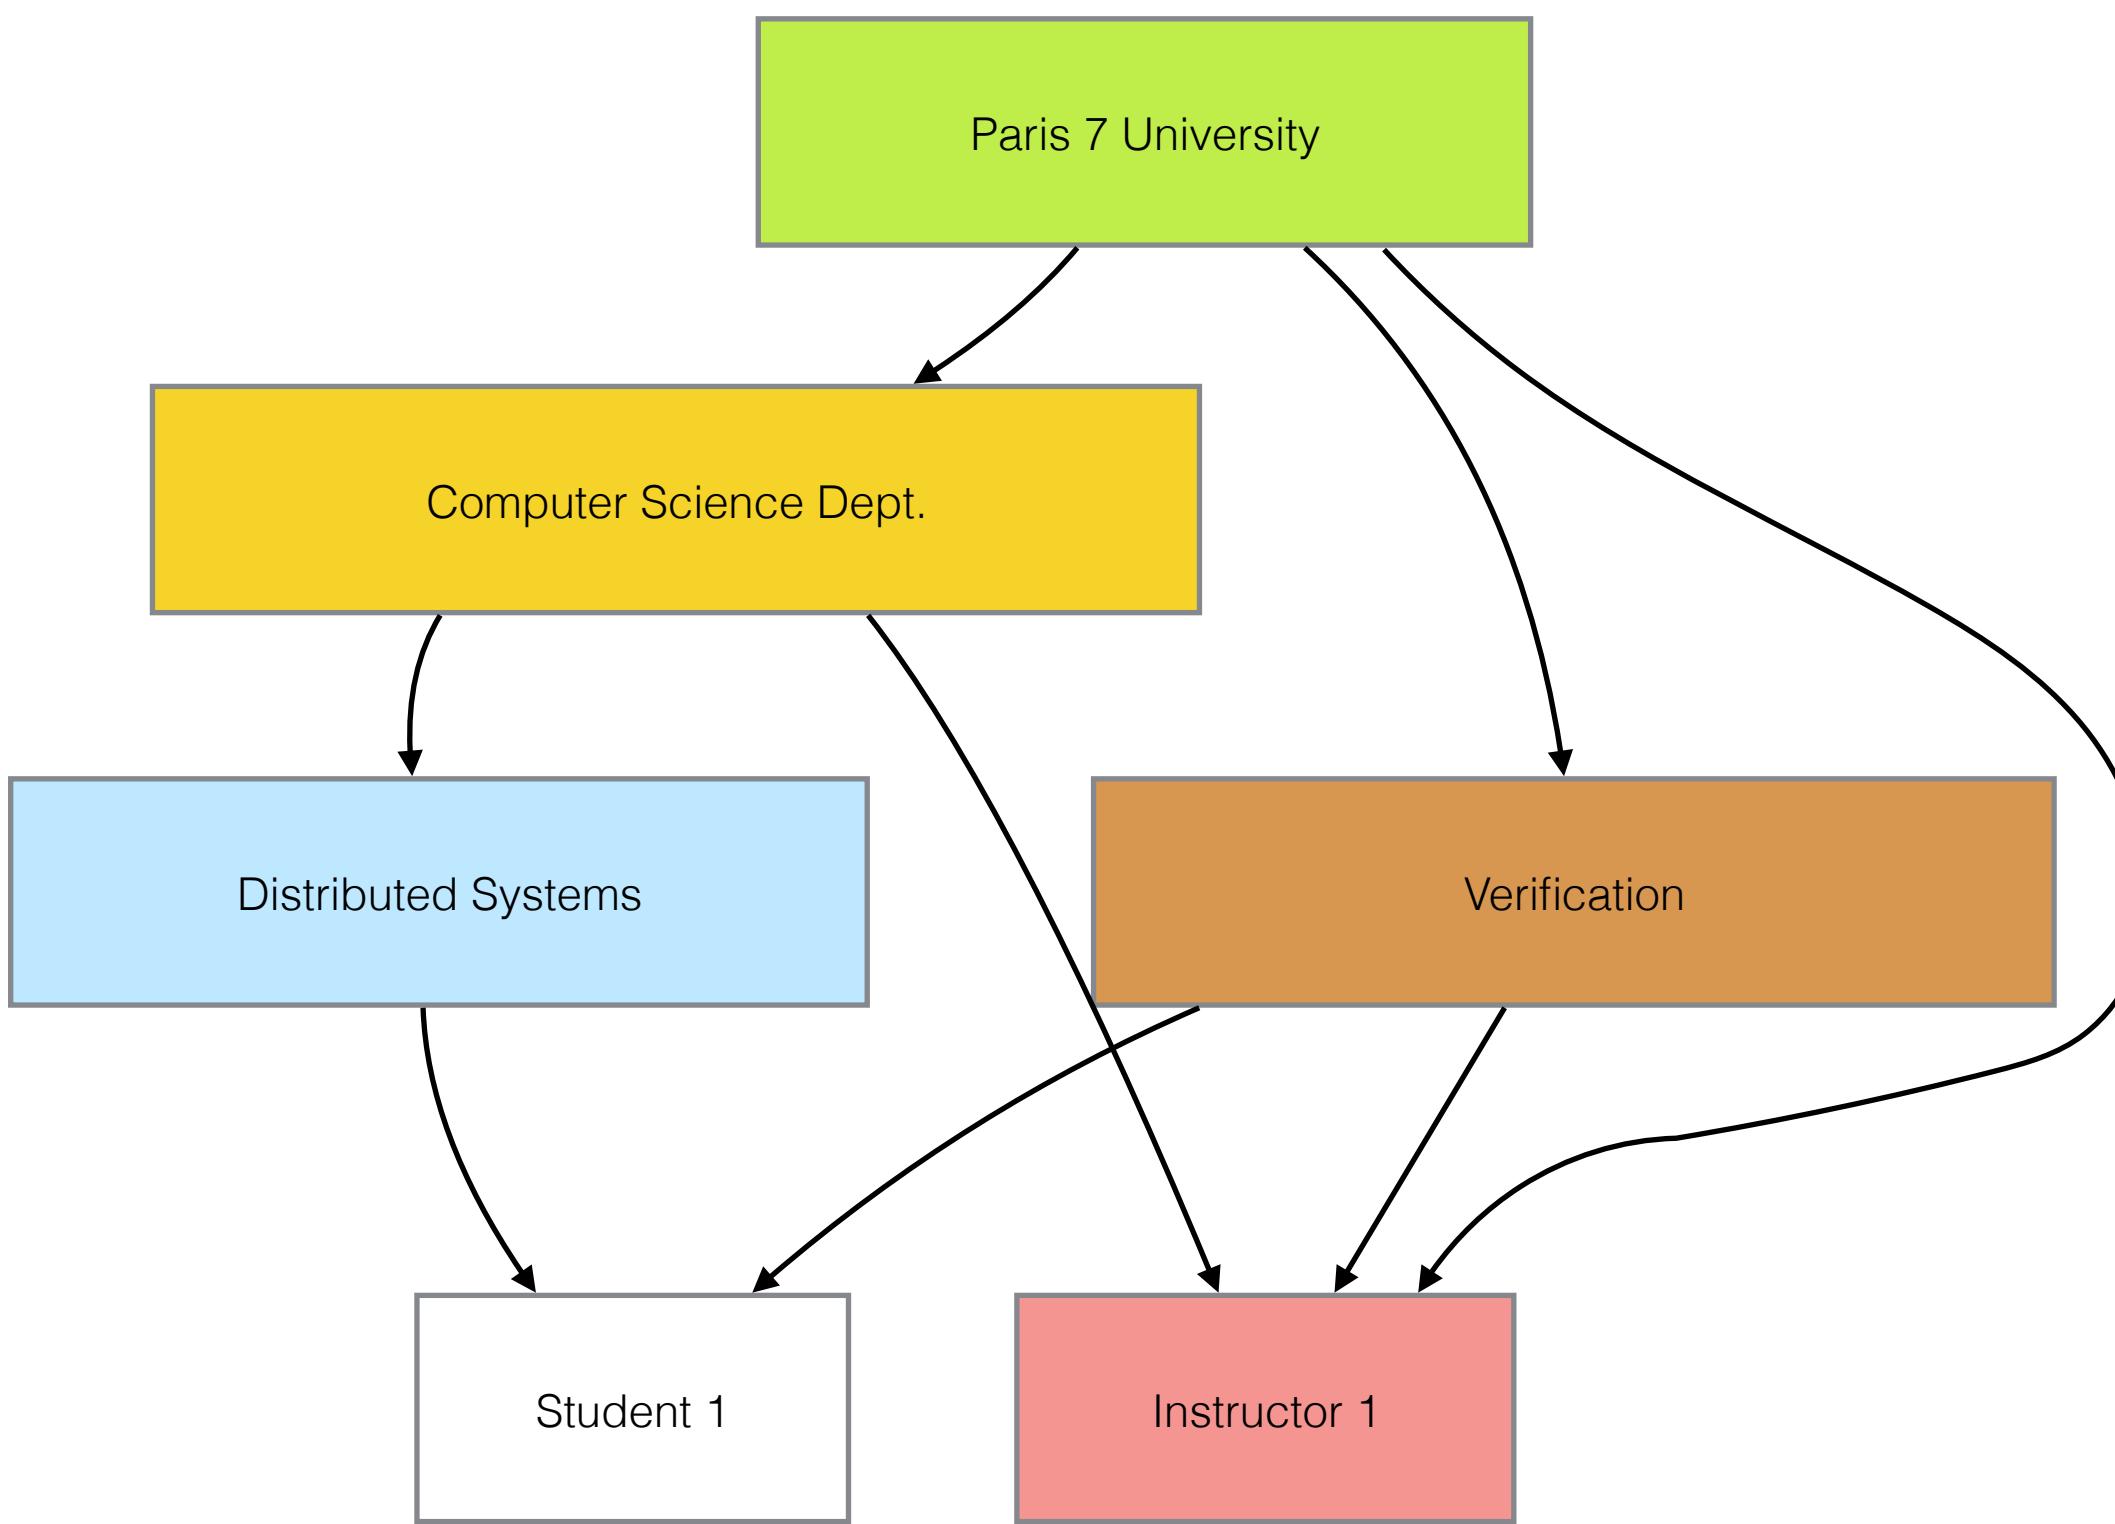
Two actors



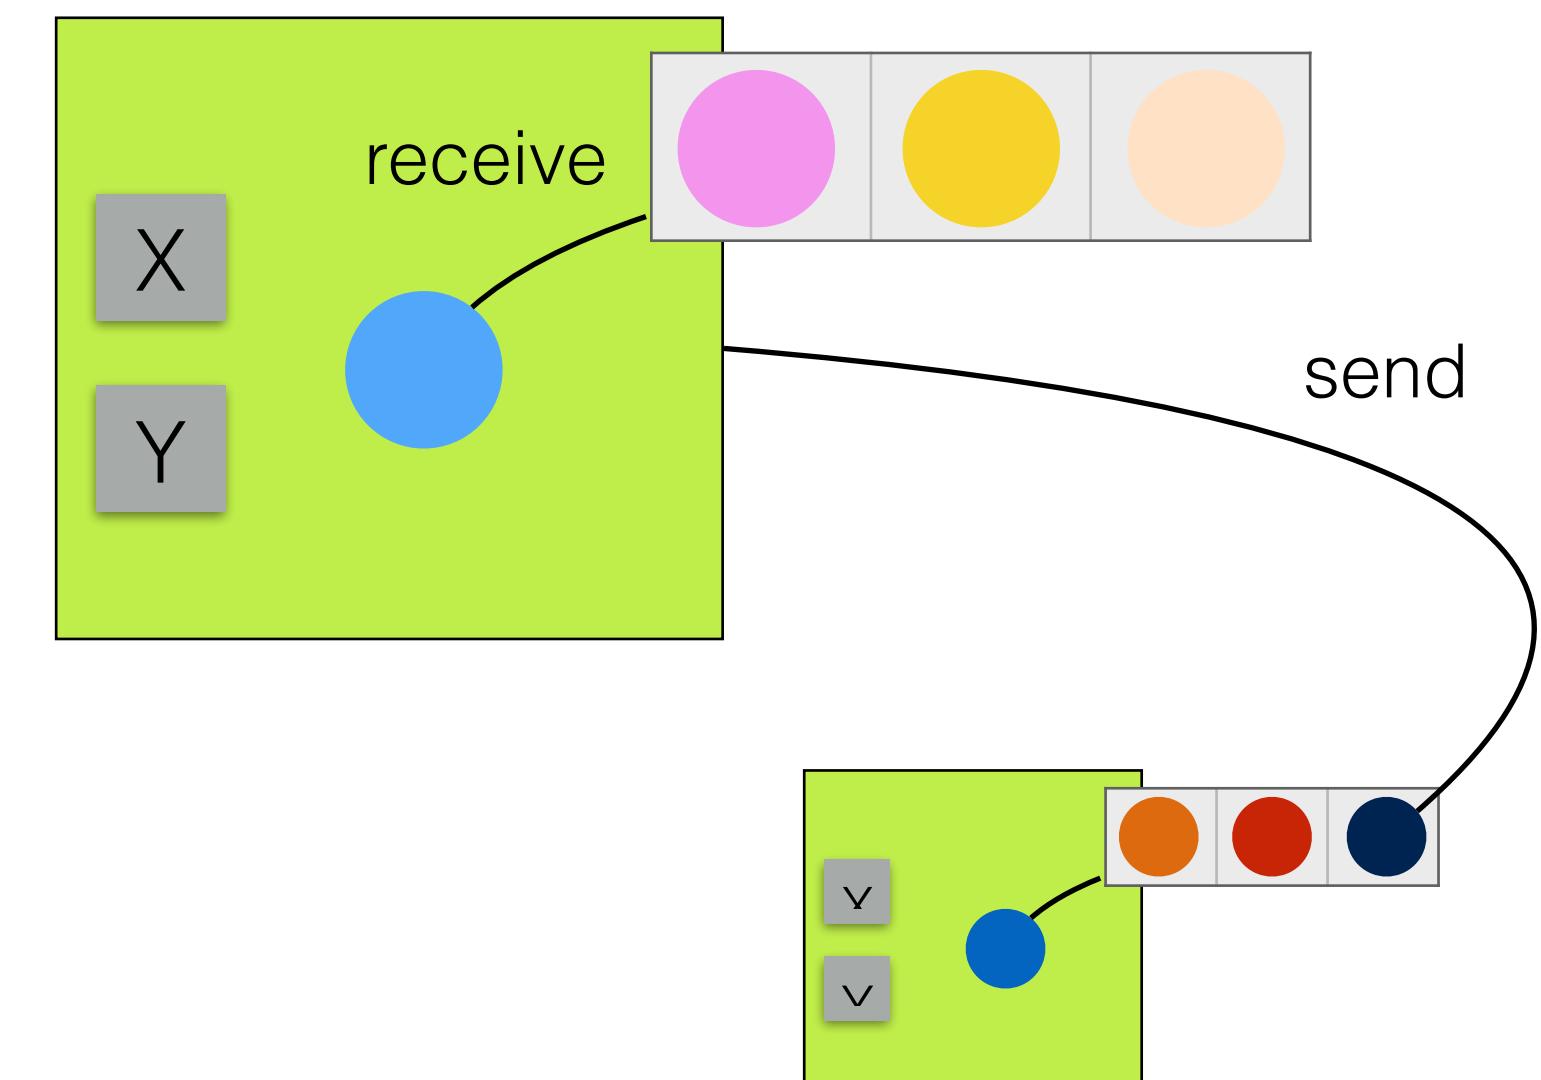
WHAT IS AN ACTOR?

- ▶ State
- ▶ Messages
- ▶ Inbox

An actor system runtime



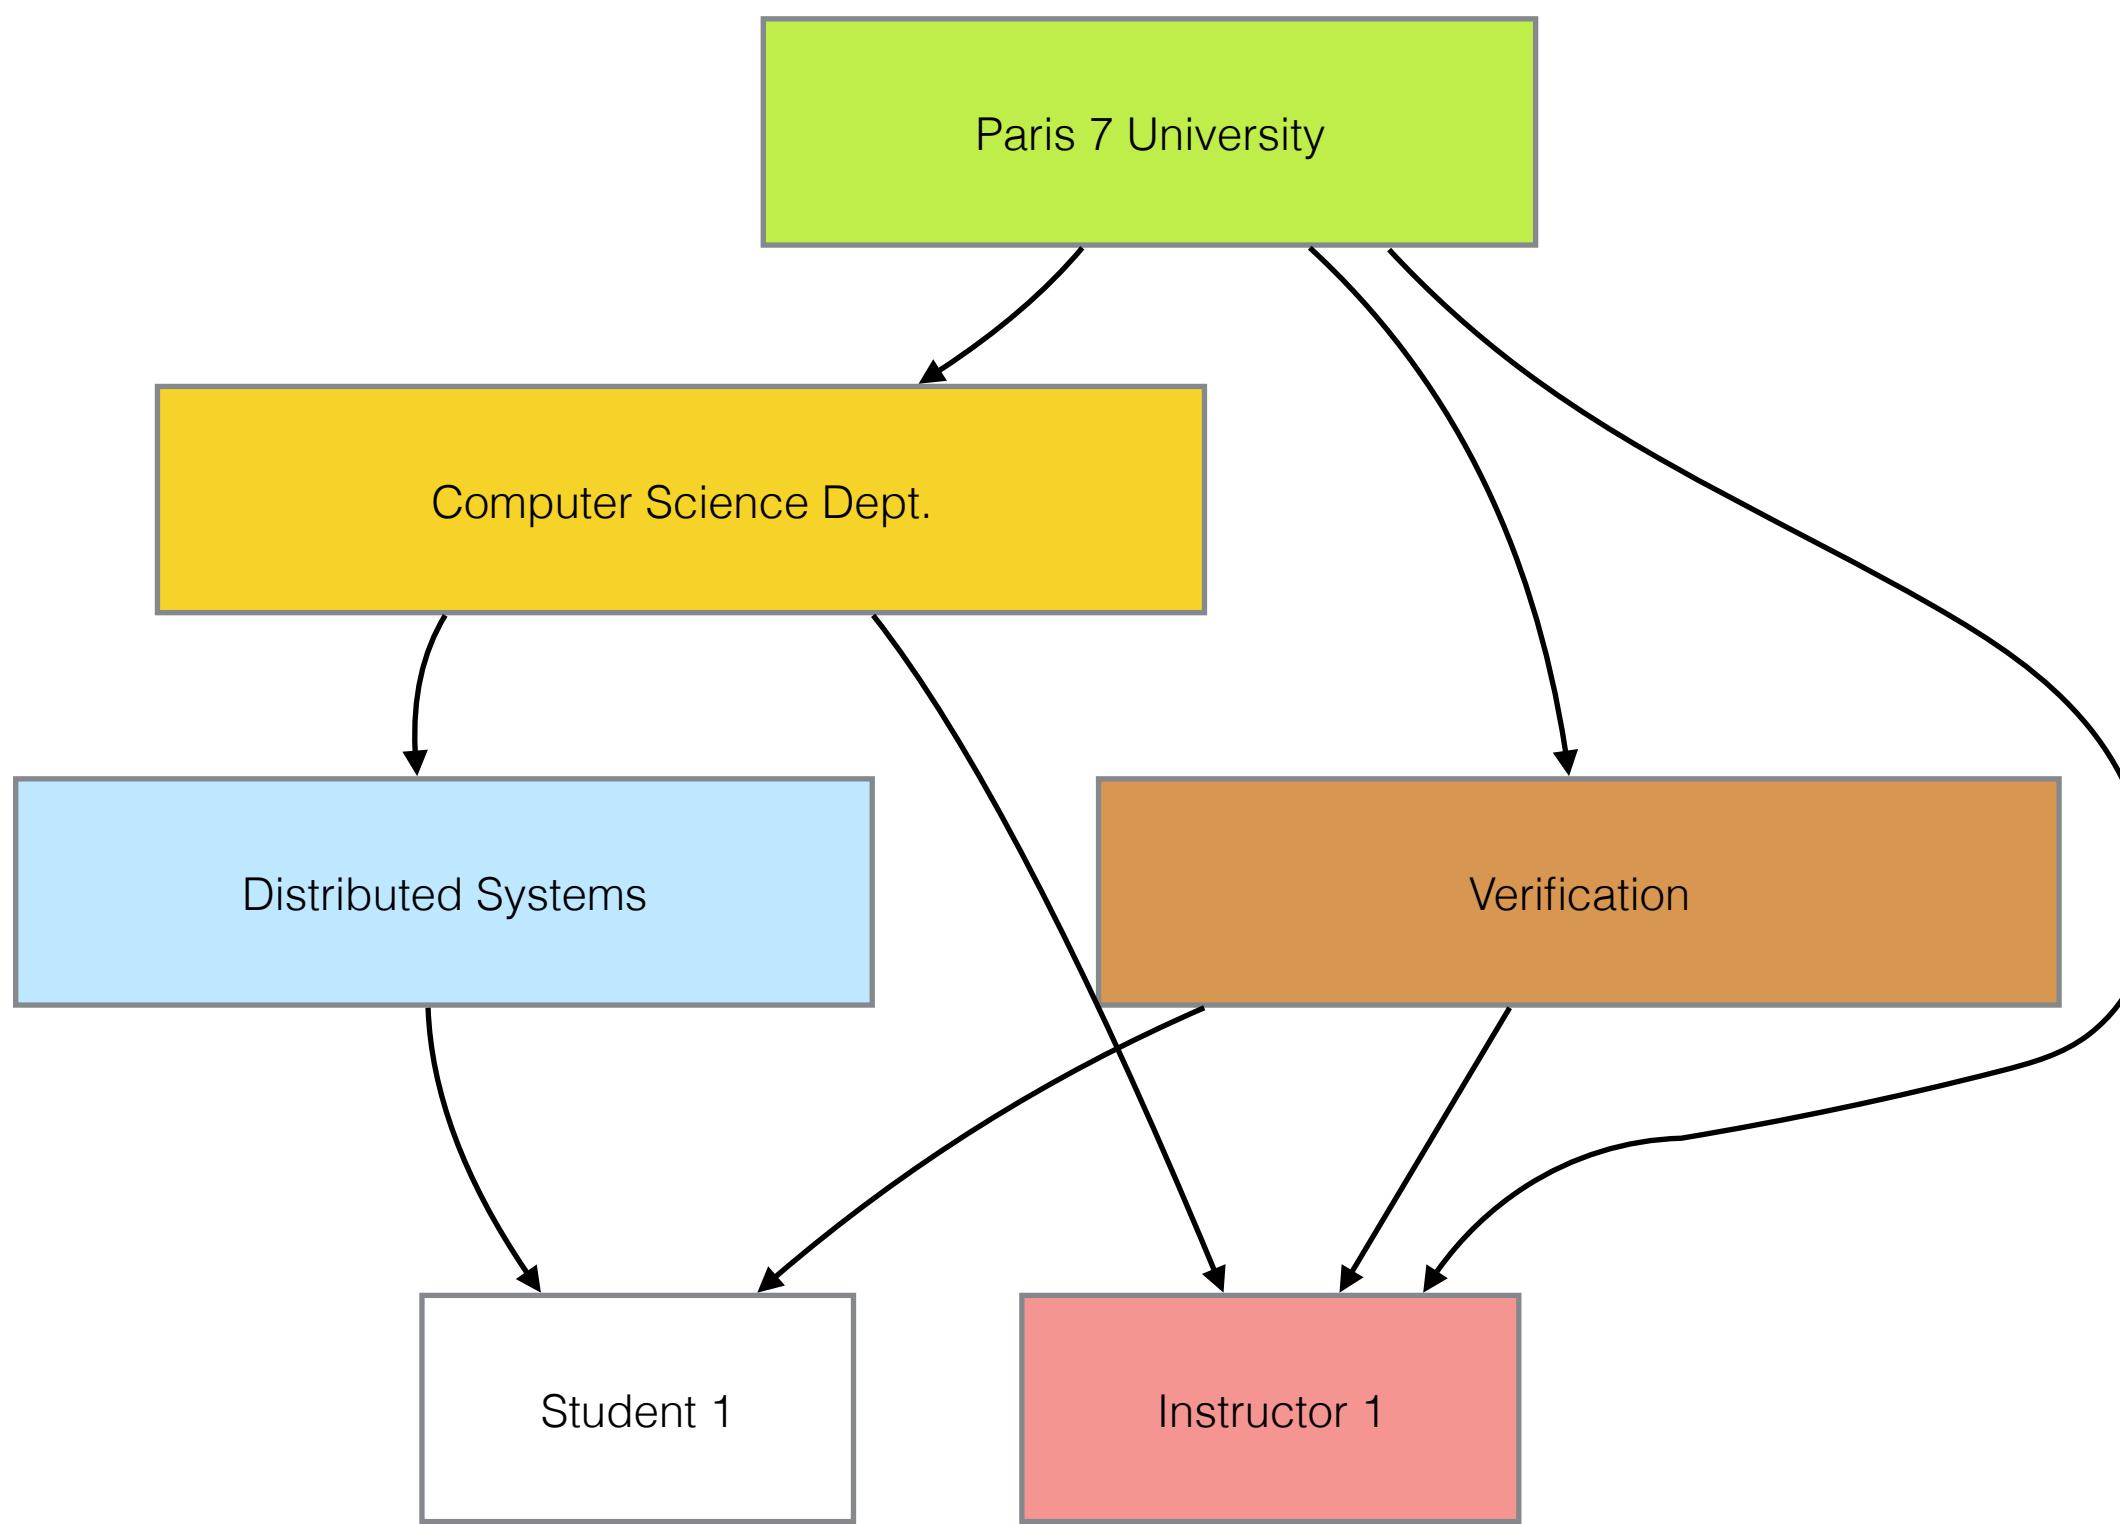
Two actors



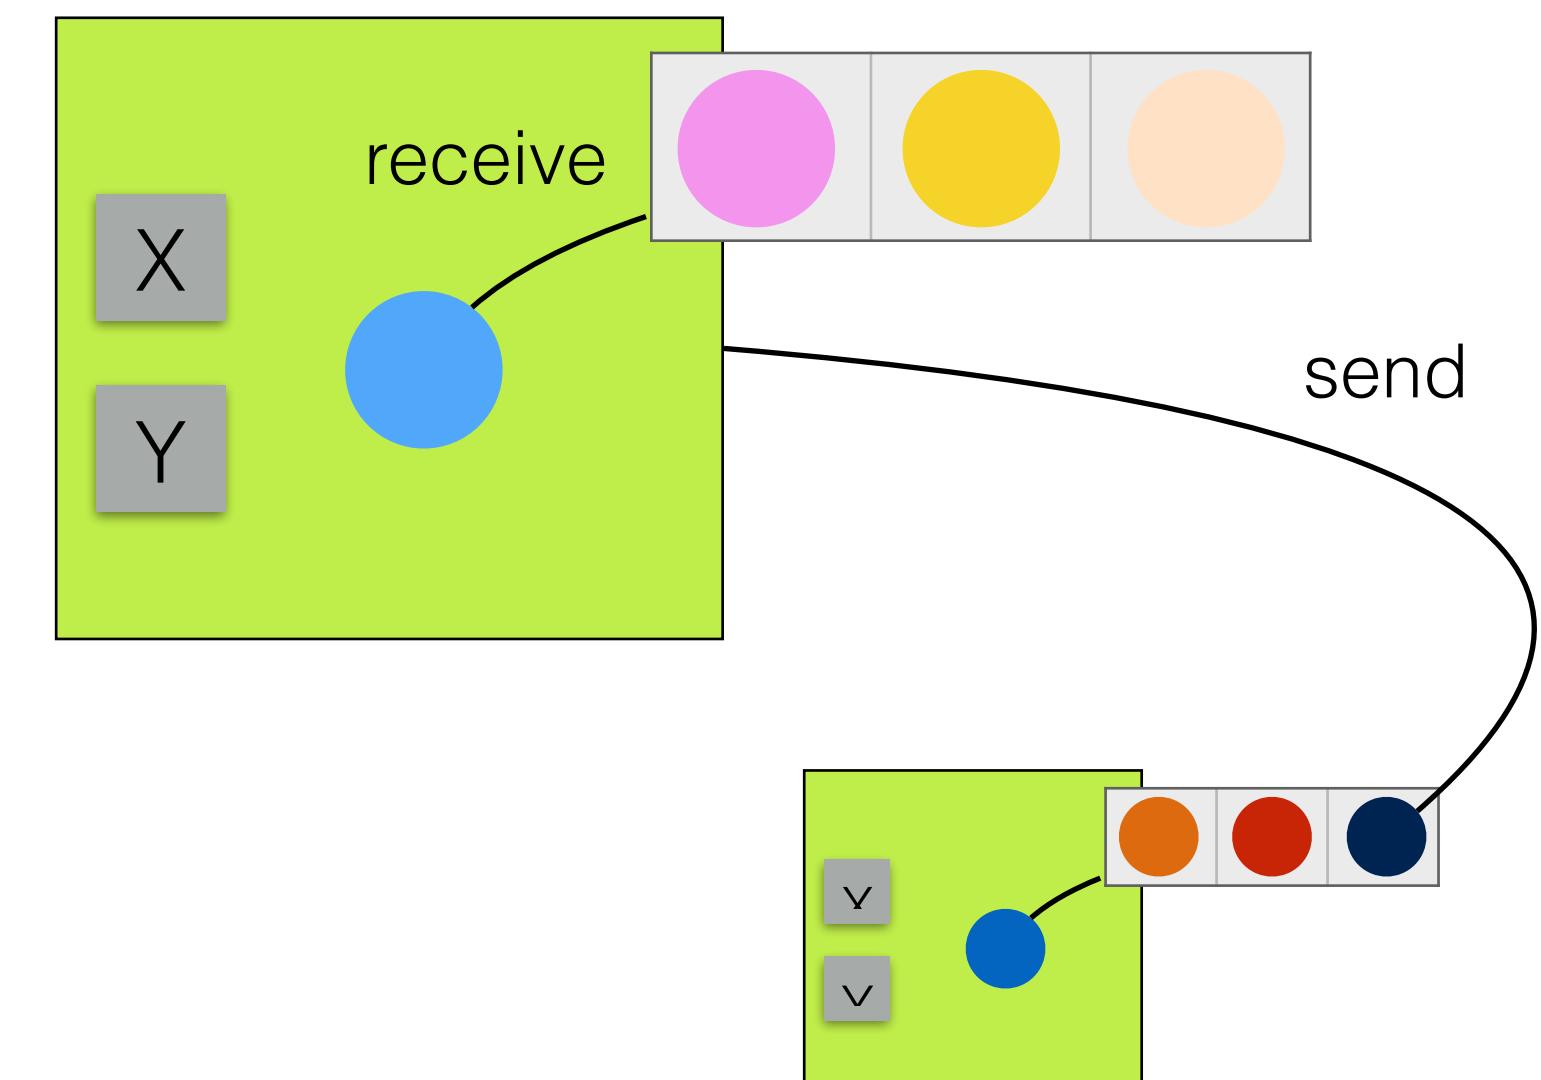
WHAT IS AN ACTOR?

- ▶ State
- ▶ Messages
- ▶ Inbox
- ▶ Supervision

An actor system runtime



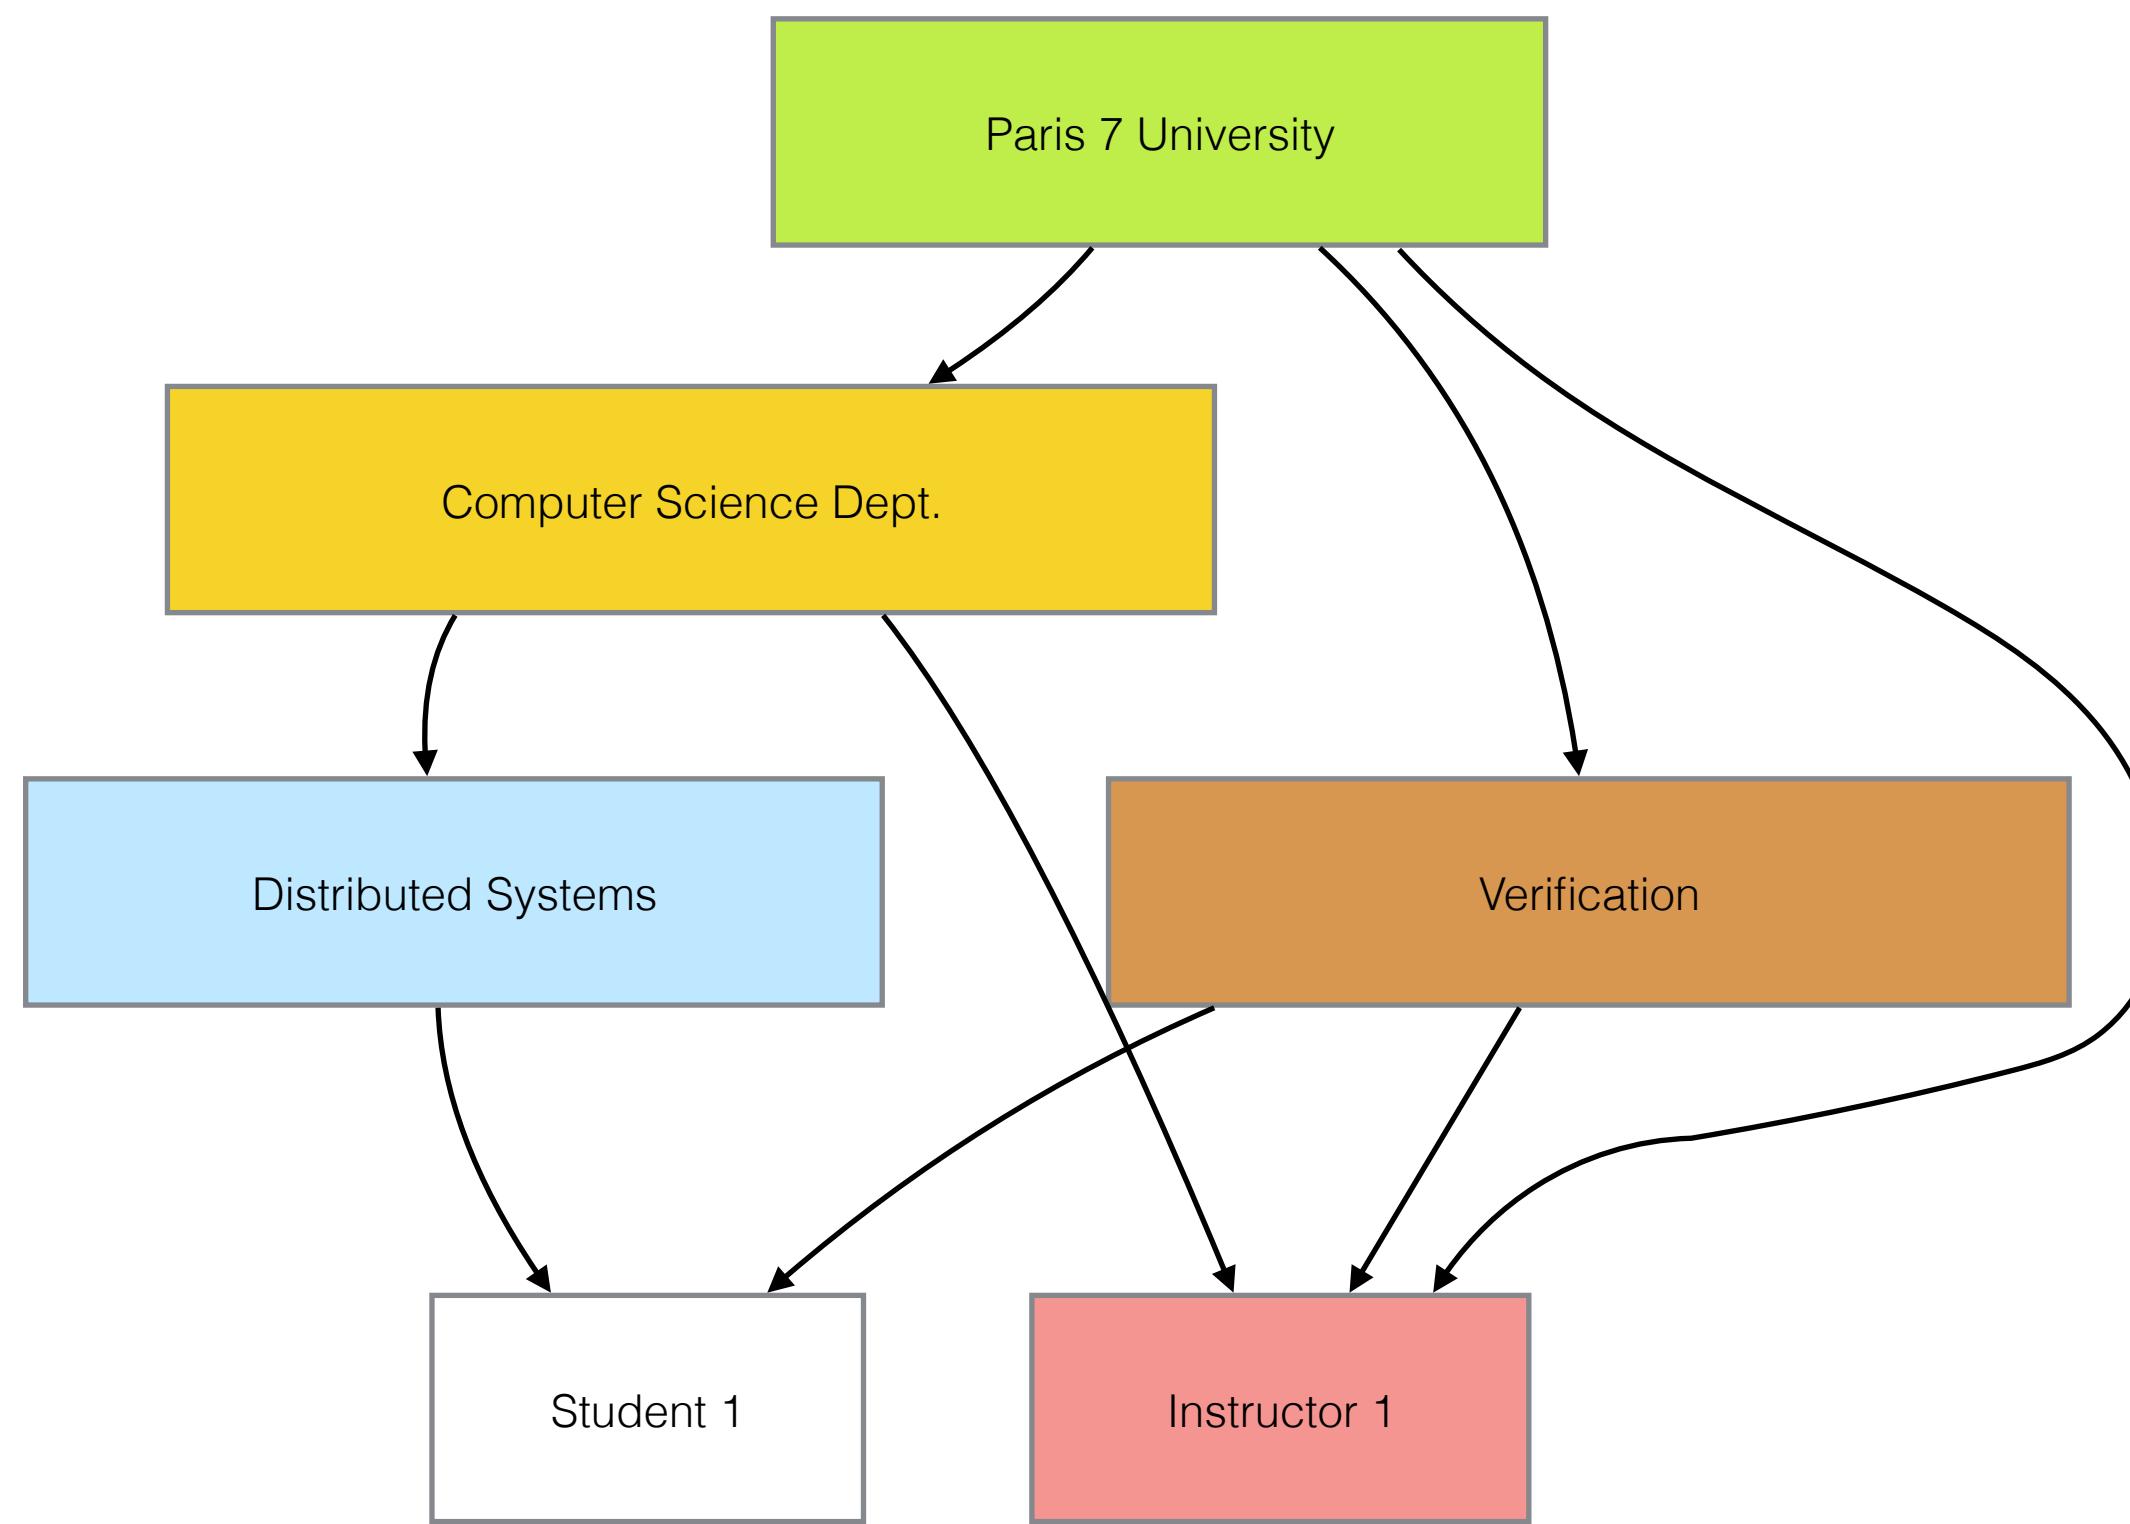
Two actors



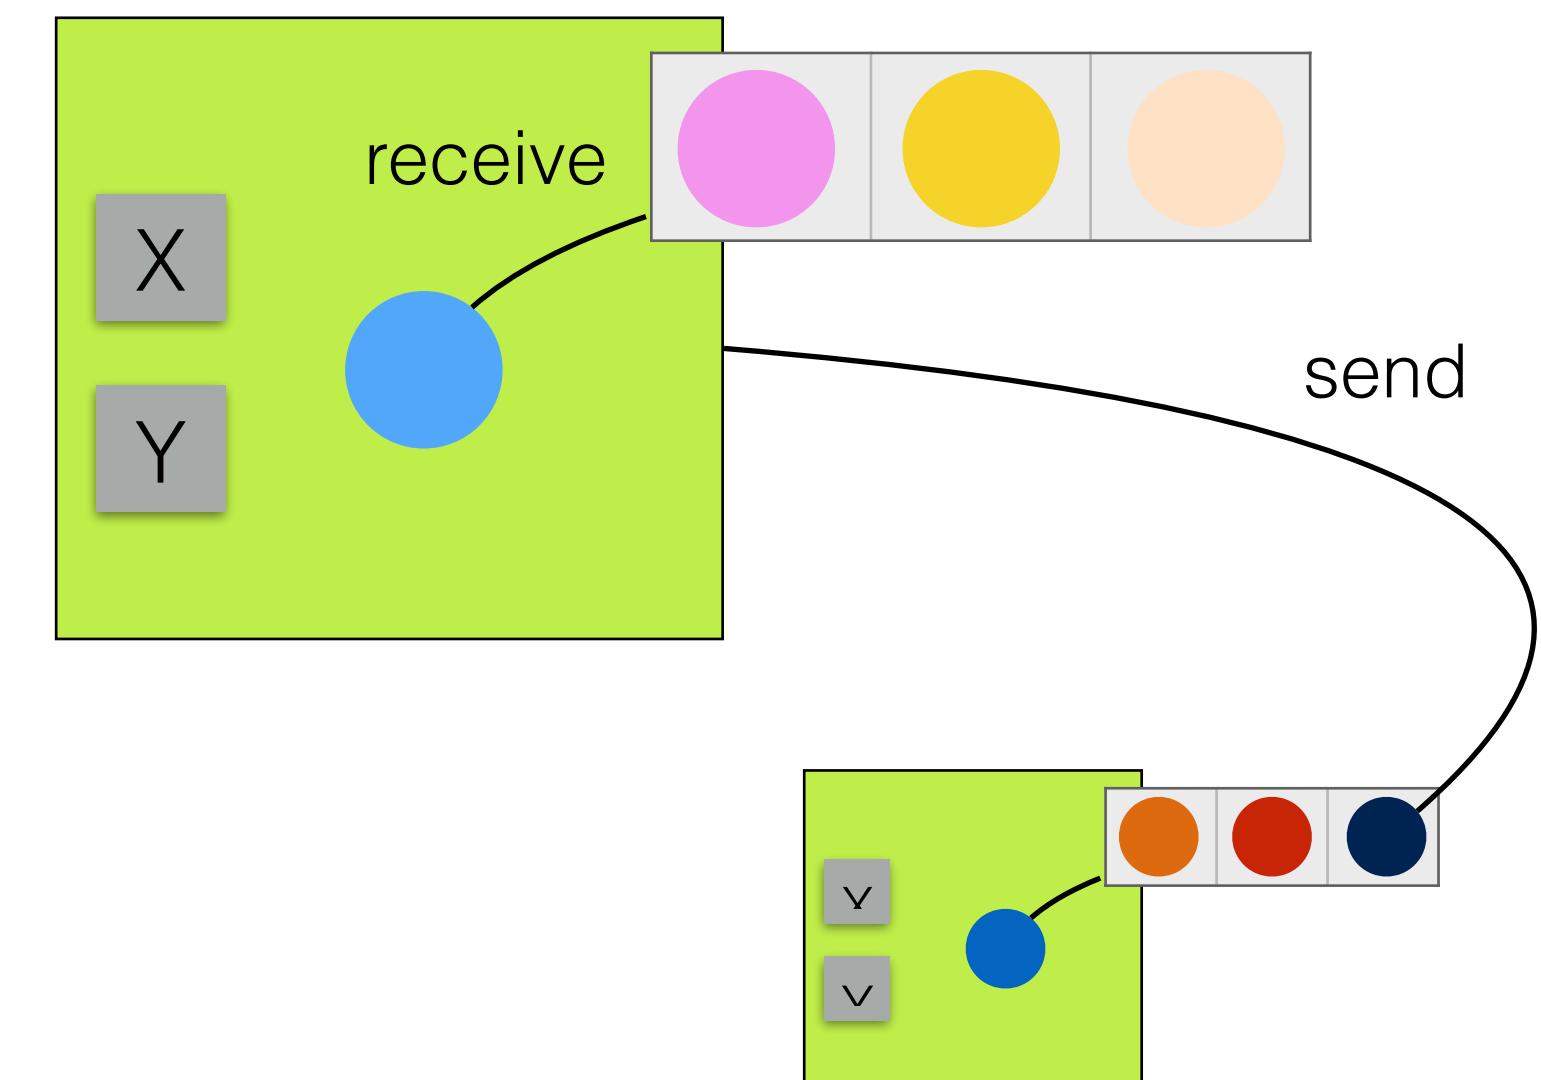
WHAT IS AN ACTOR?

- ▶ State
- ▶ Messages
- ▶ Inbox
- ▶ Supervision
- ▶ Routing (Path)

An actor system runtime

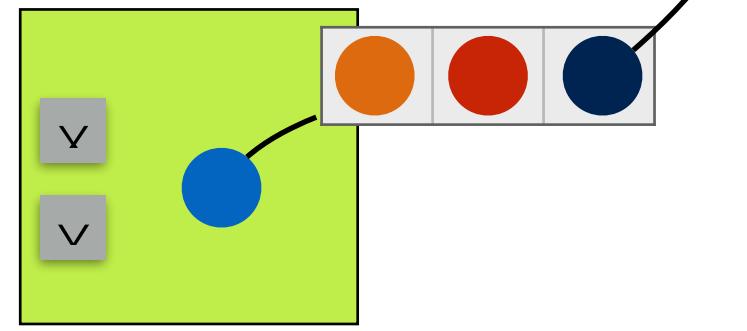
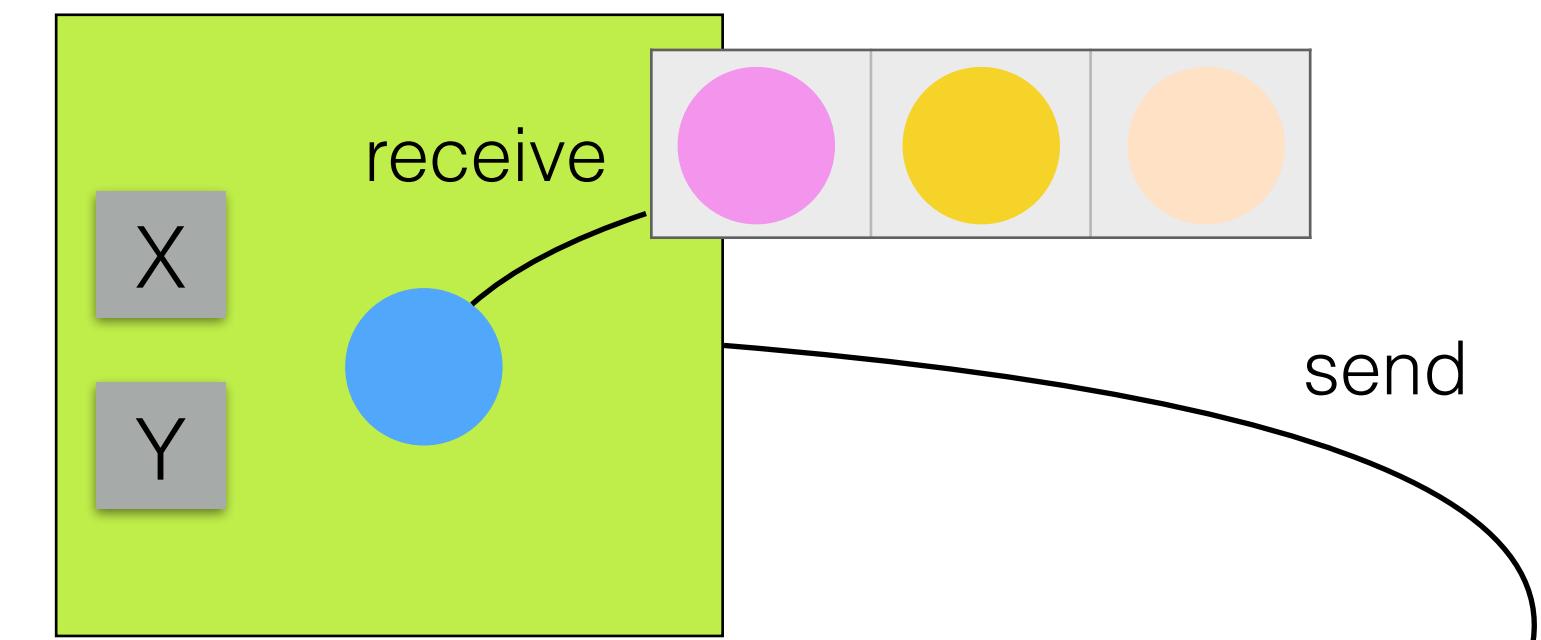


Two actors



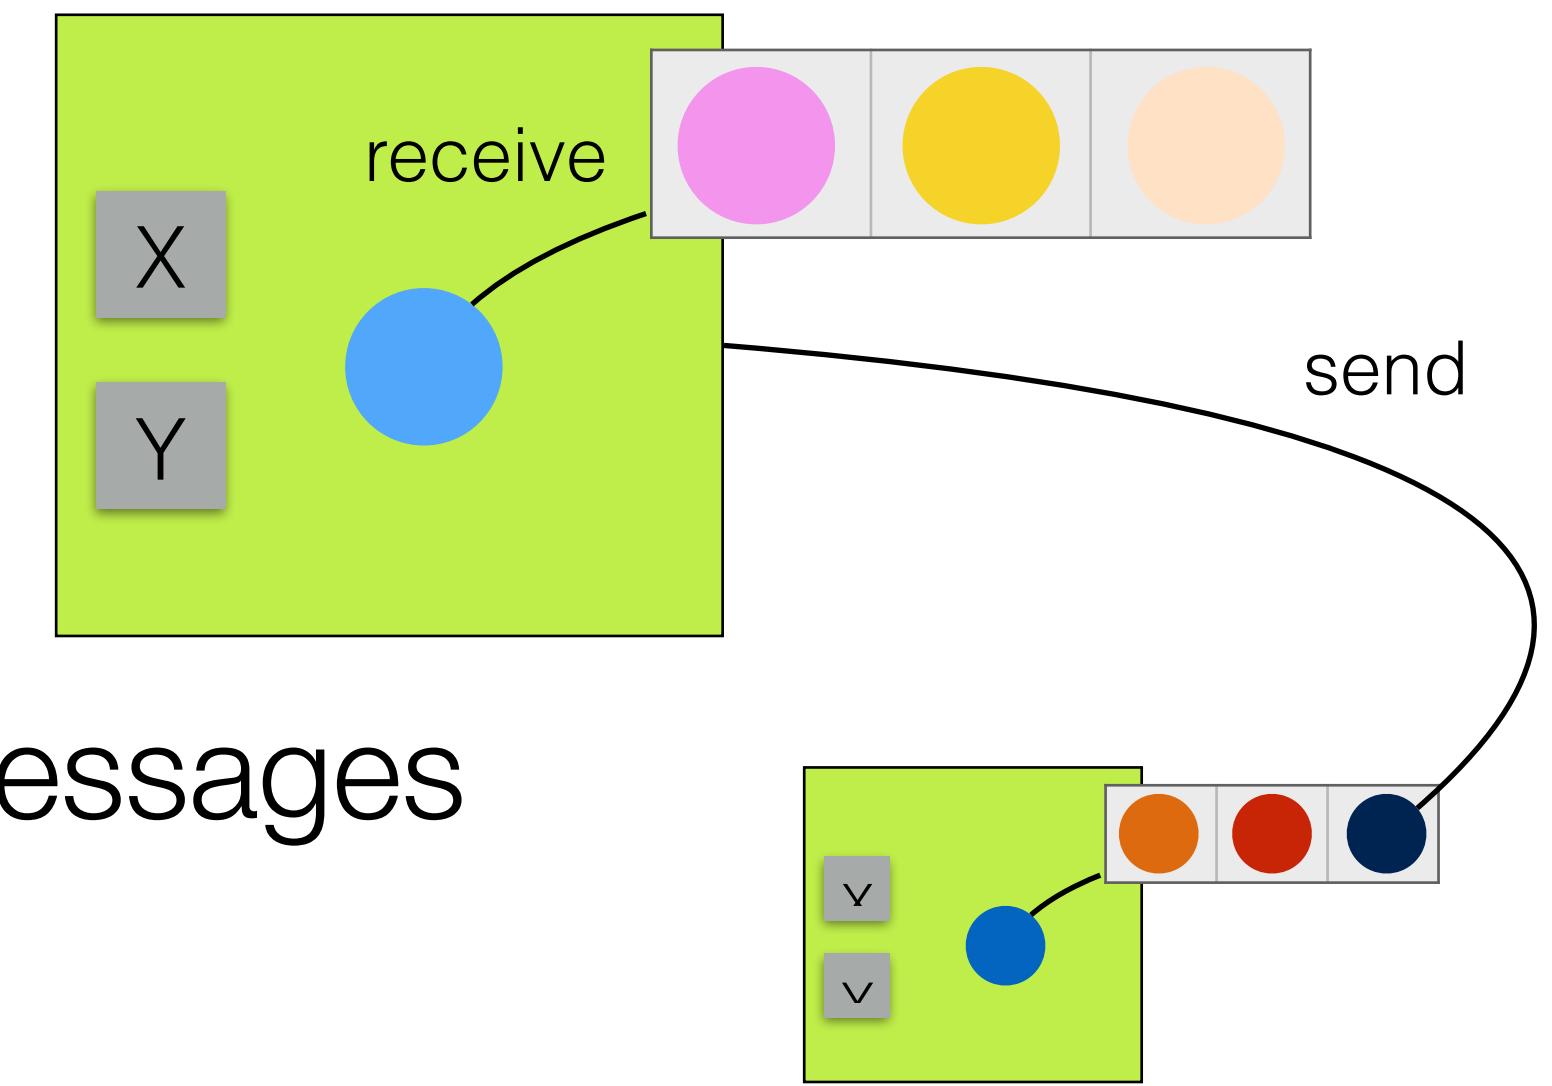
STATE AND COMPUTATION

- ▶ An class of Actors extends the Actor trait
- ▶ It can have state (like a normal class)
- ▶ It implements the receive method to respond to messages
- ▶ To send a message we use !
- ▶ BehaviorAndState.scala
- ▶ The path



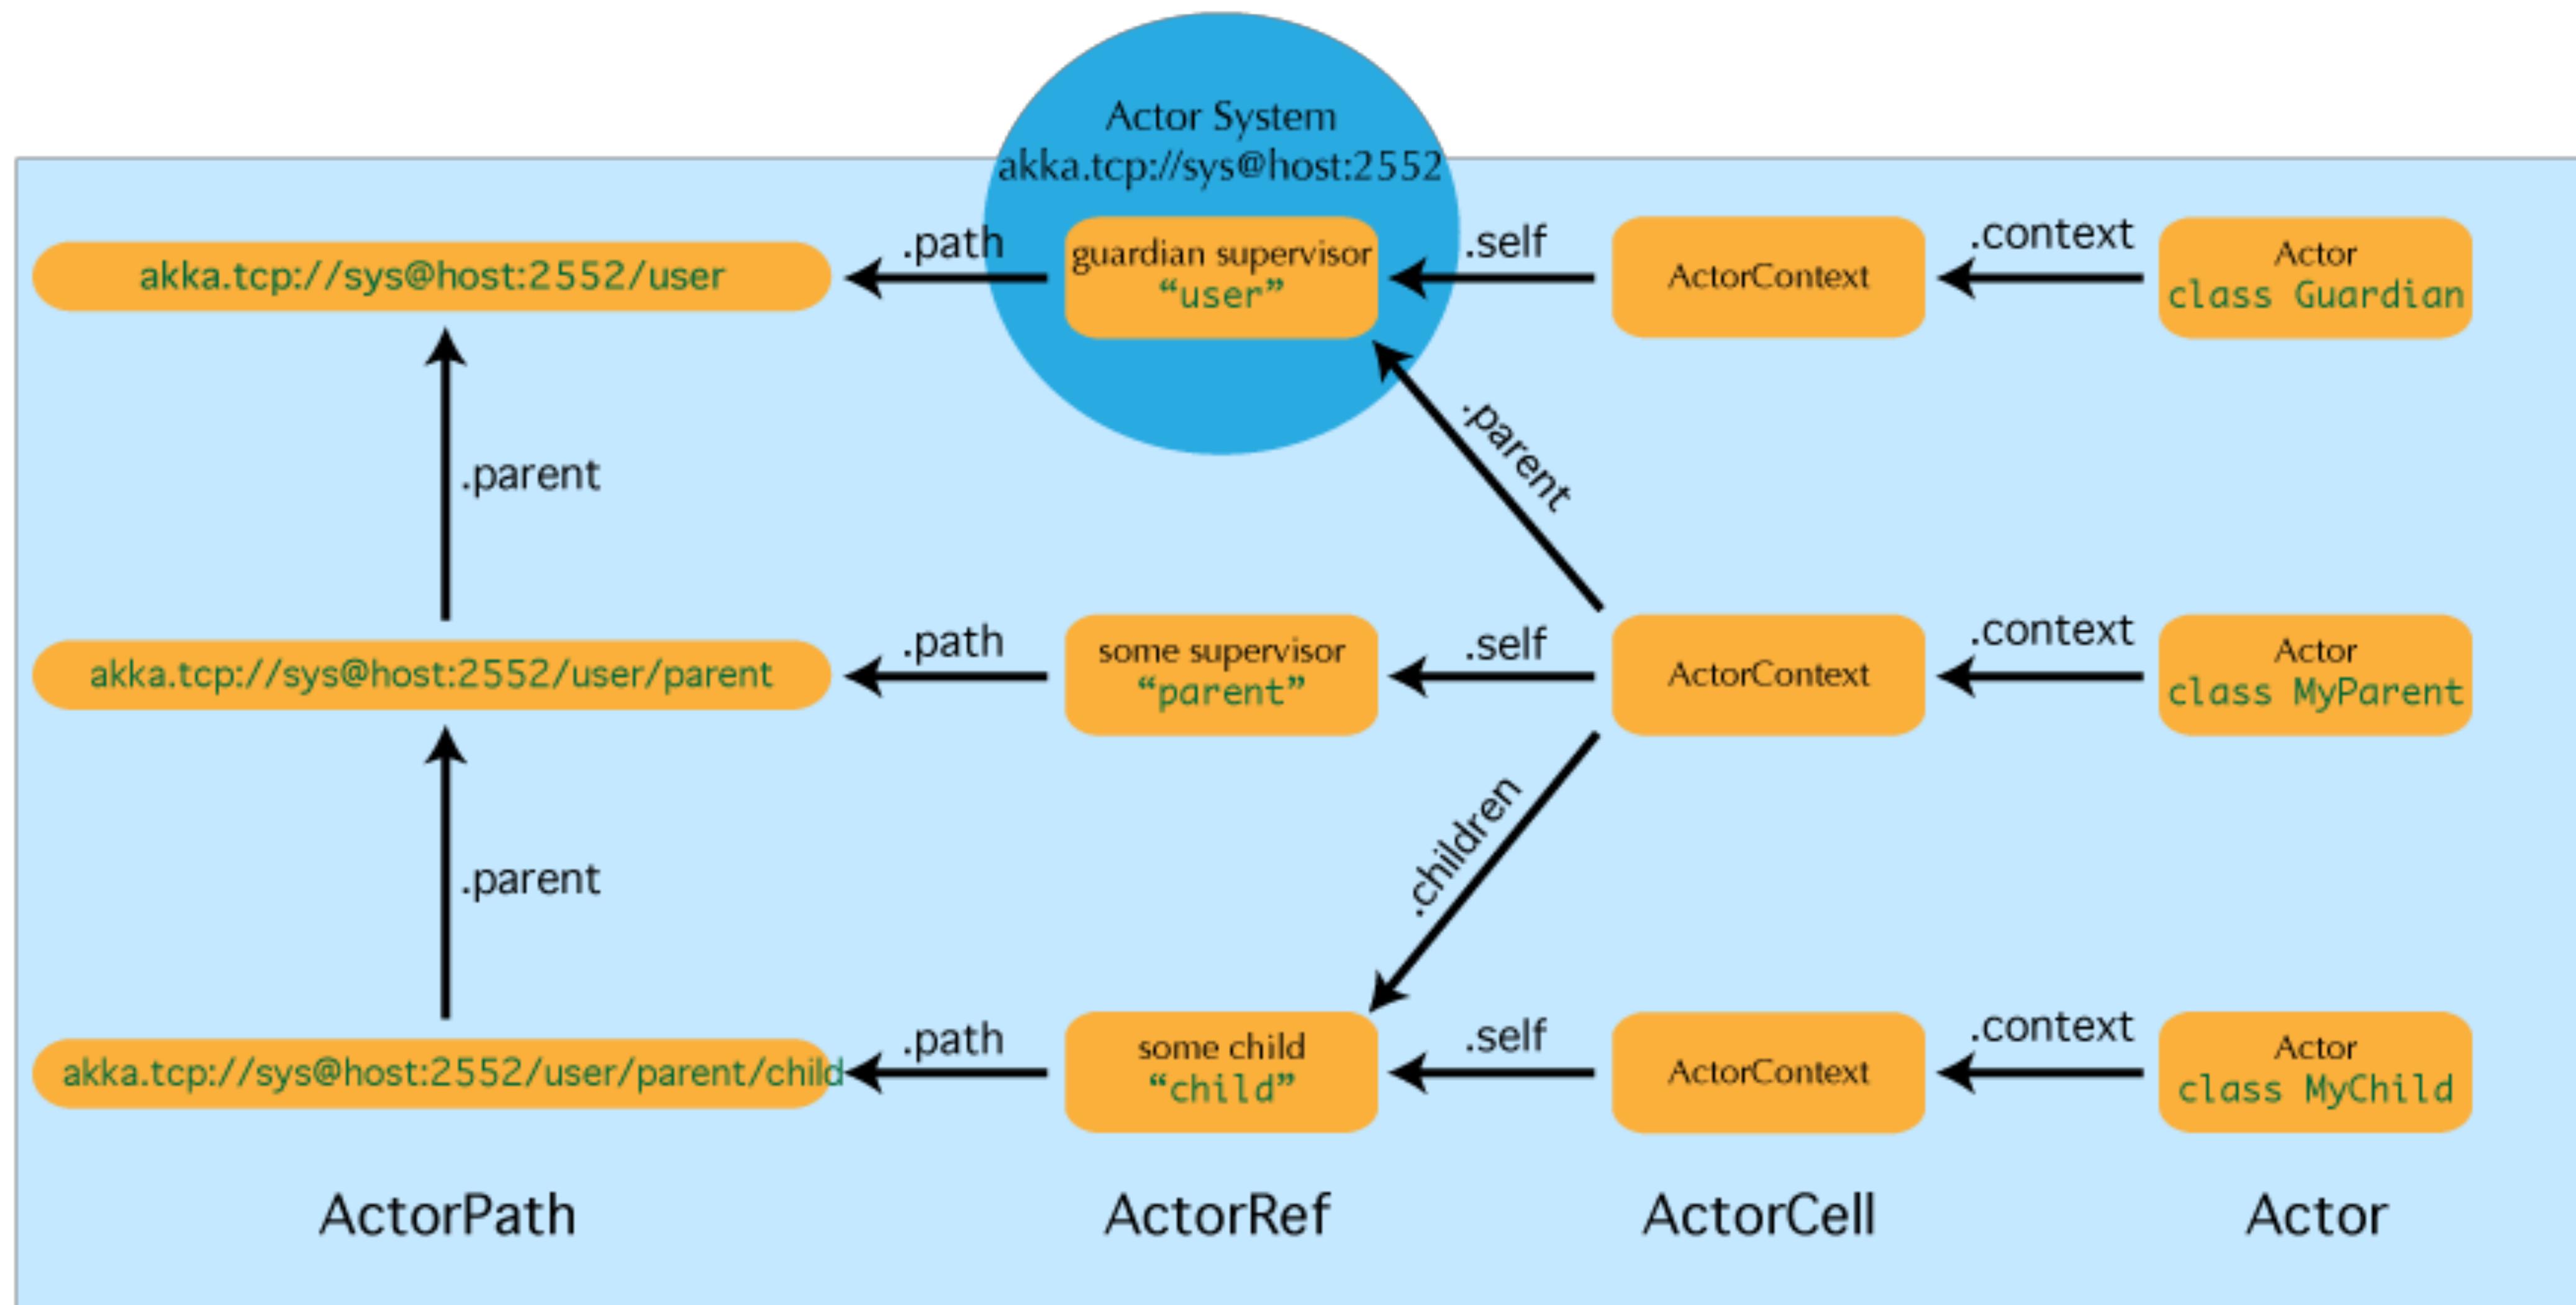
STATE AND COMPUTATION

- ▶ An class of Actors extends the Actor trait
- ▶ It can have state (like a normal class)
- ▶ It implements the receive method to respond to messages
- ▶ To send a message we use !
- ▶ BehaviorAndState.scala
- ▶ The path



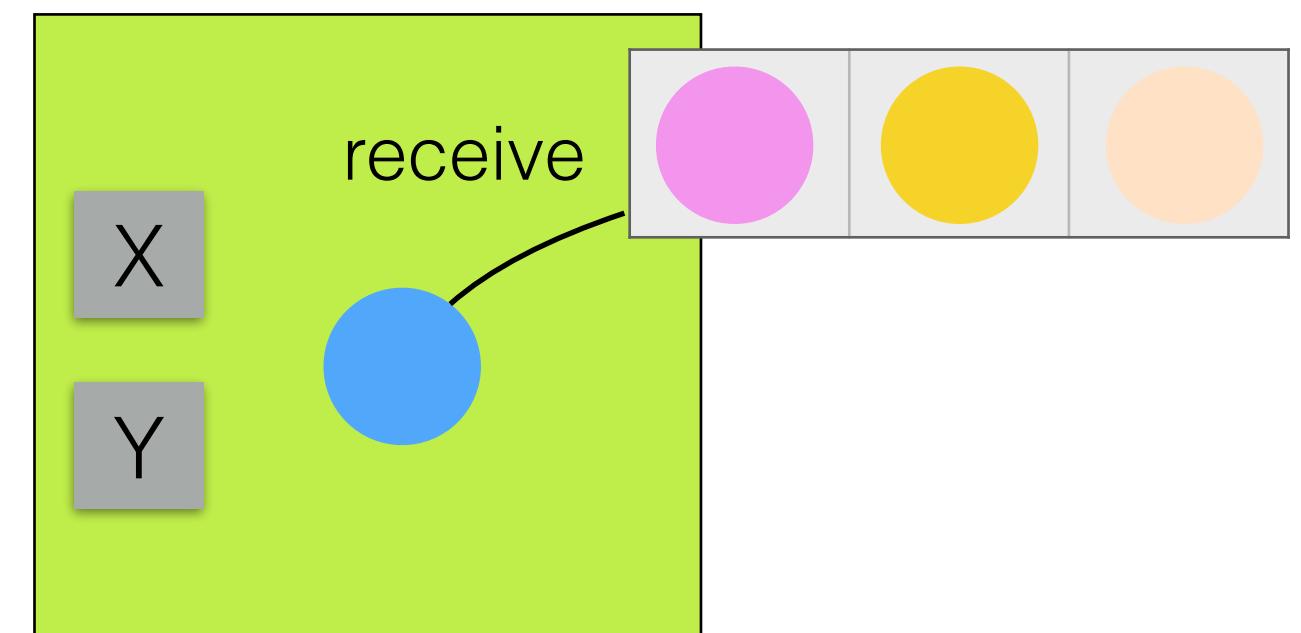
◉ **Exercise 2:** Create actors for a University Department, for Courses and for Students

PATH OF AN ACTOR



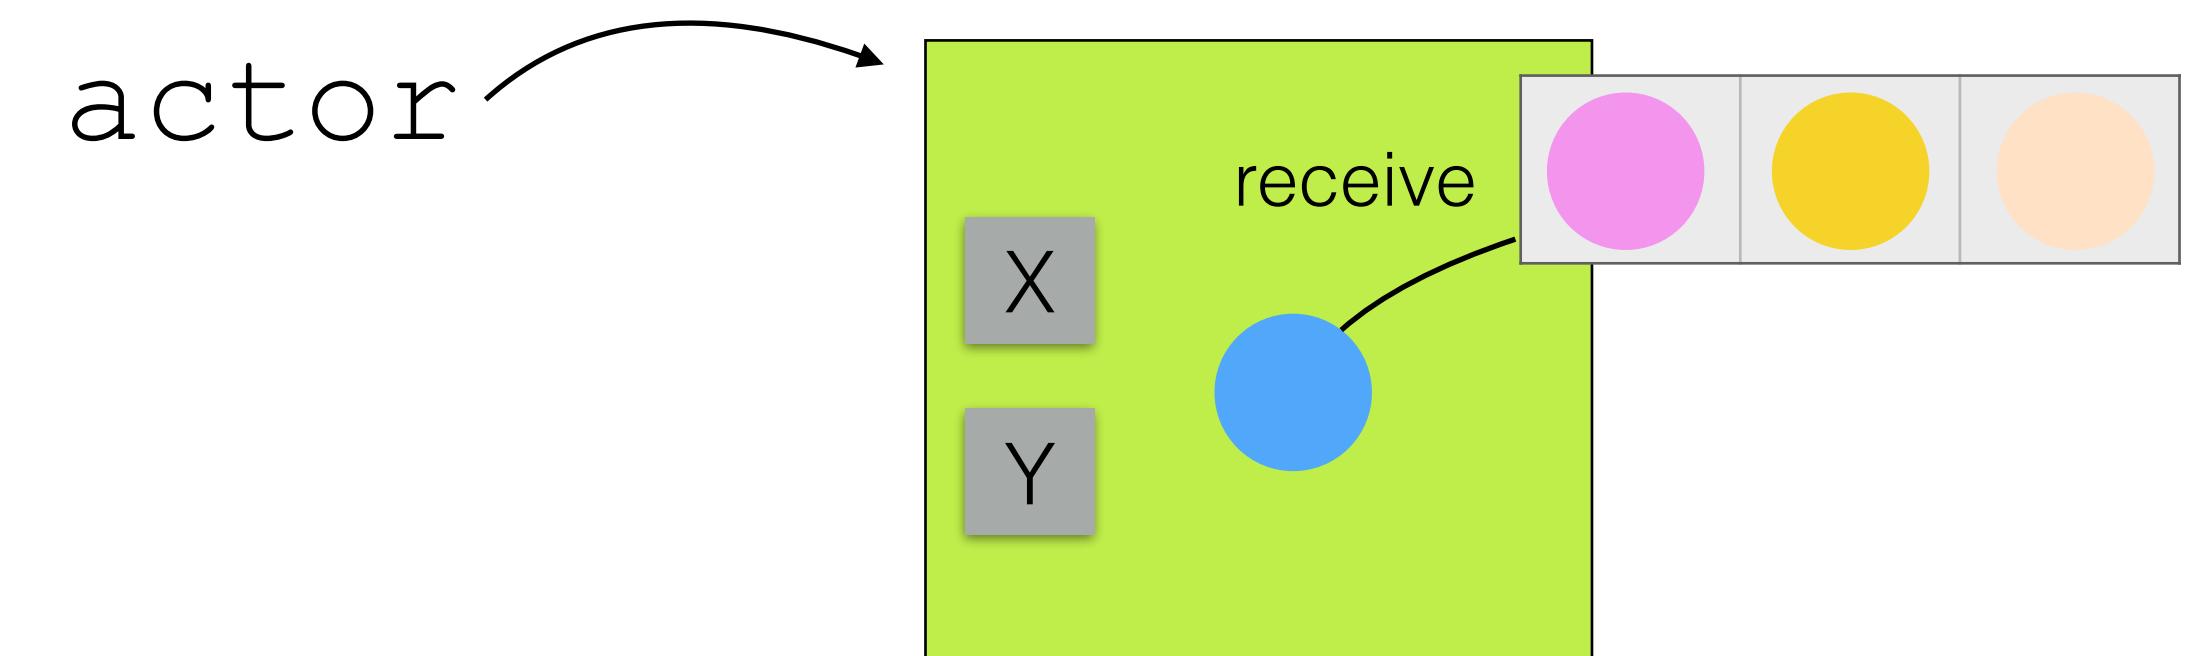
SENDING A MESSAGE TO AN ACTOR

- ▶ How can we refer to an actor?



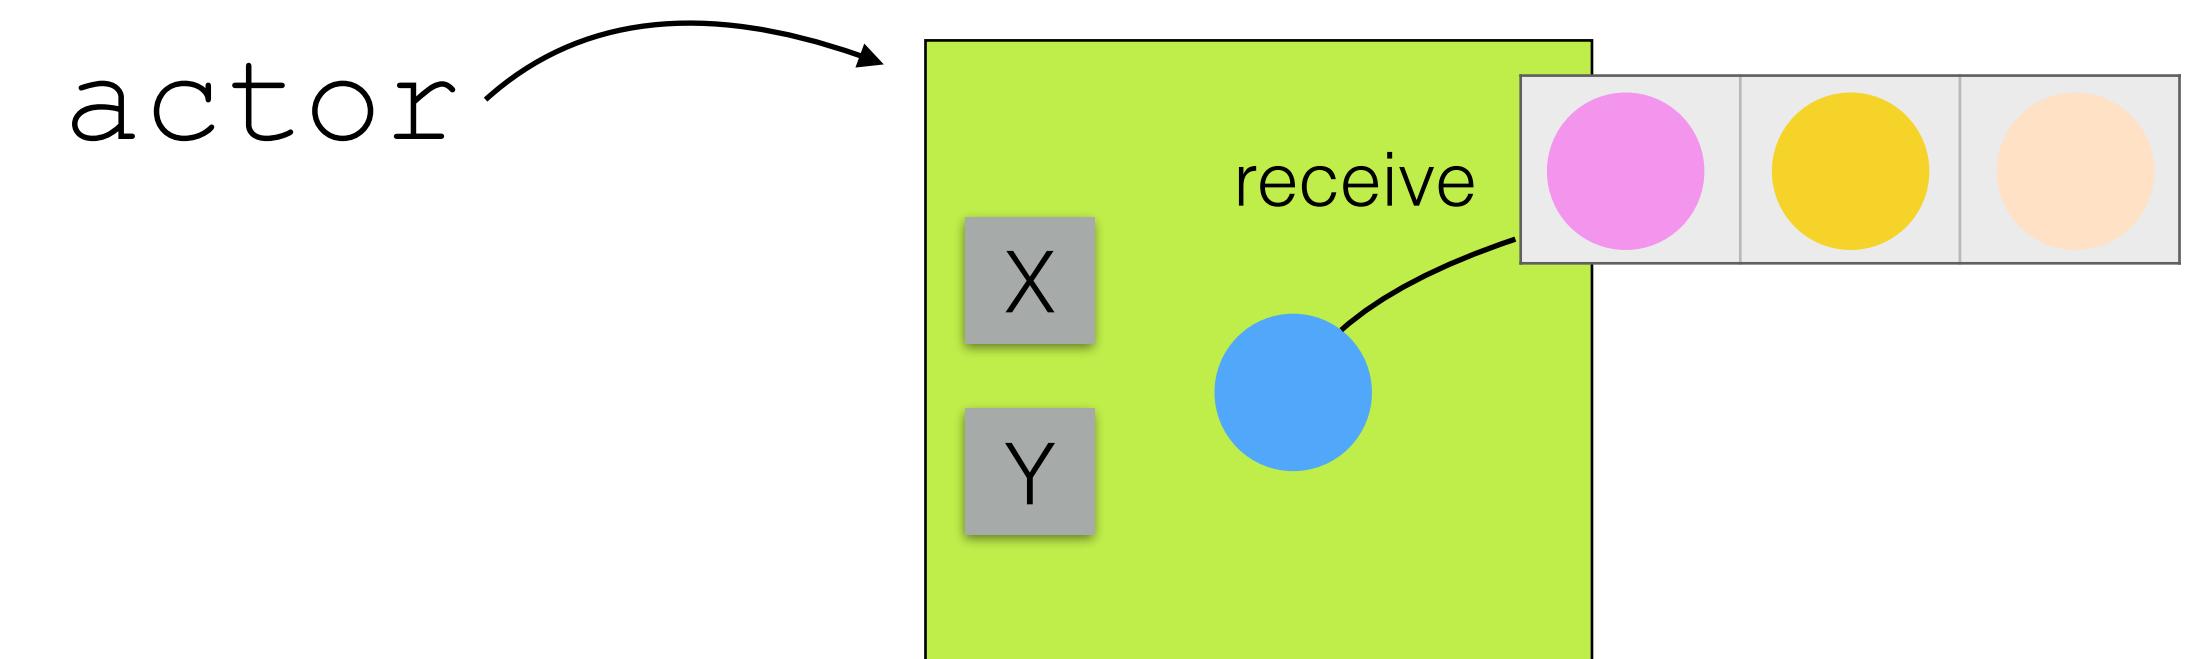
SENDING A MESSAGE TO AN ACTOR

- ▶ How can we refer to an actor?
- ▶ The creation of an actor returns an `ActorRef`, a reference to the actor



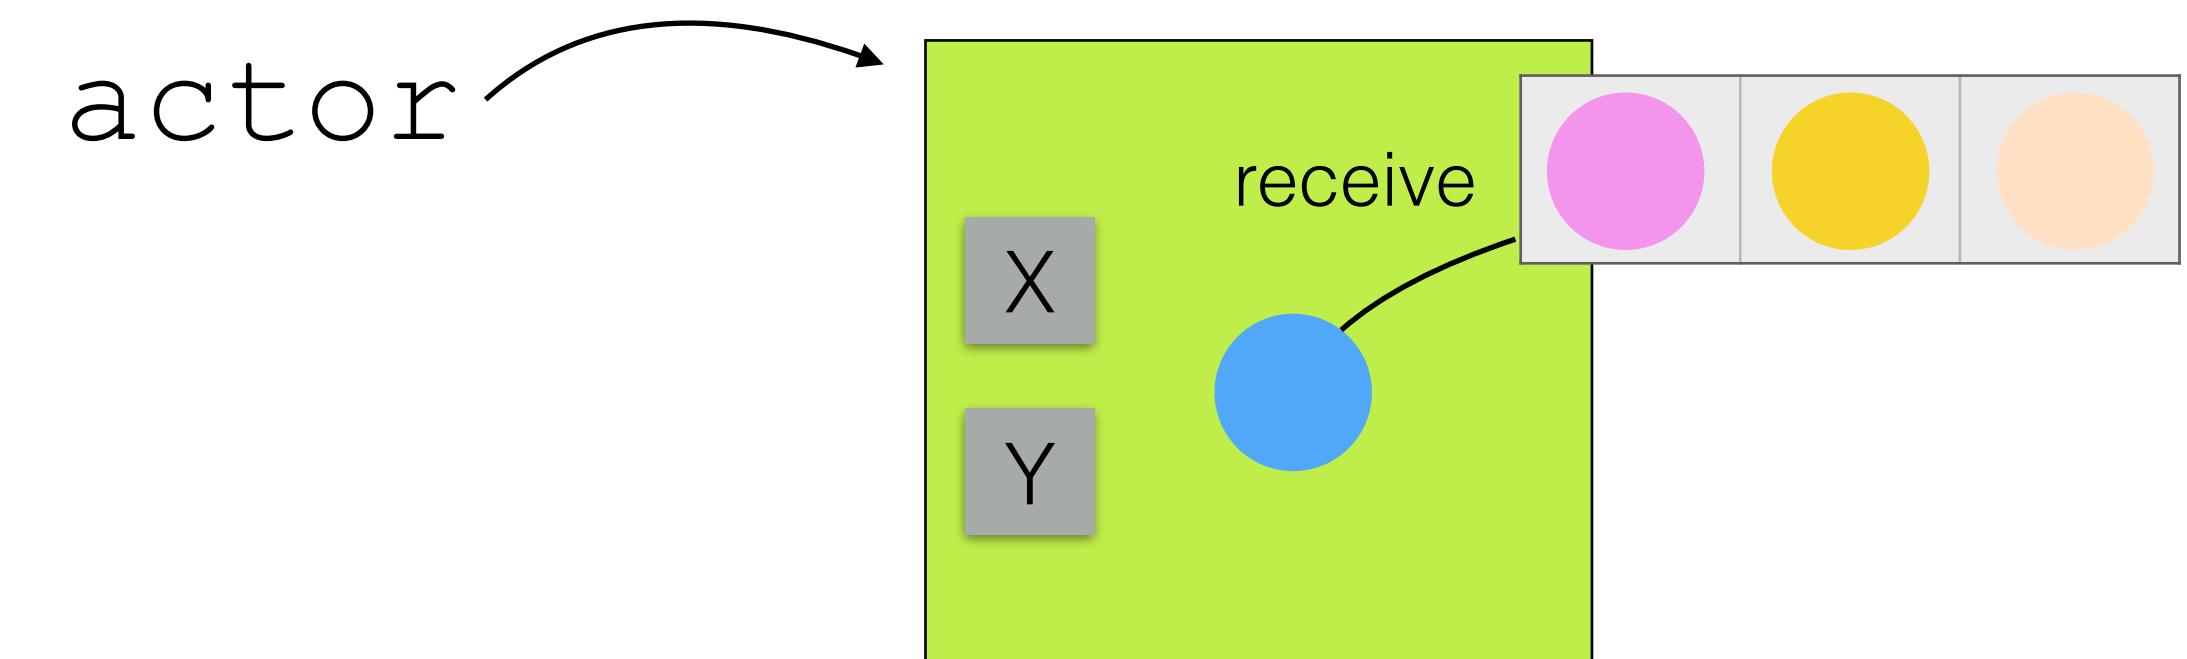
SENDING A MESSAGE TO AN ACTOR

- ▶ How can we refer to an actor?
 - ▶ The creation of an actor returns an `ActorRef`, a reference to the actor
 - ▶ We can send a message to an actor through this reference: `actor ! 1`



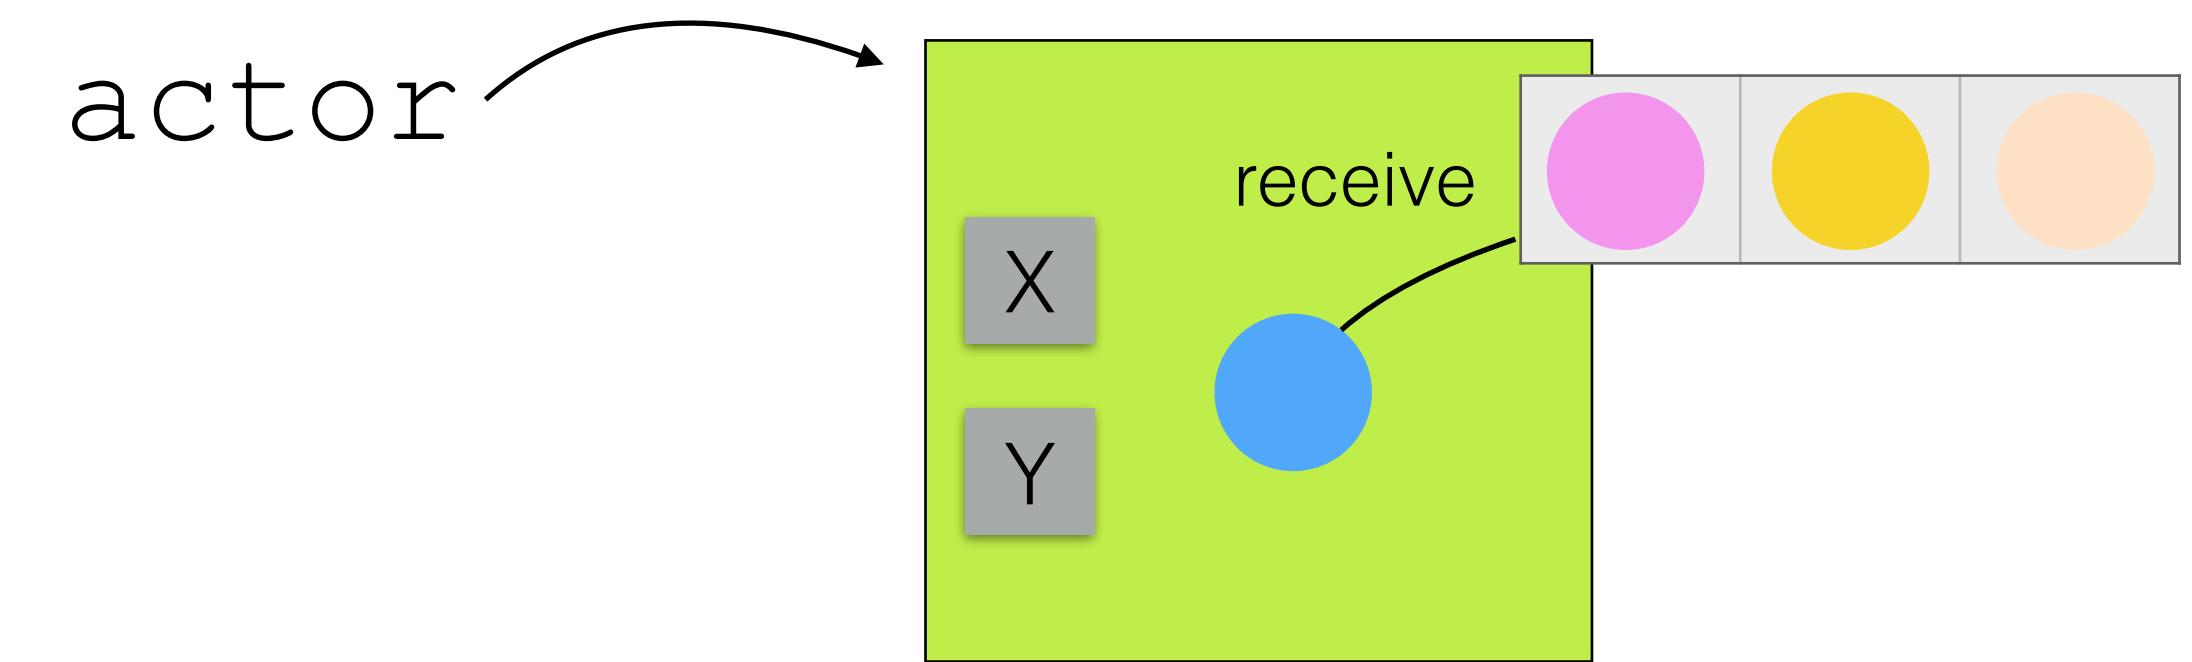
SENDING A MESSAGE TO AN ACTOR

- ▶ How can we refer to an actor?
 - ▶ The creation of an actor returns an `ActorRef`, a reference to the actor
 - ▶ We can send a message to an actor through this reference: `actor ! 1`
 - ▶ The message is dispatched to the mailbox, and executed later



SENDING A MESSAGE TO AN ACTOR

- ▶ How can we refer to an actor?
 - ▶ The creation of an actor returns an `ActorRef`, a reference to the actor
 - ▶ We can send a message to an actor through this reference: `actor ! 1`
 - ▶ The message is dispatched to the mailbox, and executed later



◉ **Exercise 3:** A small application with two departments of UMP6, a few courses and a few students. Have the students register to courses

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?
 - ▶ Each request defines a special ActorRef sender
 - ▶ The receiver sends a message to the sender

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?
 - ▶ Each request defines a special ActorRef sender
 - ▶ The receiver sends a message to the sender
- ▶ akka.pattern.ask to get answers

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?
 - ▶ Each request defines a special ActorRef sender
 - ▶ The receiver sends a message to the sender
 - ▶ akka.pattern.ask to get answers
 - ▶ The answer is not immediately present (since messages are asynchronous)
 - ▶ Therefore we must use futures

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?
 - ▶ Each request defines a special ActorRef sender
 - ▶ The receiver sends a message to the sender
 - ▶ akka.pattern.ask to get answers
 - ▶ The answer is not immediately present (since messages are asynchronous)
 - ▶ Therefore we must use futures
 - ▶ FiboActorApp.scala

GETTING AN ANSWER

- ▶ How do we get an answer from an actor?
 - ▶ Each request defines a special ActorRef sender
 - ▶ The receiver sends a message to the sender
 - ▶ akka.pattern.ask to get answers
 - ▶ The answer is not immediately present (since messages are asynchronous)
 - ▶ Therefore we must use futures
 - ▶ FiboActorApp.scala

● **Exercise 4:** Define a maximum capacity for courses and confirm or reject student registrations

COMMUNICATION BETWEEN OF ACTORS

- ▶ Actors can talk back and forth with each other
- ▶ Communication.scala

COMMUNICATION BETWEEN OF ACTORS

- ▶ Actors can talk back and forth with each other
- ▶ Communication.scala

○ **Exercise 5:** Assuming the students are registered in the department, have the registration in a course confirm with the department if the student is registered at the university

FAULT TOLERANCE

FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
- ▶ Catastrophic failures: floods, fires, etc.



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.
 - ▶ Bugs happen: ALL THE TIME!!!



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.
 - ▶ Bugs happen: ALL THE TIME!!!
- ▶ Graceful Degradation



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.
 - ▶ Bugs happen: ALL THE TIME!!!
- ▶ Graceful Degradation
 - ▶ Decreased Throughput: meh ..., some clients will have to retry



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.
 - ▶ Bugs happen: ALL THE TIME!!!
- ▶ Graceful Degradation
 - ▶ Decreased Throughput: meh ..., some clients will have to retry
 - ▶ Increased Latency: meh ..., some clients will be unhappy waiting



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
 - ▶ Catastrophic failures: floods, fires, etc.
 - ▶ Infrastructures fail: electricity, network, etc.
 - ▶ Bugs happen: ALL THE TIME!!!
- ▶ Graceful Degradation
 - ▶ Decreased Throughput: meh ..., some clients will have to retry
 - ▶ Increased Latency: meh ..., some clients will be unhappy waiting
 - ▶ *Unavailable*: oh s*** ..., clients are going away!



FAULT TOLERANCE

- ▶ Bad things happen ~~sometimes~~ all the time
- ▶ Catastrophic failures: floods, fires, etc.
- ▶ Infrastructure failures
- ▶ Bugs happen
- ▶ Graceful Degradation
 - ▶ Slow might be acceptable
 - ▶ Retries might be acceptable
 - ▶ **Unreachable is not!**
- ▶ Decreased Throughput: meh ..., some clients will have to retry
- ▶ Increased Latency: meh ..., some clients will be unhappy waiting
- ▶ *Unavailable*: oh s*** ..., clients are going away!



FAULT TOLERANCE: ARCHITECTURE

FAULT TOLERANCE: ARCHITECTURE

- ▶ Compartmentalization
 - ▶ Separation of Concerns and Functionality
 - ▶ We can harden / restart critical sub-systems

FAULT TOLERANCE: ARCHITECTURE

- ▶ Compartmentalization
 - ▶ Separation of Concerns and Functionality
 - ▶ We can harden / restart critical sub-systems
- ▶ Avoid cascading Failures

FAULT TOLERANCE: ARCHITECTURE

- ▶ Compartmentalization
 - ▶ Separation of Concerns and Functionality
 - ▶ We can harden / restart critical sub-systems
- ▶ Avoid cascading Failures
- ▶ Isolate failures from important parts of the system

FAULT TOLERANCE: ARCHITECTURE

- ▶ Compartmentalization
 - ▶ Separation of Concerns and Functionality
 - ▶ We can harden / restart critical sub-systems
- ▶ Avoid cascading Failures
- ▶ Isolate failures from important parts of the system
- ▶ Back-up / Replicate / Scale Bottlenecks

FAULT TOLERANCE: MECHANISMS

FAULT TOLERANCE: MECHANISMS

- ▶ State Machine Replication
 - ▶ Multiple nodes perform the same action
 - ▶ Easy for state-less applications
 - ▶ Requires Synchronization for stateful applications



FAULT TOLERANCE: MECHANISMS

- ▶ State Machine Replication
 - ▶ Multiple nodes perform the same action
 - ▶ Easy for state-less applications
 - ▶ Requires Synchronization for stateful applications
- ▶ State Replication
 - ▶ Requires Synchronization
 - ▶ Or weakened consistency (this will be discussed Thu. or Fri.)



FAULT TOLERANCE: MECHANISMS

- ▶ State Machine Replication
 - ▶ Multiple nodes perform the same action
 - ▶ Easy for state-less applications
 - ▶ Requires Synchronization for stateful applications
- ▶ State Replication
 - ▶ Requires Synchronization
 - ▶ Or weakened consistency (this will be discussed Thu. or Fri.)
- ▶ Component Isolation
 - ▶ Message Passing



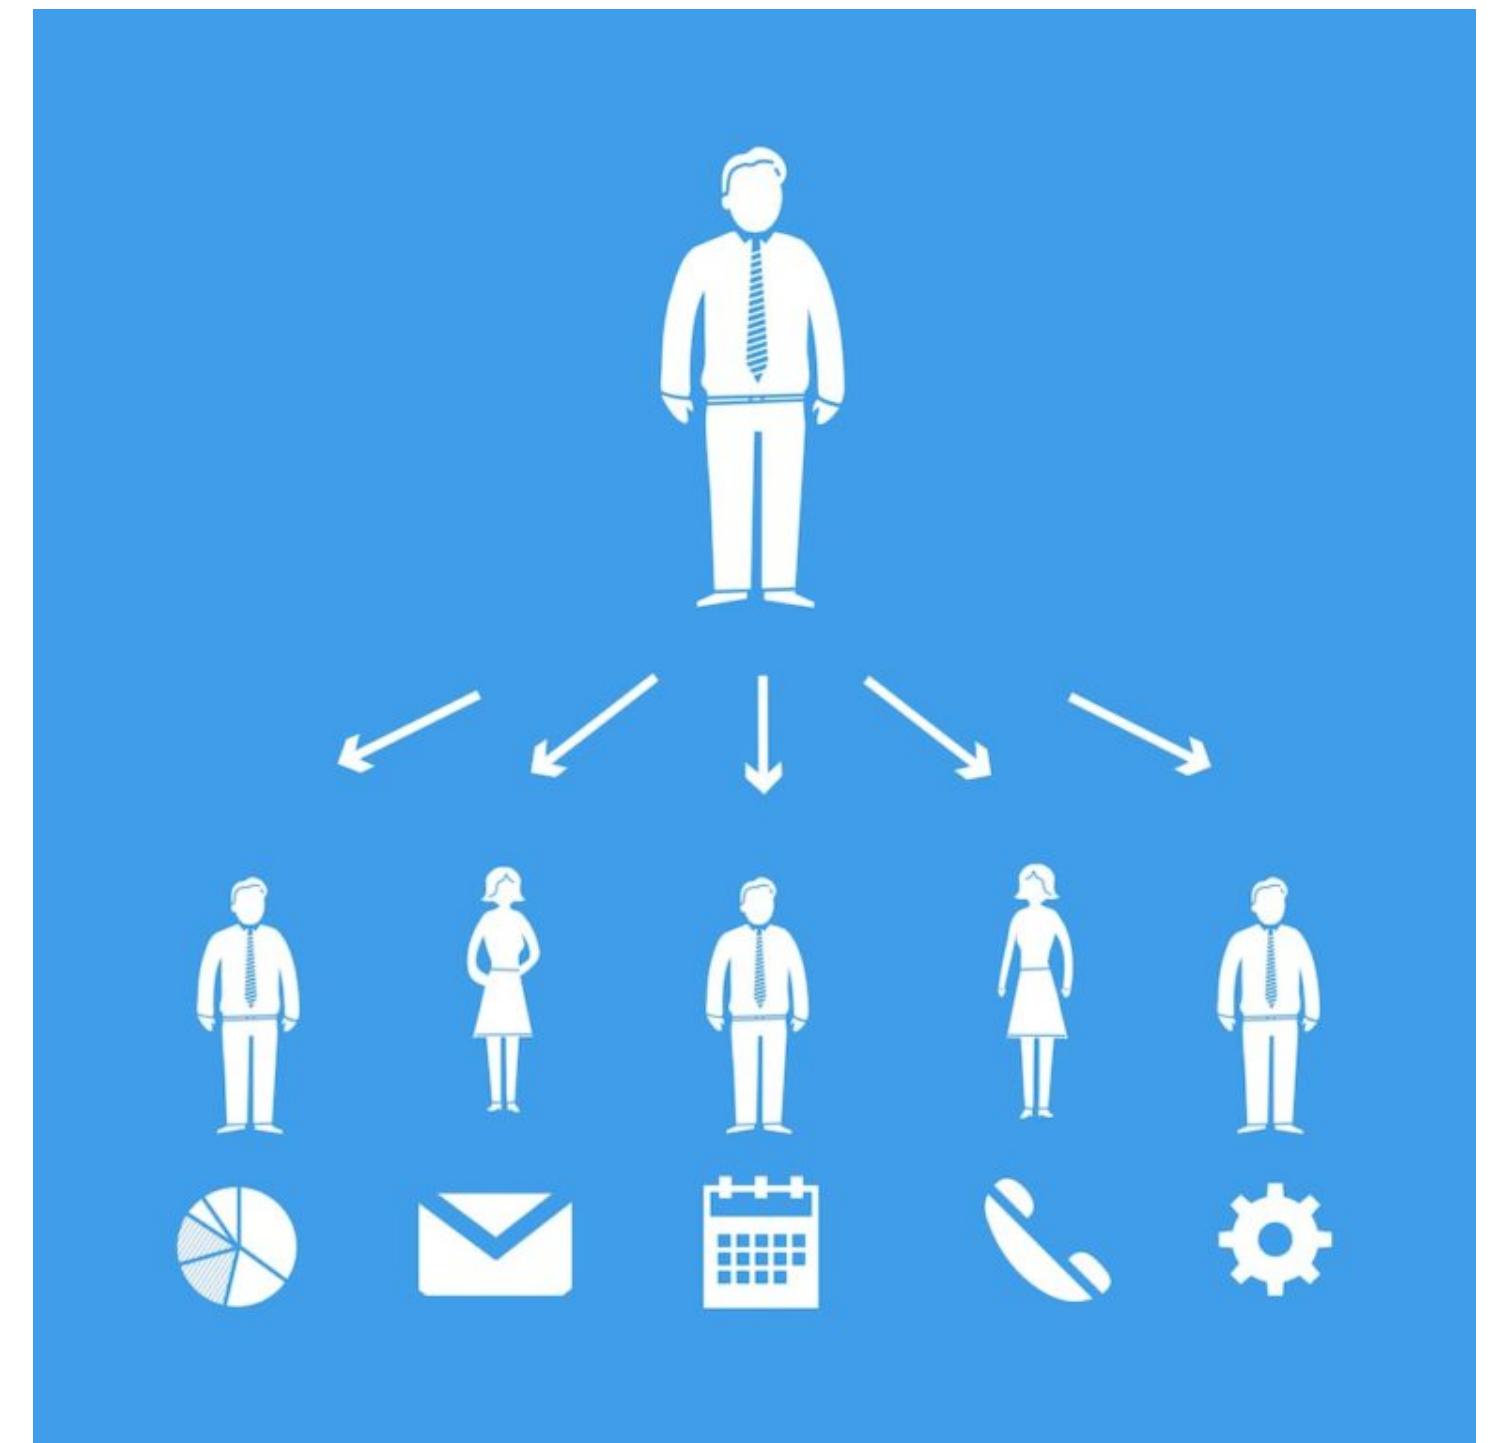
FAULT TOLERANCE: MECHANISMS

- ▶ State Machine Replication
 - ▶ Multiple nodes perform the same action
 - ▶ Easy for state-less applications
 - ▶ Requires Synchronization for stateful applications
- ▶ State Replication
 - ▶ Requires Synchronization
 - ▶ Or weakened consistency (this will be discussed Thu. or Fri.)
- ▶ Component Isolation
 - ▶ Message Passing
 - ▶ Persistent state (Snapshots)



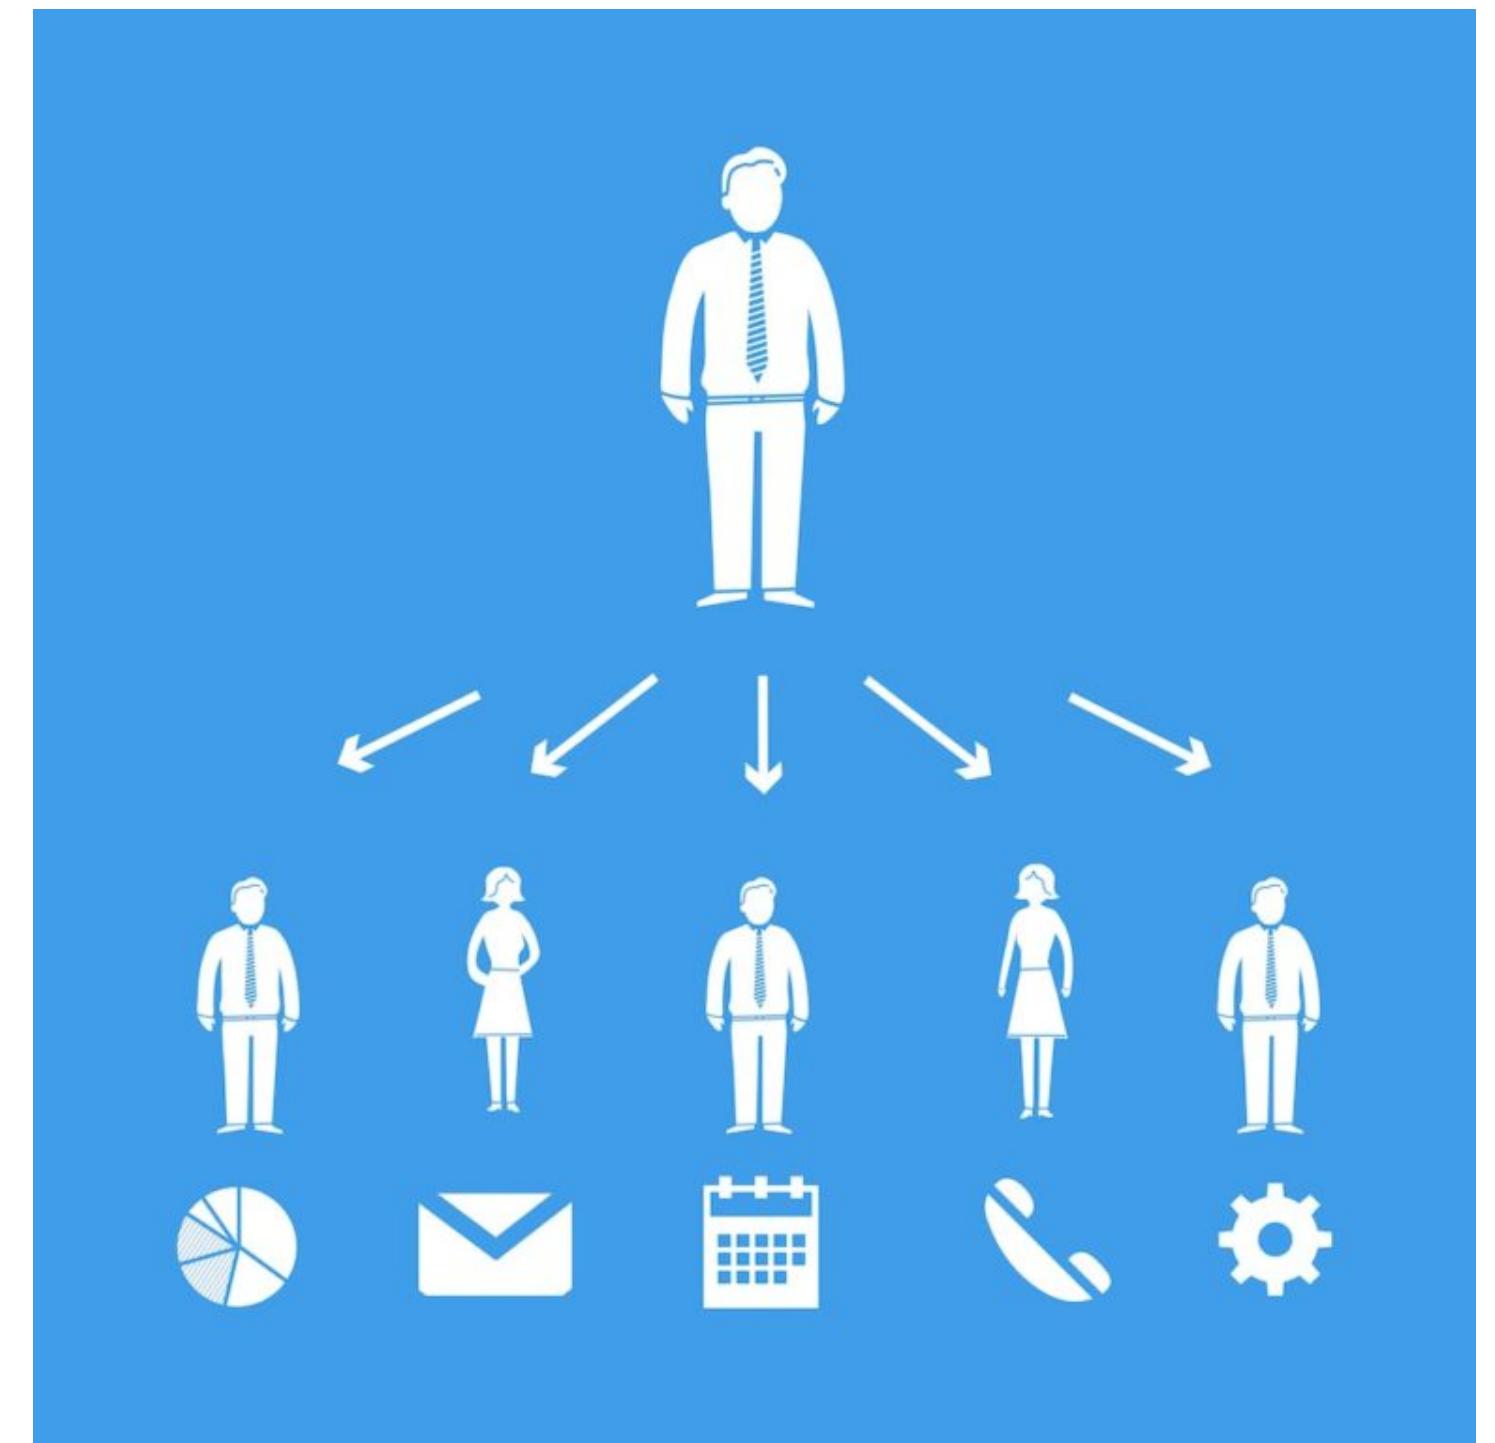
FAULT TOLERANCE: MECHANISMS

- ▶ Delegation
- ▶ Hand-over responsibility



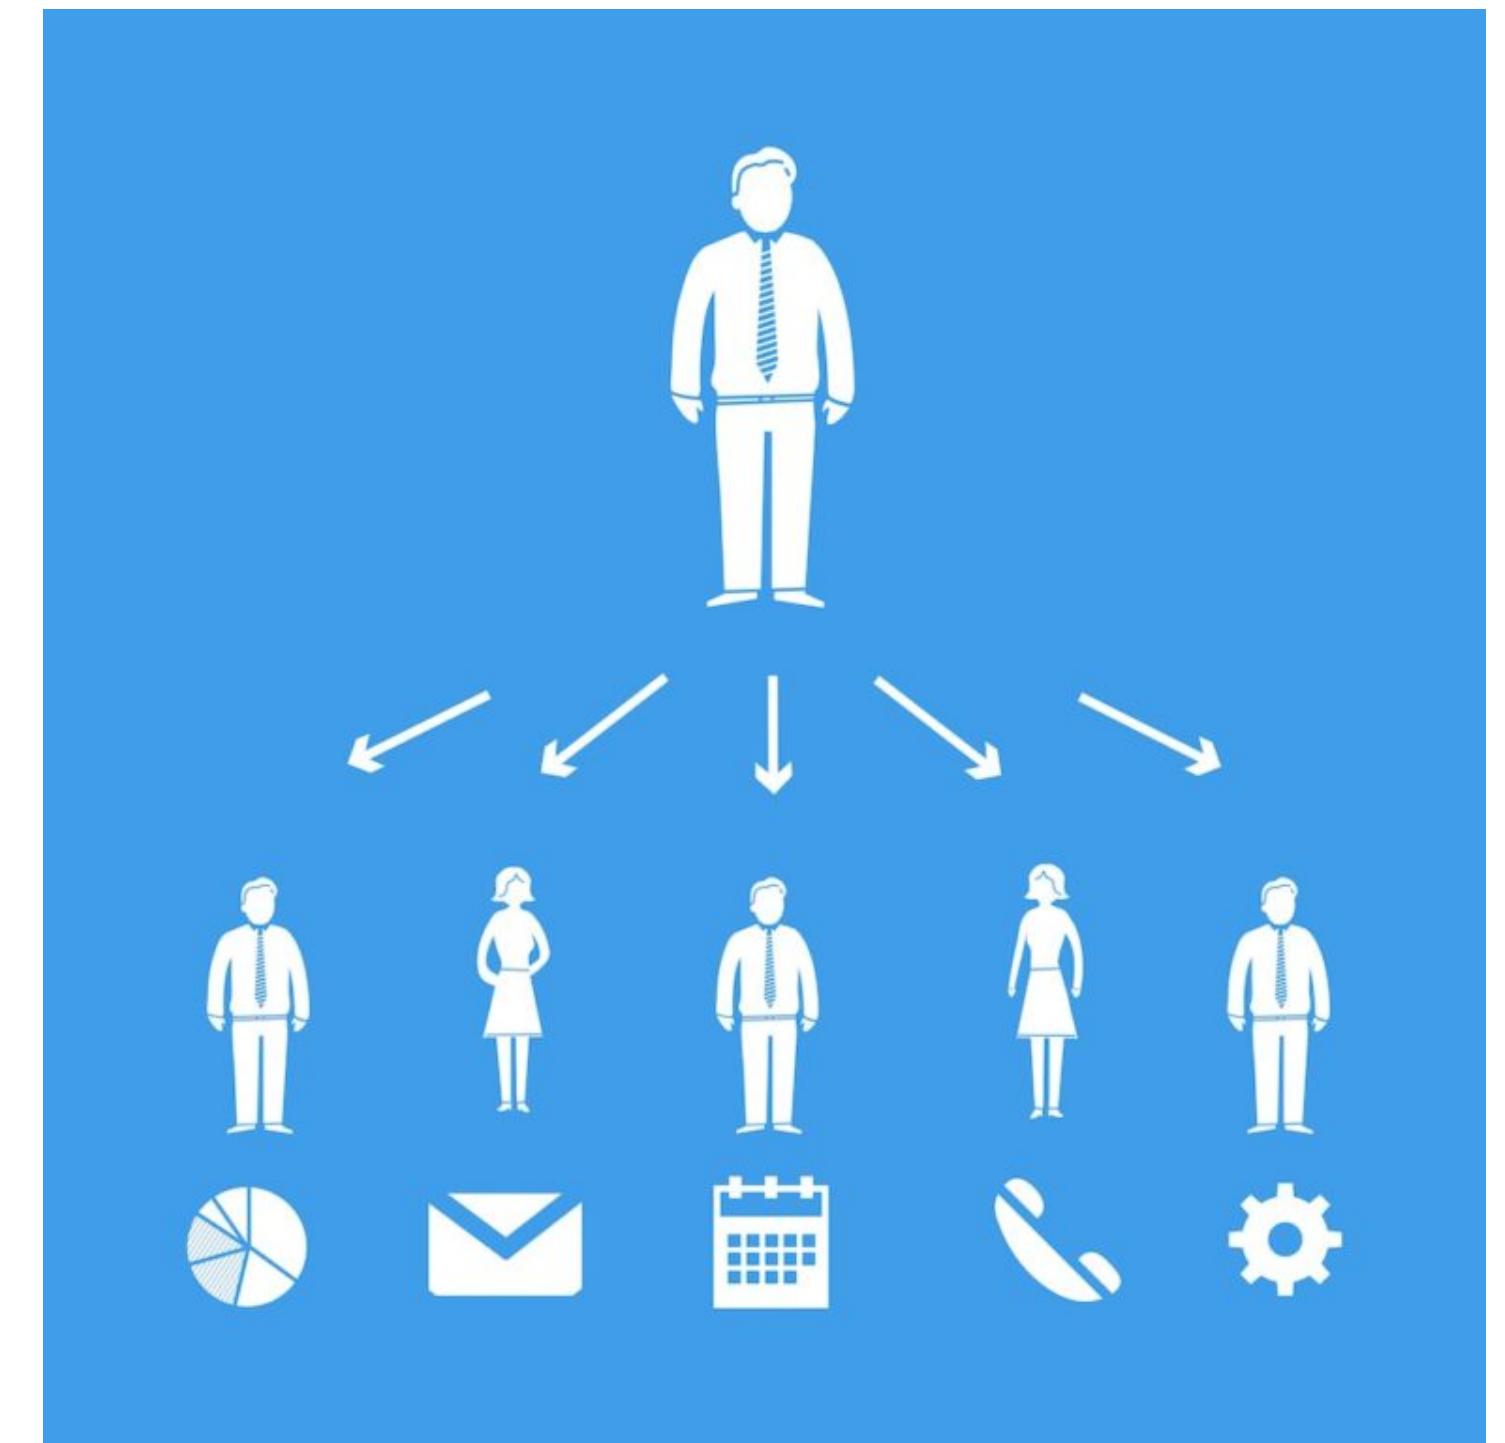
FAULT TOLERANCE: MECHANISMS

- ▶ Delegation
 - ▶ Hand-over responsibility
- ▶ Components
 - ▶ Self-contained
 - ▶ Isolated
 - ▶ Encapsulated



FAULT TOLERANCE: MECHANISMS

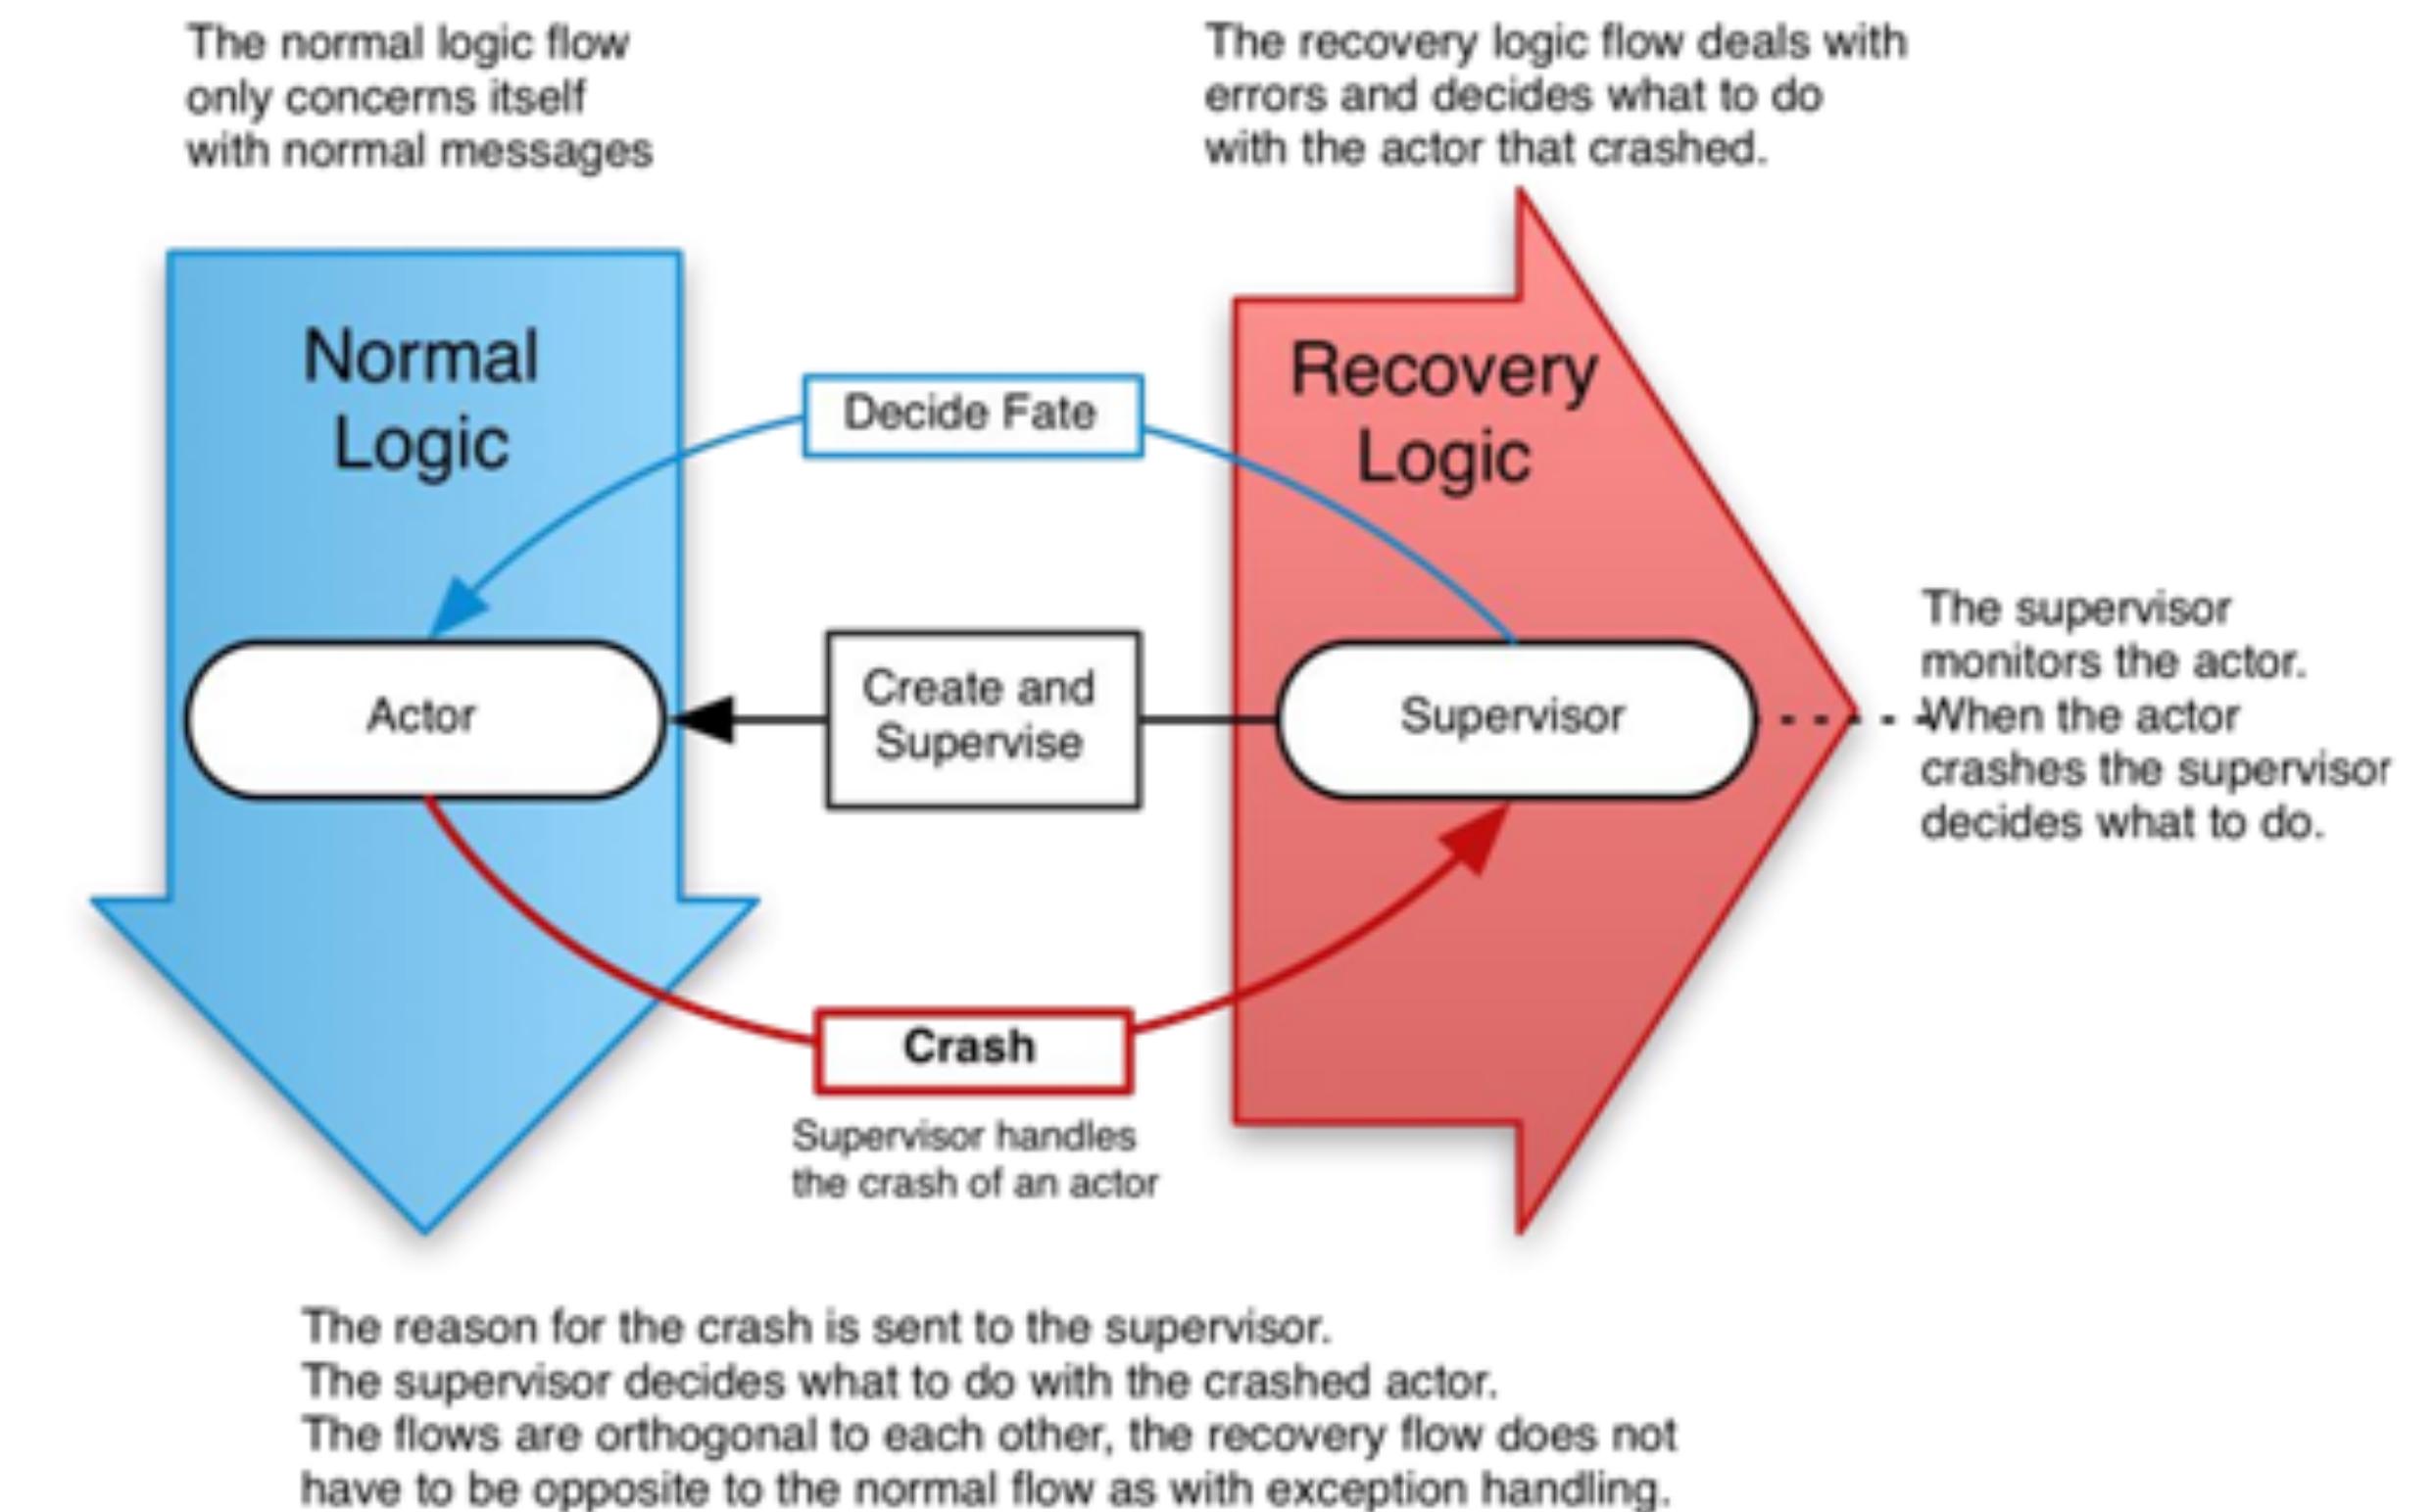
- ▶ Delegation
 - ▶ Hand-over responsibility
- ▶ Components
 - ▶ Self-contained
 - ▶ Isolated
 - ▶ Encapsulated
- ▶ Micro-services / Actors



FAULT TOLERANCE: SUPERVISION

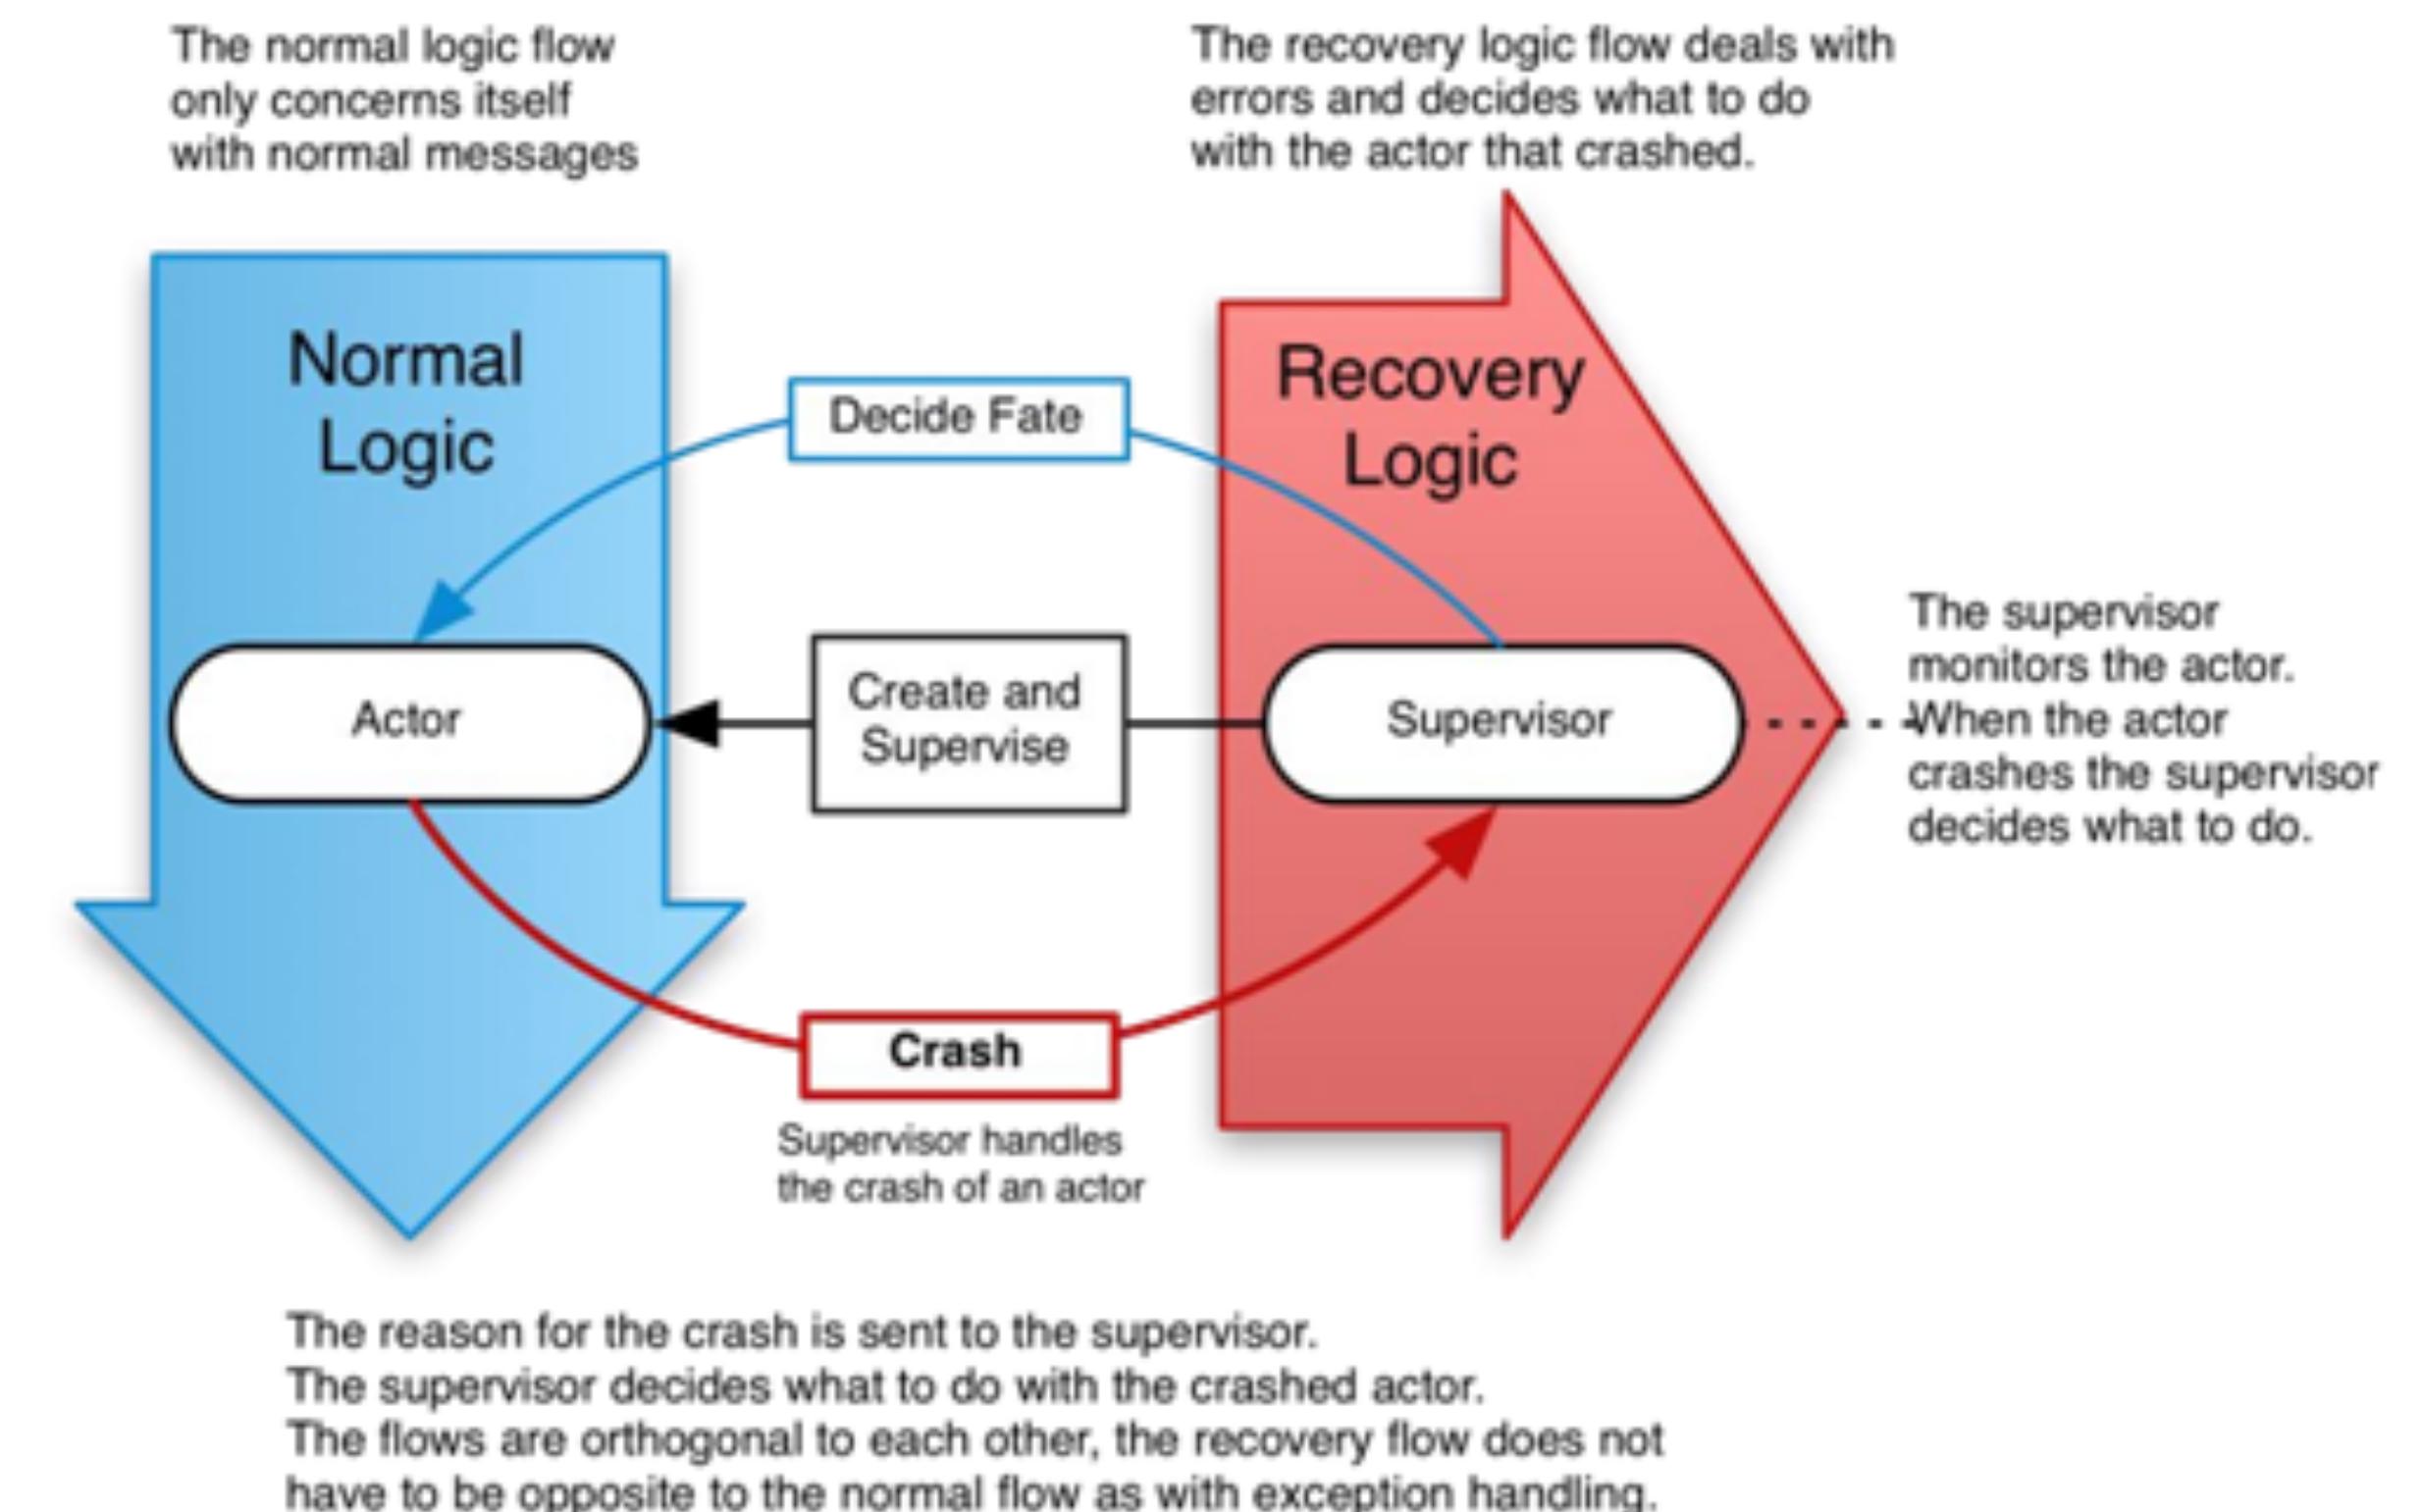
FAULT TOLERANCE: SUPERVISION

- ▶ “Let it crash”



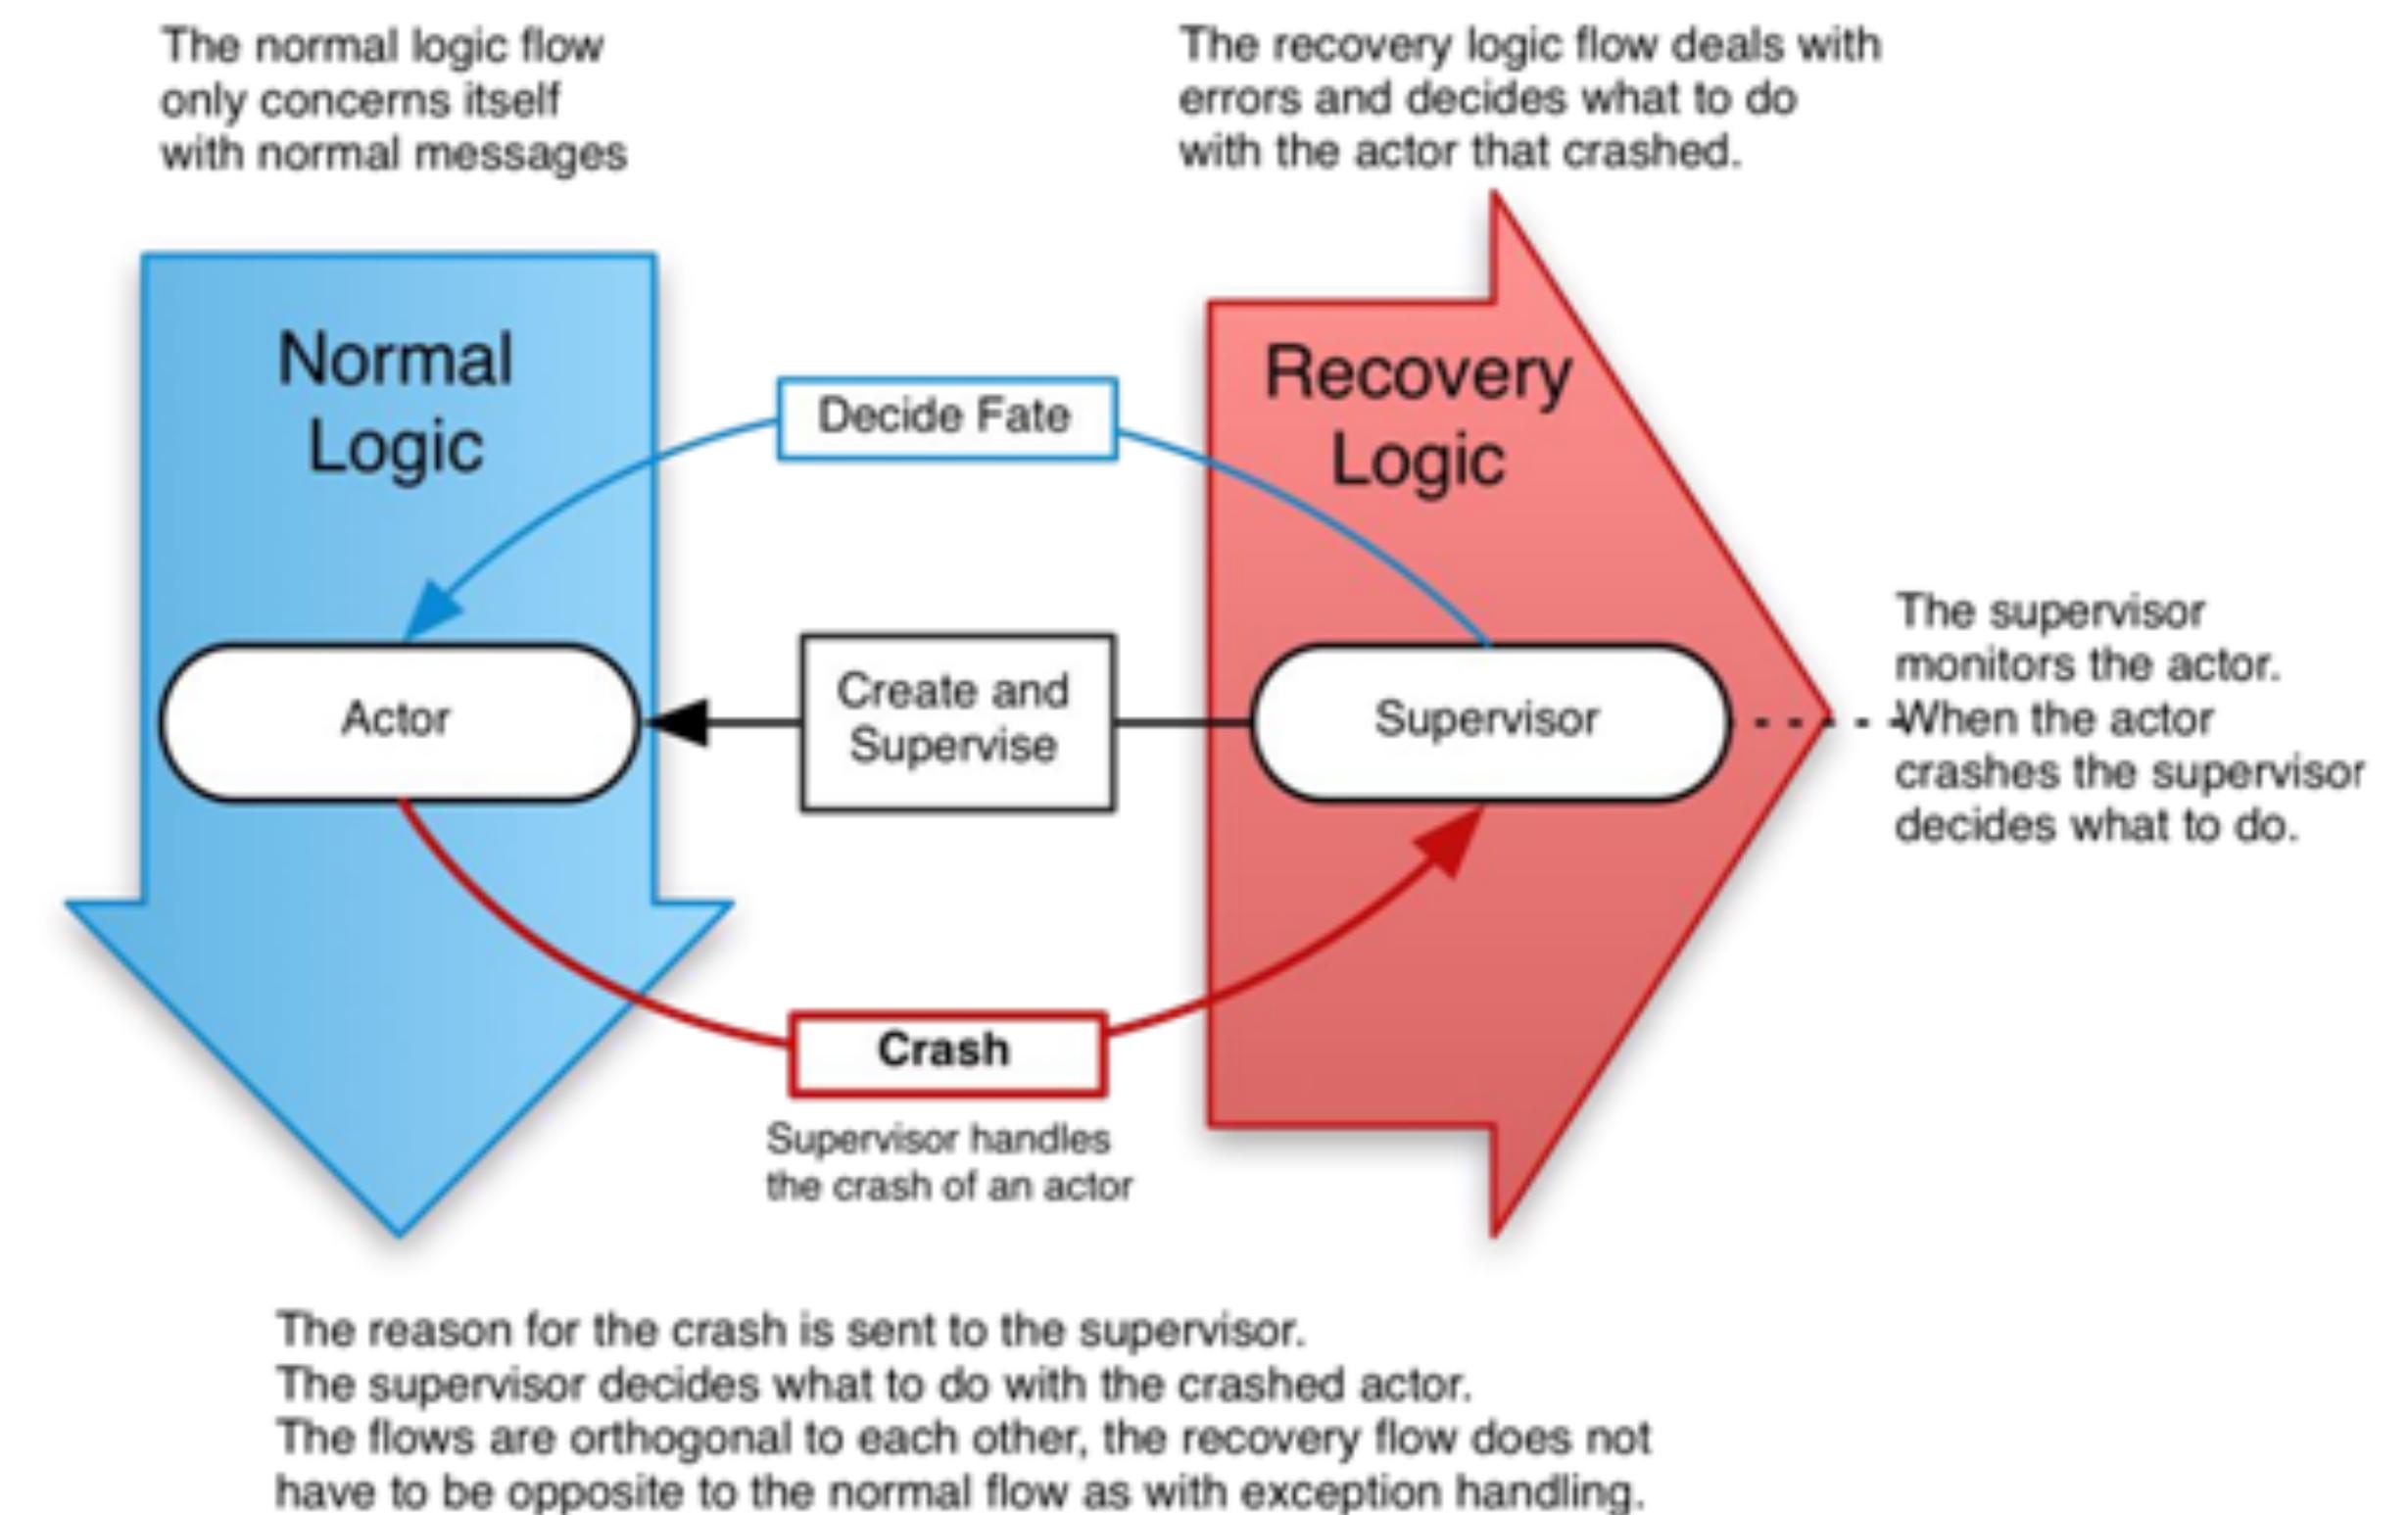
FAULT TOLERANCE: SUPERVISION

- ▶ “Let it crash”
- ▶ Parent / Child Supervision:
 - ▶ send messages
 - ▶ collect answers



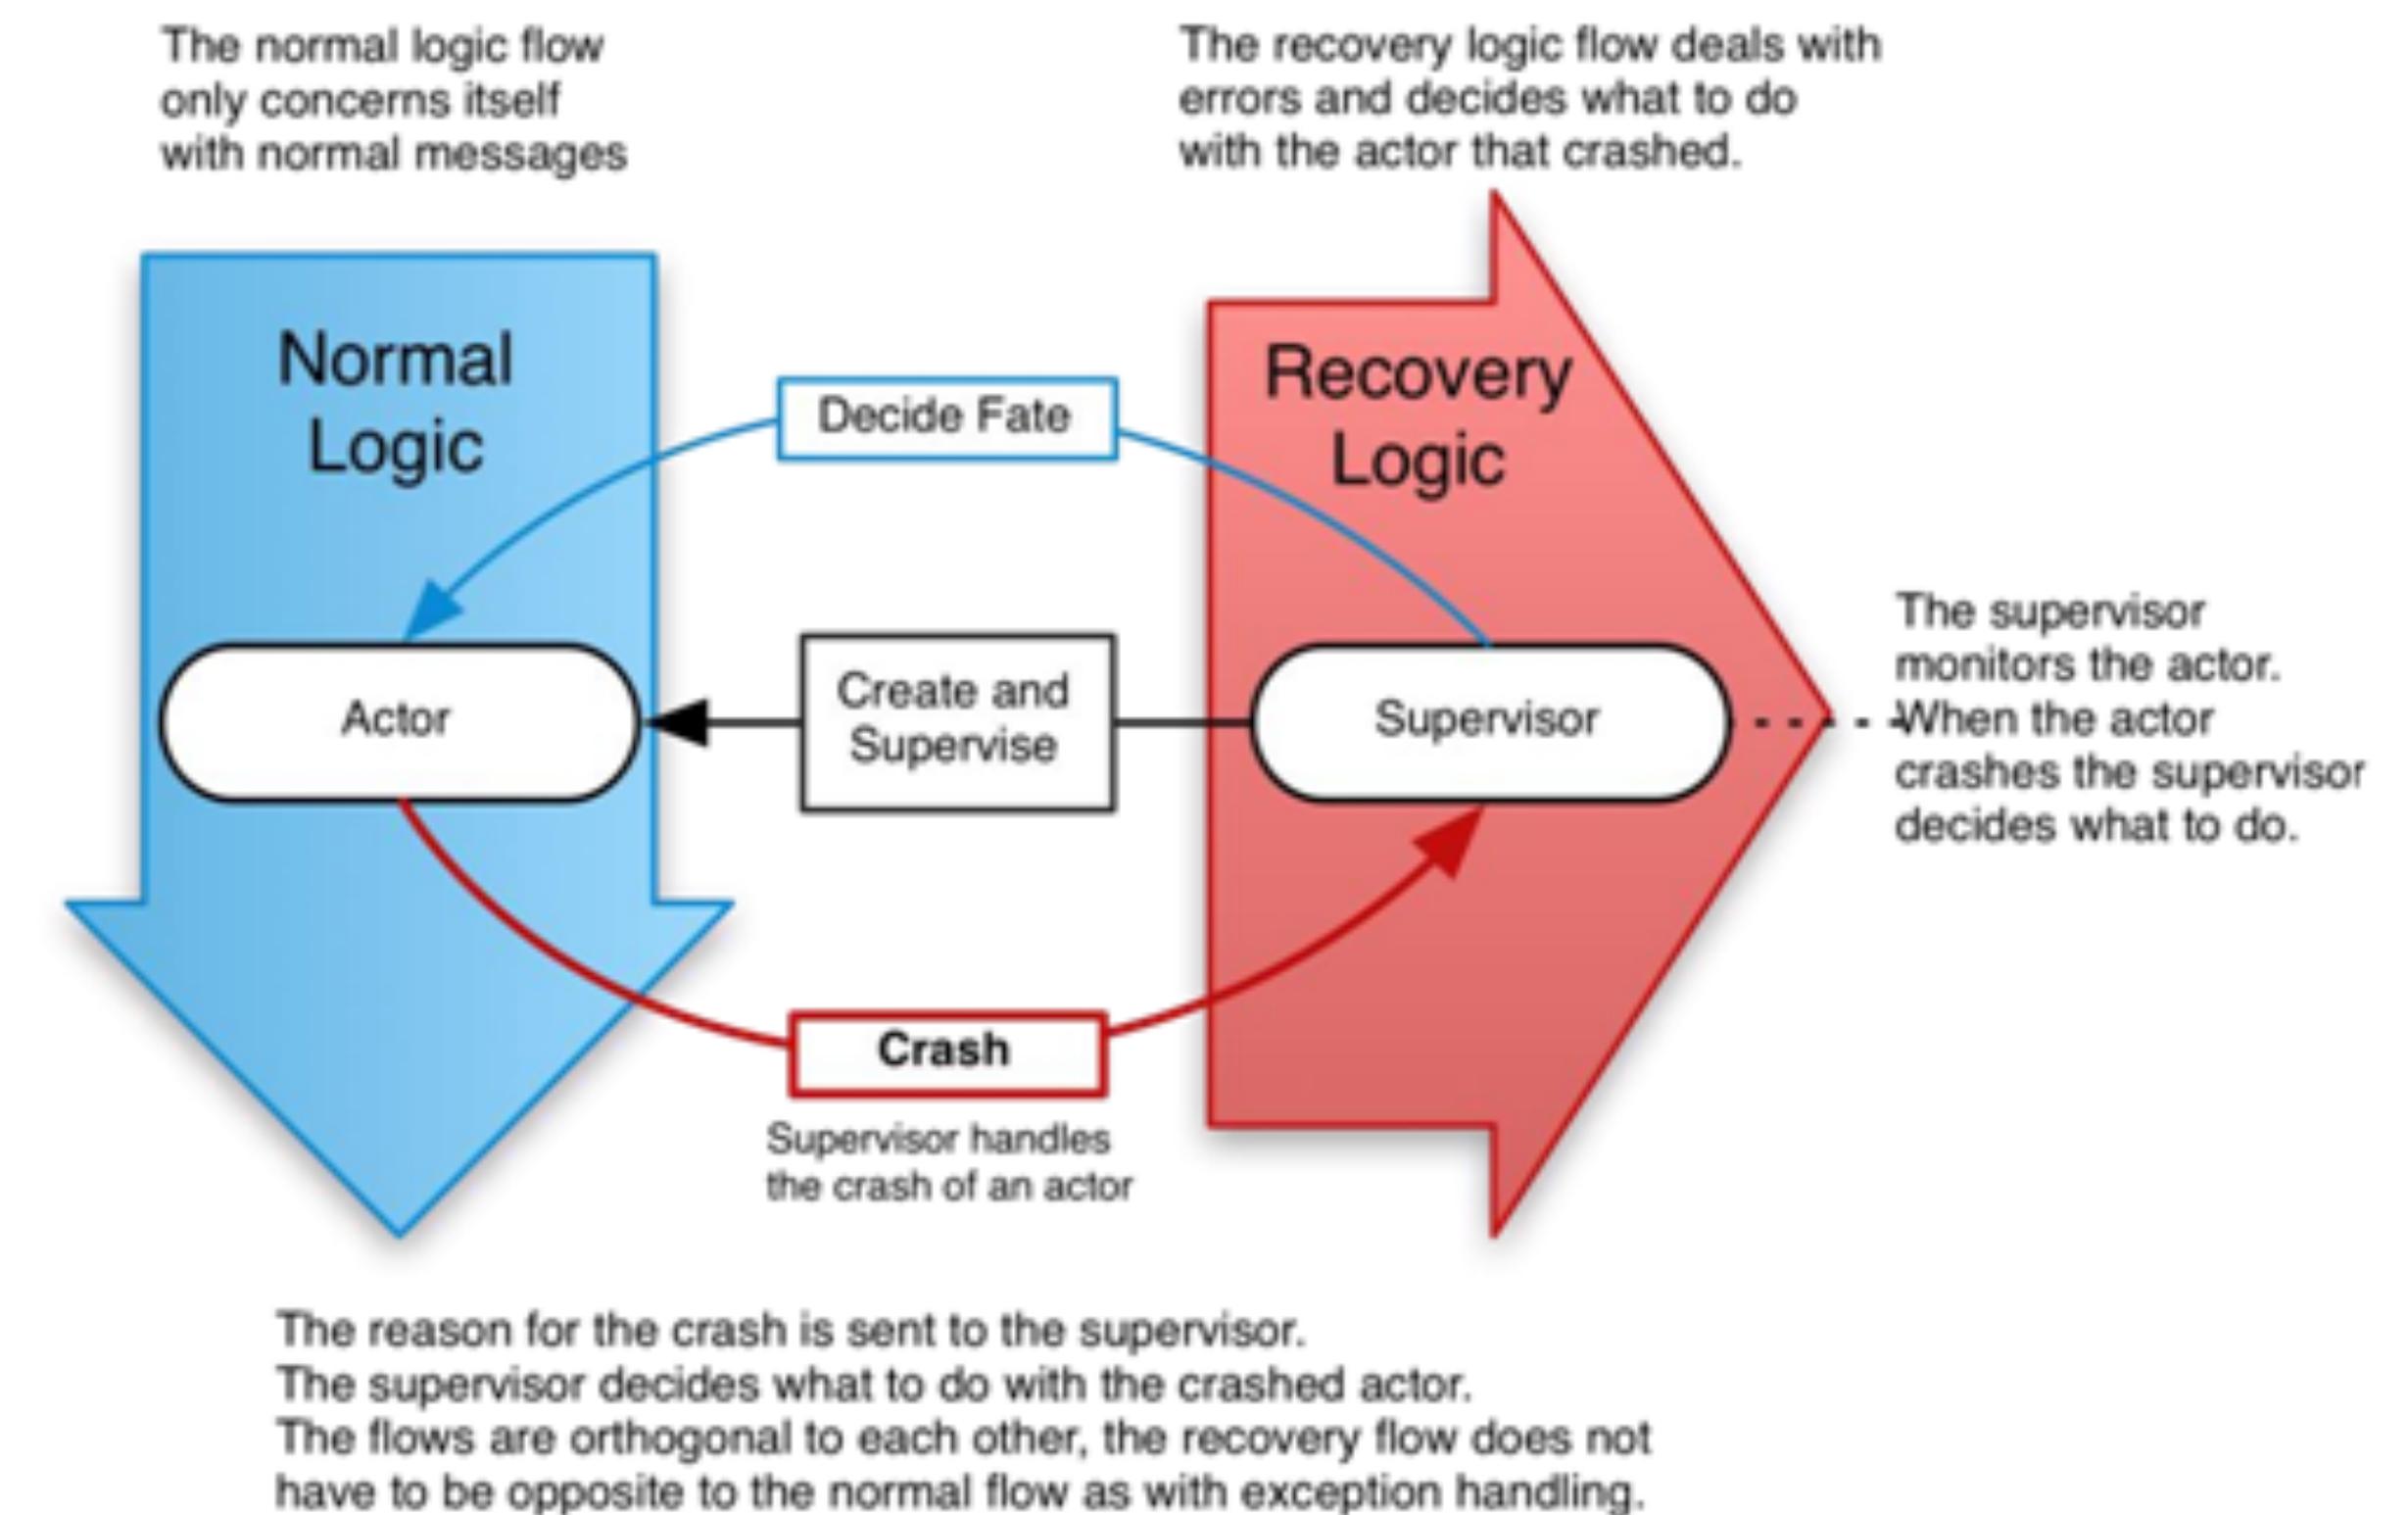
FAULT TOLERANCE: SUPERVISION

- ▶ “Let it crash”
- ▶ Parent / Child Supervision:
 - ▶ send messages
 - ▶ collect answers
- ▶ One-for-one Strategy



FAULT TOLERANCE: SUPERVISION

- ▶ “Let it crash”
- ▶ Parent / Child Supervision:
 - ▶ send messages
 - ▶ collect answers
- ▶ One-for-one Strategy
- ▶ All-for-one Strategy



FAULT TOLERANCE: SUPERVISION

- ▶ “Let it crash”
- ▶ Parent / Child Supervision:
 - ▶ send messages
 - ▶ collect answers
- ▶ One-for-one Strategy
- ▶ All-for-one Strategy
- ▶ Death-Watch

