

# Software Models

# How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs

LESLIE LAMPORT

**Abstract**—Many large sequential computers execute operations in a different order than is specified by the program. A correct execution is achieved if the results produced are the same as would be produced by executing the program steps in order. For a multiprocessor, such a correct execution by each processor does not guarantee the correct execution of the entire program. Additional conditions are given which do guarantee that a computer correctly executes multiprocess programs.

**Index Terms**—Computer design, concurrent computing, correctness, multiprocessing, parallel processing.

A high-speed processor may execute operations in a different order than is specified by the program. The correct execution is guaranteed if the processor satisfies the following condition: the result of an execution is the same as if all the operations had been executed in the order specified by the program. Consider a computer composed of several such processors sharing a common memory. The customary approach to proving the correctness of multiprocess algorithms assumes that such a computer satisfies the following condition: the result of any execution is the same as if the operations of the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the same order as they appear in the program. A multiprocessor satisfying this condition is called sequentially consistent.

# Shared Memory Consistency Models: A Tutorial \*

Sarita V. Adve<sup>†</sup> and Kourosh Gharachorloo<sup>‡</sup>

<sup>†</sup>Department of Electrical and Computer Engineering  
Rice University  
Houston, Texas 77251-1892

<sup>‡</sup>Western Research Laboratory  
Digital Equipment Corporation  
Palo Alto, California 94301-1616

Rice University ECE Technical Report 9512  
Western Research Laboratory Research Report 95/7

September 1995

## Abstract

Parallel systems that support the shared memory abstraction are becoming widely accepted in many areas of computing. Writing correct and efficient programs for such systems requires a formal specification of memory semantics, called a *memory consistency model*. The most intuitive model—sequential consistency—greatly restricts the use of many performance optimizations commonly used by uniprocessor hardware and compiler designers, thereby reducing the benefit of using a multiprocessor. To alleviate this problem, many current multiprocessors support more relaxed consistency models. Unfortunately, the models supported by various systems differ from each other in subtle yet important ways. Furthermore, precisely defining the semantics of each model often leads to complex specifications that are difficult to understand for typical users and builders of computer systems.

The purpose of this tutorial paper is to describe issues related to memory consistency models in a way that would be understandable to most computer professionals. We focus on consistency models proposed for hardware-based shared-memory systems. Many of these models are originally specified with an emphasis on the

# How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs

LESLIE LAMPORT

**Abstract**—Many large sequential computers execute operations in a different order than is specified by the program. A correct execution is achieved if the results produced are the same as would be produced by executing the program steps in order. For a multiprocessor, such guarantee the conditions are executed mult

Index Terms:  
correctness, i

A high-sp  
order than  
execution i  
condition:  
tions had b  
processor  
sider a co  
common  
proving t  
such a co  
the result  
the proc

## ABSTRACT

This paper describes the new Java memory model, which has been revised as part of Java 5.0. The model specifies the legal behaviors for a multithreaded program; it defines the semantics of multithreaded Java programs and partially determines legal implementations of Java virtual machines and compilers.

The new Java model provides a simple interface for correctly synchronized programs – it guarantees sequential consistency to data-race-free programs. Its novel contribution is requiring that the behavior of incorrectly synchronized programs be bounded by a well defined notion of causality. The causality requirement is strong enough to respect the safety and security properties of Java and weak enough to allow standard compiler and hardware optimizations. To our knowledge, other models are either too weak because they do not provide for sufficient safety/security, or are too strong because they rely on a strong notion of data and control dependences that precludes some standard compiler transformations.

For a high-level programming language such as Java, the memory model determines the transformations the compiler

Jeremy Manson and William Pugh  
Department of Computer Science  
University of Maryland, College Park  
College Park, MD  
[{jmanson, pugh}@cs.umd.edu](mailto:{jmanson, pugh}@cs.umd.edu)

## The Java Memory Model

Sarita V. Adve

Department of Computer Science  
University of Illinois at Urbana-Champaign  
Urbana-Champaign, IL  
[sadve@cs.uiuc.edu](mailto:sadve@cs.uiuc.edu)

Meanings of Programs]: Operational Semantics

General Terms: Design, Languages

Keywords: Concurrency, Java, Memory Model, Multithreading

## 1. INTRODUCTION

The memory model for a multithreaded system specifies how memory actions (e.g., reads and writes) in a program will appear to execute to the programmer, and specifically, which value each read of a memory location may return. Every hardware and software interface of a system that admits multithreaded access to shared memory requires a memory model. The model determines the transformations that the system (compiler, virtual machine, or hardware) can apply to a program written at that interface. For example, given a program in machine language, the memory model for the machine language / hardware interface will determine the optimizations the hardware can perform.

For a high-level programming language such as Java, the memory model determines the transformations the compiler

# Shared Memory Consistency Models: A Tutorial \*

Sarita V. Adve<sup>†</sup> and Kourosh Gharachorloo<sup>‡</sup>

<sup>†</sup>Department of Electrical and Computer Engineering,  
Rice University, Houston, TX 77005-1892

892

ory

616

t 9512

port 95/7

accepted in many areas of  
specification of memory  
consistency—greatly  
hardware and compiler  
problem, many current  
supported by various  
users and builders of

models in a way  
models proposed for  
emphasis on the

# The case of Java

# The case of Java

(To appear in *Concurrency: Practice and Experience*)

## The Java Memory Model is Fatally Flawed

William Pugh

Dept. of Computer Science

Univ. of Maryland, College Park

[pugh@cs.umd.edu](mailto:pugh@cs.umd.edu)

### Abstract

The Java memory model described in Chapter 17 of the Java Language Specification gives constraints on how threads interact through memory. This chapter is hard to interpret and poorly understood; it imposes constraints that prohibit common compiler optimizations and are expensive to implement on existing hardware. Most JVMs violate the constraints of the existing Java memory model; conforming to the existing specification would impose significant performance penalties.

In addition, programming idioms used by some programmers and used within Sun's Java Develop-

existing style of the specification will never be clear, and that attempts to patch the existing specification by adding new rules will make even harder to understand. If we decide to change the Java memory model, a completely new description of the memory model should be devised.

A number of terms are used in the Java memory model but not explicitly related to Java source programs nor the Java virtual machine. Some of these terms have been interpreted differently by various people. I have based my understanding of these terms on conversations with Guy Steele, Doug Lea and others.

A *variable* refers to a static variable of a loaded

# The case of Java

Jeremy Manson and William Pugh  
Department of Computer Science  
University of Maryland, College Park  
College Park, MD  
[{jmanson, pugh}@cs.umd.edu](mailto:{jmanson,pugh}@cs.umd.edu)

## ABSTRACT

This paper describes the new Java memory model, which has been revised as part of Java 5.0. The model specifies the legal behaviors for a multithreaded program; it defines the semantics of multithreaded Java programs and partially determines legal implementations of Java virtual machines and compilers.

The new Java model provides a simple interface for correctly synchronized programs – it guarantees sequential consistency to data-race-free programs. Its novel contribution is requiring that the behavior of incorrectly synchronized programs be bounded by a well defined notion of causality. The causality requirement is strong enough to respect the safety and security properties of Java and weak enough to allow standard compiler and hardware optimizations. To our knowledge, other models are either too weak because they do not provide for sufficient safety/security, or are too strong because they do not provide for sufficient optimization.

Sarita V. Adve  
Department of Computer Science  
University of Illinois at Urbana-Champaign  
Urbana-Champaign, IL  
[sadve@cs.uiuc.edu](mailto:sadve@cs.uiuc.edu)

Meanings of Programs]: Operational Semantics

General Terms: Design, Languages

Keywords: Concurrency, Java, Memory Model, Multithreading

## 1. INTRODUCTION

The memory model for a multithreaded system specifies how memory actions (e.g., reads and writes) in a program will appear to execute to the programmer, and specifically, which value each read of a memory location may return. Every hardware and software interface of a system that admits multithreaded access to shared memory requires a memory model. The model determines the transformations that the system (compiler, virtual machine, or hardware) can apply to a program written at that interface. For example, given a program in machine language, the memory model for the machine language / hardware interface will determine the

(To appear in *Concurrency: Practice and Experience*)

## The Java Memory Model is Fatally Flawed

Pugh  
Computer Science  
University of Maryland,  
College Park  
[www.cs.umd.edu](http://www.cs.umd.edu)

existing style of the specification will never be clear, and that attempts to patch the existing specification by adding new rules will make even harder to understand. If we decide to change the Java memory model, a completely new description of the memory model should be devised.

A number of terms are used in the Java memory model but not explicitly related to Java source programs nor the Java virtual machine. Some of these terms have been interpreted differently by various people. I have based my understanding of these terms on conversations with Guy Steele, Doug Lea and others.

*variable* refers to a static variable of a loaded

# The case of Java

(To appear in *Concurrency: Practice and Experience*)

## The Java Memory Model is Fatally Flawed

Pugh  
Computer Science  
University of Maryland,  
College Park  
[cs.umd.edu](http://cs.umd.edu)

## The Java Memory Model

Jeremy Manson and William Pugh  
Department of Computer Science  
University of Maryland, College Park  
College Park, MD  
[{jmanson, pugh}@cs.umd.edu](mailto:{jmanson,pugh}@cs.umd.edu)

### ABSTRACT

This paper describes the new Java memory model, which has been revised as part of Java 5.0. The model specifies the legal behaviors for a multithreaded program; it defines the semantics of multithreaded Java programs and partially determines legal implementations of Java virtual machines and compilers.

The new Java model provides a simple interface for correctly synchronized programs – it guarantees sequential consistency to data-race-free programs. Its novel contribution is requiring that the behavior of incorrectly synchronized programs be bounded by a well defined notion of causality. The causality requirement is strong enough to respect the safety and security properties of Java and weak enough to allow standard compiler and hardware optimizations. To our knowledge, other models are either too weak because they do not provide for sufficient safety/security, or are too strong because they do not provide for sufficient optimization.

Dep.  
Universit

Meaning  
General  
Keywords  
ing

### 1. Introduction

The Java memory model [1] specifies how memory will appear to threads which have different modes of operation. It is based on a partial ordering of events in time called causality. This ordering is used to determine whether one thread's actions can affect another thread's visible state.

## Java Memory Model Examples: Good, Bad and Ugly

David Aspinall and Jaroslav Ševčík

August 8, 2007

### Abstract

We review a number of illustrative example programs for the Java Memory Model (JMM) [6, 3], relating them to the original design goals and giving intuitive explanations (which can be made precise). We consider good, bad and ugly examples. The good examples are allowed behaviours in the JMM, showing possibilities for non sequentially consistent executions and reordering optimisations. The bad examples are prohibited behaviours, which are clearly ruled out by the JMM. The ugly examples are most interesting: these are tricky cases which illustrate some problem areas for the current formulation of the memory model, where the anticipated design goals are not met. For some of these we mention possible fixes, drawing on knowledge we gained while formalising the memory model in the theorem prover Isabelle [1].

# The case of Java

A collage of academic papers and abstracts related to Java's memory model, illustrating its complexity and flaws.

**Top Paper:** *The Java Memory Model is Fatally Flawed* (To appear in Concurrency: Practice and Experience)

**Authors:** Jeremy Manson and William Pugh

**Institution:** Department of Computer Science, University of Maryland, College Park, MD

**Email:** {jmanson, pugh}@cs.umd.edu

**Abstract:** This paper shows that the Java Memory Model (JMM) is fatally flawed. It presents several examples that demonstrate how the JMM fails to meet its own design goals, such as allowing as many optimizations as possible while maintaining consistency. The examples range from good cases to bad and ugly ones, illustrating the limitations of the JMM.

**Second Paper:** *On Validity of Program Transformations in the Java Memory Model*

**Authors:** Jaroslav Ševčík and David Aspinall

**Institution:** LFCS, School of Informatics, University of Edinburgh

**Abstract:** We analyse the validity of several common program transformations in multi-threaded Java, as defined by the Java Memory Model (JMM) section of Chapter 17 of the Java Language Specification. The main design goal of the JMM was to allow as many optimisations as possible. However, we find that commonly used optimisations, such as common subexpression elimination, can introduce new behaviours and violations of the JMM. In this paper, we describe several kinds of transformations and their interactions with the JMM. We provide counterexamples showing that some transformations are invalid according to the JMM.

**Third Paper:** *Java Memory Model Examples: Good, Bad and Ugly*

**Authors:** David Aspinall and Jaroslav Ševčík

**Date:** August 8, 2007

**Abstract:** This paper provides illustrative example programs for the Java Memory Model (JMM), relating them to the original design goals of the JMM (which can be made precise). We present three types of examples. The good examples are allowed by the JMM and represent possibilities for non sequentially consistent memory operations. The bad examples are prohibited by the JMM. The ugly examples represent tricky cases which illustrate some problems with the formalisation of the memory model, where the JMM has failed to meet its design goals. For some of these we mention possible improvements to the JMM, based on what we gained while formalising the memory model.

**It's not just the processor ...**

# It's not just the processor ...

## Atomic Operations

```
int x = 1;  
x++;      double x = 1.0;
```

```
monitorenter  
monitorexit
```

# It's not just the processor ...

## Atomic Operations

```
int x = 1;  
x++;      double x = 1.0
```

```
monitorenter  
monitorexit
```

## Synchronizing Operations

```
volatile int x = 0;  
thread.start();  
thread.join();  
monitorenter  
monitorexit
```

# It's not just the processor ...

## Atomic Operations

```
int x = 1;  
x++;      double x = 1.0
```

```
monitorenter  
monitorexit
```

## Compiler Optimizations

### Reordering

### Eliminations (CSE)

### Introductions

## Synchronizing Operations

```
volatile int x = 0;  
thread.start();  
thread.join();
```

```
monitorenter  
monitorexit
```

# It's not just the processor ...

## Atomic Operations

```
int x = 1;  
x++;      double x = 1.0
```

```
monitorenter  
monitorexit
```

## Compiler Optimizations

### Reordering

### Eliminations (CSE)

### Introductions

## Synchronizing Operations

```
volatile int x = 0;  
thread.start();  
thread.join();
```

```
monitorenter  
monitorexit
```

## VM Optimizations

Just in time compilation

Lock elision

Biased Locking

# Compiler Optimizations

## On Validity of Program Transformations in the Java Memory Model

Jaroslav Ševčík and David Aspinall

LFCS, School of Informatics, University of Edinburgh

**Abstract.** We analyse the validity of several common program transformations in multi-threaded Java, as defined by the Java Memory Model (JMM) section of Chapter 17 of the Java Language Specification. The main design goal of the JMM was to allow as many optimisations as possible. However, we find that commonly used optimisations, such as common subexpression elimination, can introduce new behaviours and so are invalid for Java. In this paper, we describe several kinds of transformations and explain the problems with a number of counterexamples. More positively, we also examine some valid transformations, and prove their validity. Our study contributes to the understanding of the JMM, and has the practical impact of revealing some cases where the Sun Hotspot JVM does not comply with the Java Memory Model.

### 1 Introduction

Although programmers generally assume an interleaved semantics, the Java Language Specification [11] defines more relaxed semantics, which is called the Java Memory Model [12, 19]. The reasons for having a weaker semantics become apparent from the following example:

# Compiler Optimizations

# Compiler Optimizations

```
x = y = z = 0
-----
r1 = z;
if (r1 == 1) {
    x = 1;
    y = 1;
} else {
    y = 1;
    x = 1;
}
||| r2 = x;
    r3 = y;
    if (r2==1 && r3==1) {
        z = 1;
    }
-----
r1 = r2 = r3 = 1
```

# Compiler Optimizations

```
x = y = z = 0
-----
r1 = z;
if (r1 == 1) {
    x = 1;
    y = 1;
} else {
    x = 1;
    y = 1;
}
||| r2 = x;
    r3 = y;
    if (r2==1 && r3==1) {
        z = 1;
    }
-----
r1 = r2 = r3 = 1
```

# Compiler Optimizations

x = y = z = 0

---

r1 = z;  
x = 1;  
y = 1;

r2 = x;  
r3 = y;  
**if** (r2==1 && r3==1) {  
 z = 1;  
}

---

r1 = r2 = r3 = 1

# Compiler Optimizations

x = y = z = 0

---

x = 1;  
y = 1;  
r1 = z;

r2 = x;  
r3 = y;  
**if** (r2==1 && r3==1) {  
 z = 1;  
}

---

r1 = r2 = r3 = 1

# Compiler Optimizations

x = y = z = 0

---

x = 1;  
y = 1;  
r1 = z;

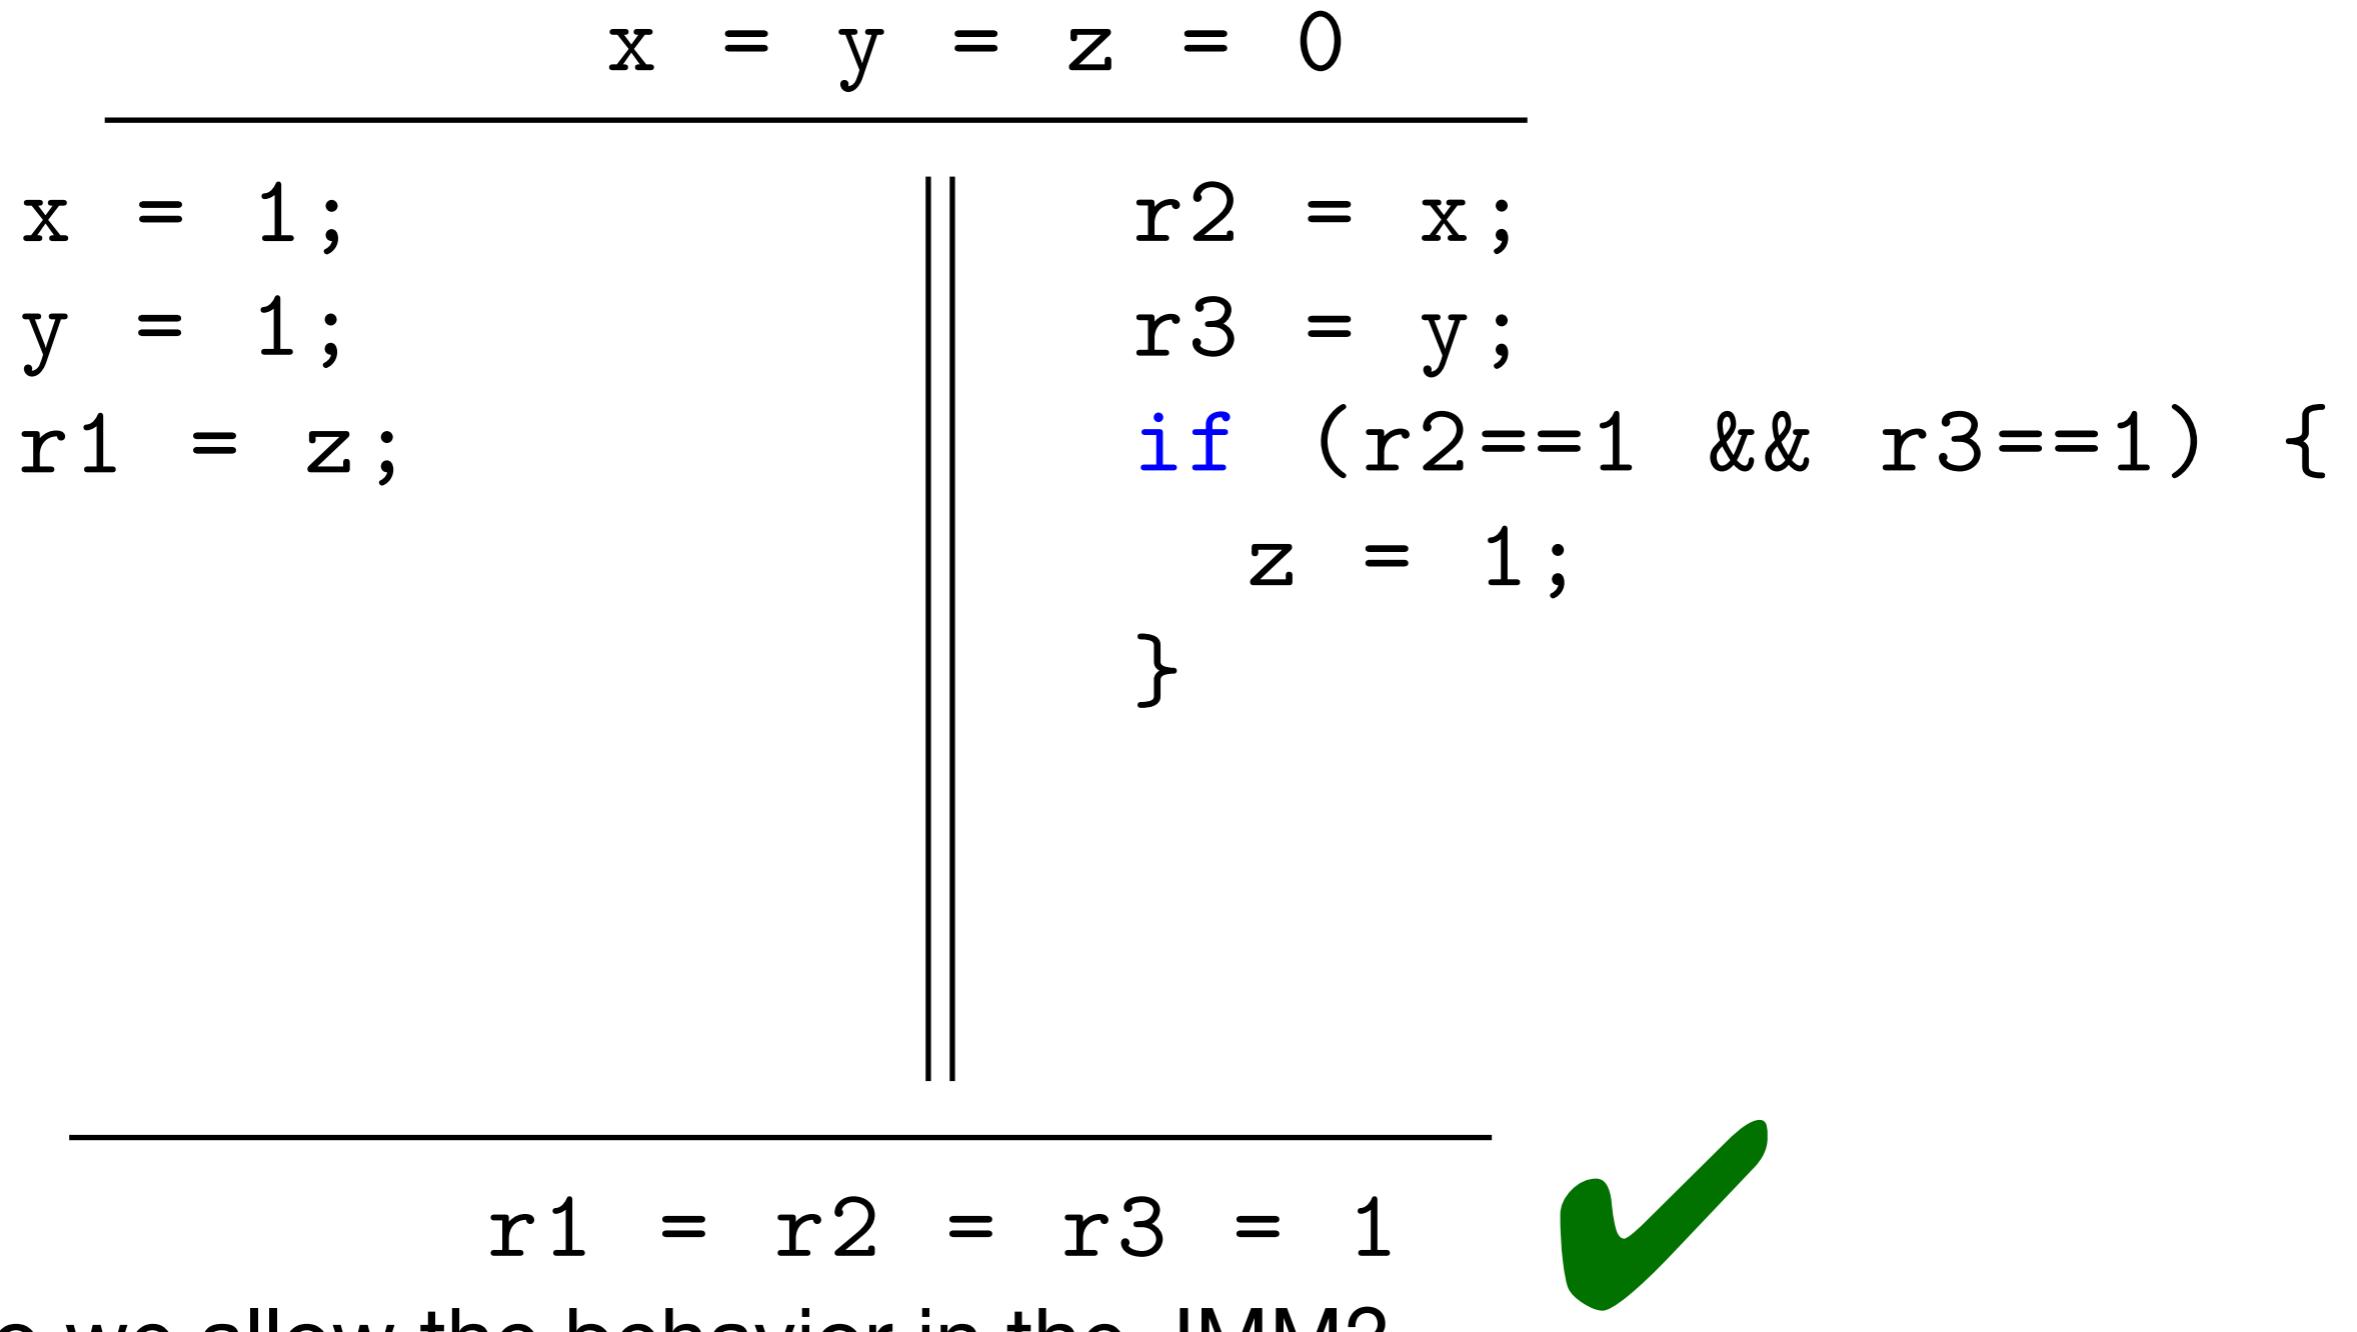
r2 = x;  
r3 = y;  
**if** (r2==1 && r3==1) {  
 z = 1;  
}

---

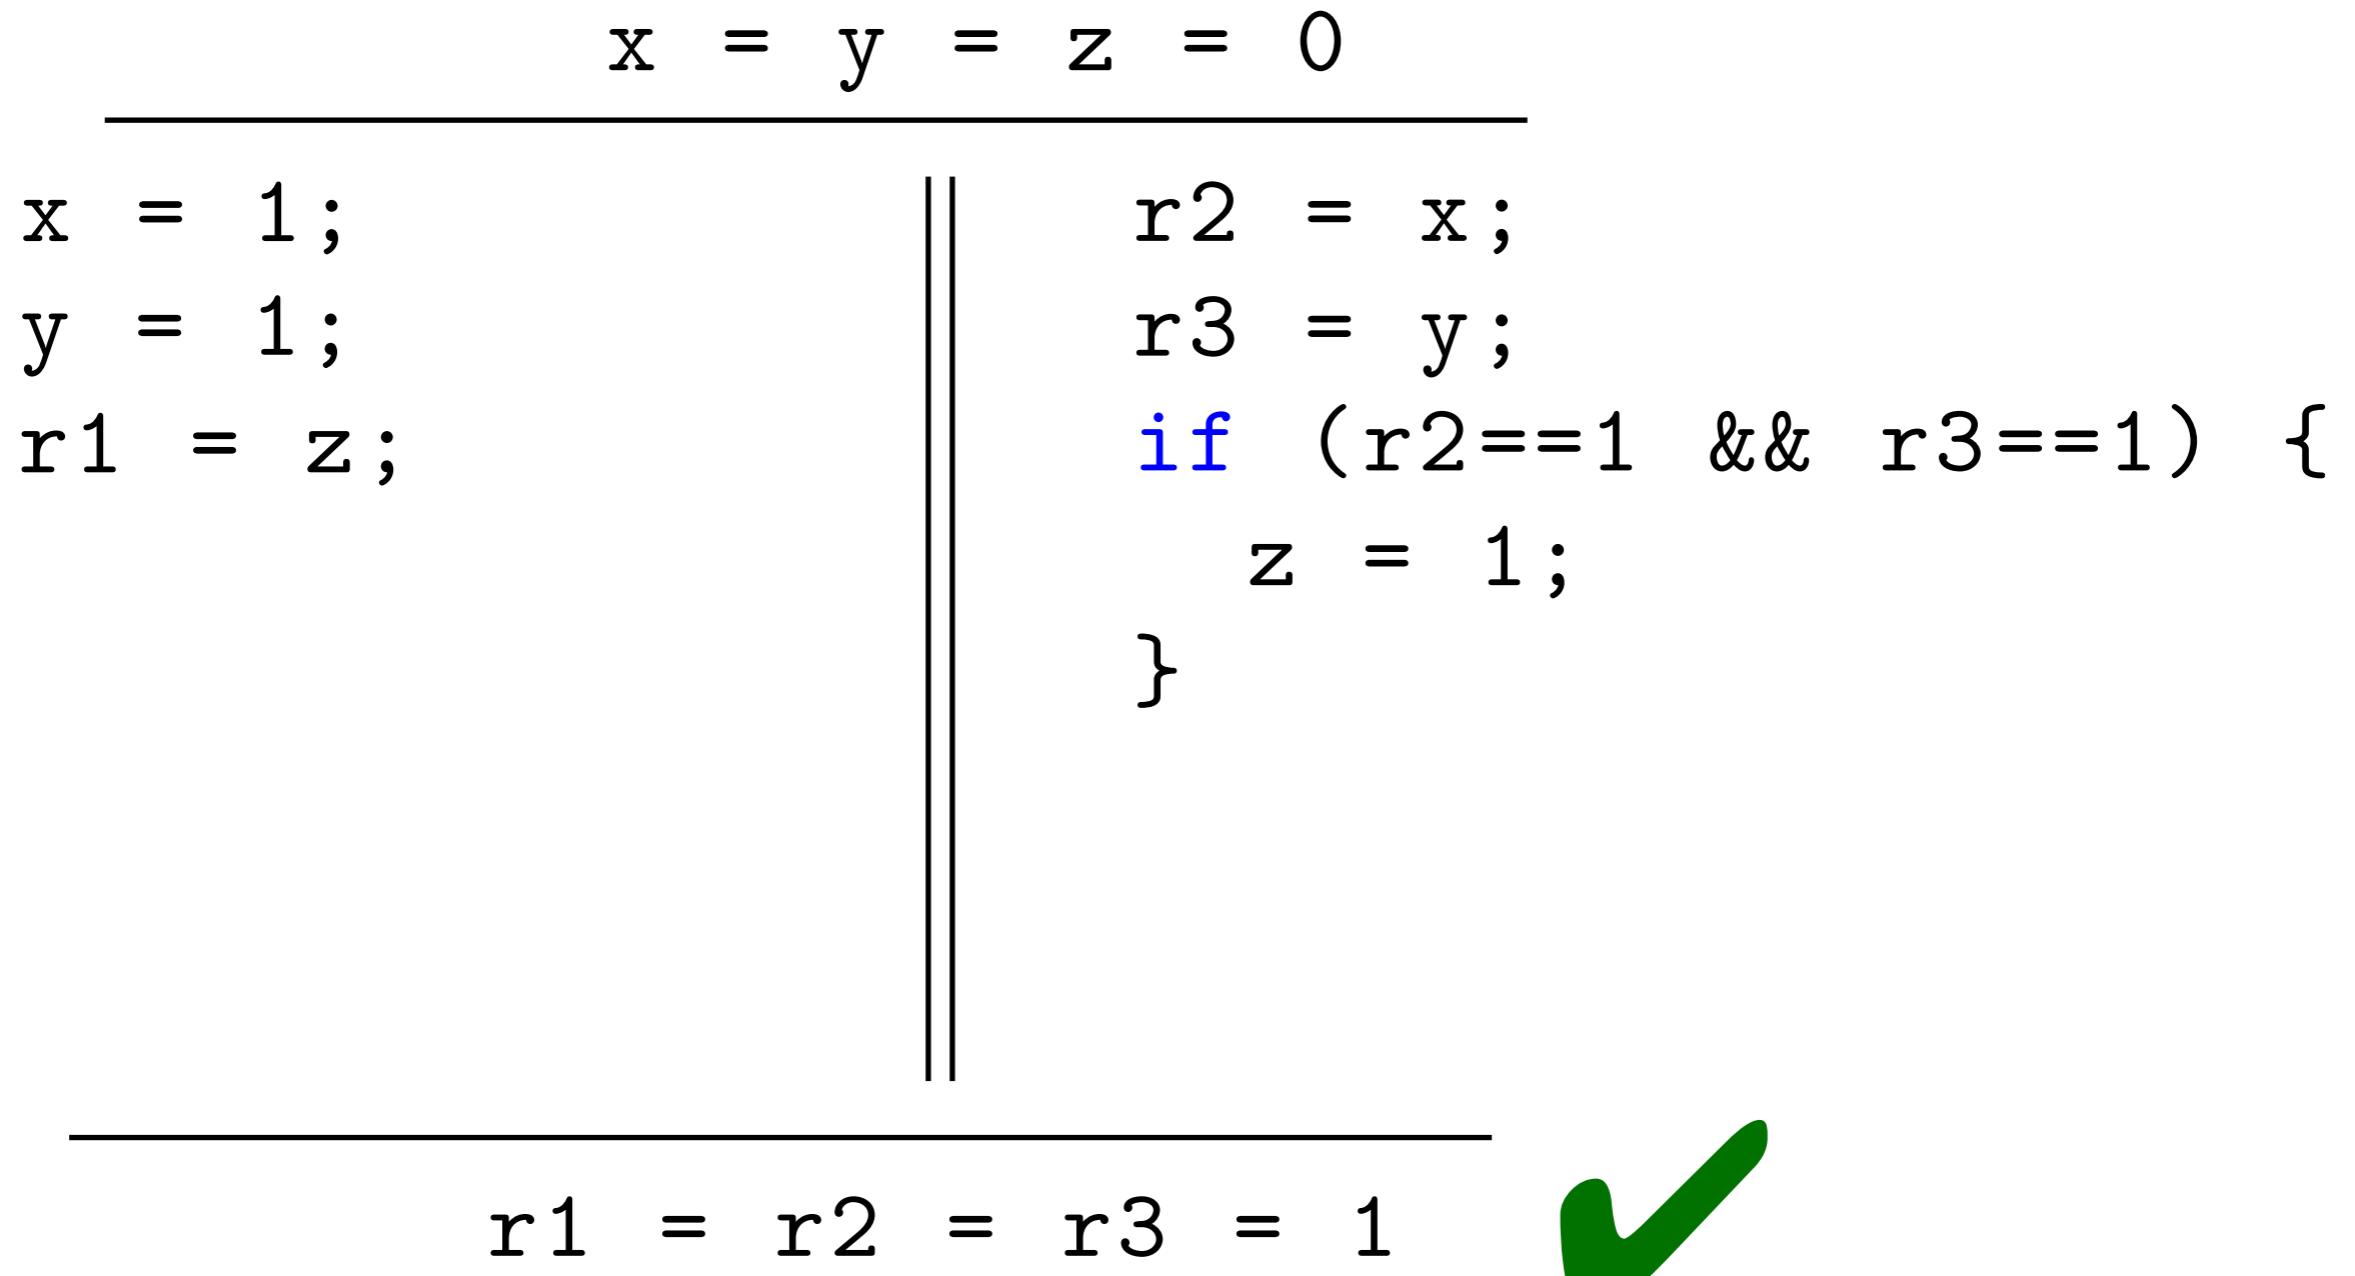
r1 = r2 = r3 = 1



# Compiler Optimizations



# Compiler Optimizations



Do we allow the behavior in the JMM?

Do we forbid these simple optimizations?

# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ \text{if}(r_1! = 0) \\ \quad \quad \quad x = r_1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ \text{if}(r_2! = 0) \\ \quad \quad \quad x = r_2; \end{array}} \quad r_1 = r_3 = 42?$$

# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ \text{if}(r_1 \neq 0) \\ \quad \quad \quad x = r_1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ \text{if}(r_2 \neq 0) \\ \quad \quad \quad x = r_2; \end{array}} \quad r_1 = r_3 = 42?$$

If we consider SC executions this program has **no data races**

# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ \text{if}(r_1 \neq 0) \\ \quad \quad \quad x = r_1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ \text{if}(r_2 \neq 0) \\ \quad \quad \quad x = r_2; \end{array}} \\ r_1 = r_3 = 42?$$

If we consider SC executions this program has **no data races**

The value 42 appears nowhere in the program

# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ \text{if}(r_1 \neq 0) \\ \quad \quad \quad x = r_1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ \text{if}(r_2 \neq 0) \\ \quad \quad \quad x = r_2; \end{array}} \\ r_1 = r_3 = 42?$$

If we consider SC executions this program has **no data races**

The value 42 appears nowhere in the program

Explicitly **forbidden** by the JMM and C++11

# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ x = (r_1 == 1) ? r_1 : 1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ y = (r_2 == 1)?r_2 : 1; \end{array}} \quad r_1 = r_2 = 1?$$

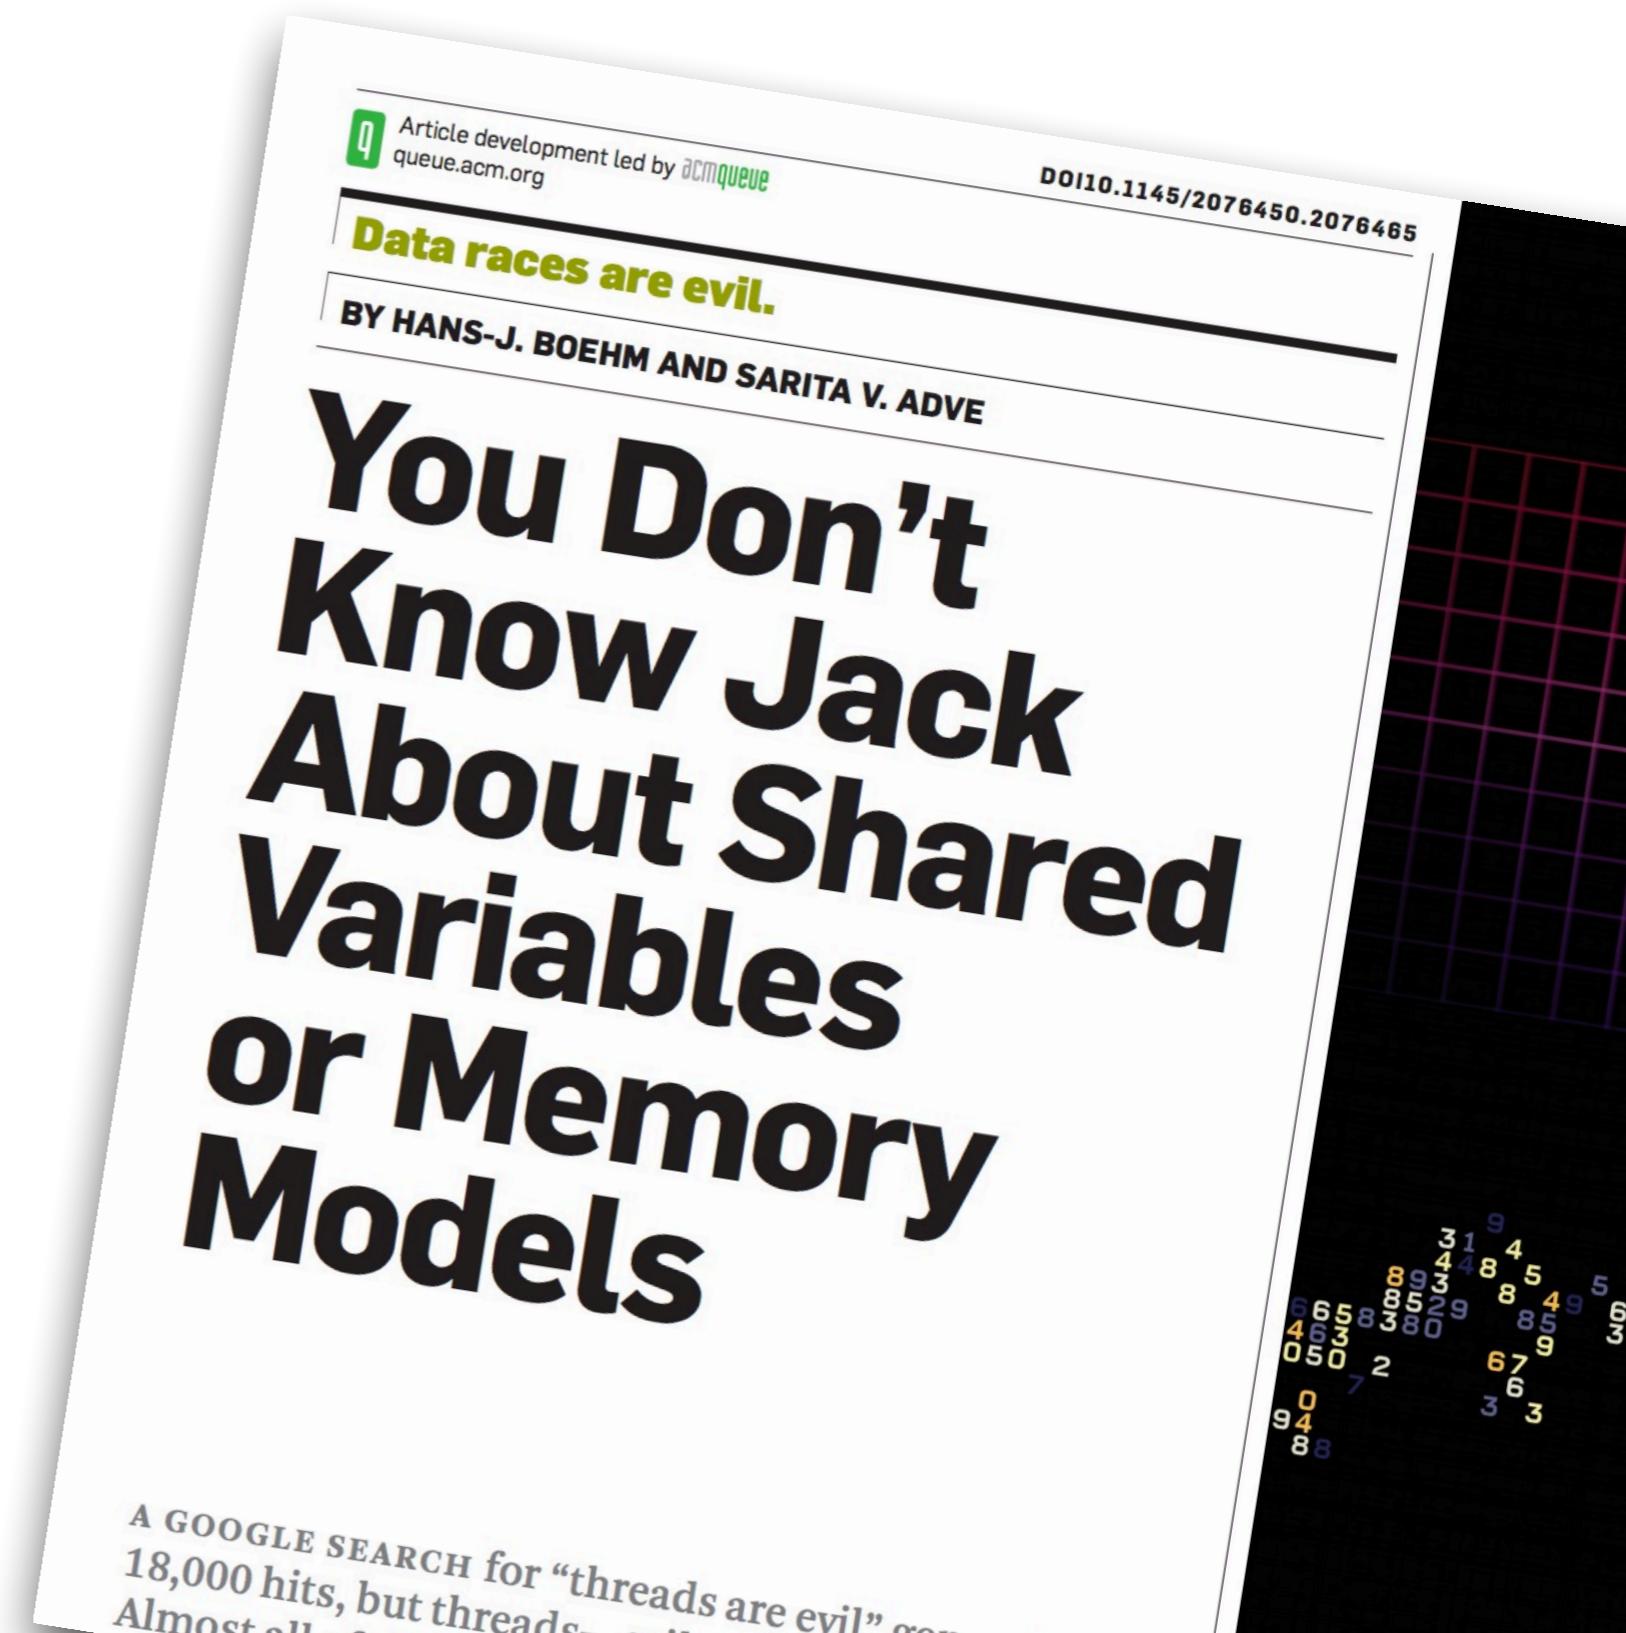
# Thin-Air-Reads

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ x = (r_1 == 1) ? r_1 : 1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ y = (r_2 == 1) ? r_2 : 1; \end{array}} \frac{}{r_1 = r_2 = 1?}$$

$\approx$

$$\frac{x = 0 \quad \& \quad y = 0}{\begin{array}{c} r_1 = y; \\ x = 1; \end{array} \quad \parallel \quad \begin{array}{c} r_2 = x; \\ y = 1; \end{array}} \frac{}{r_1 = r_2 = 1?}$$

# The fundamental Property of Memory Models



# The fundamental Property of Memory Models

# The fundamental Property of Memory Models

- Data Race
  - concurrent memory accesses on a single variable
  - at least one is a write

$$x = 1 \parallel x = 2$$
$$x = 1 \parallel r = x$$

# The fundamental Property of Memory Models

- Data Race
  - concurrent memory accesses on a single variable
  - at least one is a write

$$x = 1 \parallel x = 2$$

$$x = 1 \parallel r = x$$

- Data Race Freedom
  - Every SC execution is free of data races

# The fundamental Property of Memory Models

- Data Race
  - concurrent memory accesses on a single variable
  - at least one is a write

$$x = 1 \parallel x = 2$$

$$x = 1 \parallel r = x$$

- Data Race Freedom
  - Every SC execution is free of data races
- DRF Guarantee
  - DRF programs do not exhibit relaxed behaviors

Relaxed Executions

SC Executions

DRF Executions



Relaxed Executions  
SC Executions  
DRF Executions



Relaxed Executions  
SC Executions  
DRF Executions

- Remark: To check if a program is DRF we can use SC tools



Relaxed Executions  
SC Executions  
DRF Executions

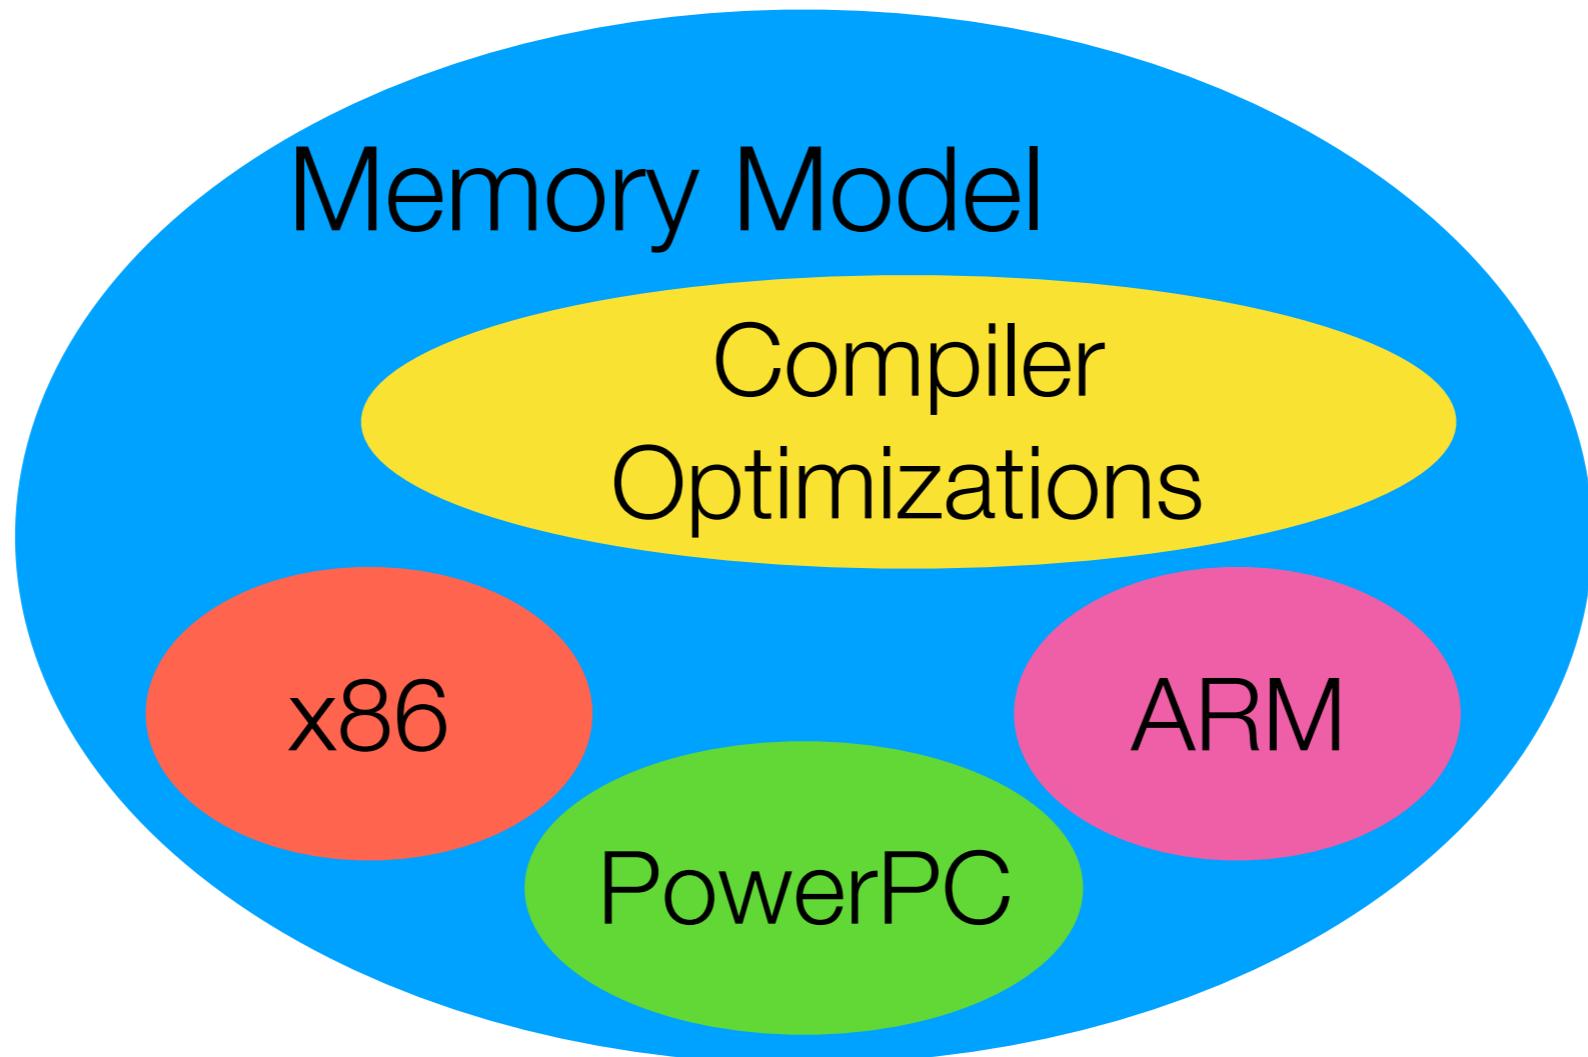
- Remark: To check if a program is DRF we can use SC tools
- Corollary: If a program is DRF we can use any SC too



Relaxed Executions  
SC Executions  
DRF Executions

- Remark: To check if a program is DRF we can use SC tools
- Corollary: If a program is DRF we can use any SC too
- Verification Strategy:
  1. Check that the program is DRF
  2. Verify the program using any SC tool

# Memory Models in Practice



# What if we define DRF to be the memory model?

- Someone has to implement the JVM
- Someone has to implement `java.util.concurrent`
- There are no complete DRF detection algorithms
- Performance critical lock-free data structures
- What happens if despite the programmers effort there are data races?

# C++ Memory Model

# C++ Memory Model

- Catch-fire semantics for data races

# C++ Memory Model

- Catch-fire semantics for data races
- Variables declared atomic cannot race
  - std::atomic<int> x;
  - Memory Ordering Decorations:
    - x.load(mo), x.store(v, mo)
    - x.store(5, memory\_order\_release)
    - x.load(memory\_order\_seq\_cst)
  - Release/Acquire: Message Passing
  - Release/Consume: Message Passing modulo dependencies
  - Relaxed: simply not a data race

# C++ Memory Model

```
bool t1;
bool t2;
int counter = 1;

t1 = true;           || t2 = true;
if (!t2)            || if (!t1)
    counter--;      || counter--;
assert(counter >= 0);
```

# C++ Memory Model

```
bool t1;
bool t2;
int counter = 1;

t1 = true;           || t2 = true;
if (!t2)            || if (!t1)
    counter--;      || counter--;
assert(counter >= 0);
```



# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

t1.store(true, relaxed);    || t2.store(true, relaxed);
if (!t2.load(relaxed))     || if (!t1.load(relaxed))
    counter--;               || counter--;

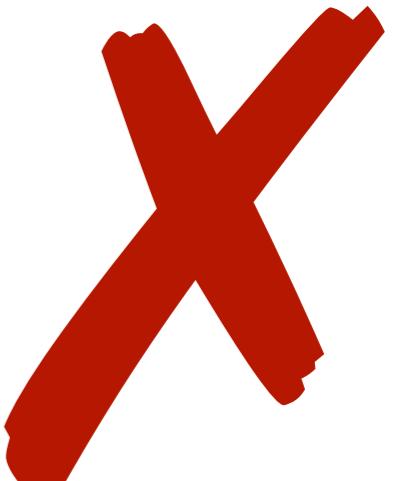
assert(counter >= 0);
```

# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

t1.store(true, relaxed);    || t2.store(true, relaxed);
if (!t2.load(relaxed))     || if (!t1.load(relaxed))
    counter--;               counter--;

assert(counter >= 0);
```



# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

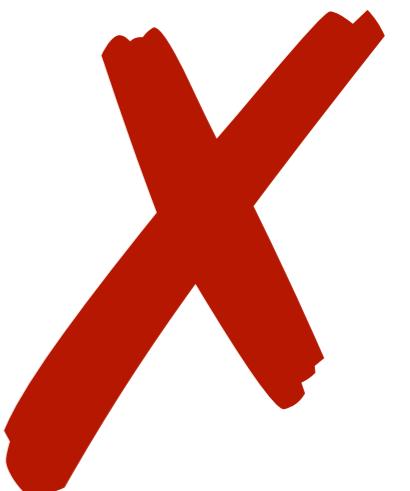
t1.store(true, release);    || t2.store(true, release);
if (!t2.load(acquire))    || if (!t1.load(acquire))
    counter--;            || counter--;

assert(counter >= 0);
```

# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

t1.store(true, release);    || t2.store(true, release);
if (!t2.load(acquire))     || if (!t1.load(acquire))
    counter--;
assert(counter >= 0);
```



# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

t1.store(true, seq_cst);    || t2.store(true, seq_cst);
if (!t2.load(seq_cst))     || if (!t1.load(seq_cst))
    counter--;             || counter--;

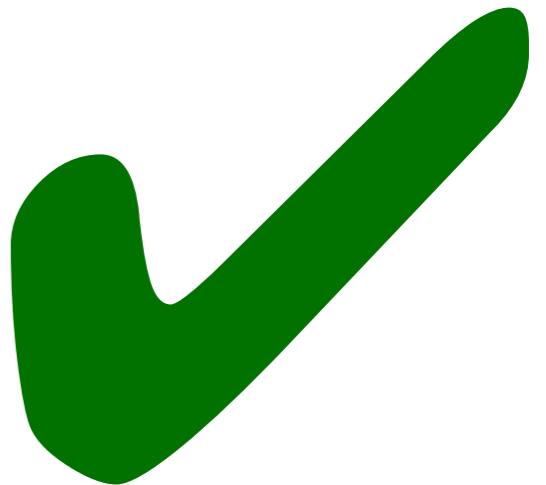
assert(counter >= 0);
```

# C++ Memory Model

```
atomic<bool> t1;
atomic<bool> t2;
int counter = 1;

t1.store(true, seq_cst);    || t2.store(true, seq_cst);
if (!t2.load(seq_cst))     || if (!t1.load(seq_cst))
    counter--;             || counter--;

assert(counter >= 0);
```



# C++11 demo

CPPMem

# The Java Memory Model (JMM)

# The JMM by Intimidation

$x = 0$

```
lock( $\ell$ );
 $x++;$     ||   lock( $\ell$ );
unlock( $\ell$ );      unlock( $\ell$ );    lock( $\ell$ );
 $x++;$     ||    $x++;$     ||   while(1) print( $x$ );
```

0 0 0 1 1 2 2 2 3 3 ...

# The JMM by Intimidation

$x = 0$

```
lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||      lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||      lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||  while(1) print( $x$ );
```

0 0 0 1 1 2 2 2 3 3 ...

$\approx$

```
lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||      lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||      lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||  while(1)
    print(random(0,3));
```

0 3 1 3 2 0 0 1 2 0 0 ...

# The JMM by Intimidation

$x = 0$

```
lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||          lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||          lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||  while(1) print( $x$ );
```

0 0 0 1 1 2 2 2 3 3 ...

$\approx$

```
lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||          lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||          lock( $\ell$ );
 $x++;$ 
unlock( $\ell$ );
||  while(1)
      print(random(0,3));
```

0 3 1 3 2 0 0 1 2 0 0 ...

Of course, no reasonable implementation would do that!

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
so: $1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
so: $1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	
sw: $4 \xrightarrow{sw} 5 \vee 8 \xrightarrow{sw} 1$	

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
so: $1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	
sw: $4 \xrightarrow{sw} 5 \vee 8 \xrightarrow{sw} 1$	
hb: $1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4 \xrightarrow{hb} 5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \vee$	
hb: $5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \xrightarrow{hb} 1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4$	

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
so: $1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	
sw: $4 \xrightarrow{sw} 5 \vee 8 \xrightarrow{sw} 1$	
hb: $1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4 \xrightarrow{hb} 5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \vee$	
hb: $5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \xrightarrow{hb} 1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4$	

Theorem:

The JMM respects the DRF guarantee



# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

$x = y = 0$	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
$po: 1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
$so: 1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	
$sw: 4 \xrightarrow{sw} 5 \vee 8 \xrightarrow{sw} 1$	
$hb: 1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4 \xrightarrow{hb} 5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \vee$	
$hb: 5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \xrightarrow{hb} 1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4$	

Theorem:

The JMM respects the DRF guarantee



Theorem:

Independent non-synchronizing statements can always be reordered

# JMM in a Nutshell

- happens-before is enforced through locks or volatile variables
- reads can see any non-hb related write, or
- any immediate predecessor in hb

x = y = 0	
1: lock m	5: lock m
2: r1 = x	6: r2 = y
3: y = r1	7: x = r2
4: unlock m	8: unlock m
po: $1 \xrightarrow{po} 2 \xrightarrow{po} 3 \xrightarrow{po} 4; 5 \xrightarrow{po} 6 \xrightarrow{po} 7 \xrightarrow{po} 8$	
so: $1 \xrightarrow{so} 4 \xrightarrow{so} 5 \xrightarrow{so} 8 \vee 5 \xrightarrow{so} 8 \xrightarrow{so} 1 \xrightarrow{so} 4$	
sw: $4 \xrightarrow{sw} 5 \vee 8 \xrightarrow{sw} 1$	
hb: $1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4 \xrightarrow{hb} 5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \vee$	
hb: $5 \xrightarrow{hb} 6 \xrightarrow{hb} 7 \xrightarrow{hb} 8 \xrightarrow{hb} 1 \xrightarrow{hb} 2 \xrightarrow{hb} 3 \xrightarrow{hb} 4$	

Theorem:

The JMM respects the DRF guarantee



Theorem:

Independent non-synchronizing statements can always be reordered



Sevcik, Aspinall 2008

# JMM Roach Motel Semantics

Increasing happens-before is always safe

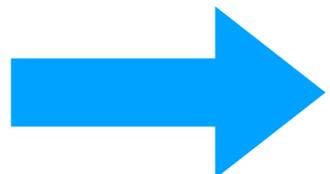
```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```



# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```



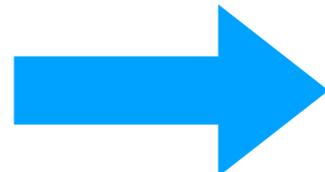
```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```



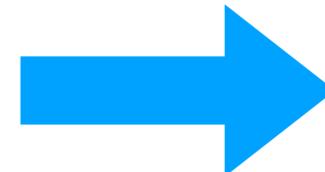
# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```



```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```



```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```



# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
v = 1;  
x = 1;
```



# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
v = 1;  
x = 1;
```

```
x = 1;  
v = 1;
```



# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
v = 1;  
x = 1;
```

```
x = 1;  
v = 1;
```

```
r1 = x;  
r2 = v;
```



# JMM Roach Motel Semantics

Increasing happens-before is always safe

```
x = 1;  
synchronized(this) {  
    y = 1;  
    r1 = x;  
}  
r2 = y;
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
synchronized(this) {  
    x = 1;  
    y = 1;  
    r1 = x;  
    r2 = y;  
}
```

```
v = 1;  
x = 1;  
  
r1 = x;  
r2 = v;
```



# Compiler Optimizations

Transformation	SC	JMM	DRF
Trace-preserving transformations	✓	✓	✓
Reordering normal memory accesses	✗	✗(✓)	✓
Redundant read after read elimination	✓*	✗	✓
Redundant read after write elimination	✓*	✓	✓
Irrelevant read elimination	✓	✓	✓
Irrelevant read introduction	✓	✗	✗
Redundant write before write elimination	✓*	✓	✓
Redundant write after read elimination	✓*	✗	✓
Roach-motel reordering	✓	✗	✓
External action reordering	✗	✗	✓

✓ – correct, ✗ – incorrect, ✓\* – correct only for adjacent memory accesses.

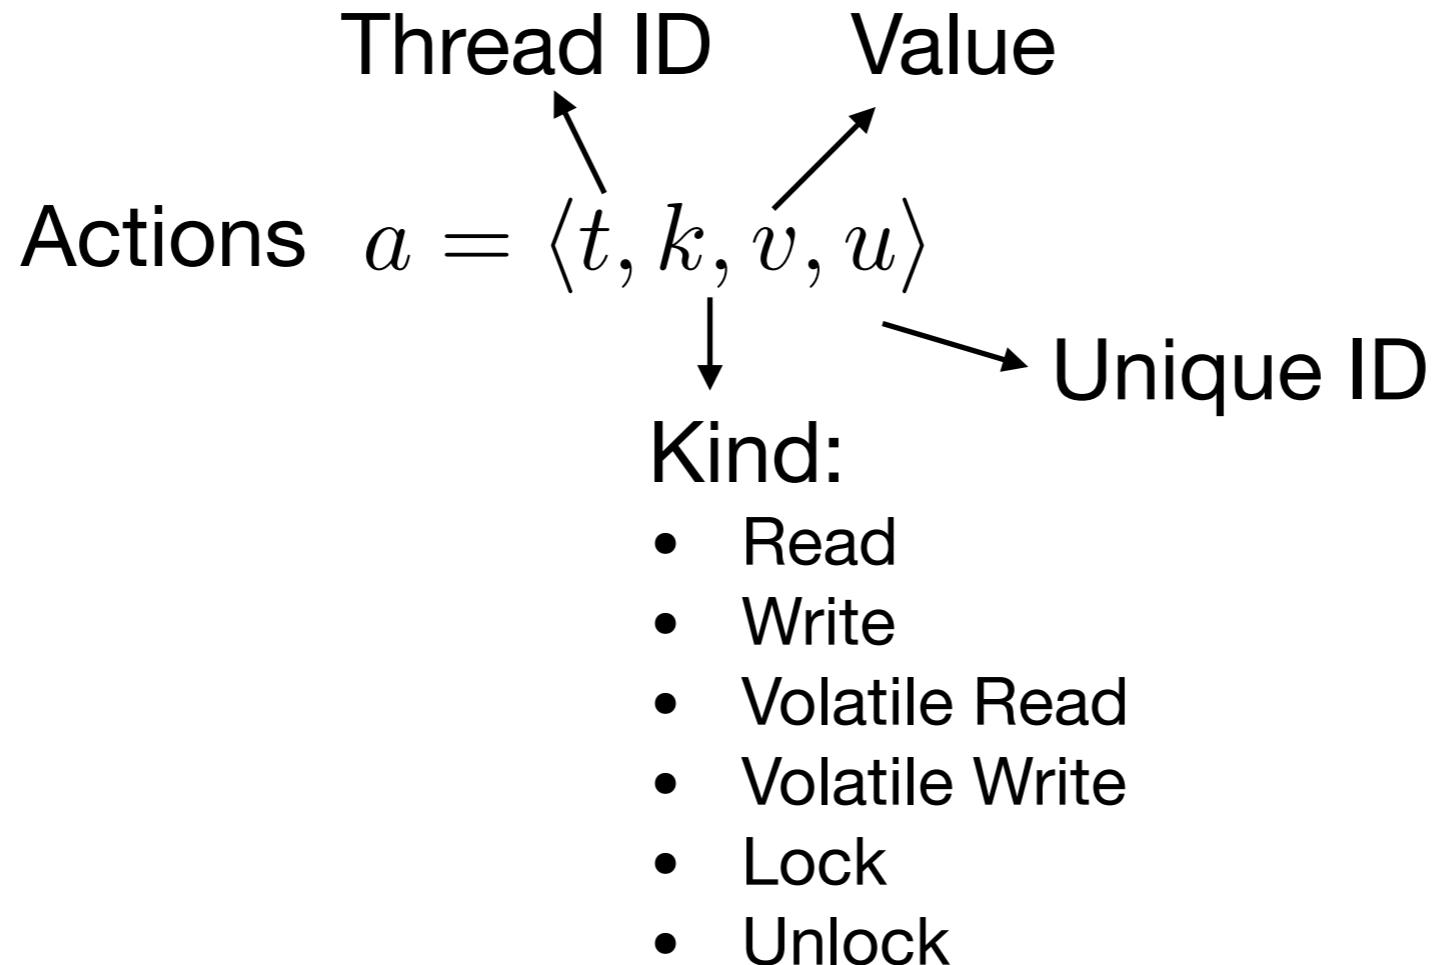
# The Formal JMM

Adve, Manson, Pugh'05

Actions  $a = \langle t, k, v, u \rangle$

# The Formal JMM

Adve, Manson, Pugh'05



# The Formal JMM

Adve, Manson, Pugh'05

**Actions**  $a = \langle t, k, v, u \rangle$

**Executions**  $E = \langle P, A, \xrightarrow{po}, \xrightarrow{so}, W, V, \xrightarrow{sw}, \xrightarrow{hb} \rangle$

# The Formal JMM

Adve, Manson, Pugh'05

Actions  $a = \langle t, k, v, u \rangle$

Program  
↑  
Executions  $E = \langle P, A, \xrightarrow{po}, \xrightarrow{so}, W, V, \xrightarrow{sw}, \xrightarrow{hb} \rangle$

# The Formal JMM

Adve, Manson, Pugh'05

Actions  $a = \langle t, k, v, u \rangle$

Program      Actions  
↑                ↑  
Executions  $E = \langle P, A, \xrightarrow{po}, \xrightarrow{so}, W, V, \xrightarrow{sw}, \xrightarrow{hb} \rangle$

# The Formal JMM

Adve, Manson, Pugh'05

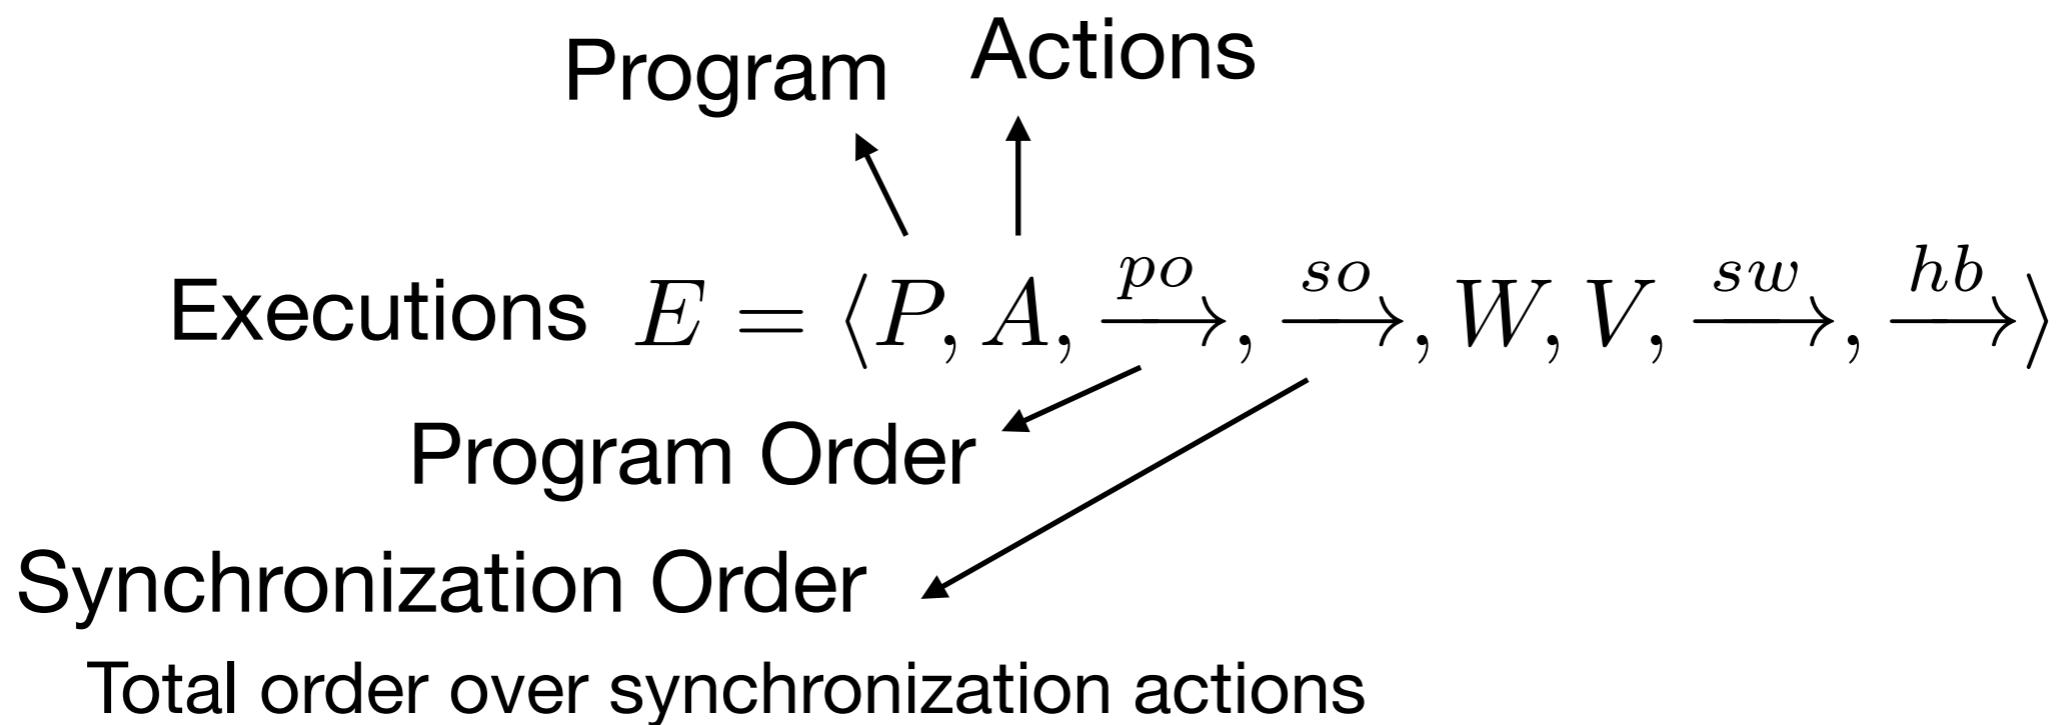
Actions  $a = \langle t, k, v, u \rangle$

Program      Actions  
↑               ↑  
Executions  $E = \langle P, A, \xrightarrow{po}, \xrightarrow{so}, W, V, \xrightarrow{sw}, \xrightarrow{hb} \rangle$   
Program Order    ↖

# The Formal JMM

Adve, Manson, Pugh'05

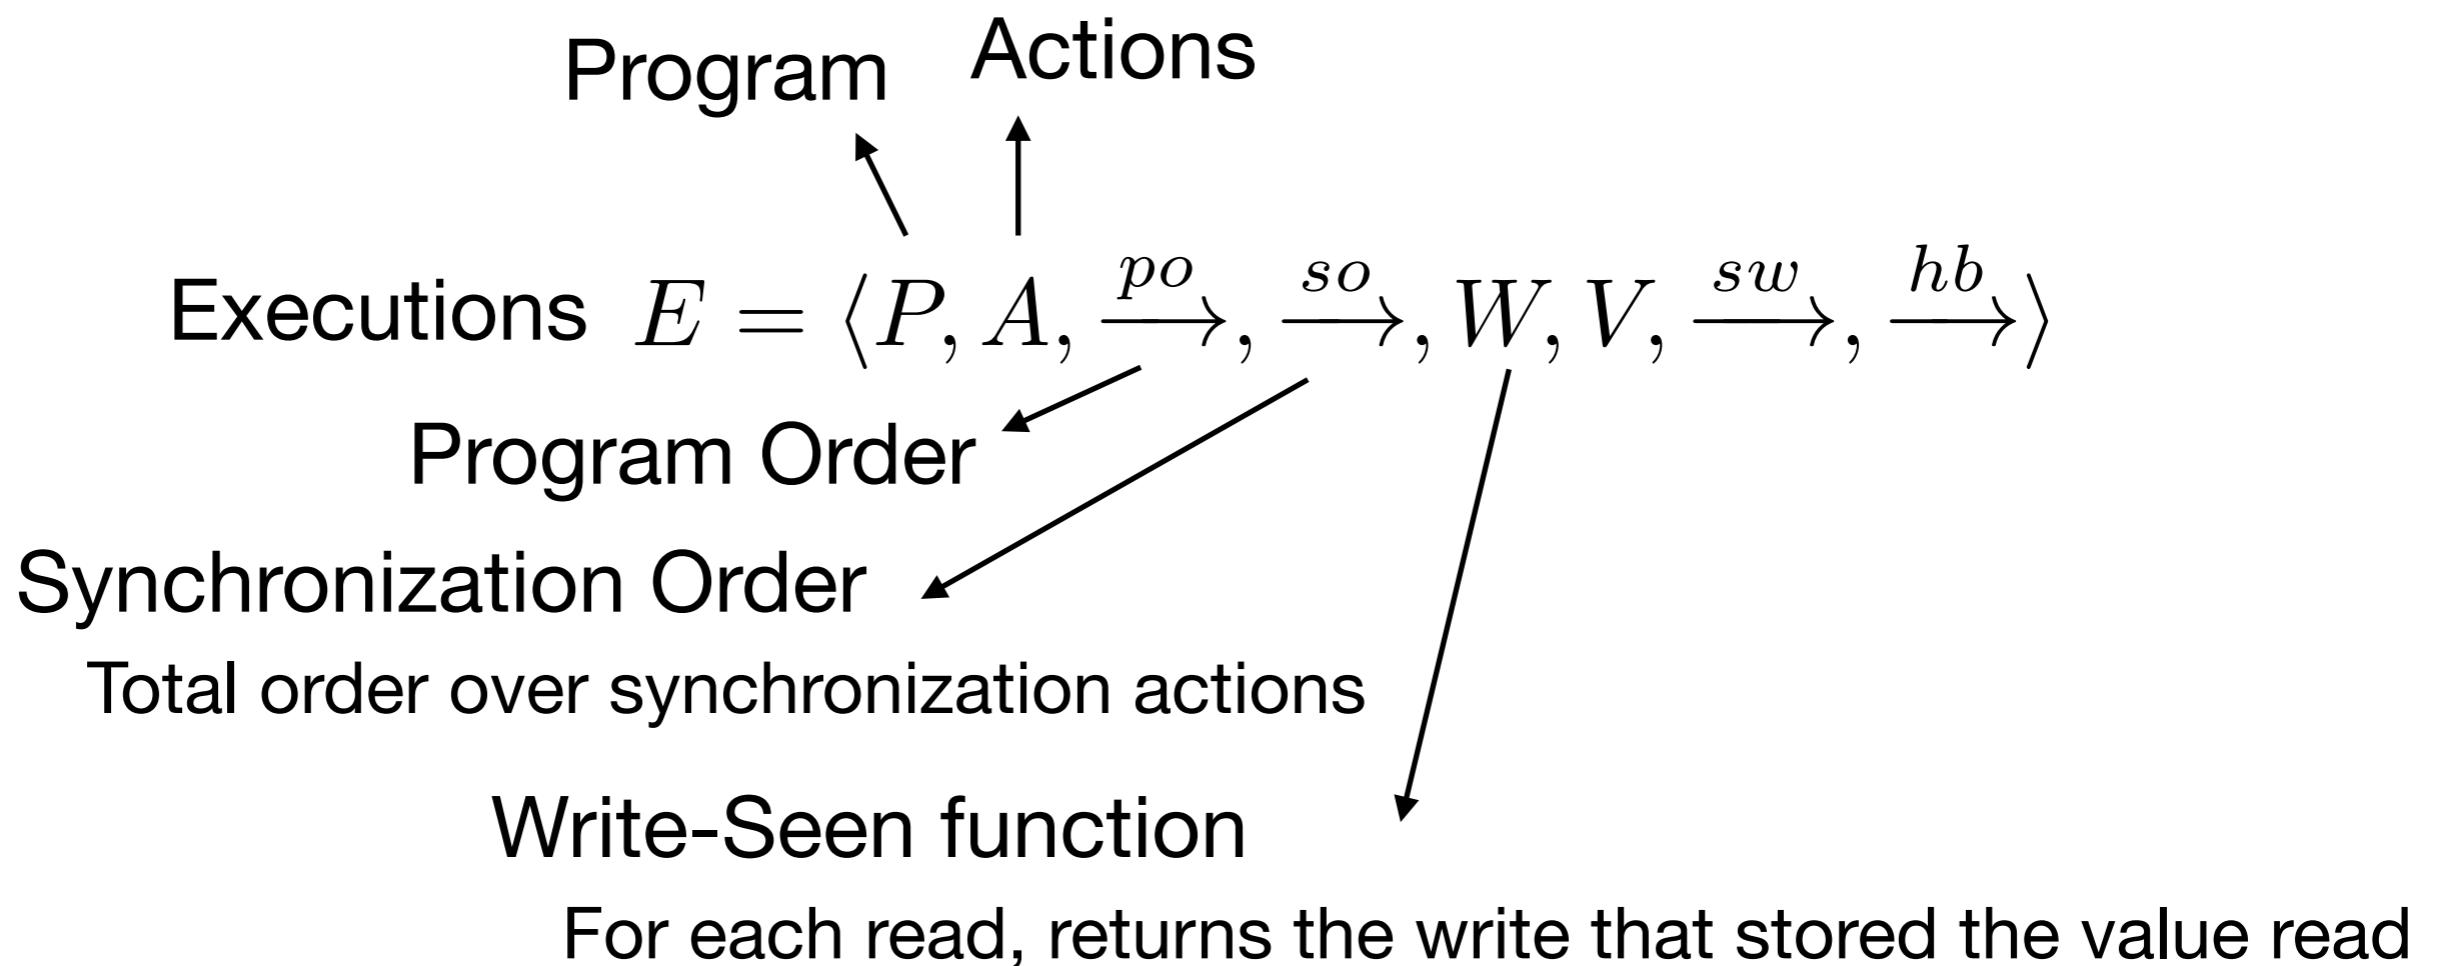
Actions  $a = \langle t, k, v, u \rangle$



# The Formal JMM

Adve, Manson, Pugh'05

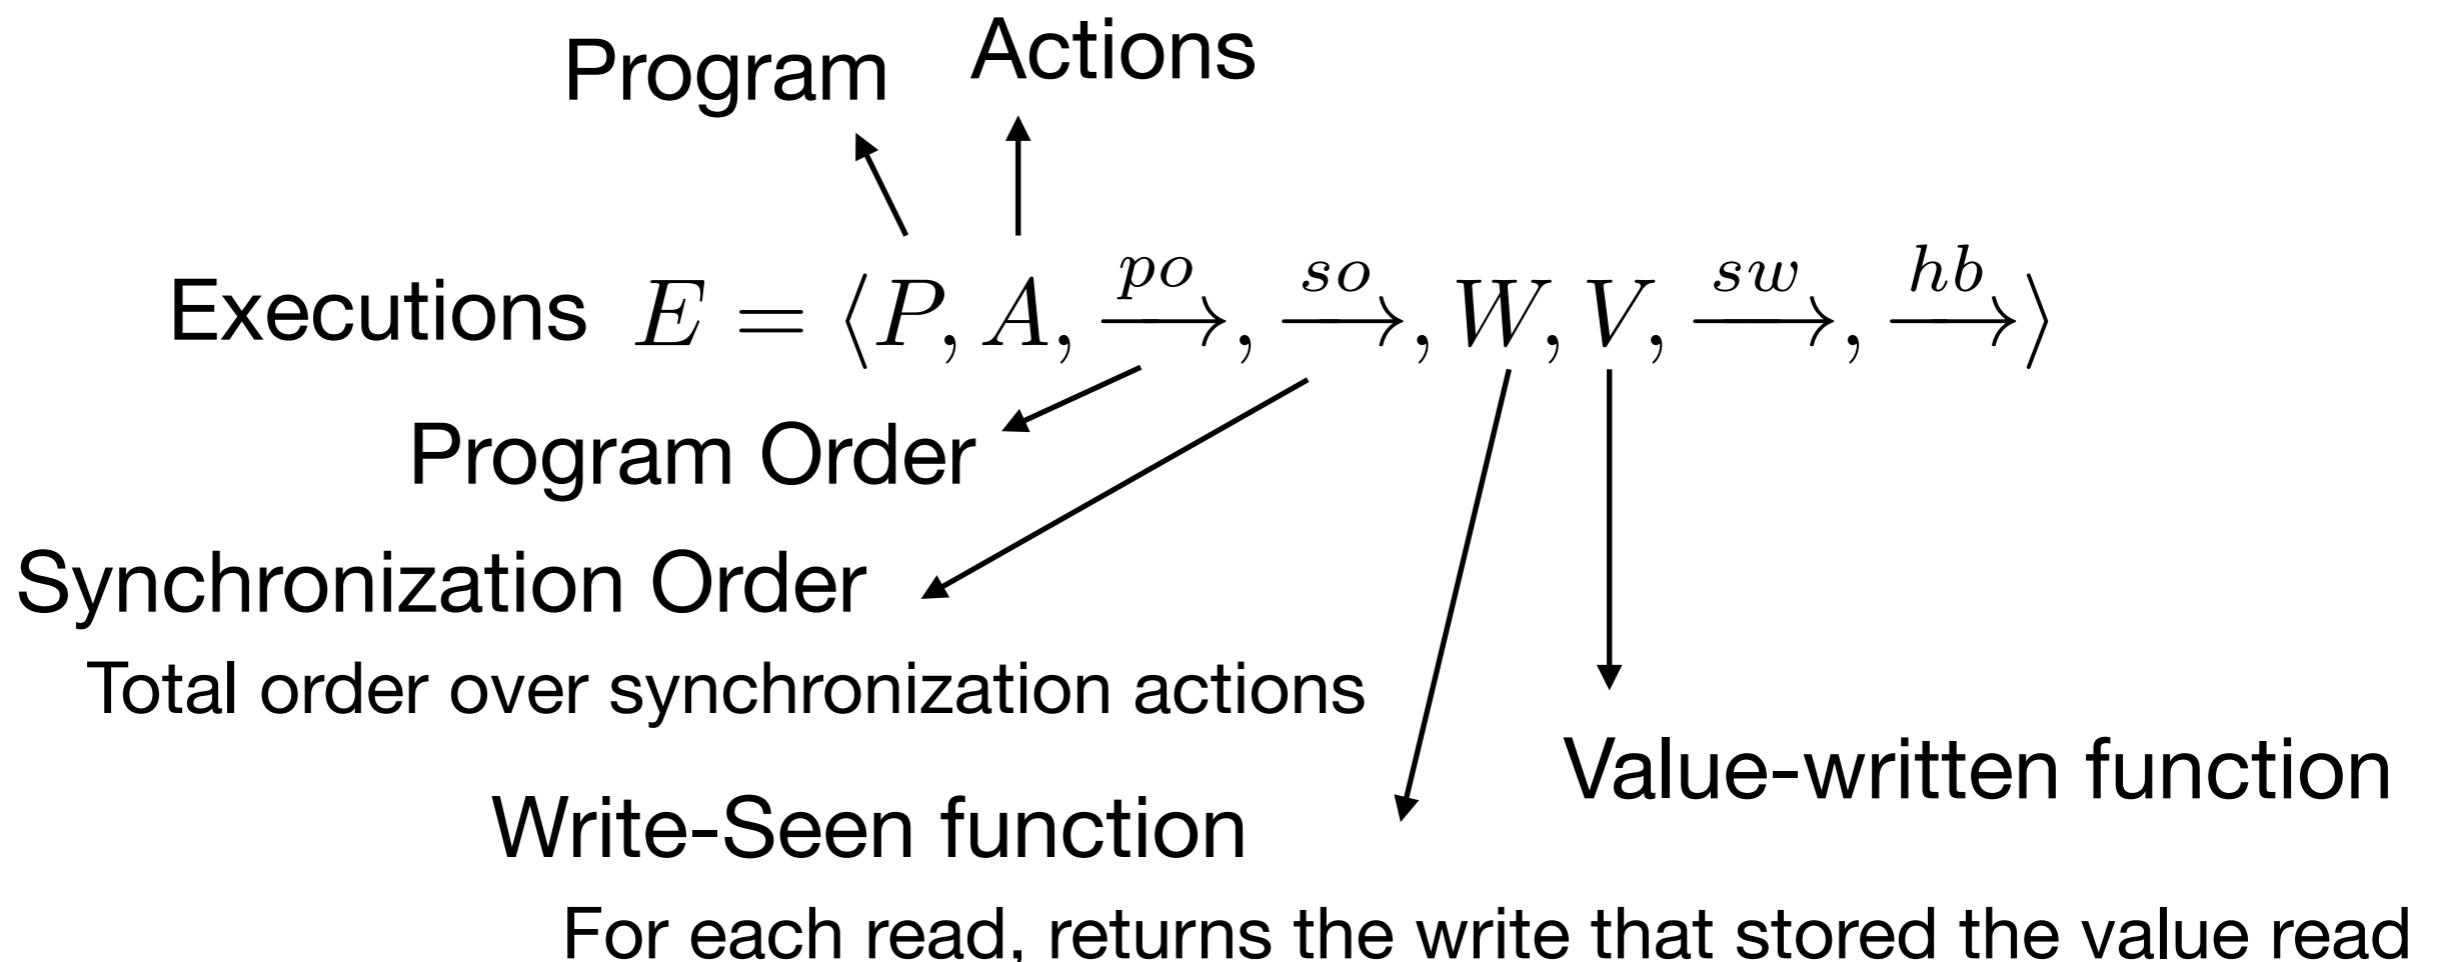
Actions  $a = \langle t, k, v, u \rangle$



# The Formal JMM

Adve, Manson, Pugh'05

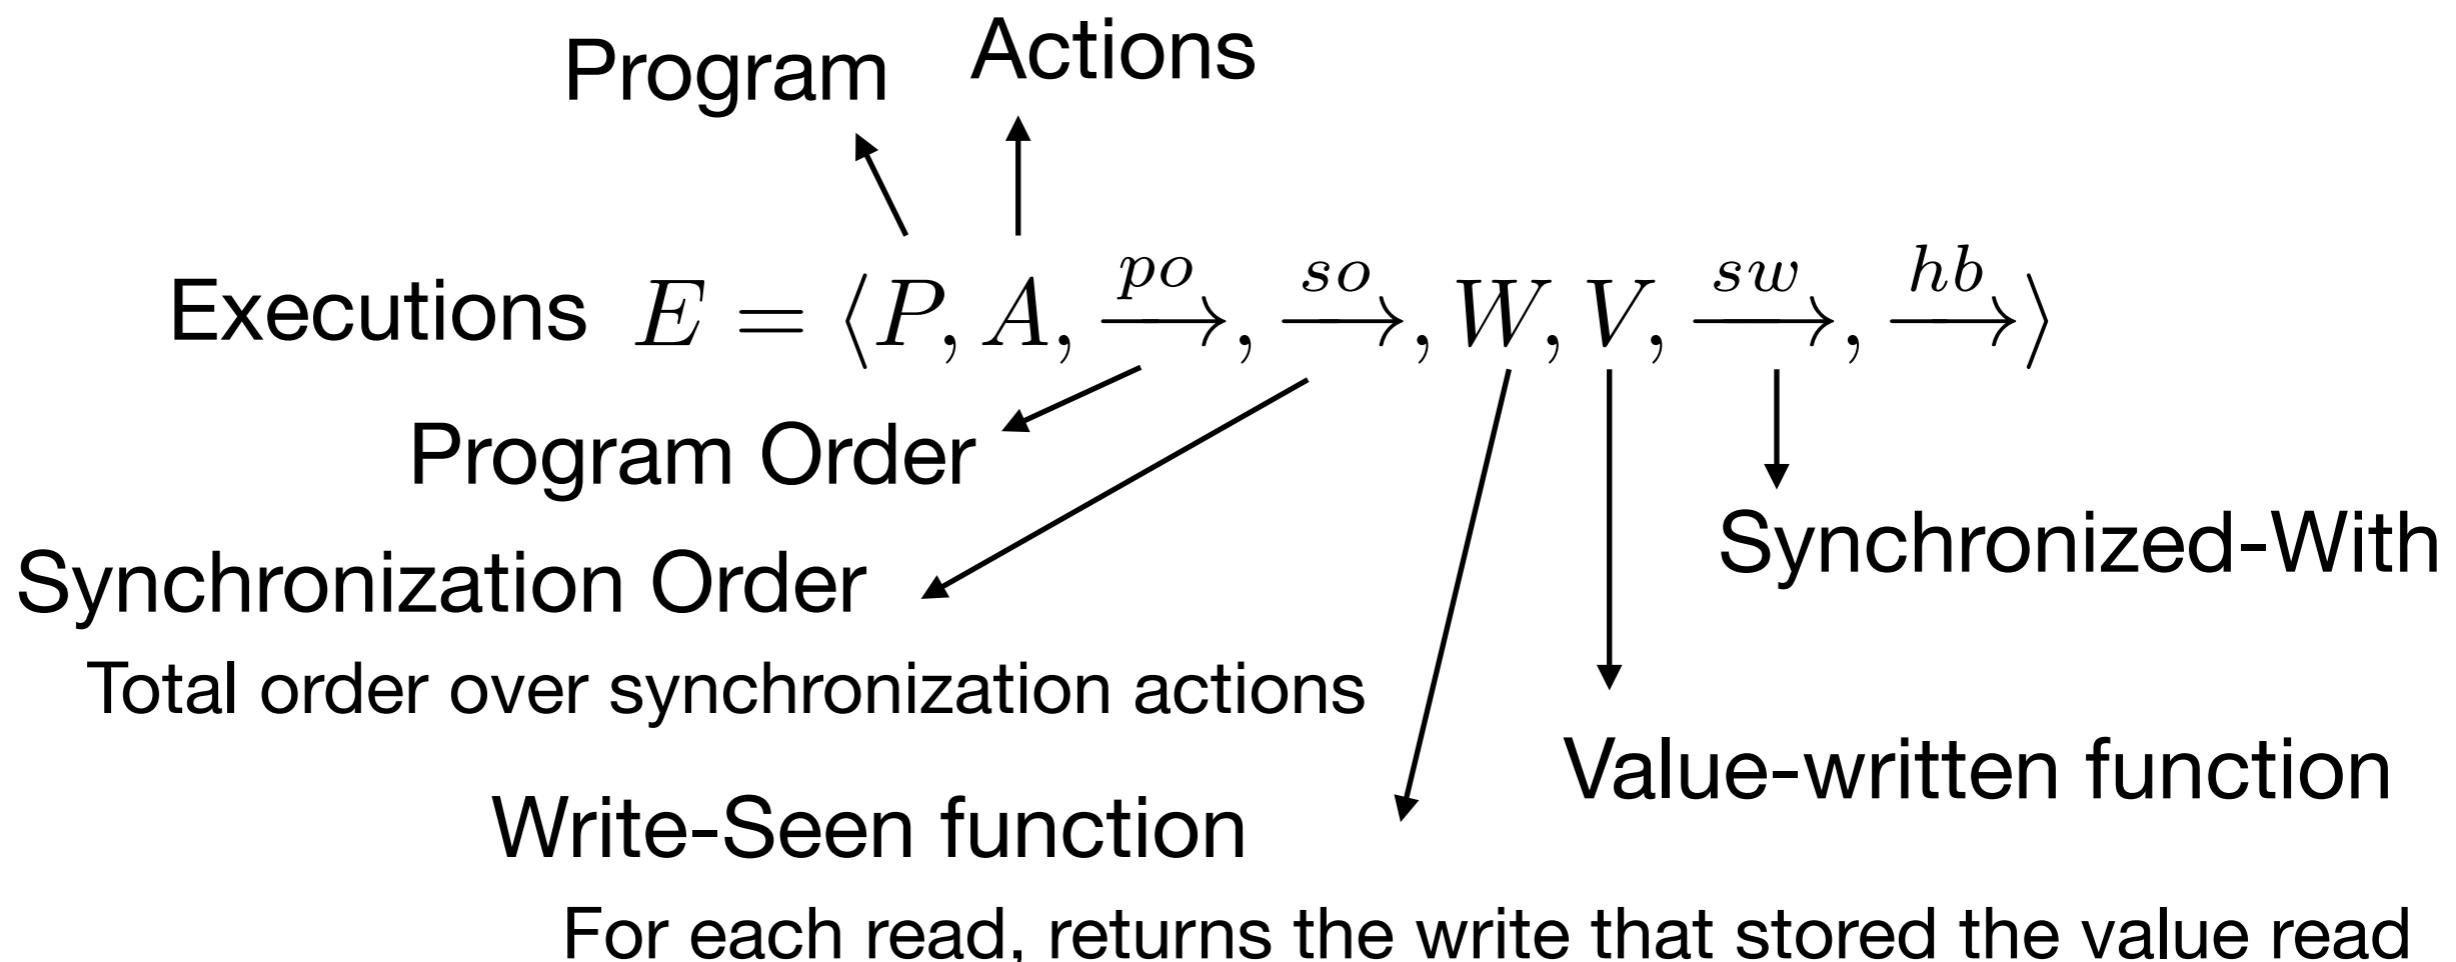
Actions  $a = \langle t, k, v, u \rangle$



# The Formal JMM

Adve, Manson, Pugh'05

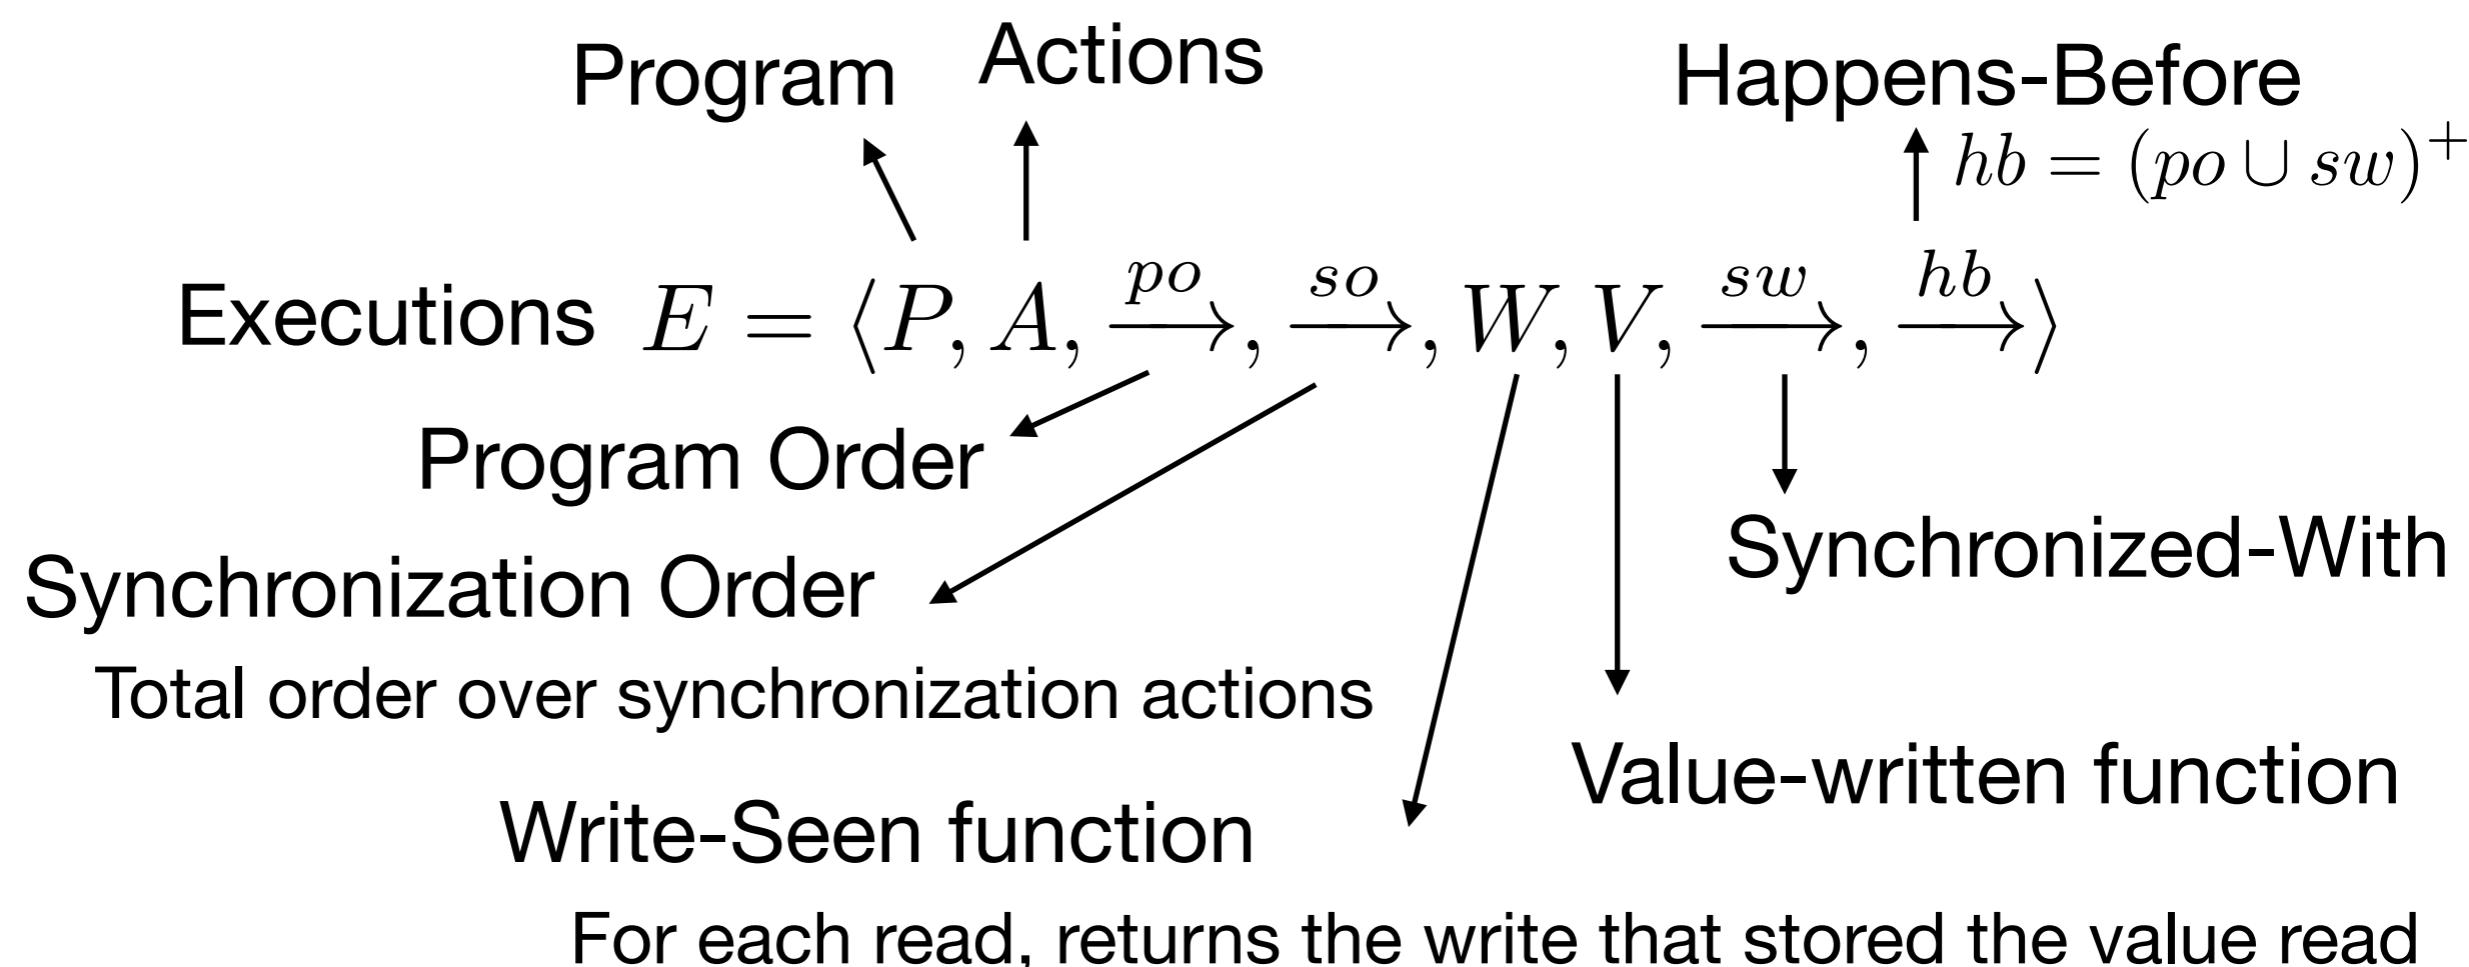
Actions  $a = \langle t, k, v, u \rangle$



# The Formal JMM

# Adve, Manson, Pugh'05

## Actions $a = \langle t, k, v, u \rangle$



# The Formal JMM

Adve, Manson, Pugh'05

## Well-Formedness

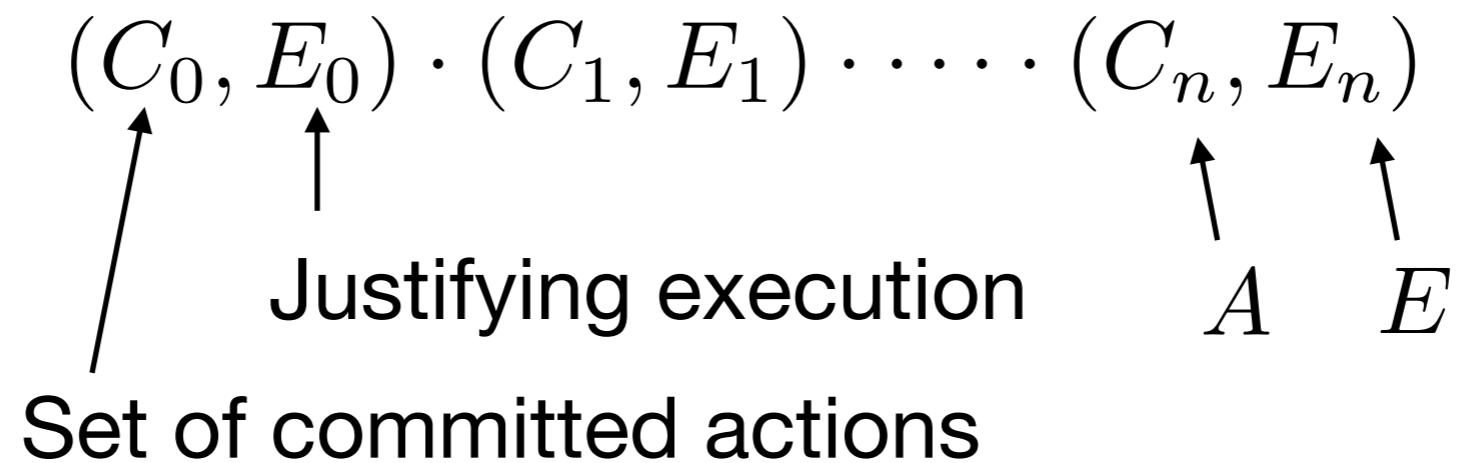
- for every read  $r$ ,  $V(W(r)) = r.v$
- $\xrightarrow{po}$  is consistent with  $\xrightarrow{so}$  ( $\xrightarrow{hb}$  is a partial order)
- lock are consistent with mutual exclusion
- the semantics of single threads is respected
- synchronization is consistent
- happens-before is consistent

# The Formal JMM

Adve, Manson, Pugh'05

Causality: we need to justify each of the actions of the execution

Commit actions of  $E$  step by step



# The Formal JMM

Adve, Manson, Pugh'05

1.  $C_i \subseteq A_i$
2.  $\xrightarrow{hb_i} |_{C_i} = \xrightarrow{hb} |_{C_i}$
3.  $\xrightarrow{so_i} |_{C_i} = \xrightarrow{so} |_{C_i}$
4.  $V_i|_{C_i} = V|_{C_i}$
5.  $\forall r \in Reads(A_i - C_{i-1}) : W_i(r) \xrightarrow{hb_i} r$
6.  $W_i|_{C_{i-1}} = W|_{C_{i-1}}$
7.  $\forall r \in Reads(C_i - C_{i-1}) : \{W_i(r), W(r)\} \subseteq C_{i-1}$

# The Formal JMM

Adve, Manson, Pugh'05

1.  $C_i \subseteq A_i$
2.  $\xrightarrow{hb_i} |_{C_i} = \xrightarrow{hb} |_{C_i}$
3.  $\xrightarrow{so_i} |_{C_i} = \xrightarrow{so} |_{C_i}$
4.  $V_i|_{C_i} = V|_{C_i}$
5.  $\forall r \in Reads(A_i - C_{i-1}) : W_i(r) \xrightarrow{hb_i} r$
6.  $W_i|_{C_{i-1}} = W|_{C_{i-1}}$
7.  $\forall r \in Reads(C_i - C_{i-1}) : \{W_i(r), W(r)\} \subseteq C_{i-1}$

Causality Test Cases

# Operational JMM?

## Generative Operational Semantics for Relaxed Memory Models\*

Radha Jagadeesan, Corin Pitcher, and James Riely

School of Computing, DePaul University

**Abstract.** The specification of the Java Memory Model (JMM) is phrased in terms of acceptors of execution sequences rather than the standard generative view of operational semantics. This creates a mismatch with language-based techniques, such as simulation arguments and proofs of type safety.

We describe a semantics for the JMM using standard programming language techniques that captures its full expressivity. For data-race-free programs, our model coincides with the JMM. For lockless programs, our model is more expressive than the JMM. The stratification properties required to avoid causality cycles are derived, rather than mandated in the style of the JMM.

The JMM is arguably non-canonical in its treatment of the interaction of data races and locks as it fails to validate roach-motel reorderings and various peephole optimizations. Our model differs from the JMM in these cases. We develop a theory of simulation and use it to validate the legality of the above optimizations in any program context.

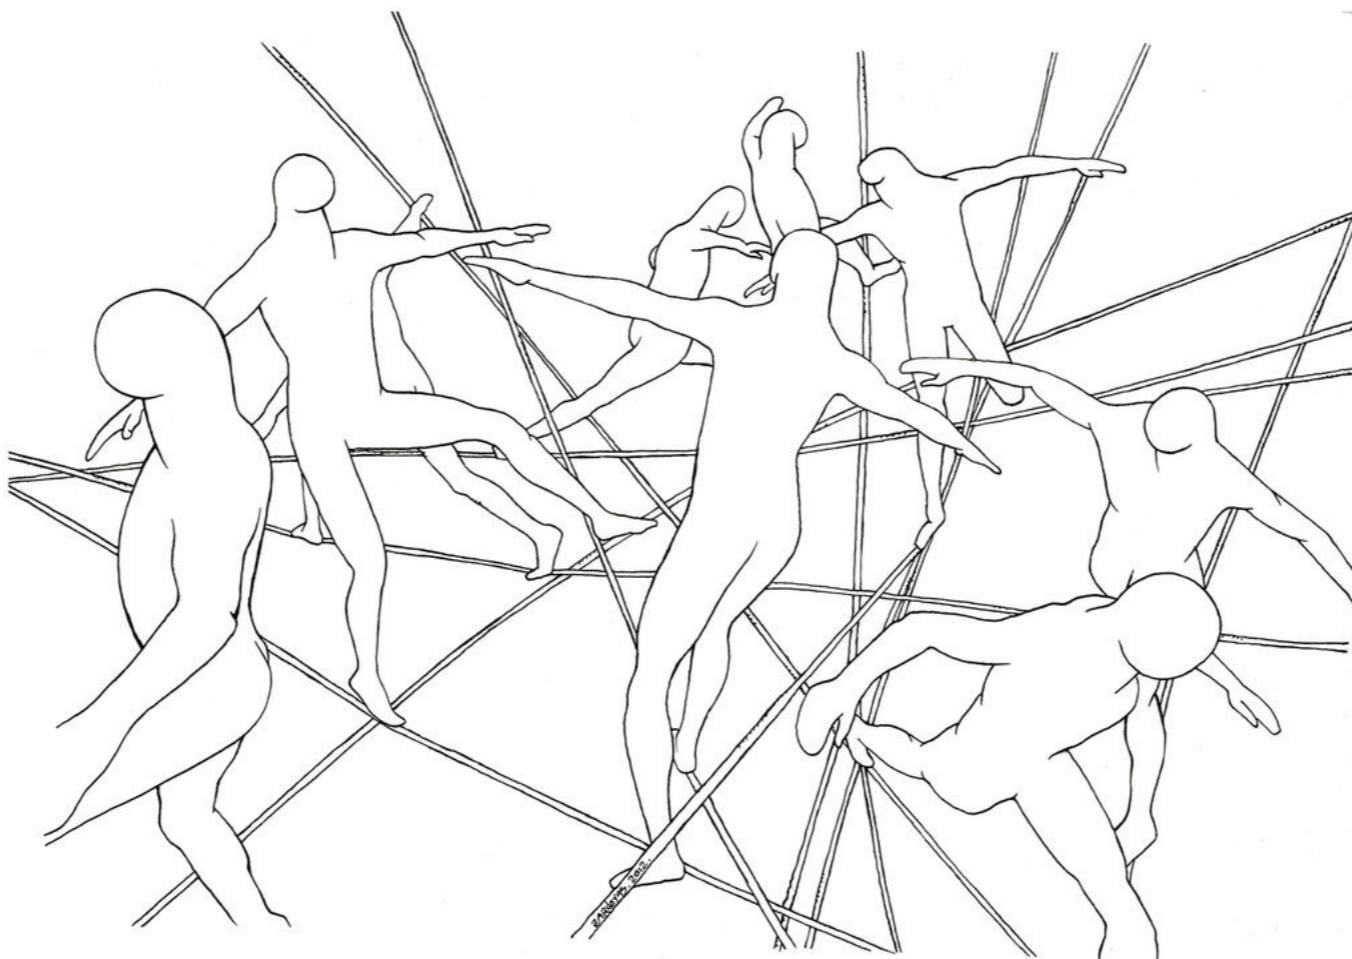
### 1 Introduction

In the context of shared memory imperative programs, Sequential Consistency (SC) (Lamport 1979) enforces a global total order on memory operations that includes the program order of each individual thread in the program. SC may be realized by a traditional interleaving semantics where shared memory is represented as a map from locations to values. It has been observed that SC disables compiler optimizations such as reordering of independent statements. Despite arguments that SC does not impair efficiency (Kamil et al. 2005), this observation and others have motivated a body of work on relaxed memory models; Adve and Gharachorloo (1996) provide a tutorial introduction with detailed bibliography.

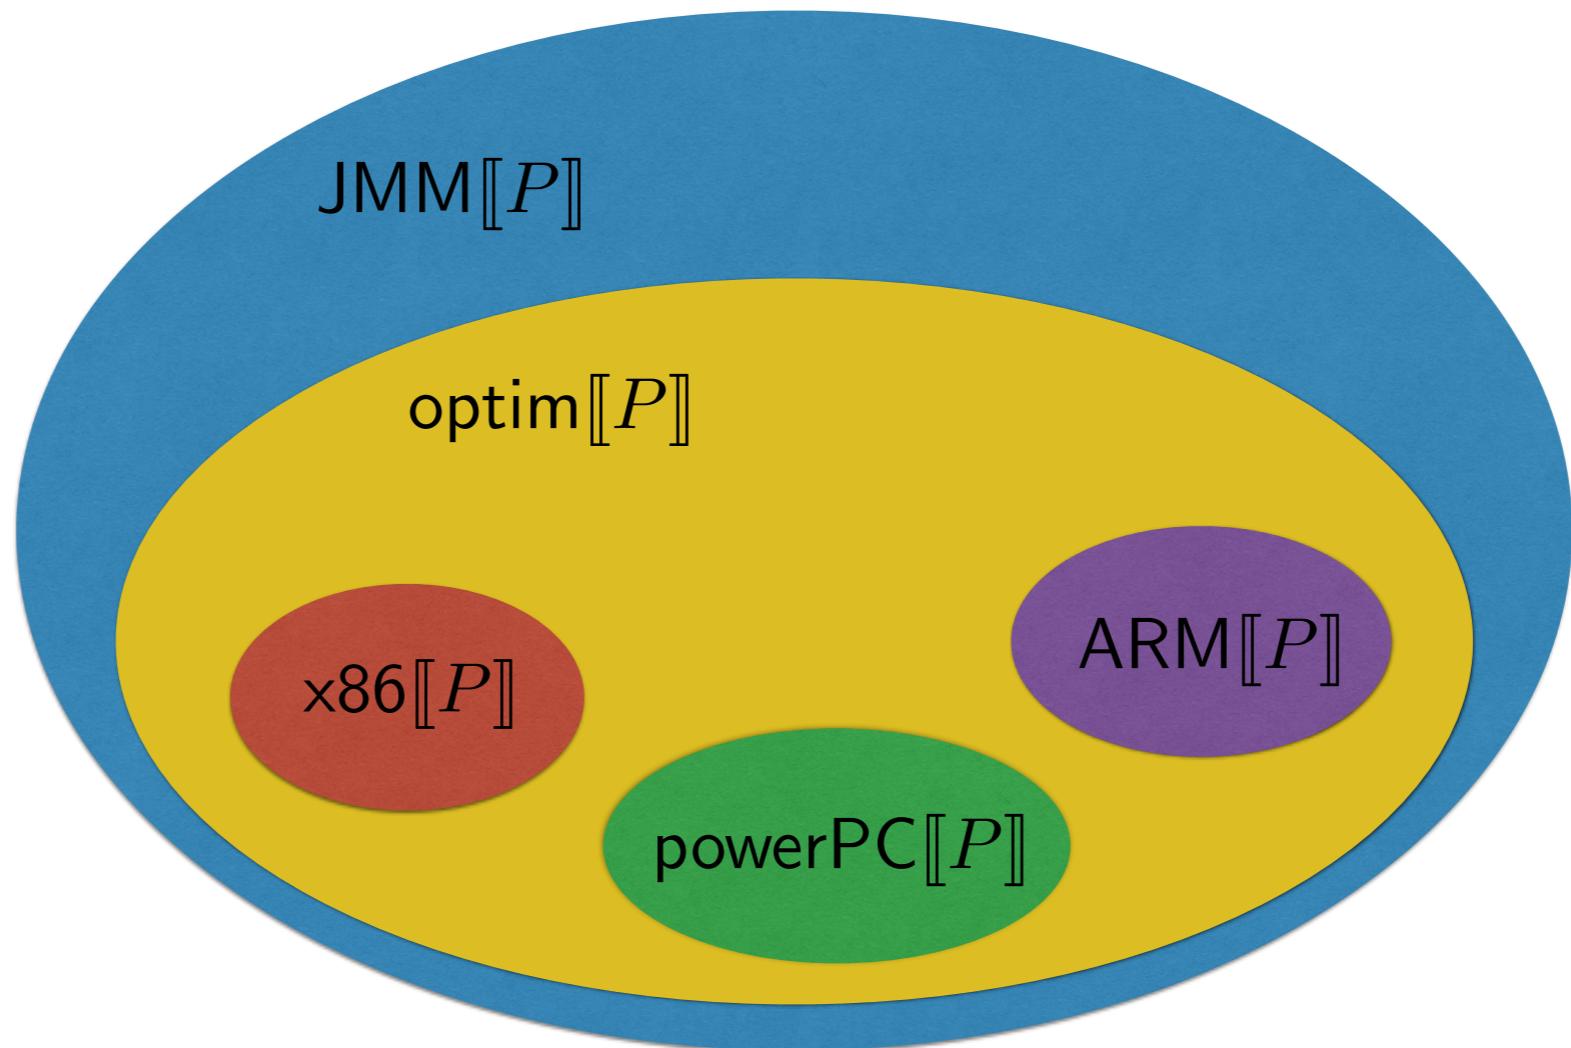
A first (conceptual, if not chronological) step in generalizing SC is to consider the Data Race Free (DRF) models. Informally, a program is DRF if no execution of the program contains a write operation in which a write happens concurrently with another operation on the same location. DRF provides a more refined abstraction than SC, and it preserves the programmer view of computation co-

- Partial solutions
- Remains elusive

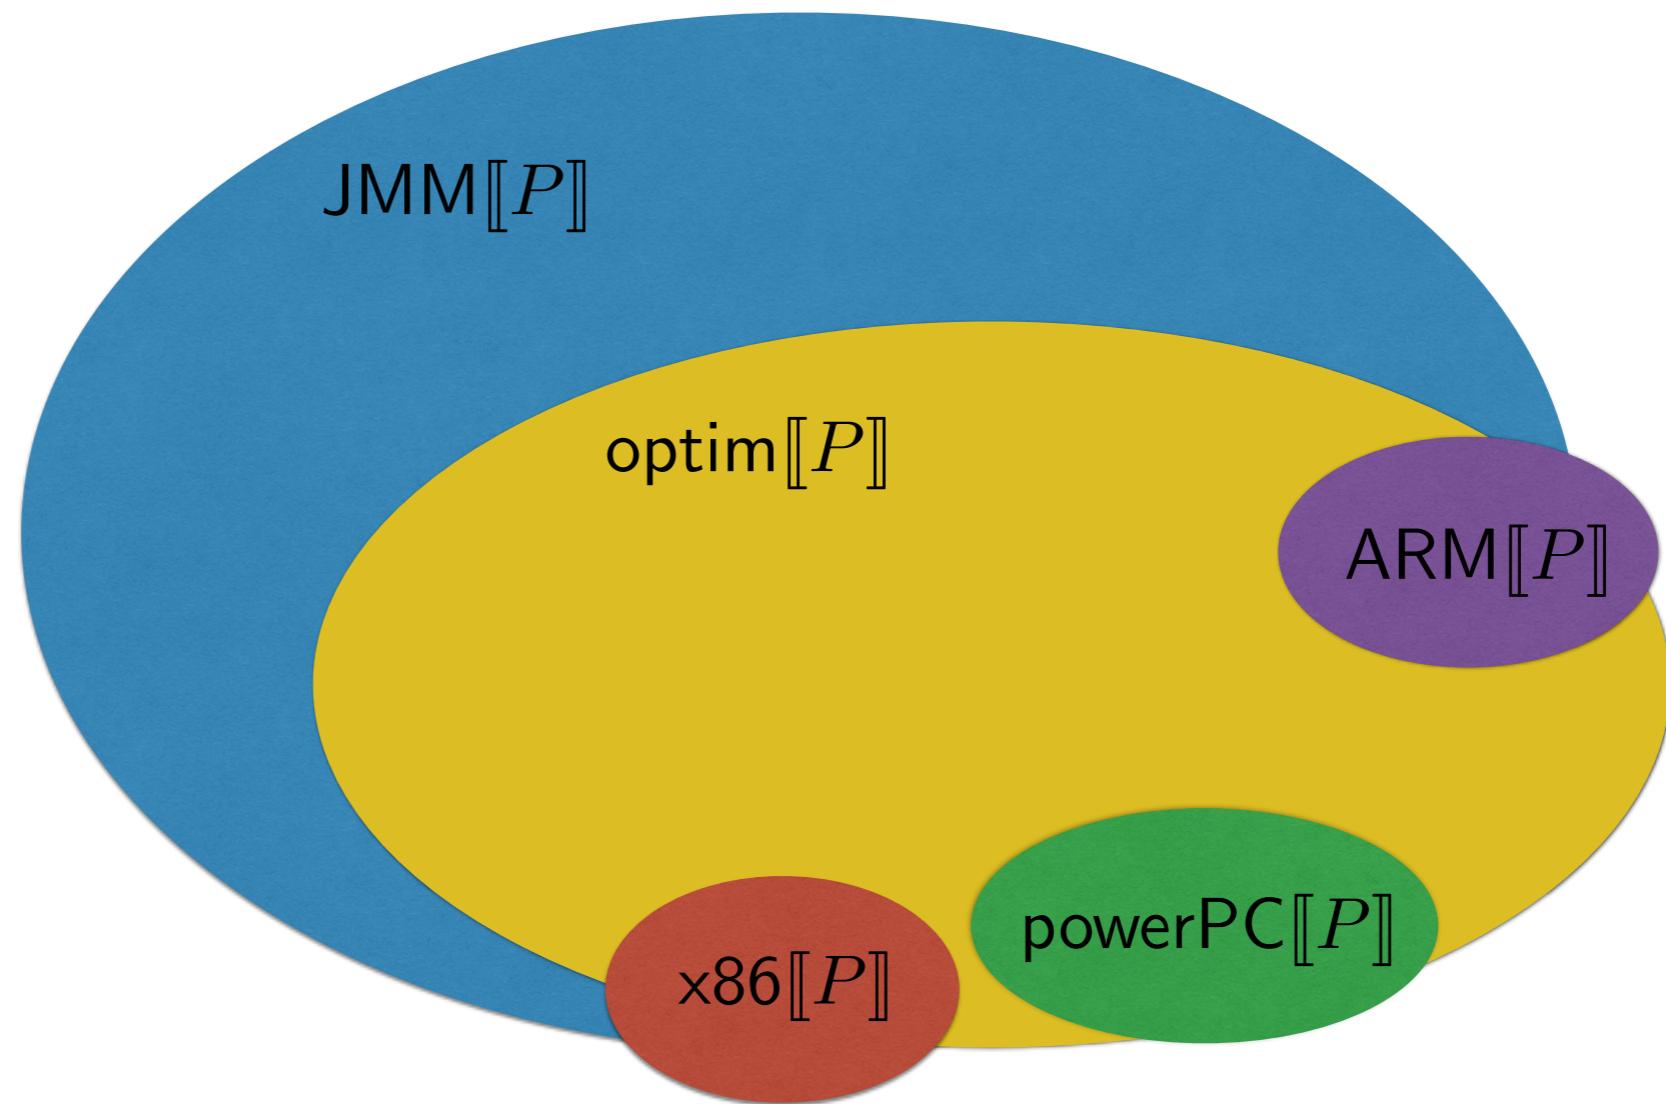
# Implementing the JMM



Given a Source Java Program  $P$  we expect



Given a Source Java Program  $P$  we expect



# The JSR-133 Cookbook for Compiler Writers

<http://g.oswego.edu/dl/jmm/cookbook.html>

## The JSR-133 Cookbook for Compiler Writers

by [Doug Lea](#), with help from members of the [JMM mailing list](#).

[dl@cs.oswego.edu](mailto:dl@cs.oswego.edu).

**Preface:** Over the 10+ years since this was initially written, many processor and language memory model specifications and issues have become clearer and better understood. And many have not. While this guide is maintained to remain accurate, it is incomplete about some of these evolving details. For more extensive coverage, see especially the work of Peter Sewell and the [Cambridge Relaxed Memory Concurrency Group](#)

This is an unofficial guide to implementing the new [Java Memory Model \(JMM\)](#) specified by [JSR-133](#). It provides at most brief backgrounds about why various rules exist, instead concentrating on their consequences for compilers and JVMs with respect to instruction reorderings, multiprocessor barrier instructions, and atomic operations. It includes a set of recommended recipes for complying to JSR-133. This guide is "unofficial" because it includes interpretations of particular processor properties and specifications. We cannot guarantee that the interpretations are correct. Also, processor specifications and implementations may change over time.

### Reorderings

For a compiler writer, the JMM mainly consists of rules disallowing reorderings of certain instructions that access fields (where "fields" include array elements) as well as monitors (locks).

### Volatile and Monitors

The main JMM rules for volatiles and monitors can be viewed as a matrix with cells indicating that you cannot reorder instructions associated with particular sequences of bytecodes. This table is not itself the JMM specification; it is just a useful way of viewing its main consequences for compilers and runtime systems.

Can Reorder	2nd operation			
	1st operation	Normal Load Normal Store	Volatile Load MonitorEnter	Volatile Store MonitorExit
Normal Load				No
Normal Store				
Volatile Load				

# The JSR-133 Cookbook for Compiler Writers

- A collection of informal arguments about what is possible under the JMM (Java) – *Roach-Motel Semantics* for Volatiles and Locks

Can Reorder	2nd operation			
	1st operation	Normal Load Normal Store	Volatile Load MonitorEnter	Volatile Store MonitorExit
Normal Load Normal Store				No
Volatile Load MonitorEnter	No	No	No	
Volatile store MonitorExit		No	No	

# The JSR-133 Cookbook for Compiler Writers

- A collection of informal “recipes” indicating how disallow unwanted *reorderings* by adding generic barrier instructions

Required barriers	2nd operation			
1st operation	Normal Load	Normal Store	Volatile Load MonitorEnter	Volatile Store MonitorExit
Normal Load				LoadStore
Normal Store				StoreStore
Volatile Load MonitorEnter	LoadLoad	LoadStore	LoadLoad	LoadStore
Volatile Store MonitorExit			StoreLoad	StoreStore

# The JSR-133 Cookbook for Compiler Writers

- A final table indicating how to implement each of the generic barriers in the different architectures

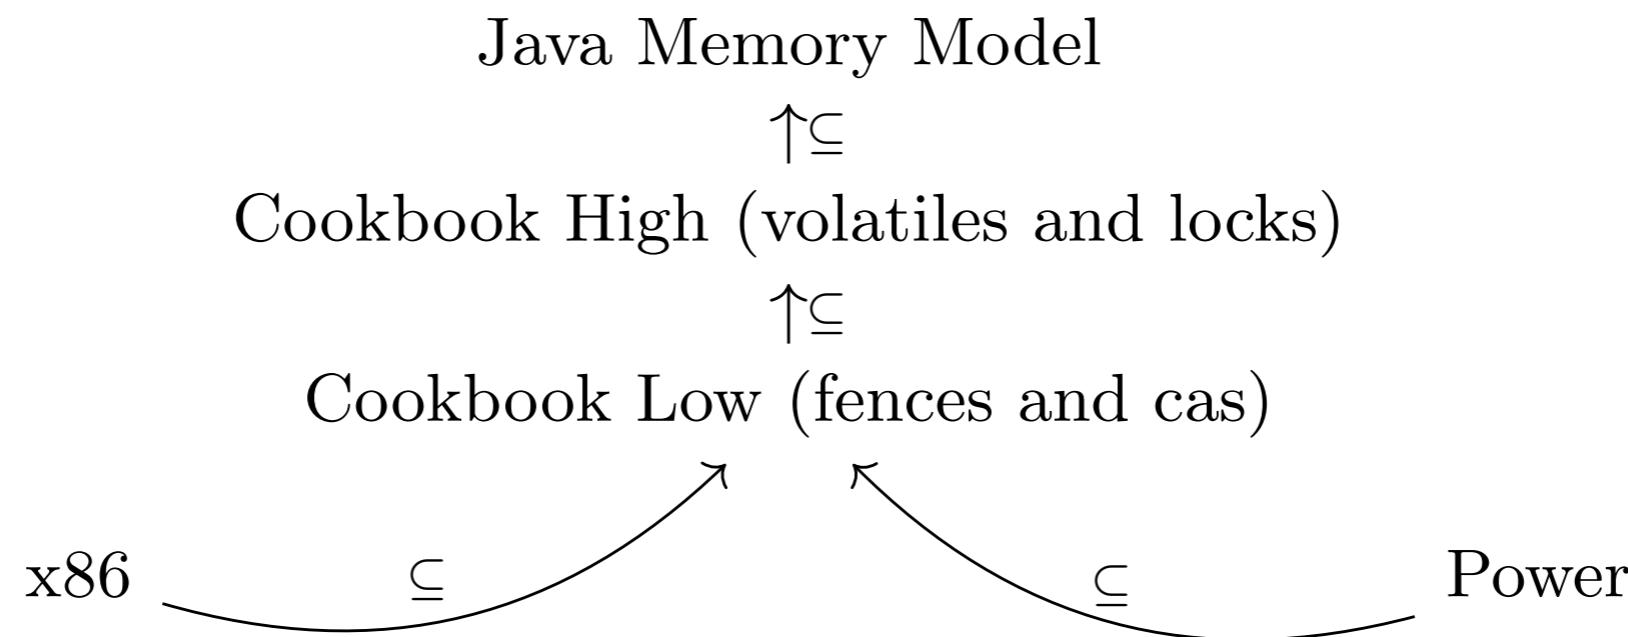
Processor	LoadStore	LoadLoad	StoreStore	StoreLoad	Data dependency orders loads?	Atomic Conditional	Other Atomics	Atomics provide barrier?
sparc-TSO	no-op	no-op	no-op	membar (StoreLoad)	yes	CAS: casa	swap, ldstub	full
x86	no-op	no-op	no-op	mfence or cpuid or locked insn	yes	CAS: cmpxchg	xchg, locked insn	full
ia64	<i>combine with st.rel or ld.acq</i>	ld.acq	st.rel	mf	yes	CAS: cmpxchg	xchg, fetchadd	target + acq/rel
arm	dmb (see below)	dmb (see below)	dmbsync	dmb	indirection only	LL/SC: ldrex/strex		target only
ppc	lwsync (see below)	lwsync (see below)	lwsync	hwsync	indirection only	LL/SC: ldarx/stwcx		target only
alpha	mb	mb	wmb	mb	no	LL/SC: idx_1/stx_c		target only
pa-risc	no-op	no-op	no-op	no-op	yes	<i>build from ldcw</i>	ldcw	(NA)

# The JSR-133 Cookbook for Compiler Writers

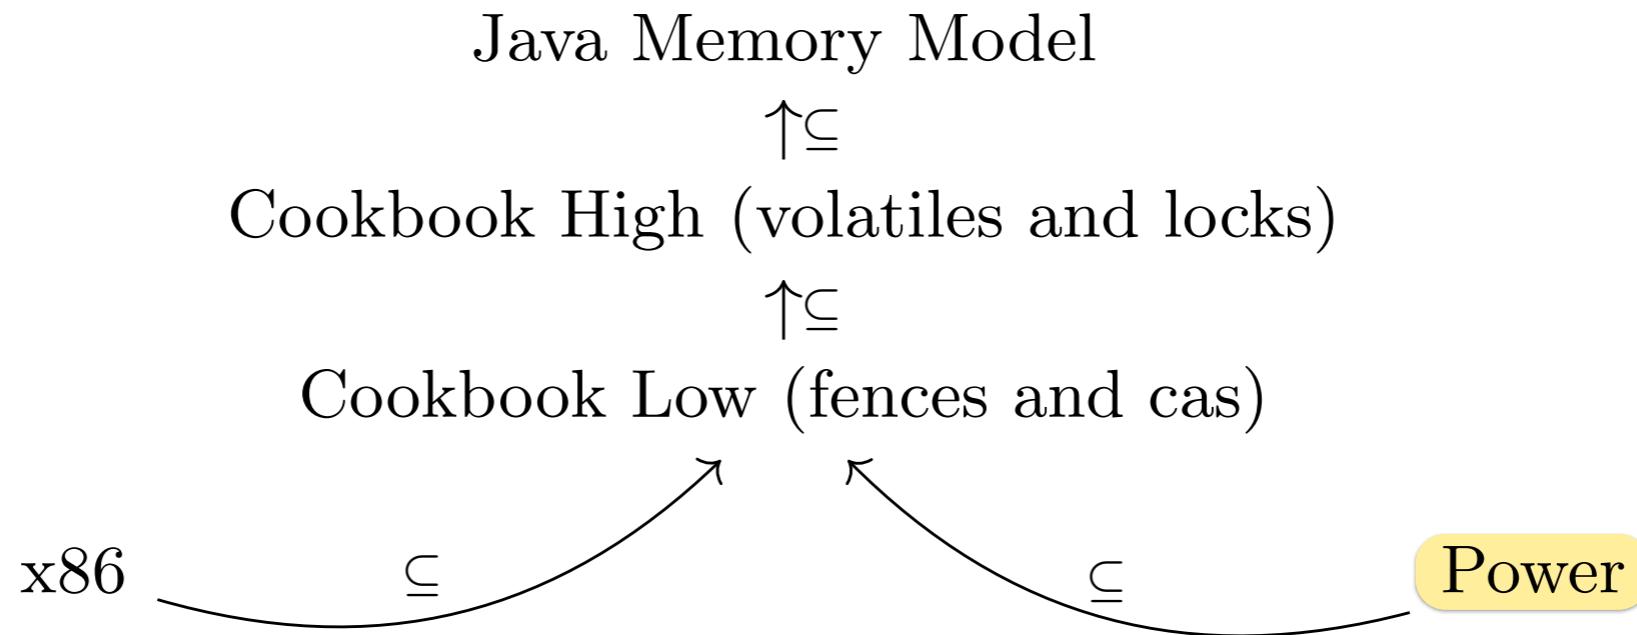
- A final table indicating how to implement each of the generic barriers in the different architectures

Processor	LoadStore	LoadLoad	StoreStore	StoreLoad	Data dependency orders loads?	Atomic Conditional	Other Atomics	Atomics provide barrier?
sparc-TSO	no-op	no-op	no-op	membar (StoreLoad)	yes	CAS: casa	swap, ldstub	full
x86	no-op	no-op	no-op	mfence or cpuid or locked insn	yes	CAS: cmpxchg	xchg, locked insn	full
ia64	<i>combine with st.rel or ld.acq</i>	ld.acq	st.rel	mf	yes	CAS: cmpxchg	xchg, fetchadd	target + acq/rel
arm	dmb (see below)	dmb (see below)	dmb-st	dmb	indirection only	LL/SC: ldrex/strex		target only
ppc	lwsync (see below)	lwsync (see below)	lwsync	hwsync	indirection only	LL/SC: ldarx/stwcx		target only
alpha	mb	mb	wmb	mb	no	LL/SC: ldx_l/stx_c		target only
pa-risc	no-op	no-op	no-op	no-op	yes	<i>build from ldcw</i>	ldcw	(NA)

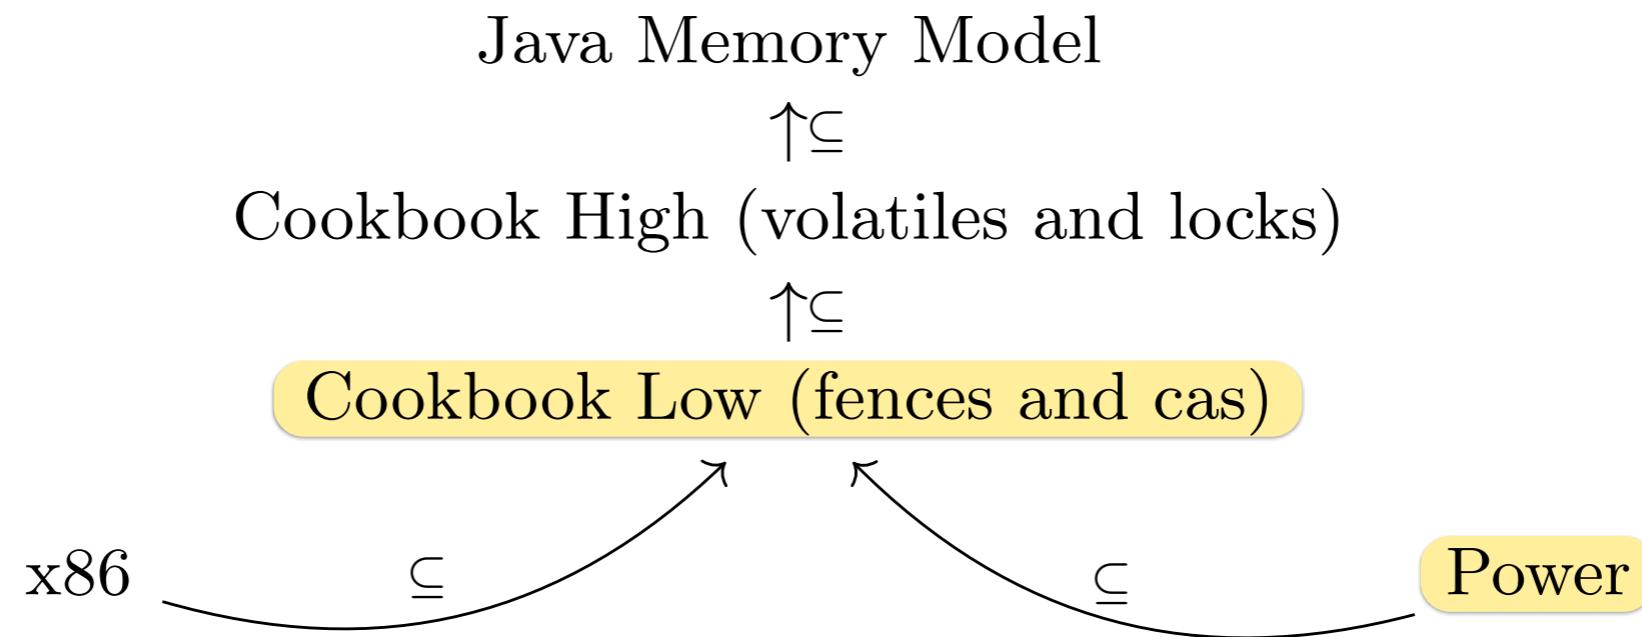
# A Semantical Framework to Verify the Cookbook



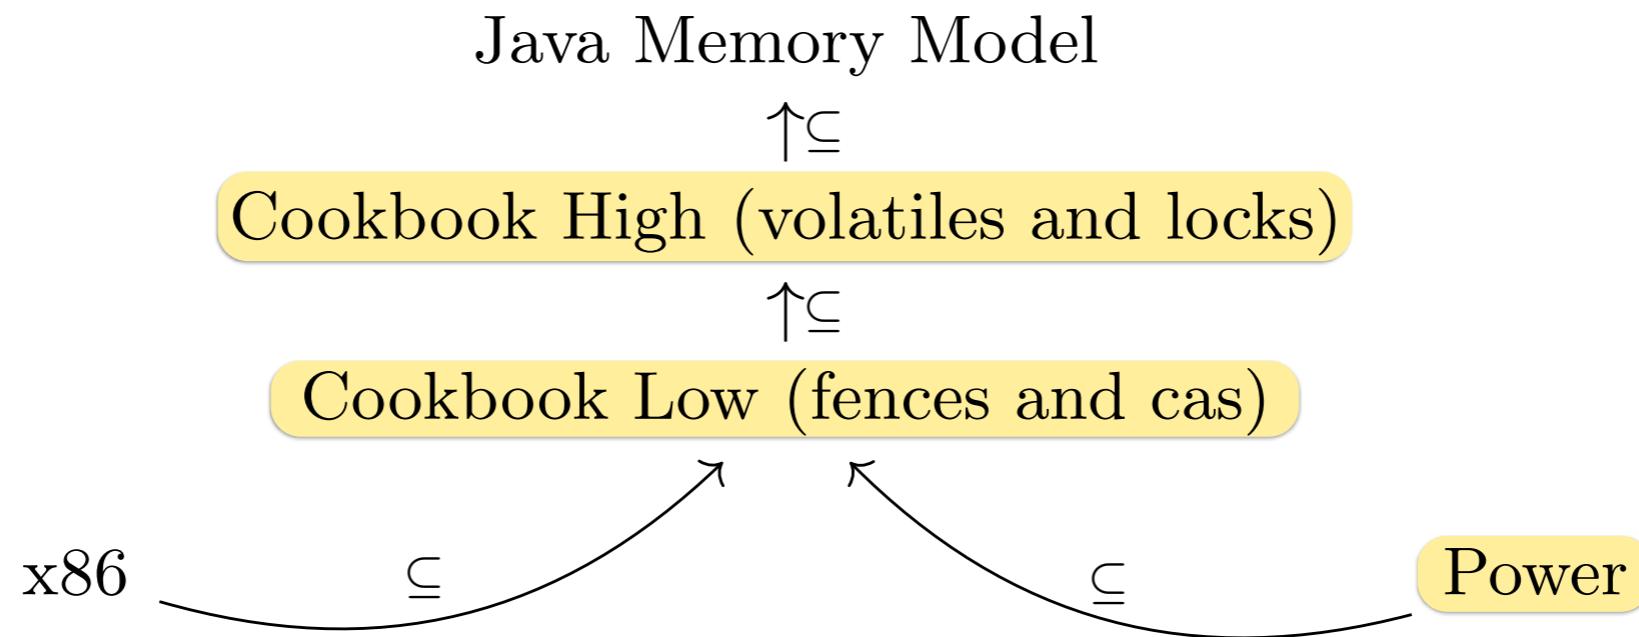
# A Semantical Framework to Verify the Cookbook



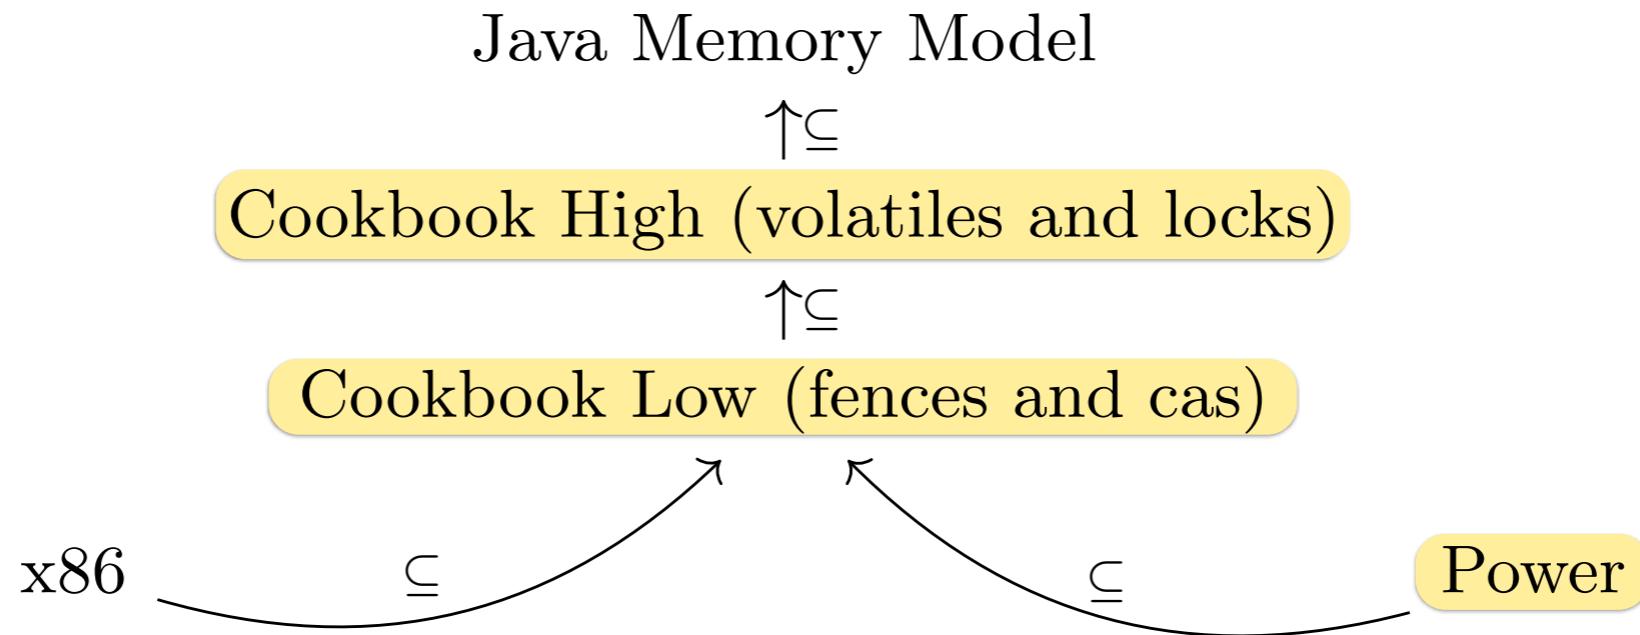
# A Semantical Framework to Verify the Cookbook



# A Semantical Framework to Verify the Cookbook

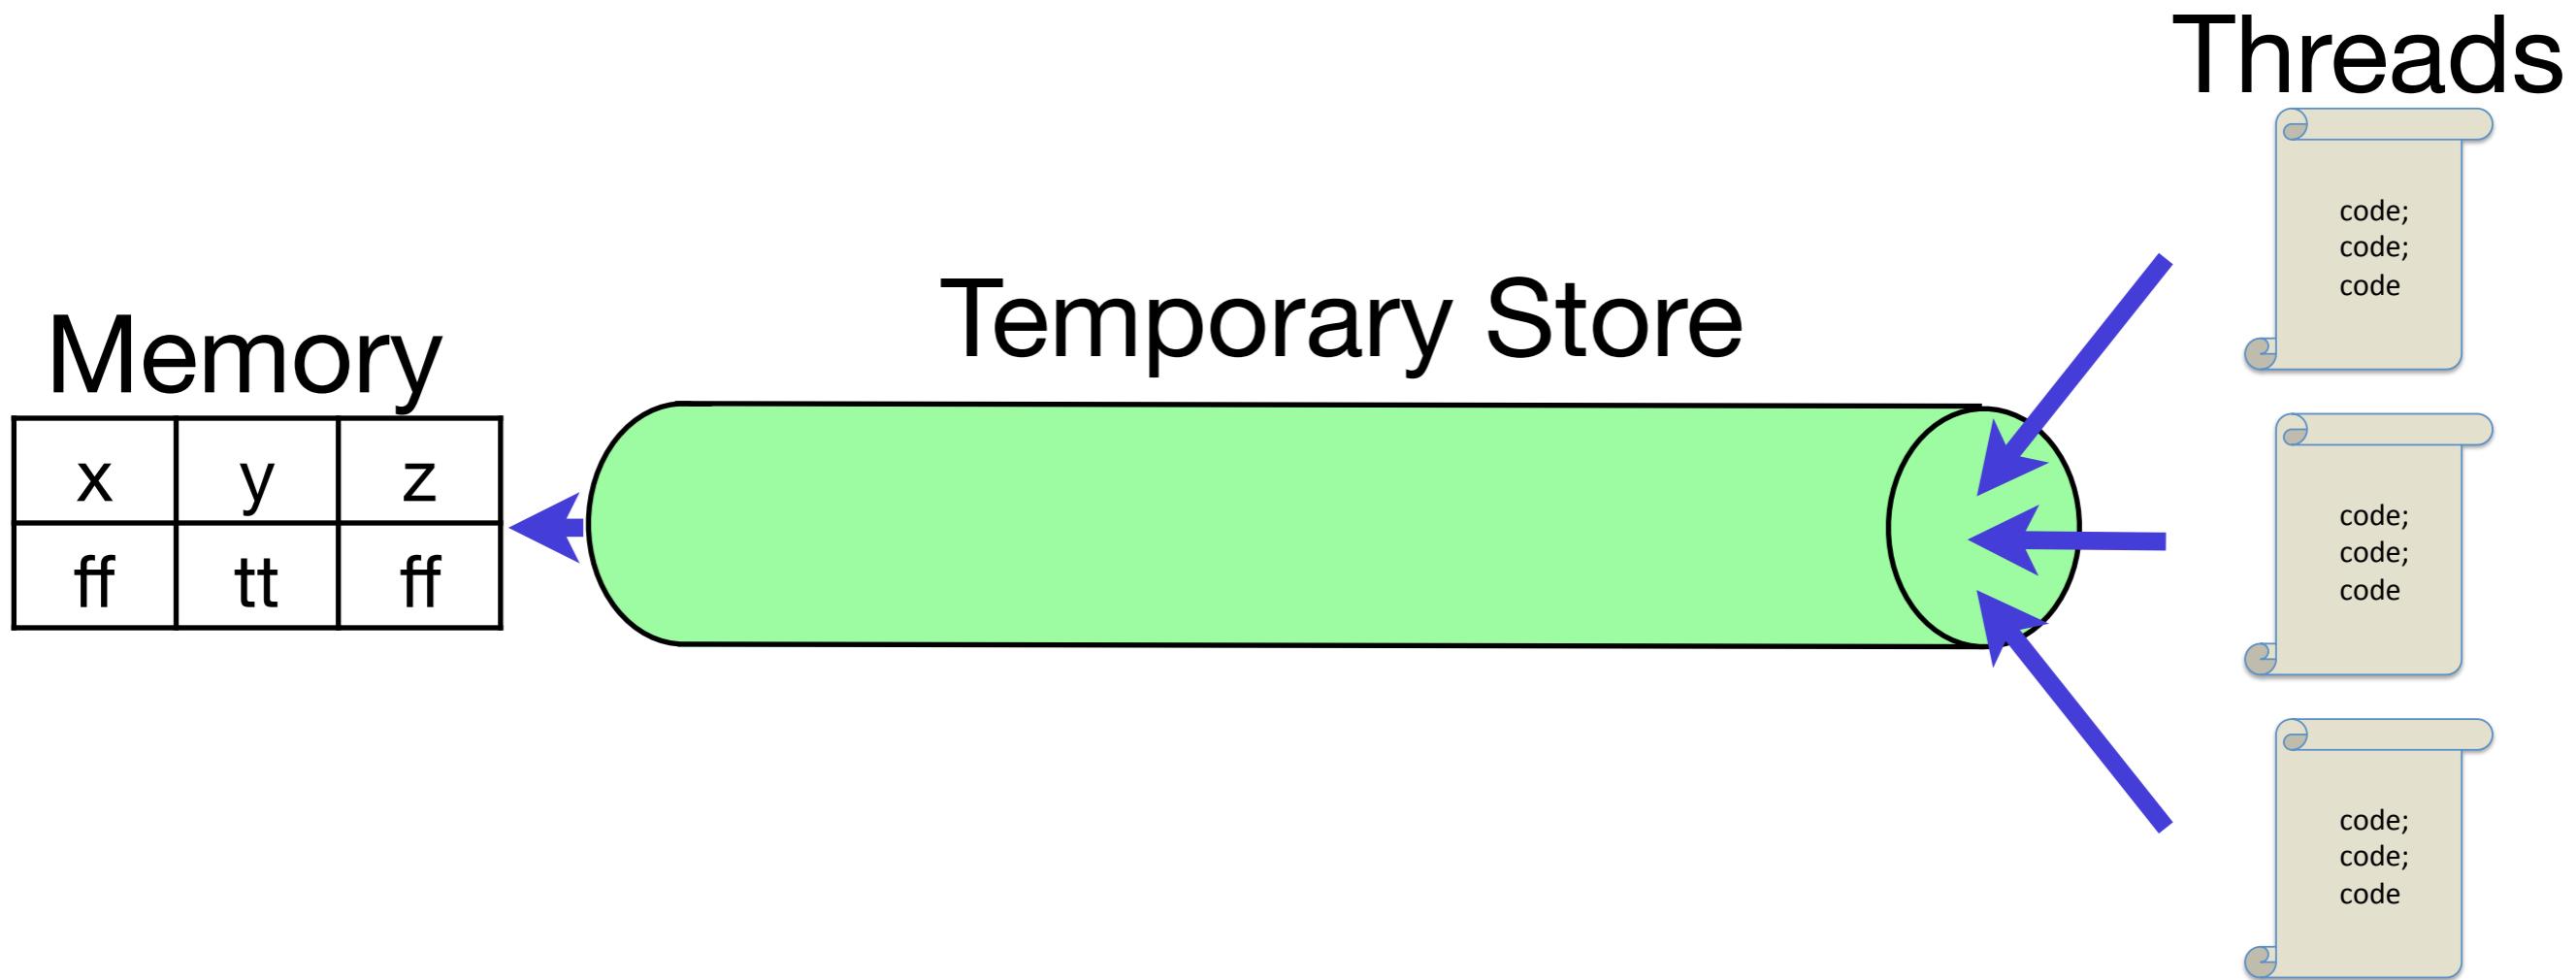


# A Semantical Framework to Verify the Cookbook



- We propagate the semantics of Power for *Normal Variables*
- We Give SC semantics to *Volatiles* and *Locks*

# Use the Pipeline Framework



- Encode the different architectures: TSO, PPC
- Encode Cookbook Low and Cookbook High
- Prove simulation at every step
- Prove that Cookbook High is a correct implementation of JMM

# Power Barriers

- **sync:** Enforces ordering between all operations
  - In our pipeline no action can bypass a sync action of the same thread
  - It requires global consensus
- **lwsync:** It enforces ordering between all operations except: **wr lwsync rd**
  - It does not enforce consensus

$$\left[ \begin{array}{l} x := 1; \end{array} \right] \parallel \left[ \begin{array}{l} y := 1; \end{array} \right] \parallel \left[ \begin{array}{l} r_1 := x; \\ \text{lwsync}; \\ r_2 := y; \end{array} \right] \parallel \left[ \begin{array}{l} r_3 := y; \\ \text{lwsync}; \\ r_4 := x; \end{array} \right]$$

$r_1 = r_3 = 0 \ \& \ r_2 = r_4 = 1$  is allowed

# The JSR-133 Cookbook for Compiler Writers

- A final table indicating how to implement each of the generic barriers in the different architectures

Processor	LoadStore	LoadLoad	StoreStore	StoreLoad	Data dependency orders loads?	Atomic Conditional	Other Atomics	Atomics provide barrier?
sparc-TSO	no-op	no-op	no-op	membar (StoreLoad)	yes	CAS: casa	swap, ldstub	full
x86	no-op	no-op	no-op	mfence or cpuid or locked insn	yes	CAS: cmpxchg	xchg, locked insn	full
ia64	<i>combine with st.rel or ld.acq</i>	ld.acq	st.rel	mf	yes	CAS: cmpxchg	xchg, fetchadd	target + acq/rel
arm	dm <sup>b</sup> <i>(see below)</i>	dm <sup>b</sup> <i>(see below)</i>	dm <sup>b</sup> -st	dm <sup>b</sup>	indirection only	LL/SC: ldrex/strex		target only
ppc	lwsync <i>(see below)</i>	lwsync <i>(see below)</i>	lwsync	hw <sup>s</sup> ync	indirection only	LL/SC: ldarx/stwcx		target only
alpha	mb	mb	wmb	mb	no	LL/SC: idx_l/stx_c		target only
pa-risc	no-op	no-op	no-op	no-op	yes	<i>build from ldcw</i>	ldcw	(NA)

# IRIW with Volatiles

$v_0 = v_1 = 0$  and  $v_0, v_1$  are volatile

$[v_0 := 1;] \parallel [v_1 := 1;] \parallel \begin{bmatrix} r_1 := v_0; \\ r_2 := v_1; \end{bmatrix} \parallel \begin{bmatrix} r_3 := v_1; \\ r_4 := v_0; \end{bmatrix}$

$r_1 = r_3 = 0 \ \& \ r_2 = r_4 = 1$  is allowed

# IRIW with Volatiles

$v_0 = v_1 = 0$  and  $v_0, v_1$  are volatile  
 $[v_0 := 1;] \parallel [v_1 := 1;] \parallel [r_1 := v_0; r_2 := v_1;] \parallel [r_3 := v_1; r_4 := v_0;]$   
 $r_1 = r_3 = 0 \& r_2 = r_4 = 1$  is allowed

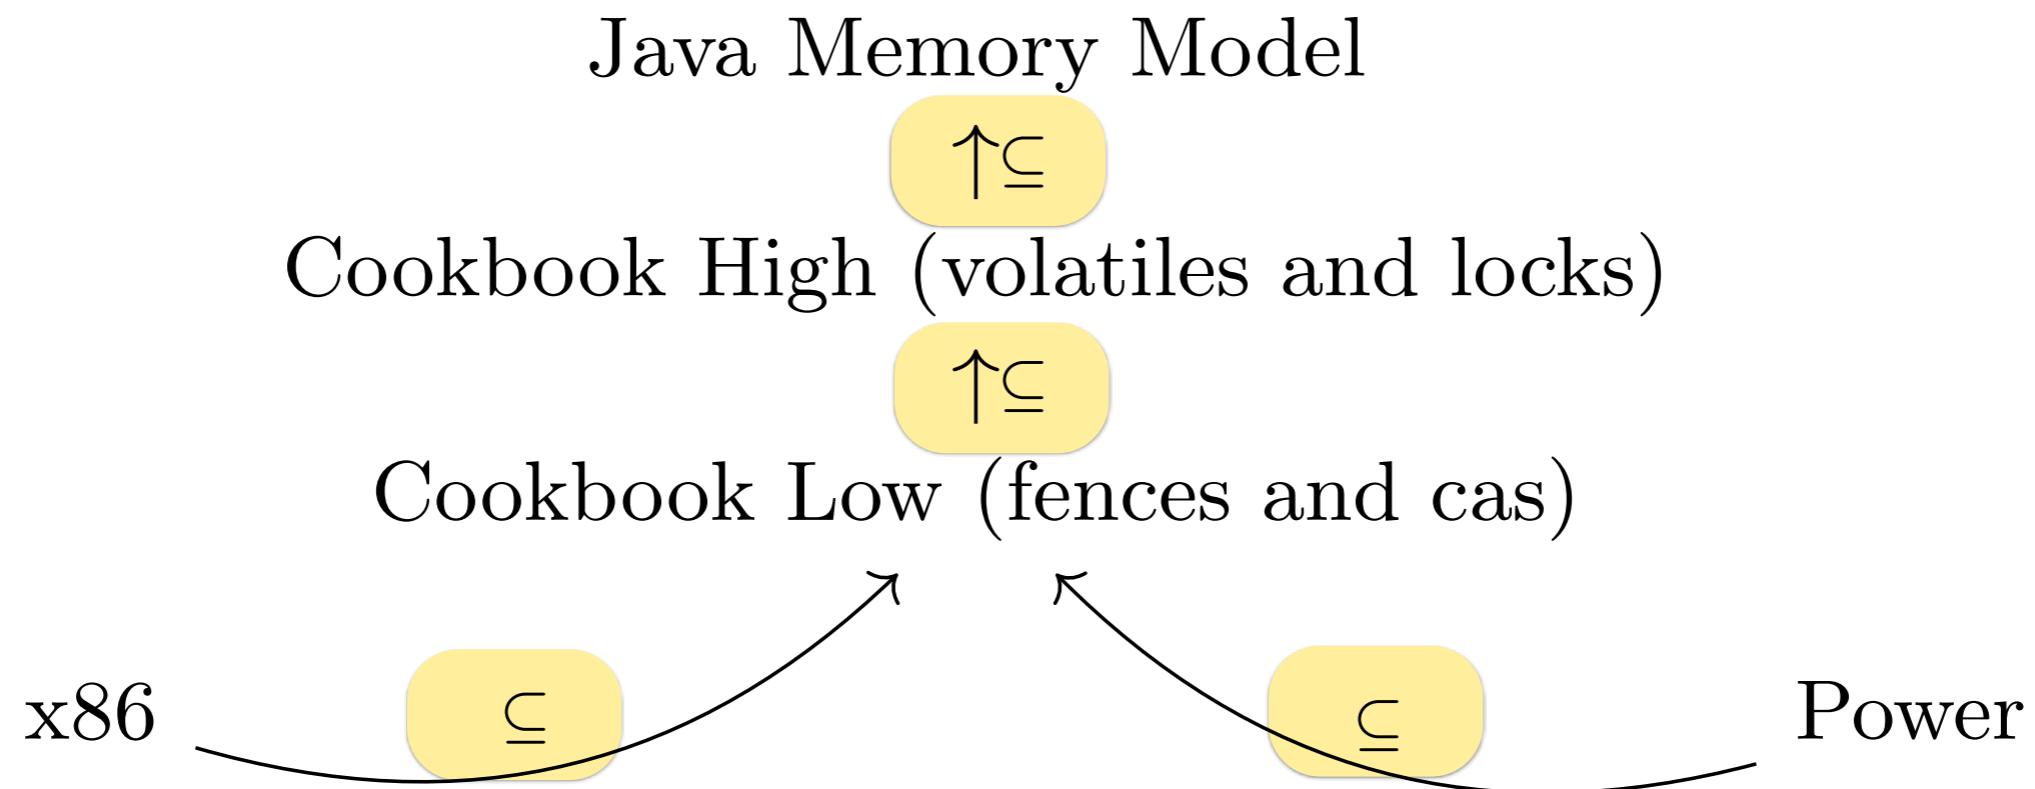
Fix

PPC(LoadLoad)	=	sync
PPC(LoadStore)	=	lwsync
PPC(StoreStore)	=	lwsync
PPC(StoreLoad)	=	sync

# Surprises

- IRIW with volatiles is reproducible on a 32-core Power7 machine
  - With two different VM's
- We found bugs in the implementation of volatiles of the Fiji ahead of time compiler
- Many implementers of JVM's follow the Cookbook

# Guarantees



**End of day 2**